


A large, red, textured, five-pointed star shape is centered on a dark blue background. The star has a rough, organic texture, resembling a biological or geological form. It is slightly tilted and has a subtle gradient of red, with some darker areas in the recesses and lighter areas on the points.

Lighting and Shading


Bo Zhu



School of Interactive Computing
Georgia Institute of Technology



NVIDIA Marbles at Night: https://www.youtube.com/watch?v=NgcYLvlp_k





Motivational Video: Art Lesson for Kids: Shading a Sphere





The Illusion of Light on a Sphere

- Light is how we see and shading informs us of the light within a scene.
- Light values are called tints and dark values are called shades.
- Contrast deals with difference.





Steps of Shading a Sphere

- Draw a **hollow circle** with a light line.
- Choose a **light source** for this sphere.
- Begin **shading** by starting opposite the light source and move inwards.
- As you move towards the light source, **fill the circle** with less and less pressure, heavier on the dark area, lighter on the light area.

How do we let computers to shade a sphere?

Shading a surface with math and GLSL




The background of the slide features five reflective spheres arranged in a circle. Starting from the top left and moving clockwise, the spheres are green, blue, red, cyan, and yellow. Each sphere has a bright, circular highlight on its upper-left surface, representing a light source. The spheres are semi-transparent, allowing the others behind them to be partially visible.

We are more ambitious than that

We also want to put specular highlights on the surface




to render shining objects

like this!



Video Credit: Logan Chang, Dartmouth 24'

Today's Study Goal: Shading a Bunny with GLSL!



Study Plan

- Vertex Normal
- Barycentric Interpolation
- Lights
- Shading Model
 - Ambient
 - Diffuse
 - Specular
 - Put Things Together: Phong Shading
- Shader Implementation

Vertex Normal


Quick Recap: Cross Product

- Geometric definition:
- $\mathbf{a} \times \mathbf{b} = |\mathbf{a}||\mathbf{b}|\sin\theta$
- right-hand rule to define direction
- Coordinate formula

$$\mathbf{a} \times \mathbf{b} := \begin{bmatrix} a_y b_z - a_z b_y \\ a_z b_x - a_x b_z \\ a_x b_y - a_y b_x \end{bmatrix}$$


related to determinant:

$$\begin{vmatrix} \mathbf{x} & \mathbf{y} & \mathbf{z} \\ a_x & a_y & a_z \\ b_x & b_y & b_z \end{vmatrix}$$





Triangle Normal

$$\vec{N} = (\vec{AB} \times \vec{AC})$$
$$\vec{N} = \vec{N} / |\vec{N}|$$



We can define a normal vector for each triangle



- Face normals: same normal for all points in face
 - geometrically correct, but faceted look



Not great for rendering a smooth surface!

Or, we can calculate a normal vector for each vertex

- Vertex normal: store normal at vertices, interpolate in face
 - geometrically “inconsistent”, but smooth look




Then how to calculate a vertex normal?

Vertex Normal


The normal vector of a vertex is the average of the normal vectors of its incident triangles

Implementation Idea: Sum all normal vectors from incident triangles, and then normalize

$$\vec{N}_v = \sum_i \vec{N}_i / \left| \sum_i \vec{N}_i \right|$$




Practice: Calculate Vertex Normal



Q: What is the normal of vertex C?


Formula of cross product:

$$\mathbf{a} \times \mathbf{b} := \begin{bmatrix} a_y b_z - a_z b_y \\ a_z b_x - a_x b_z \\ a_x b_y - a_y b_x \end{bmatrix}$$



Visualize vertex normal on a mesh?

- A nice visual trick is to draw the **xyz** of a vertex normal as **rgb** colors
- We interpolate colors within each triangle using barycentric interpolation scheme




Next question: How do we interpolate with a triangle?

Triangle Interpolation

Barycentric Interpolation

Triangle Color


- Every pixel inside a triangle projected onto the image plane requires (R,G,B) values
- If the triangle has a uniform color, one can simply color every interior pixel accordingly
- However, if a triangle has multiple colors, we need to specify the color for every pixel
- Since the pixel locations aren't known ahead of time, we need a way of specifying color throughout the interior of the triangle



Interpolation


- Convert discrete values into a continuous function by filling in all the “in between” values

Linear Interpolation $y(t) = (1 - t) y_1 + t y_2$



Barycentric Interpolation


2D/3D Edges




$$\mathbf{p} = \alpha_0 \mathbf{p}_0 + \alpha_1 \mathbf{p}_1$$
$$\alpha_0 + \alpha_1 = 1$$

2D/3D Triangles

$$\mathbf{p} = \alpha_0 \mathbf{p}_0 + \alpha_1 \mathbf{p}_1 + \alpha_2 \mathbf{p}_2$$
$$\alpha_0 + \alpha_1 + \alpha_2 = 1$$



Geometric Approach to Calculate Barycentric Weights



$$\alpha_0 = \frac{\text{area}(p \ p_1 \ p_2)}{\text{area}(p_0 \ p_1 \ p_2)}$$

$$\alpha_1 = \frac{\text{area}(p_0 \ p \ p_2)}{\text{area}(p_0 \ p_1 \ p_2)}$$

$$\alpha_2 = \frac{\text{area}(p_0 \ p_1 \ p)}{\text{area}(p_0 \ p_1 \ p_2)}$$

$$\text{area}(A, B, C) = \frac{|AB \times AC|}{2}$$


Algebraic Approach to Calculate Barycentric Weights

$$\alpha_0 \mathbf{p}_0 + \alpha_1 \mathbf{p}_1 + \alpha_2 \mathbf{p}_2 = \mathbf{p}$$

$$\alpha_0 \begin{pmatrix} x_0 \\ y_0 \end{pmatrix} + \alpha_1 \begin{pmatrix} x_1 \\ y_1 \end{pmatrix} + (1 - \alpha_0 - \alpha_1) \begin{pmatrix} x_2 \\ y_2 \end{pmatrix} = \begin{pmatrix} x \\ y \end{pmatrix}$$

Solve a 2x2 matrix equation:


$$\begin{pmatrix} x_0 - x_2 & x_1 - x_2 \\ y_0 - y_2 & y_1 - y_2 \end{pmatrix} \begin{pmatrix} \alpha_0 \\ \alpha_1 \end{pmatrix} = \begin{pmatrix} x \\ y \end{pmatrix} - \begin{pmatrix} x_2 \\ y_2 \end{pmatrix}$$



Key Idea: If we interpolate the three vertices of a triangle using barycentric interpolation, we should be able to obtain the position of the current interpolation point.

Interpolating Vertex Data

- Interpolate colors
 - $R = \alpha_0 R_0 + \alpha_1 R_1 + \alpha_2 R_2$
 - $G = \alpha_0 G_0 + \alpha_1 G_1 + \alpha_2 G_2$
 - $B = \alpha_0 B_0 + \alpha_1 B_1 + \alpha_2 B_2$
- Interpolate normal vectors
 - $N = \alpha_0 N_0 + \alpha_1 N_1 + \alpha_2 N_2$
- Interpolate texture coordinates $u = \alpha_0 u_0 + \alpha_1 u_1 + \alpha_2 u_2$
 - $v = \alpha_0 v_0 + \alpha_1 v_1 + \alpha_2 v_2$




This step happens in the rendering pipeline between the vertex shader stage and the fragment shader stage --- but it's not programmable!

- We don't need to implement our barycentric interpolation, all the vertex->fragment interpolation within each triangle happen automatically.

Lights

Different Types of Light

- Ambient light
- Directional light
- Point light



Ambient Lights


- Ambient lighting models a constant illumination independent of the incident light angle
- Useful for adding some light in the shadowed regions that would otherwise potentially be completely black



Directional lights


- Light is emitted from infinity in **one direction** \vec{d}
- Simulate distance lighting, e.g. sun, same at all surface points \vec{p}
- Light direction: $\vec{l} = -\vec{d}$

Note the negative sign here:
we typically use vector \vec{l} to
denote the vector pointing
from a surface point to the
light source




Point lights

- Light is emitted equally from point \vec{s} in **all directions**
- Lighting different at each surface point \vec{p}
- Light direction: $\vec{l} = (\vec{s} - \vec{p}) / |\vec{s} - \vec{p}|$





Again, here \vec{l} is defined as a vector pointing from a surface point to the light source



Shading Model


What is a shading model

- A shading model simulate the way light interacts with object surfaces
 - They define how lights and shades appear on the surface of the object
- A few key properties we need to consider when devising a shading model:
 - Type of lights
 - Surface material properties
 - Light-material interaction
 - View dependency



Mathematically, how do we describe a shading model?


- For a surface point, we typically consider the relation among **eye**, **light source**, and **surface normal** to calculate a light model



Let's start with a simple **ambient model**

Why do we distinguish k_a and I_a ?

- We use k_a for material property and I_a for light property.
- The final color is the **component-wise multiplication** of light color and material color.



light source
ambient intensity

↓

$$L_a = k_a I_a$$

↑ ↑

ambient color material ambient
 coefficient

Each color component is treated in separate!

Visual Intuition

- You see some object in red, it could be because
 - (1) the object's surface material is red, or
 - (2) the object is casted by some red light

It is difficult to distinguish these two cases visually



Component-wise Multiplication

- Both k_a and I_a are vec3
- We conduct multiplication for each component, i.e.:



$$L_a = k_a I_a \quad \longrightarrow \quad L_a = (k_{ar} * I_{ar}, k_{ag} * I_{ag}, k_{ab} * I_{ab})$$

- Component-wise multiplication is neither dot nor cross product
- It can be simply implemented with * operator in GLSL, e.g.,


```
vec3 a=vec3(1., 0.5, 0.3);  
vec3 b=vec3(0.2, 0.4, 0.2);  
vec3 c=a*b;          ////c=vec3(0.2, 0.2, 0.06);
```

Next, let's make it a bit more complicated by adding a **Diffuse Model**

- The diffuse model models a very rough surface with lots of microfacets that randomly reflect incoming light in **every** outgoing direction
- The shading does not depend on the position of the camera/viewer
- The shading still depends on the position of the light
 - Light sources on the top or on the side will produce different colors on the object's surface
- This is a good model for ideally diffuse/dull/matte surfaces, such as chalk




Let's illustrate a diffusive surface with simple geometry




Also called ideal **Lambertian** reflection

Lambert's Cosine Law


- Light source position matters!



Top face of cube receives a certain amount of light



Top face of 60° rotated cube intercepts half the light





In general, light per unit area is proportional to

$$\cos\theta = \vec{l} \cdot \vec{n}$$


Visual Intuition

- Objects will appear brighter when illuminated by a light source directly, and darker when the light hit them at an angle.



A light source doesn't produce color if it is on the back

- Light source position matters!
 - Needs to check if a light source is on the backside of a face




We can express this check together with Lambert's cosine law as

$$\max(0, \cos\theta) = \max(0, \vec{l} \cdot \vec{n})$$

Put everything together in one expression

- Diffusion (Lambertian) shading model:



light source
diffuse intensity

check if light is
on the back
and calculate its
contribution


$$L_d = k_d I_d \max(0, \vec{l} \cdot \vec{n})$$

diffuse color

material diffuse
coefficient


Key Takeaways for Lambertian Shading

- Light source location matters!
- Observer location does not matter
- Need to check if the light source is behind the surface (by checking the dot product of surface normal and light direction)




Our Next Goal: **Specular Highlights**

- These are effectively just “blurry” reflections of the light source




Let's illustrate a specular surface with simple geometry



Light intensity depends on view direction:

brighter around the mirror direction and darker elsewhere

Specular Highlights (Phong)



The larger the index value, the smaller the highlight region

light specular contribution according to the view

Power index controlling the size of the highlight region

light source specular intensity

shininess

$$L_s = k_s I_s \max(0, \vec{v} \cdot \vec{r})^p$$

specular color


material specular coefficient

$\cos \alpha$

Indicates how close the mirror direction and view direction are together

How to calculate the mirror reflection vector?

- Mirror - ideally smooth and specular surface
- Law of reflection



$$\theta_i = \theta_o \quad \textcircled{1} \quad \vec{r} = -\vec{l} + 2 \left(\vec{l} \cdot \vec{n} \right) \vec{n}$$

Attention: The definition of l is opposite to our convention here (**you should use $-l$ as input**)


2

`vec3 reflect(vec3 l, vec3 N)`

For a given incident vector l and surface normal N **reflect** returns the reflection direction calculated as $l - 2.0 * \text{dot}(N, l) * N$.

N should be normalized in order to achieve the desired result.


- Incident light direction \vec{l} , surface normal vector \vec{n} , and reflected light direction \vec{r} are all coplanar
- In GLSL we implement reflection vector with built-in **reflect()**



Key Takeaways for Specular Shading

- Both light source location and viewer location matter!
- Observe brighter color when closer to the mirror direction
- The shininess power controls the size of the highlight region

Putting Everything Together: Phong Shading!





Ambient

Diffuse

Specular

Final Image



$$\mathbf{L}_{Phong} = \sum_{j \in lights} \underbrace{(k_a \mathbf{I}_a^j)}_{\text{Ambient}} + \underbrace{k_d \mathbf{I}_d^j \max(0, \vec{l}^j \cdot \vec{n})}_{\text{Diffuse}} + \underbrace{k_s \mathbf{I}_s^j \max(0, \vec{v} \cdot \vec{r}^j)^p}_{\text{Specular}}$$




Phong Shading

invented by Dr. Bui Tuong Phong in 1973

One of the greatest shading models in the history of computer graphics!



Phong Shading has enabled us
to render shining objects easily
with several lines of code!



Summary: Equations

- Ambient shading:



$$L_{Ambient} = \sum_{j \in lights} (k_a I_a^j)$$

- Lambertian shading:


$$L_{Lambertian} = \sum_{j \in lights} (k_a I_a^j + k_d I_d^j \max(0, \vec{l}^j \cdot \vec{n}))$$

- Phong shading:

$$L_{Phong} = \sum_{j \in lights} (k_a I_a^j + k_d I_d^j \max(0, \vec{l}^j \cdot \vec{n}) + k_s I_s^j \max(0, \vec{v} \cdot \vec{r}^j)^p)$$



Summary (Geometric Picture)



3 points
4 vectors
2 angles

Summary (Data Types)

- On the material side

```
vec3 ka;  
vec3 kd;  
vec3 ks;  
float shininess;
```

- On the light side

```
vec3 la;  
vec3 ld;  
vec3 ls;
```

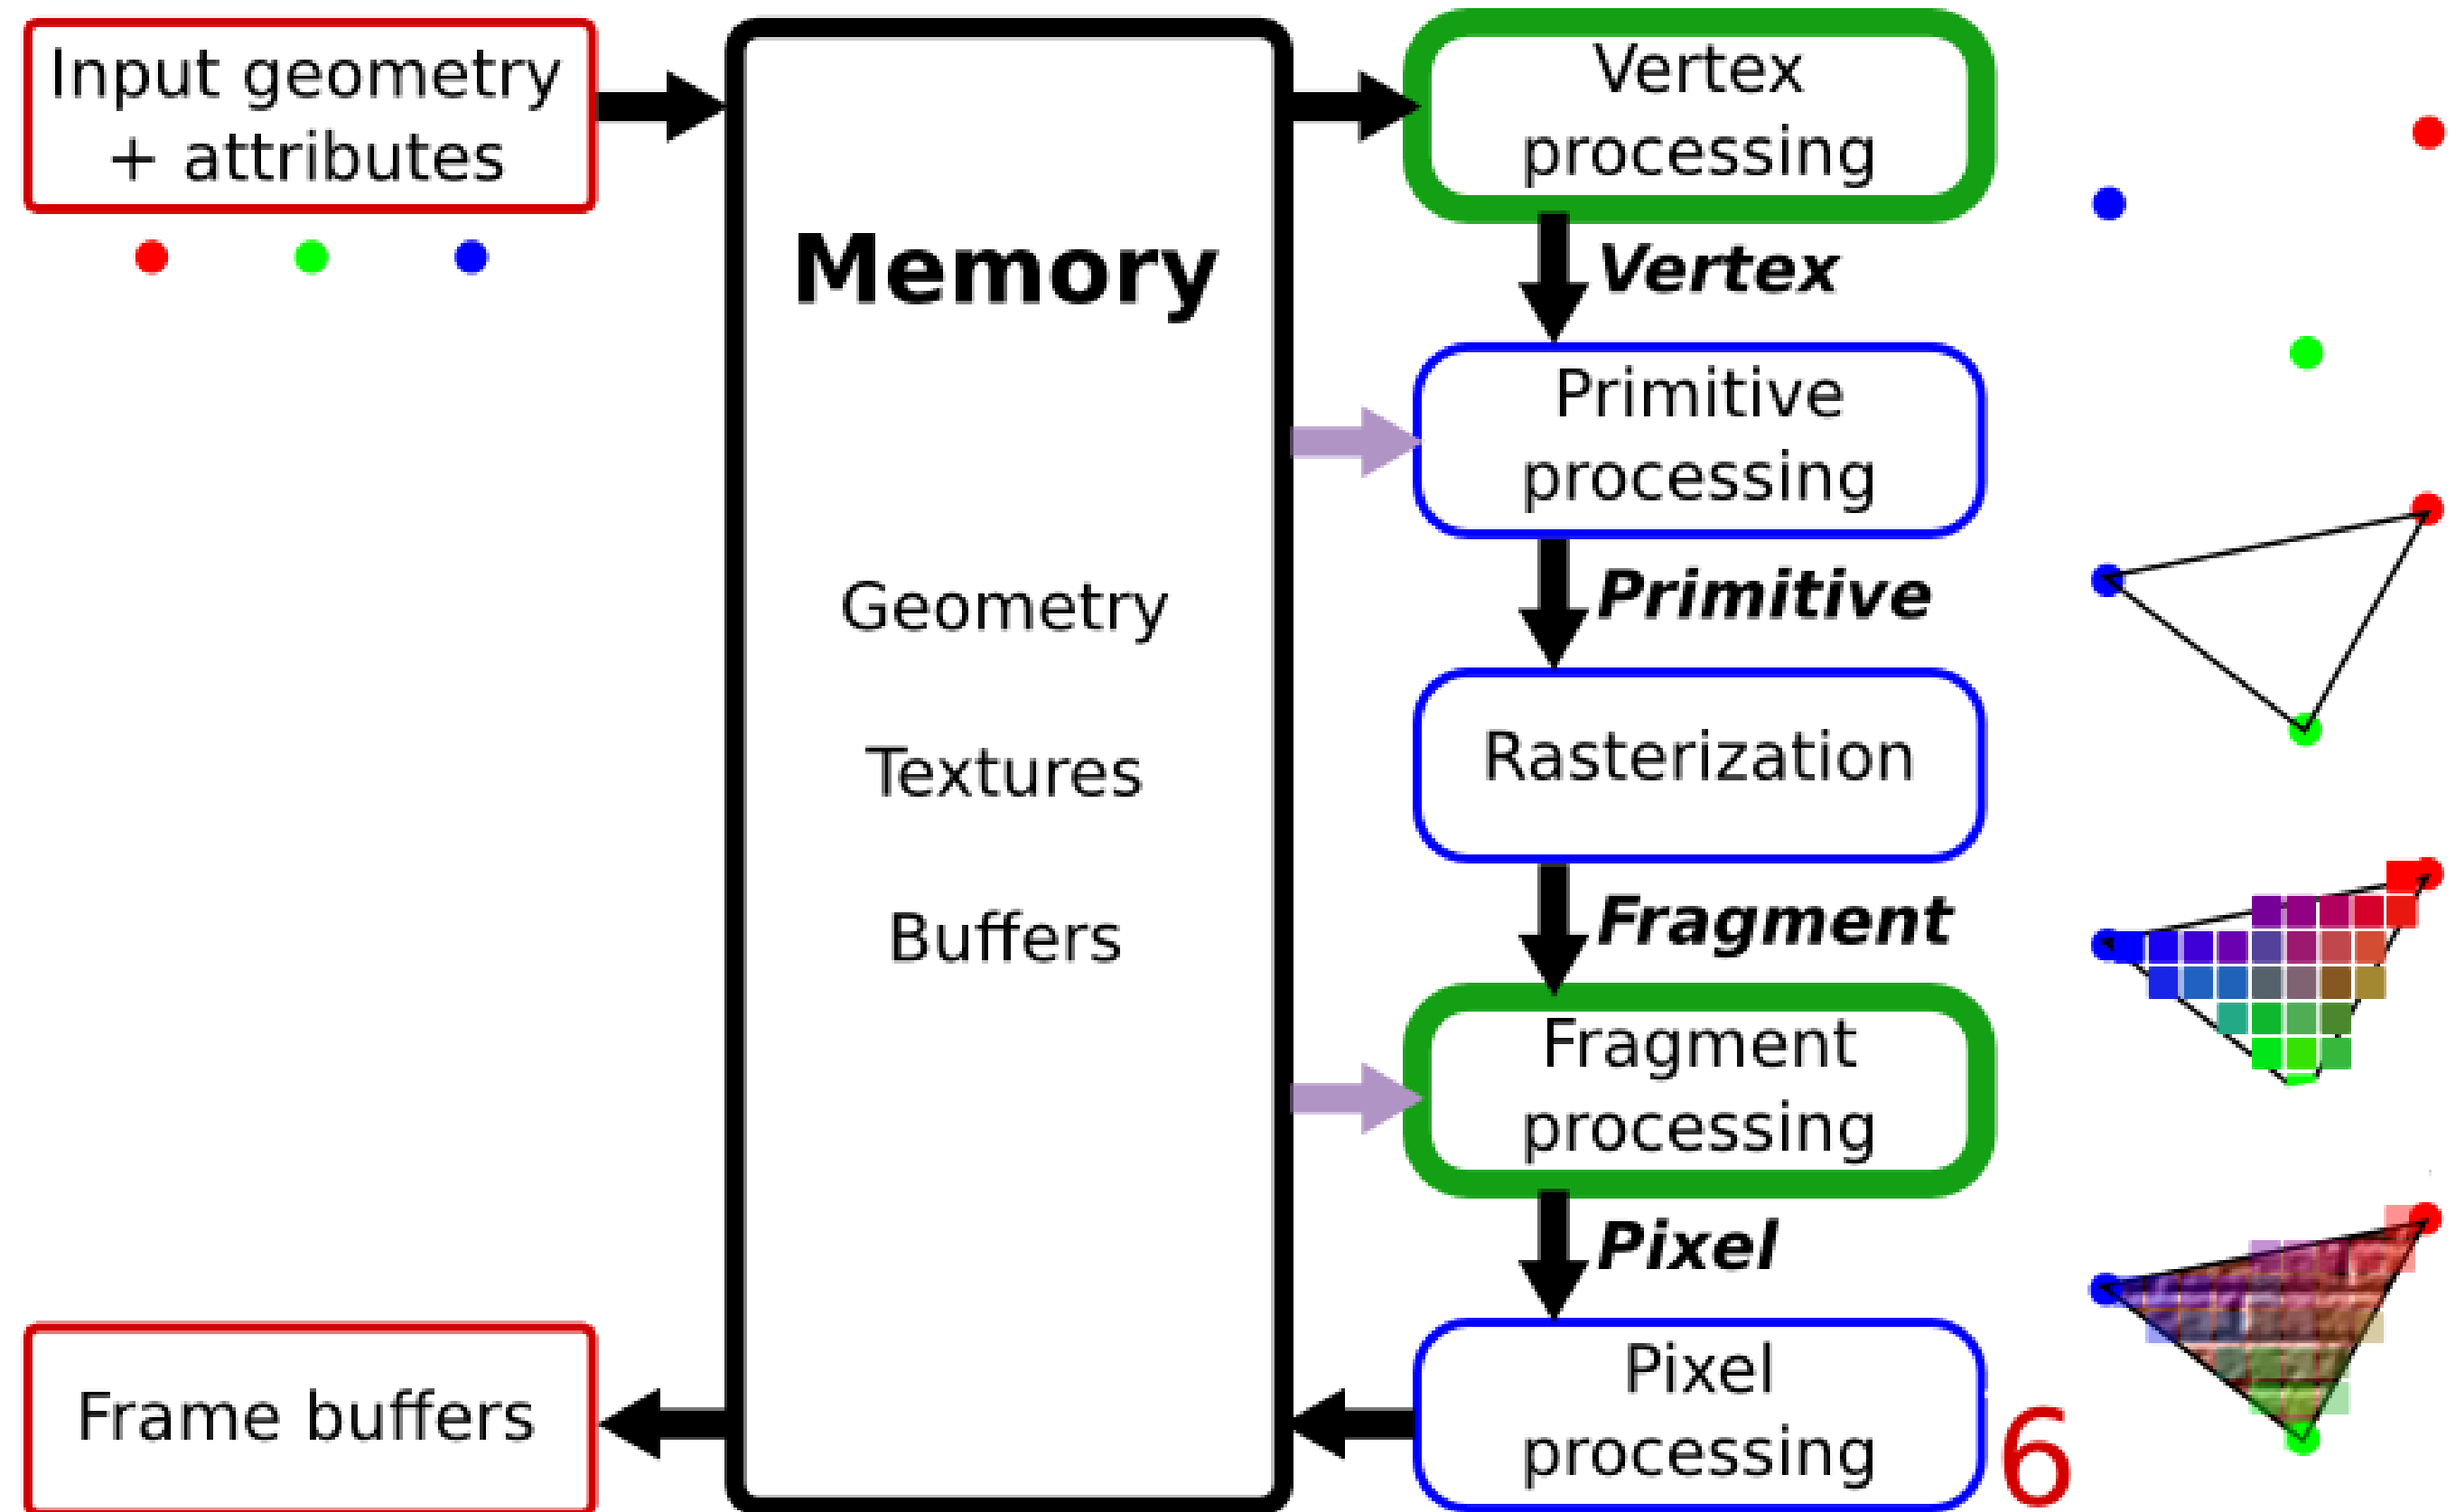
- On the geometry side


```
vec3 s;  
vec3 p;  
vec3 e;  
vec3 n;
```


GLSL Implementation

Quick Recap:

Modern Graphics Pipeline






Data Flows in Pipelines

- Consider each variable as a component of a pipeline that channels data from the CPU through the Vertex Shader and Fragment Shader, and ultimately to the Screen.
- The transfer of data within each stage is facilitated by the use of 'in' and 'out' qualifiers to denote incoming and outgoing data, respectively.
- The continuity of data between stages is maintained by using the same variable name, thus allowing for a seamless data flow from one stage to the next within the pipeline.

We need barycentric interpolation when data flows from vertex shader to fragment shader

- Interpolate colors
 - $R = \alpha_0 R_0 + \alpha_1 R_1 + \alpha_2 R_2$
 - $G = \alpha_0 G_0 + \alpha_1 G_1 + \alpha_2 G_2$
 - $B = \alpha_0 B_0 + \alpha_1 B_1 + \alpha_2 B_2$
- Interpolate normal vectors
 - $N = \alpha_0 N_0 + \alpha_1 N_1 + \alpha_2 N_2$
- Interpolate z-buffer depth values
 - $z = \alpha_0 z_0 + \alpha_1 z_1 + \alpha_2 z_2$
- Interpolate texture coordinates $u = \alpha_0 u_0 + \alpha_1 u_1 + \alpha_2 u_2$
 - $v = \alpha_0 v_0 + \alpha_1 v_1 + \alpha_2 v_2$



Data Flow in Shaders

- **Vertex shader:**

in vec4 pos;

in vec4 normal;

out vec4 vtx_pos;


out vec4 vtx_normal;

...

void main()

{...}

**Processed and
interpolated data**



Fragment shader:

in vec4 vtx_pos;

in vec4 vtx_normal;

out vec4 frag_color;

...

void main()

{...}


- Variables in the same data channel must have the name;
- out in vertex shader, and in in fragment shader.

The Simplest Vertex Shader

```
mat4 pvm;  
in vec4 pos;  
void main()  
  
{  
    gl_Position=pvm*vec4(pos.xyz,1.f);  
}
```

The Simplest Fragment Shader

```
void main()  
{  
    frag_color=vec4(1.f,0.f,0.f,1.f);  
}
```




Another Example of Vertex Shader

```
mat4 pvm;  
in vec4 pos;  
in vec4 normal;  
out vec3 vtx_normal;  
void main()  
{  
    gl_Position=pvm*vec4(pos.xyz,1.f);  
    vtx_normal=normal.xyz;          ////assuming model matrix is identity  
}
```



Example Fragment Shader: Use Normal as Color

```
in vec3 vtx_normal;  
void main()  
{  
    vec3 norm=normalize(vtx_normal.xyz);  
    frag_color=vec4(norm,1.f);  
}
```



Example: Pseudocode for Phong Shading

```
in vec3 vtx_normal;  
void main()  
{  
    ...  
    vec3 norm=normalize(vtx_normal.xyz);  
    vec3 color=phong(base_color, light_dir, norm, shininess, specular);  
    frag_color=vec4(color, 1.f);  
}
```



Your implementation

Where should we put the shading model?

1. Vertex shading:

Compute illumination/shading **per vertex** in the **vertex stage**,
and **interpolate shaded color** across triangles

2. Fragment shading:

Compute illumination/shading **per pixel** in the **fragment stage**

Pipeline for Vertex Shading

- **Vertex stage** (input: position, color, and normal / vtx)
 - transform position and normal (object to eye space)
 - **compute shaded color per vertex**
 - transform position (eye to screen space)
- **Rasterizer**
 - interpolated parameters: r, g, b color
- **Fragment stage** (output: color)
 - write to pixel


**Gouraud
Shading**
shading and then
interpolation



Put on Vertex Shader: Gouraud Shading

- Can apply Gouraud shading to any illumination model
 - it's just a method to interpolate per-vertex colors
- Results are not so good for fast-varying shading
 - e.g. problems with any highlights smaller than a triangle

II.31 Shutterbug. Gouraud shaded polygons with specular reflection (Sections 14.4.6.2.5). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)




Put on Fragment Shader: Per-pixel (Phong) Shading

- Get higher quality by interpolating shading parameters (e.g. normal), instead of shaded color
- Evaluate the illumination model per pixel rather than per vertex (and normalizing the normal first)
 - in pipeline, this means we are moving illumination from the vertex processing stage to the fragment processing stage

Pipeline for Fragment Shading

- Vertex stage (input: position, color, and normal / vtx)
 - transform position and normal (object to eye space)
 - transform position (eye to screen space)
 - pass through material color
- Rasterizer
 - interpolated parameters: r, g, b color; x, y, z normal
- Fragment stage (output: color, z')
 - **compute shading using interpolated color and normal**
 - write to pixel

Per-vertex vs Per-pixel Shading




Phong



Toon



Flat



Gouraud



Live Demo: GLSL Shading



https://gitlab.com/boolzhu/cs345l-computer-graphics-starter-code/-/tree/master/tutorials/tutorial_lighting?ref_type=heads