

CS345 I: Noise

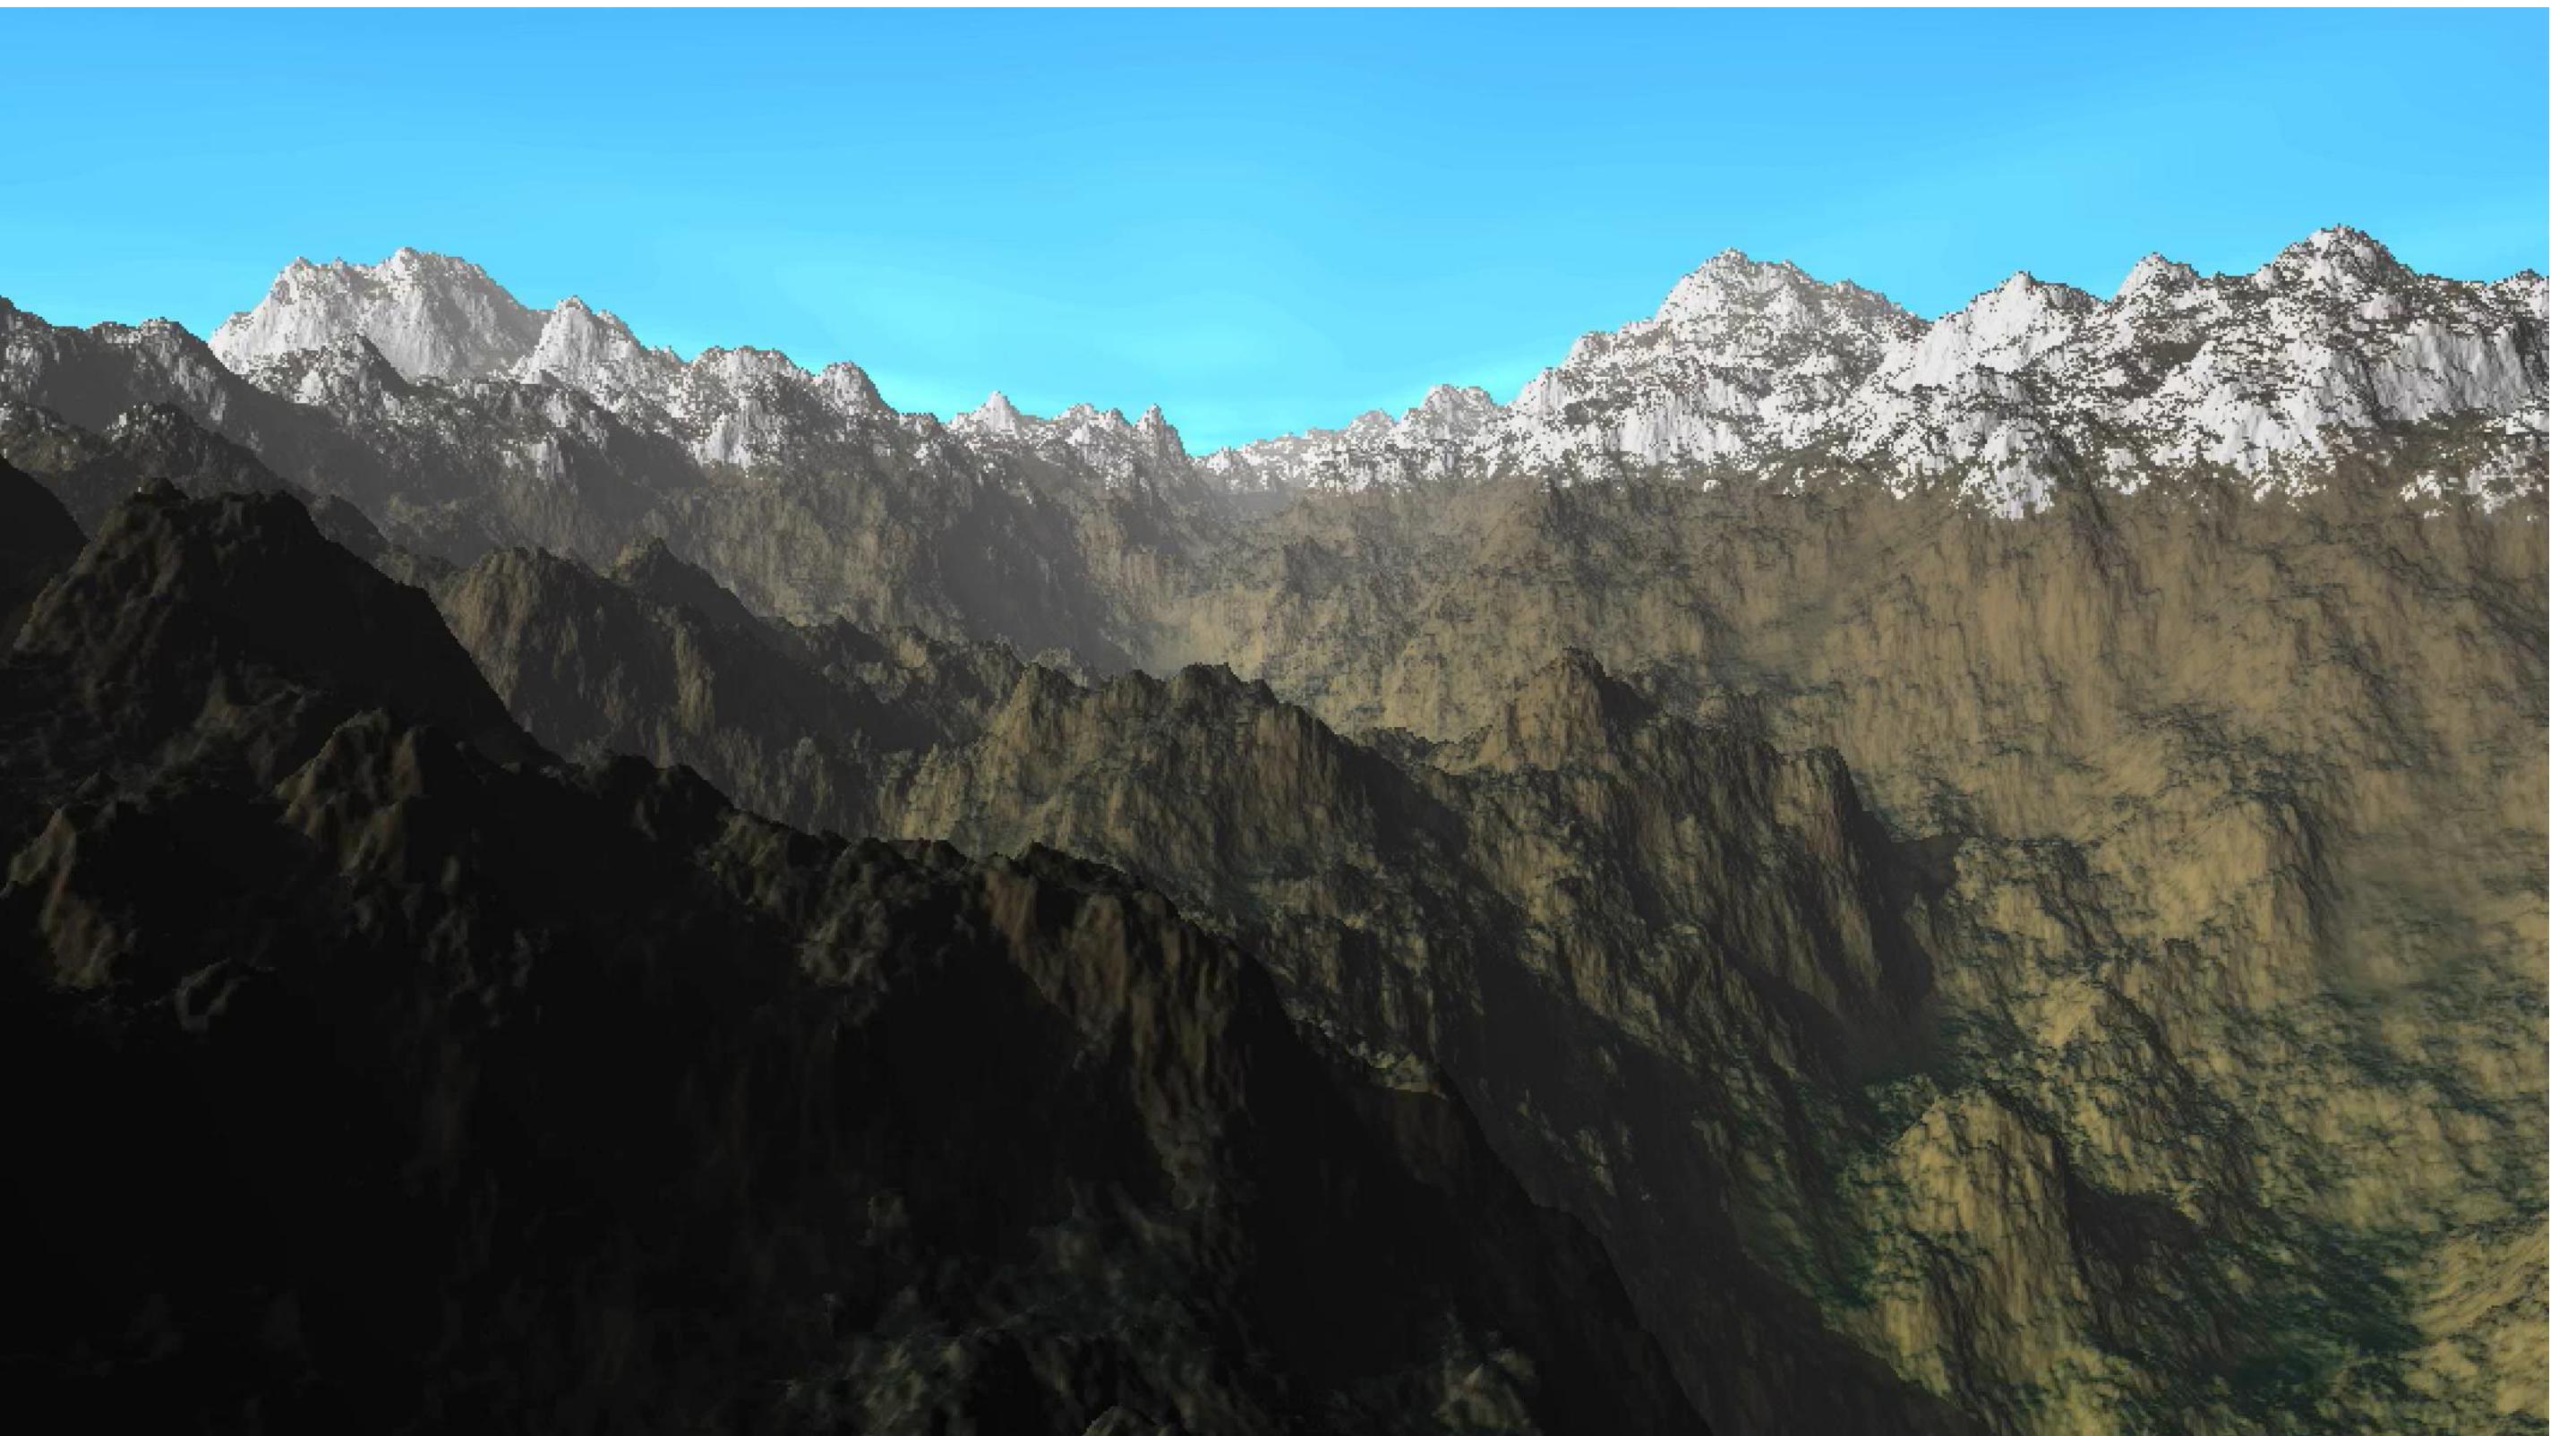
Bo Zhu

School of Interactive Computing
Georgia Institute of Technology

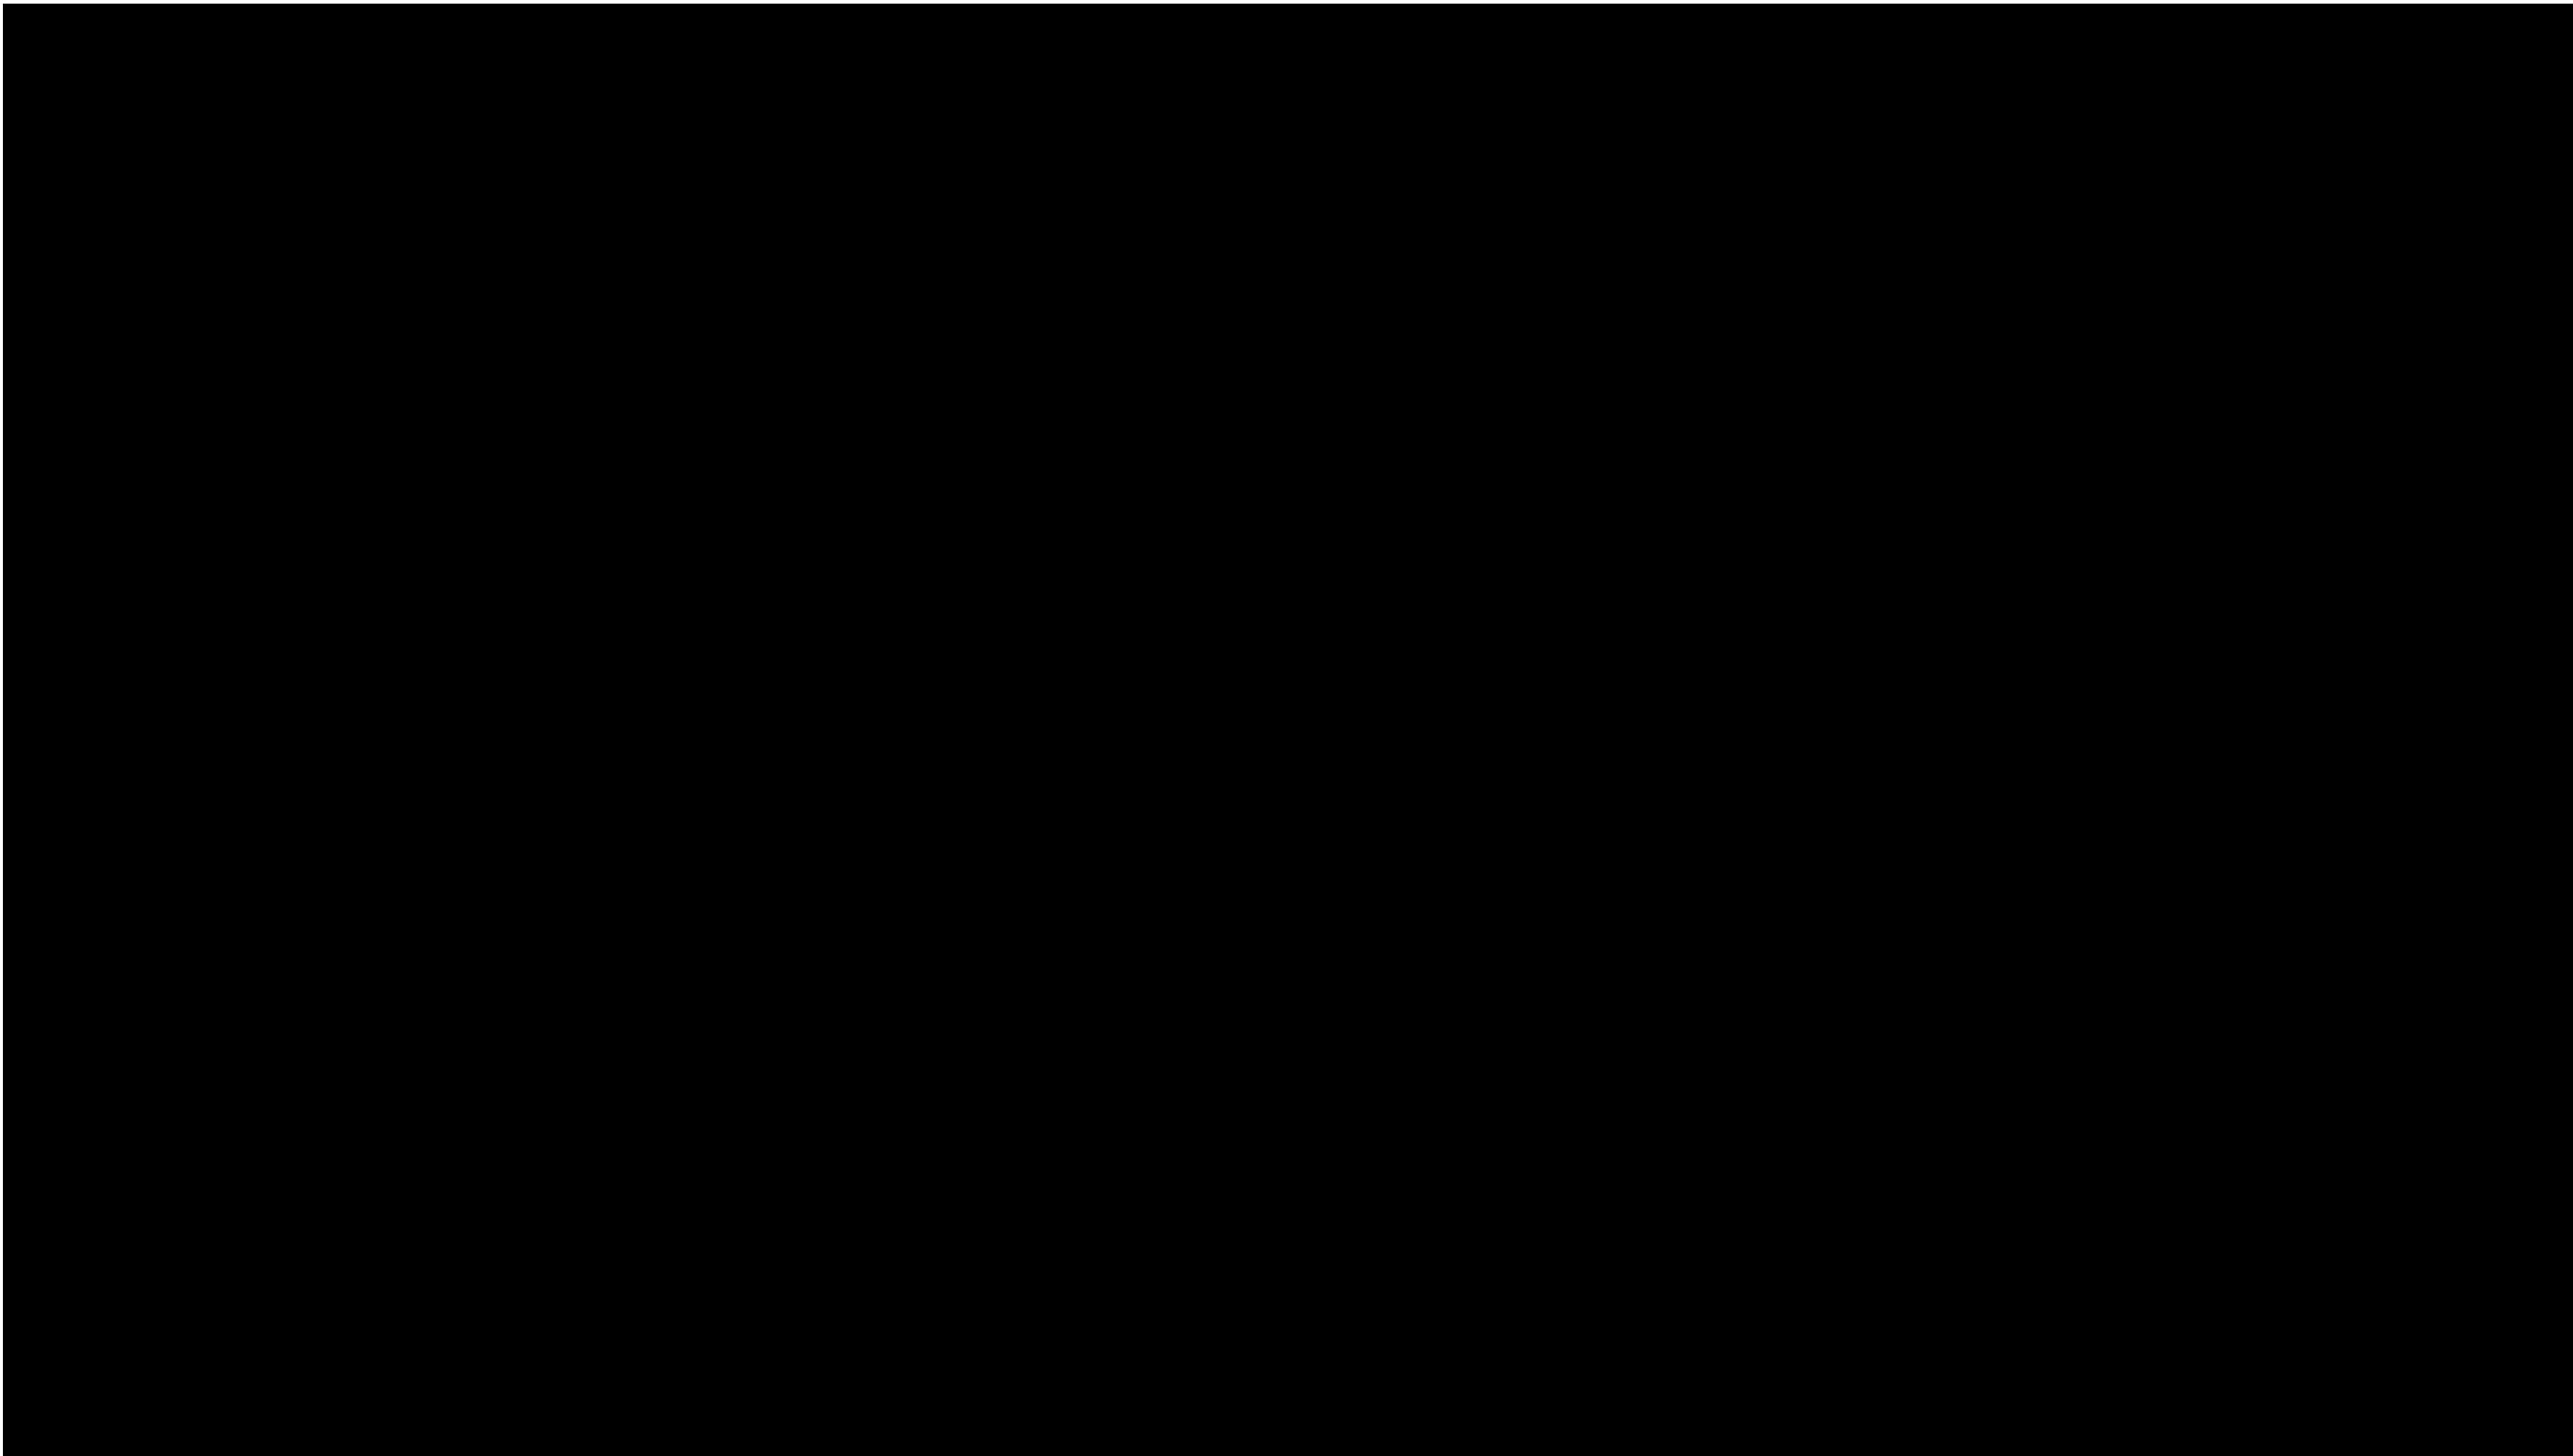
Motivational Video: Infinite Terrain Generation with Procedural Modeling

by Evan Z. Muscatel,
Dartmouth CS 21'

Inventor of our A6 ☺



Motivational Video: Houdini Procedural Modeling Demo



https://www.youtube.com/watch?v=CE7H_P8q6QY



Motivational Video:
WonderJourney:
Going from
Anywhere to
Everywhere



<https://kovenyu.com/wonderjourney/>



Motivational Video: OpenAI Sora:World Generation Model

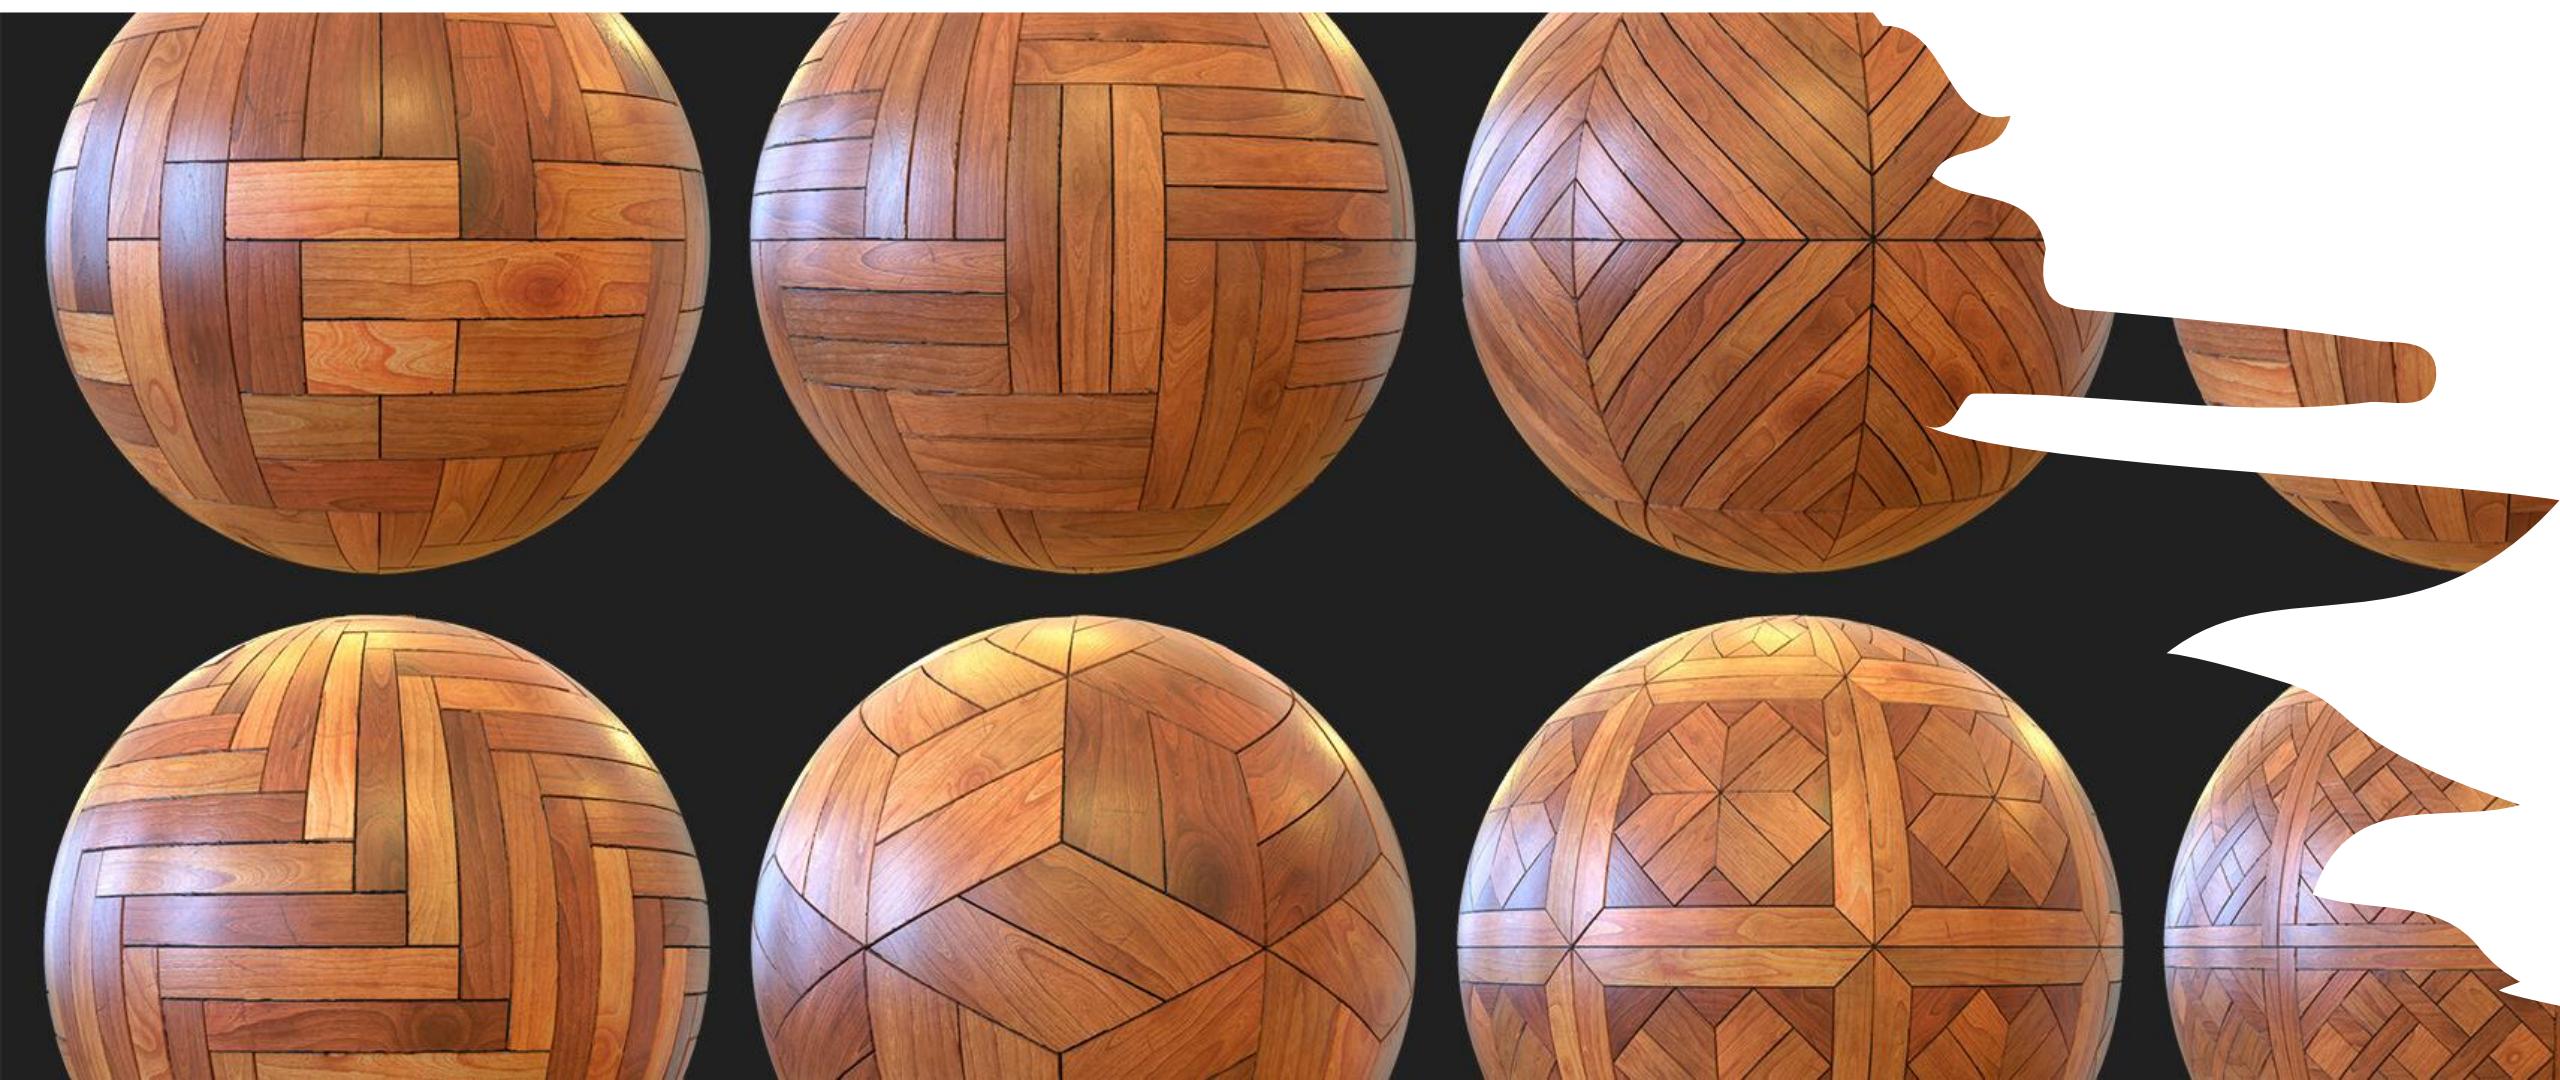
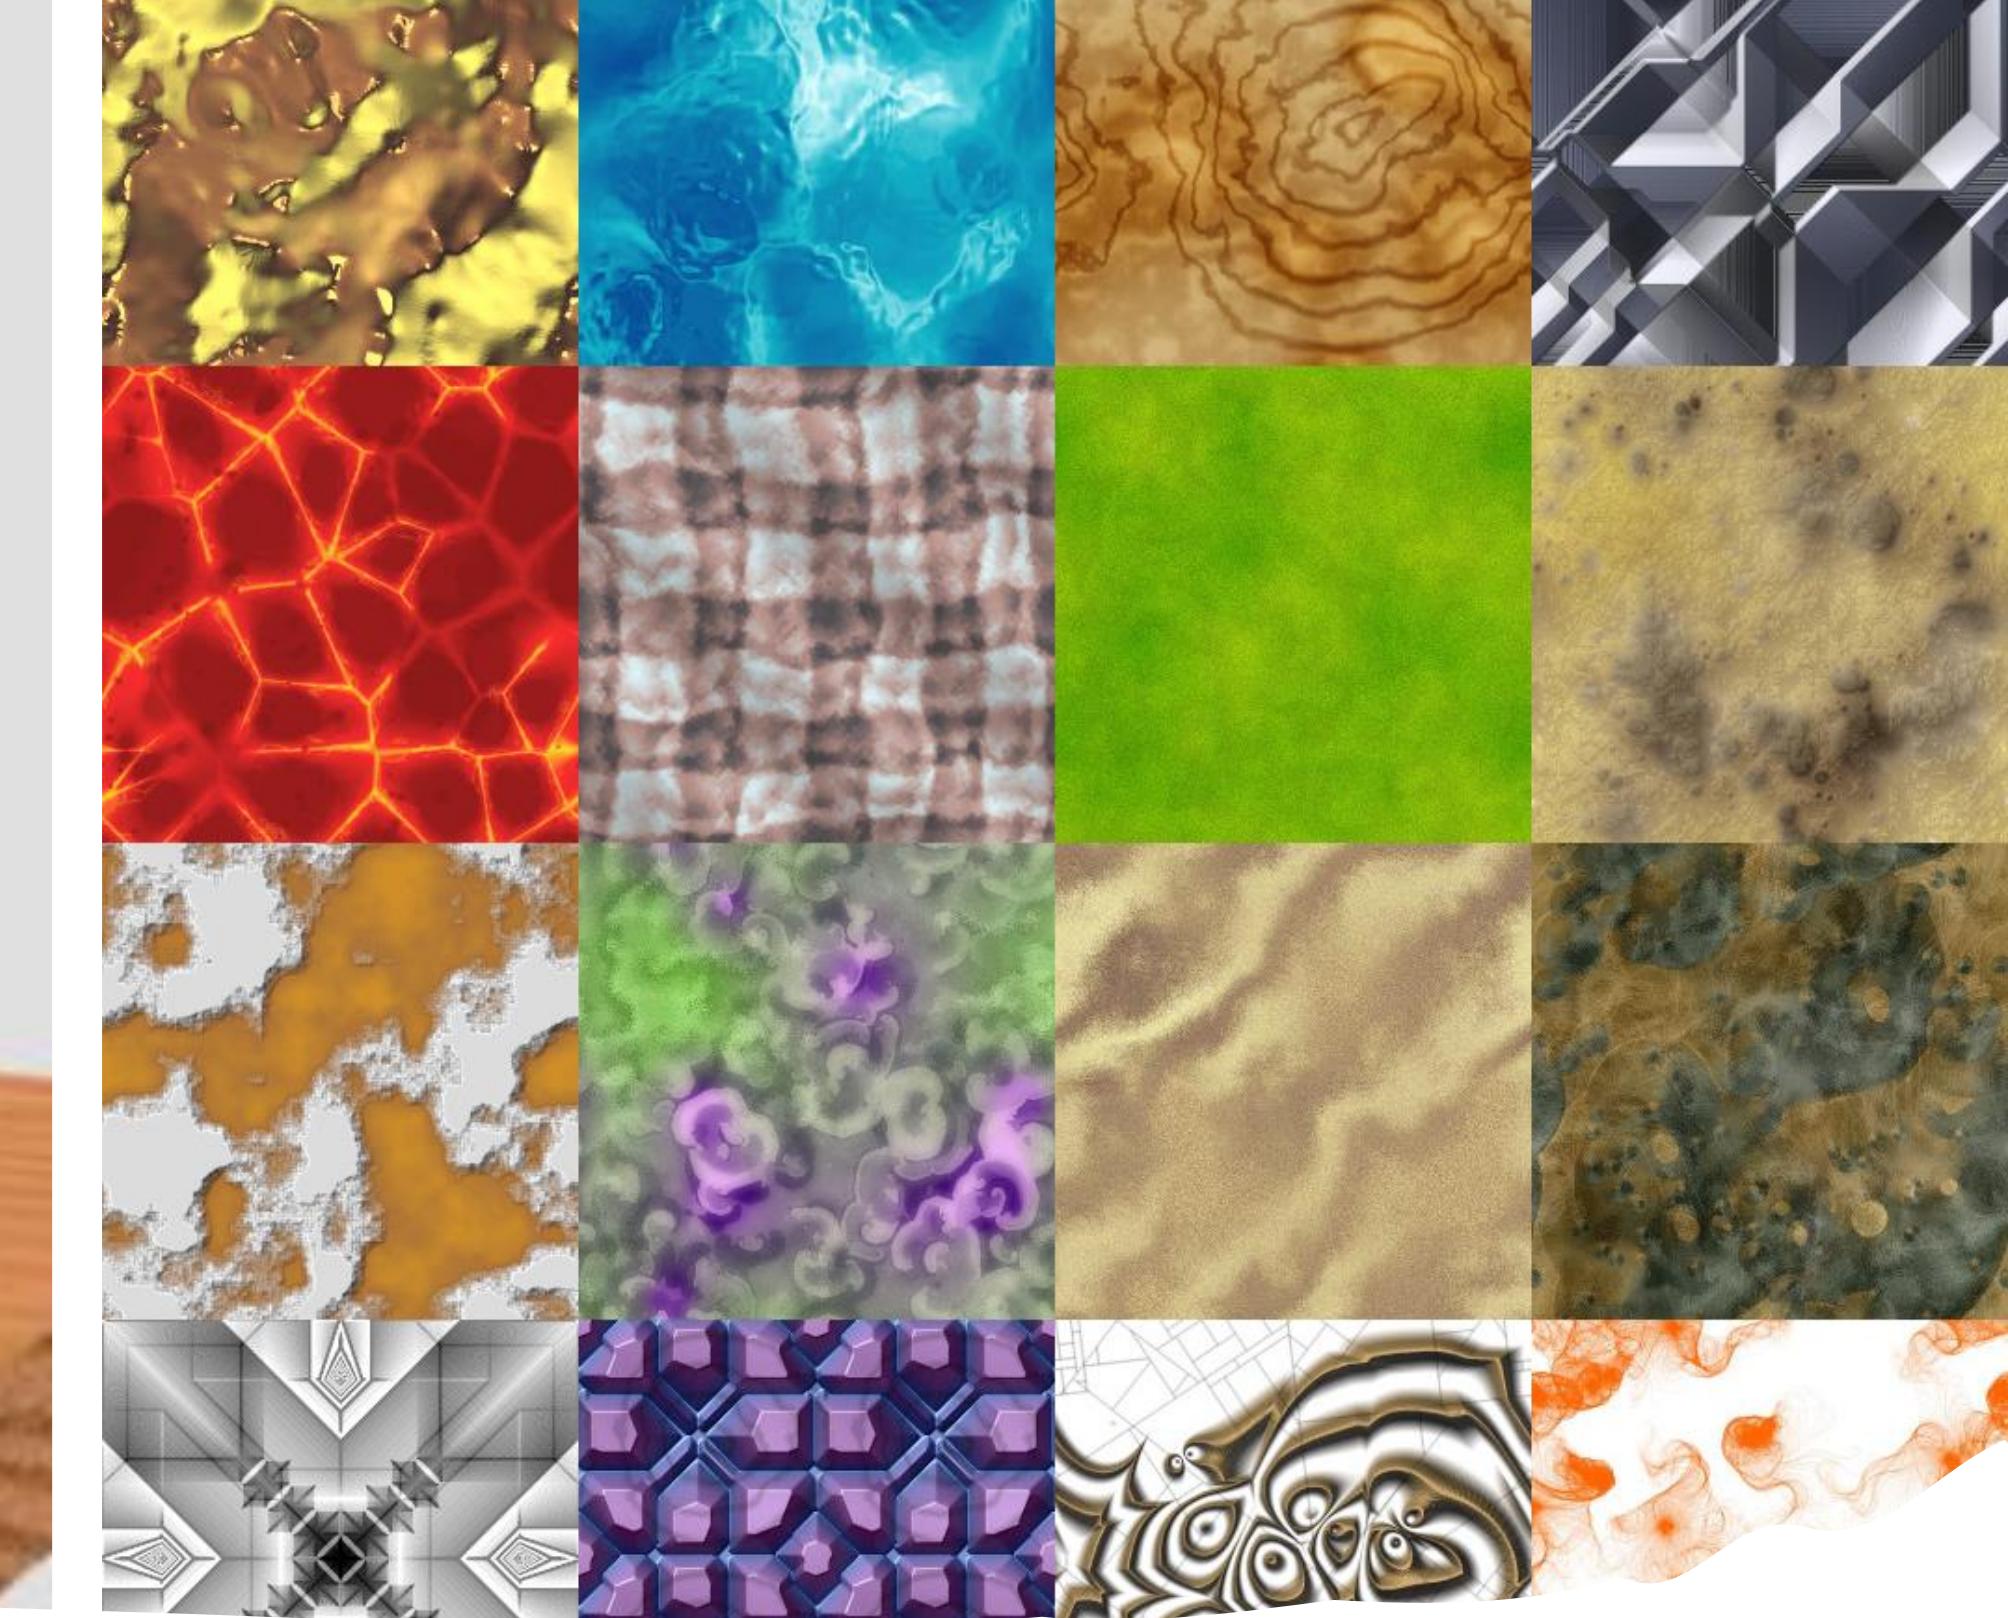


<https://www.youtube.com/watch?v=anmuklFtu8U>



A scenic landscape featuring a dense forest of tall evergreen trees in the foreground. In the middle ground, there is a calm lake reflecting the surrounding environment. In the background, there are majestic mountains under a clear blue sky with a few wispy clouds.

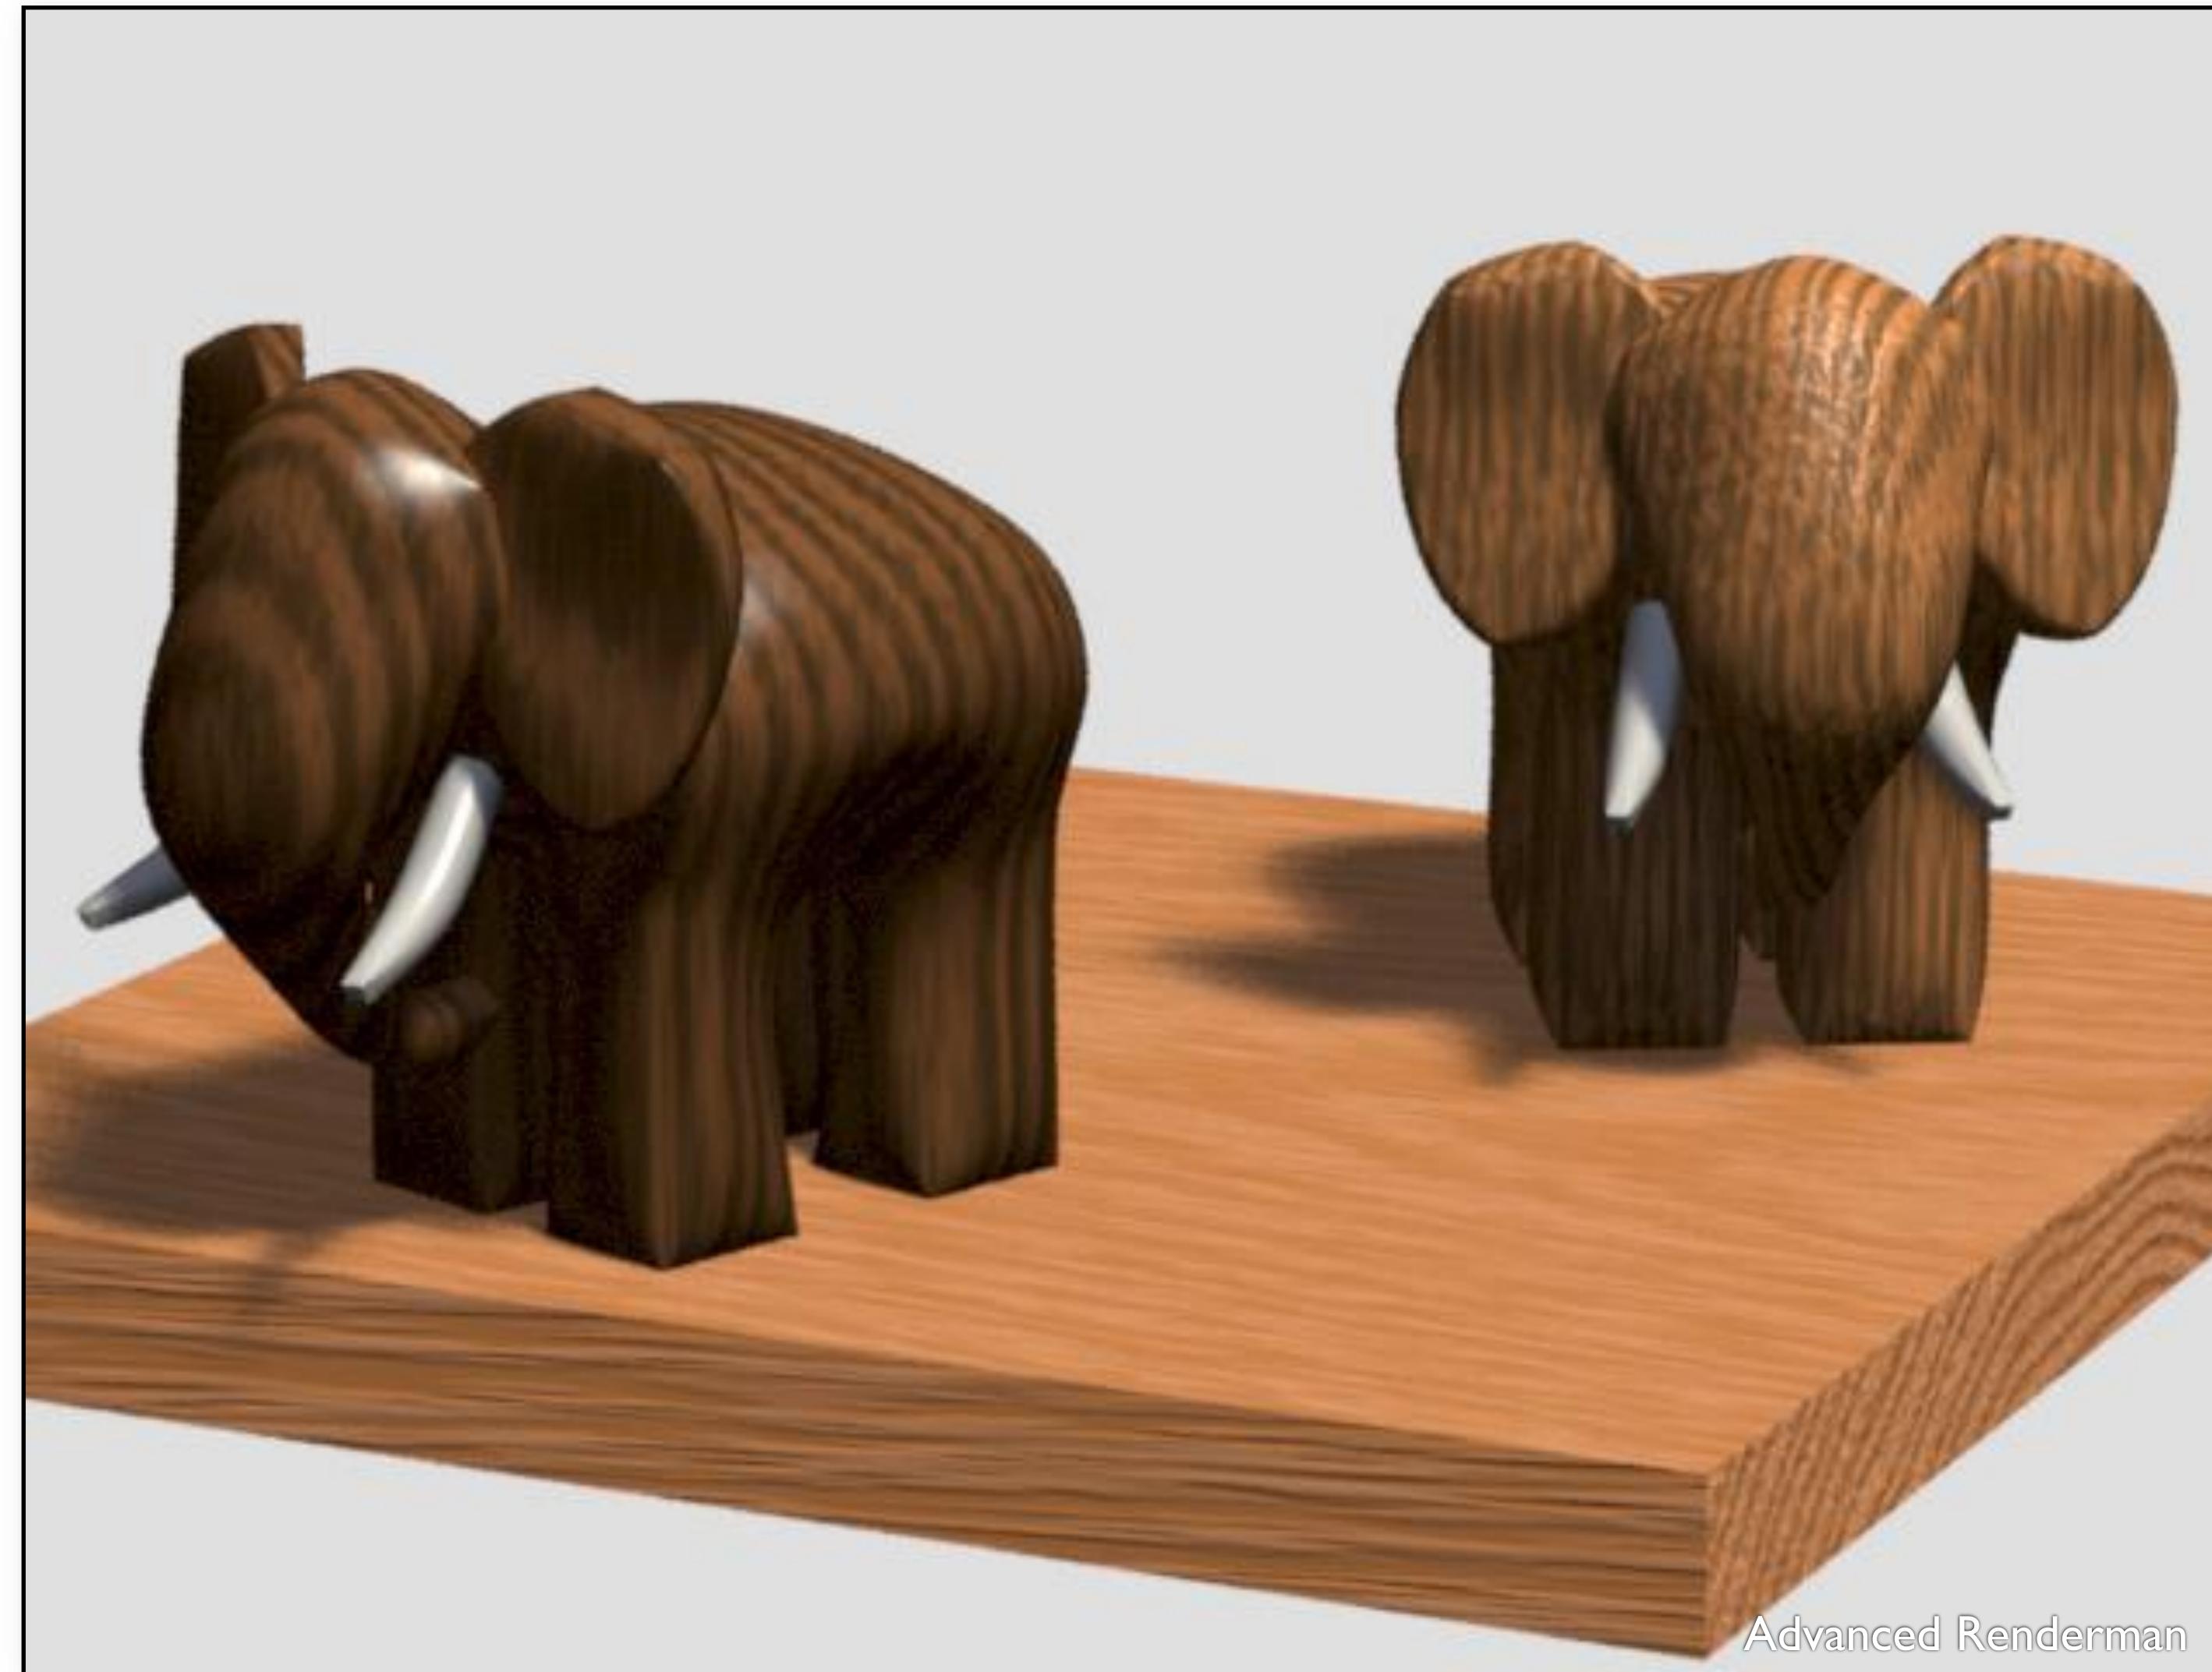
We want to Define Rules
to Generate the World



Pioneering Exploration in Graphics:
**Generating Texture
Images with Simple Math**

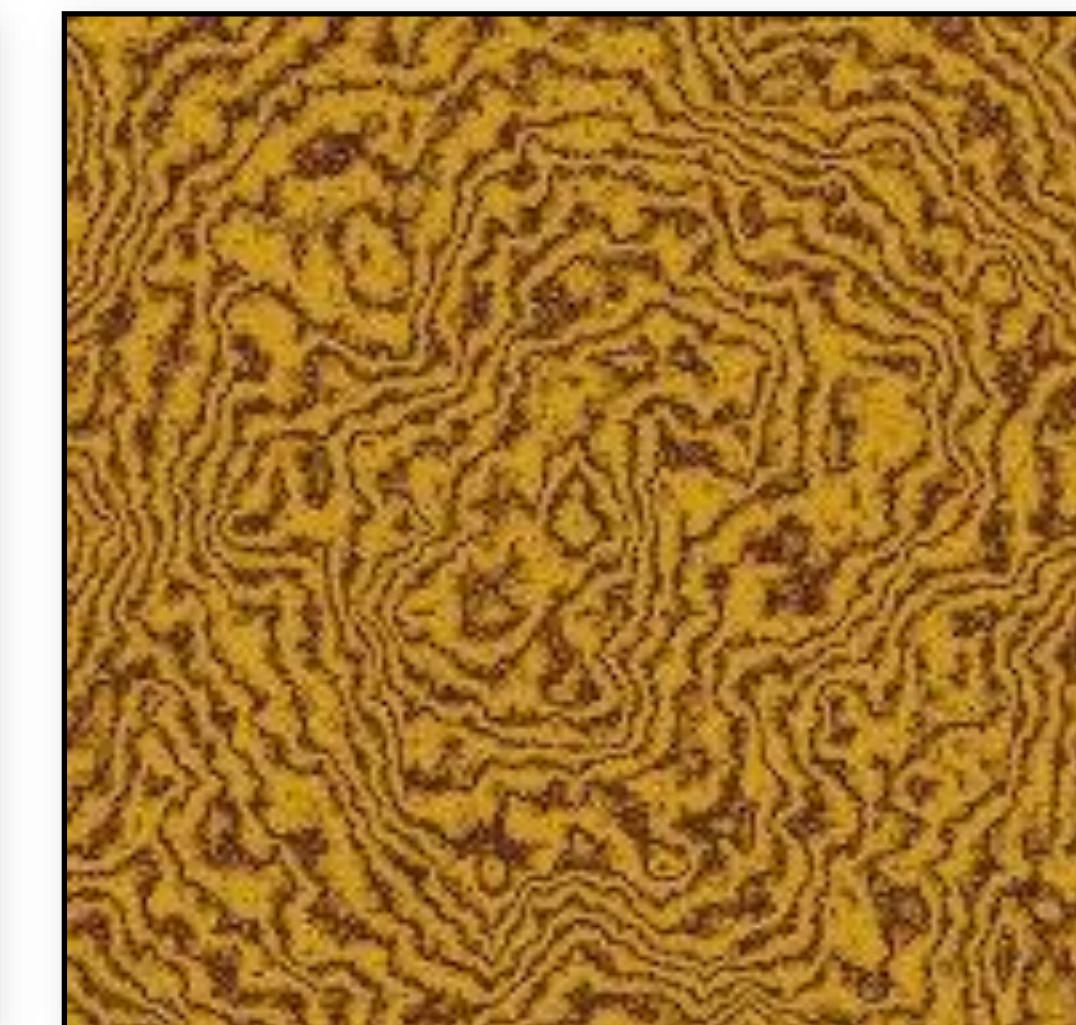
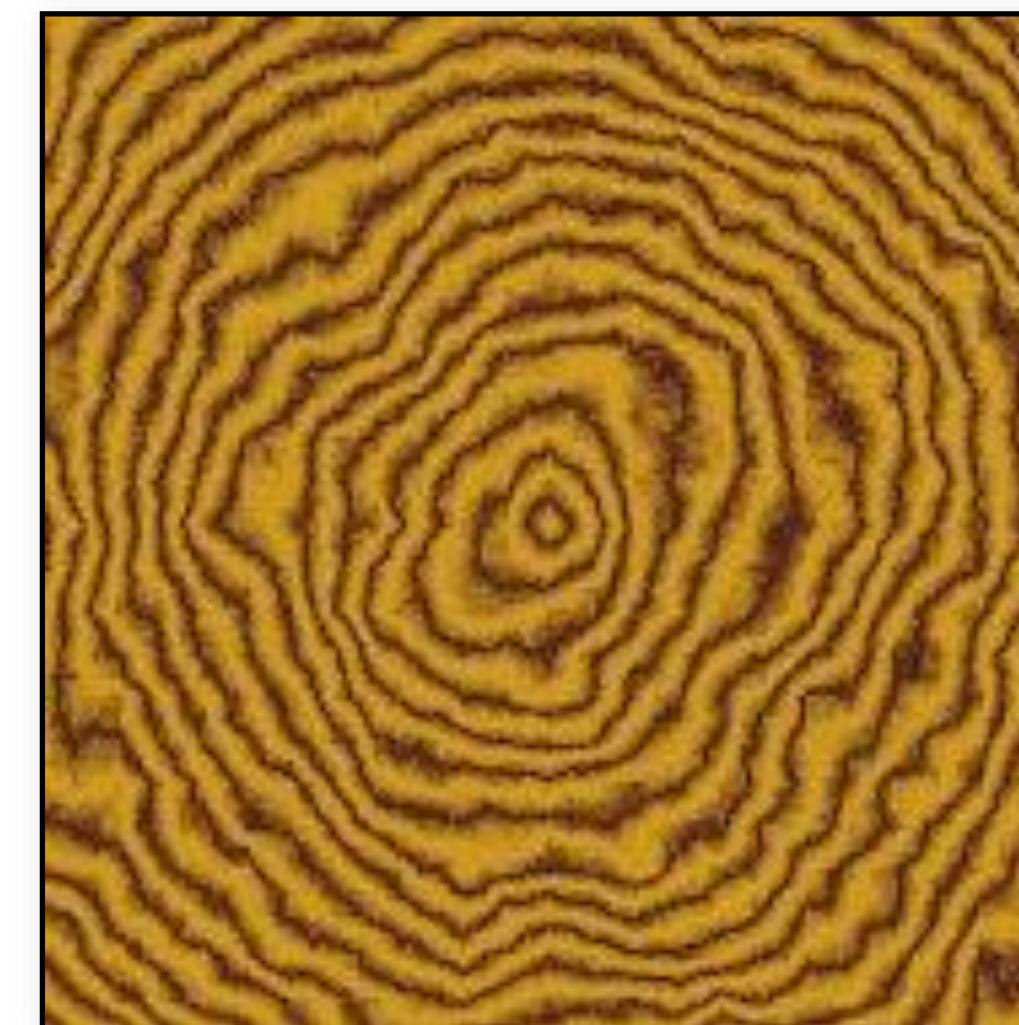
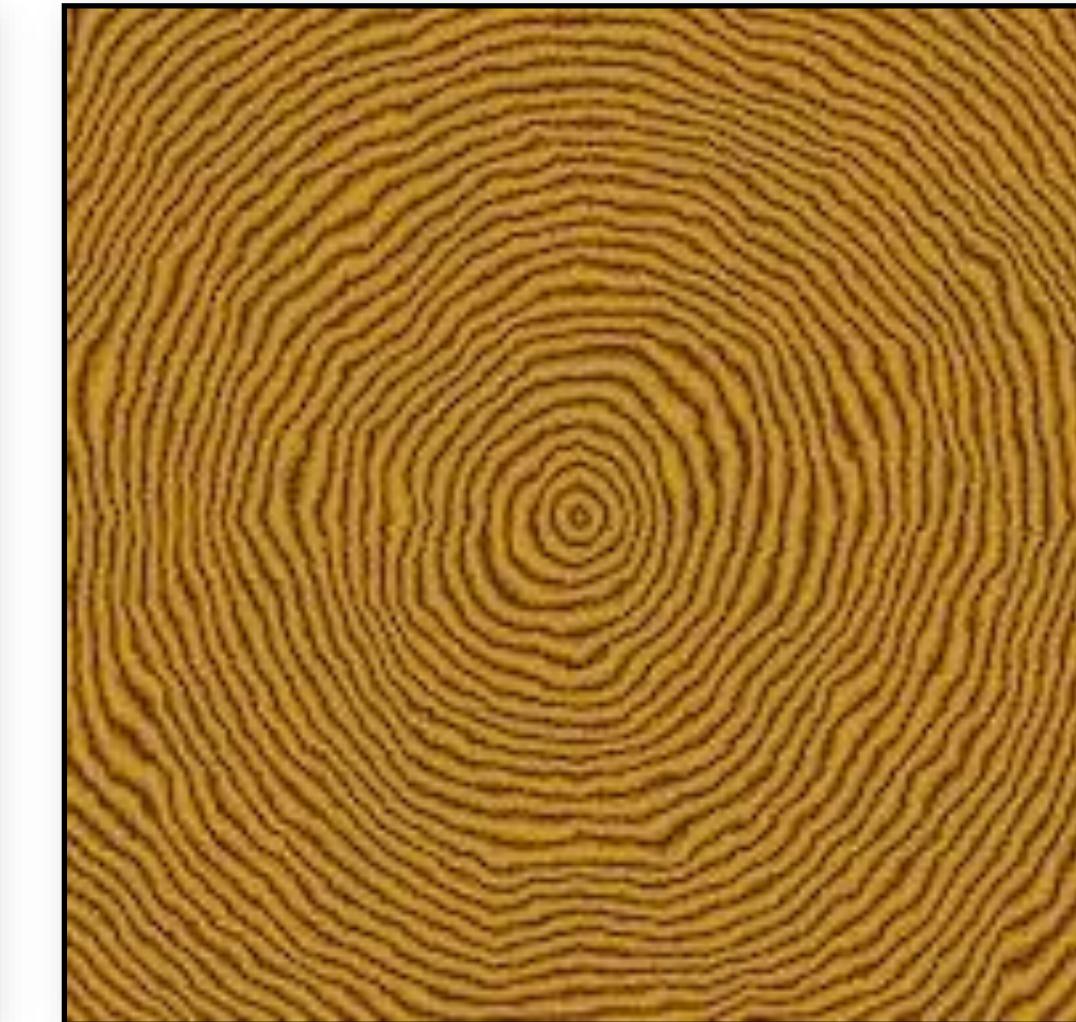
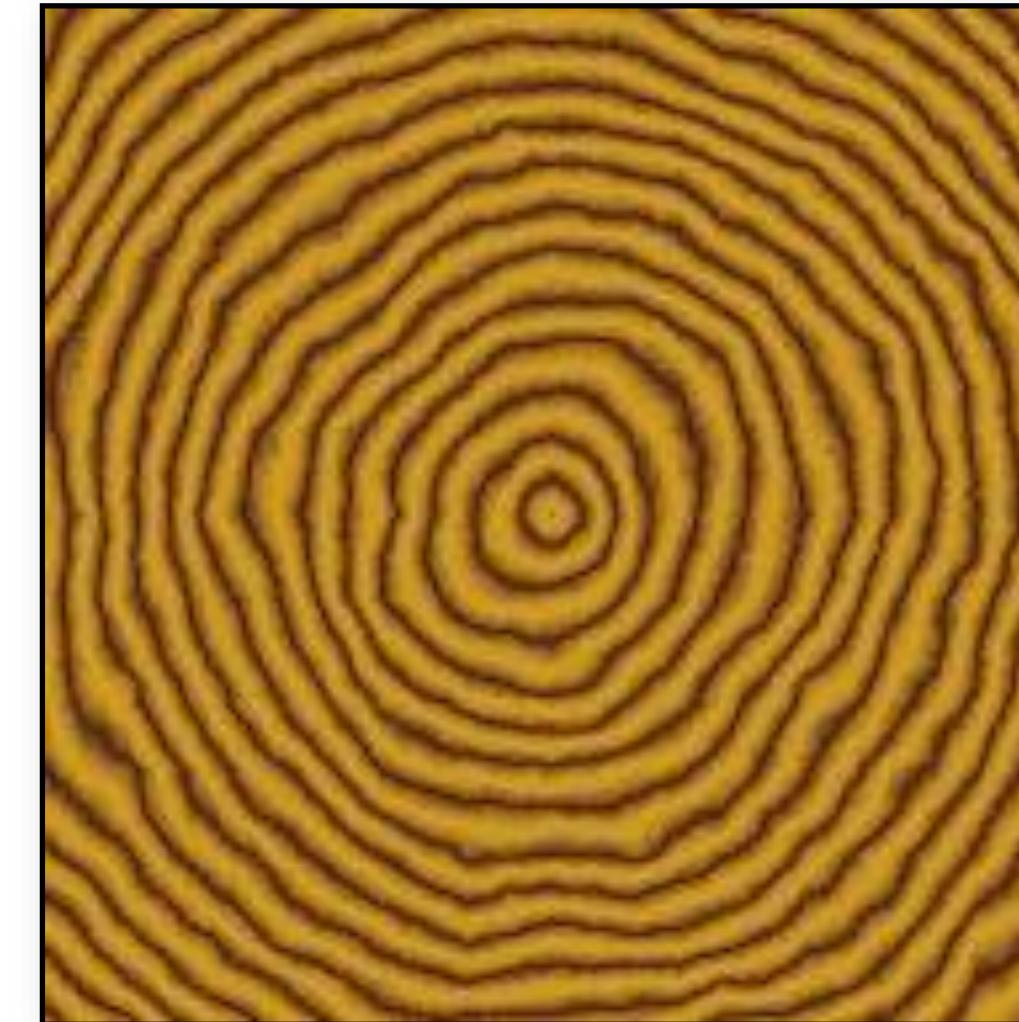
Wood

$$(1 + \sin(\sqrt{p_x^2 + p_y^2}) + fBm(p))) / 2$$



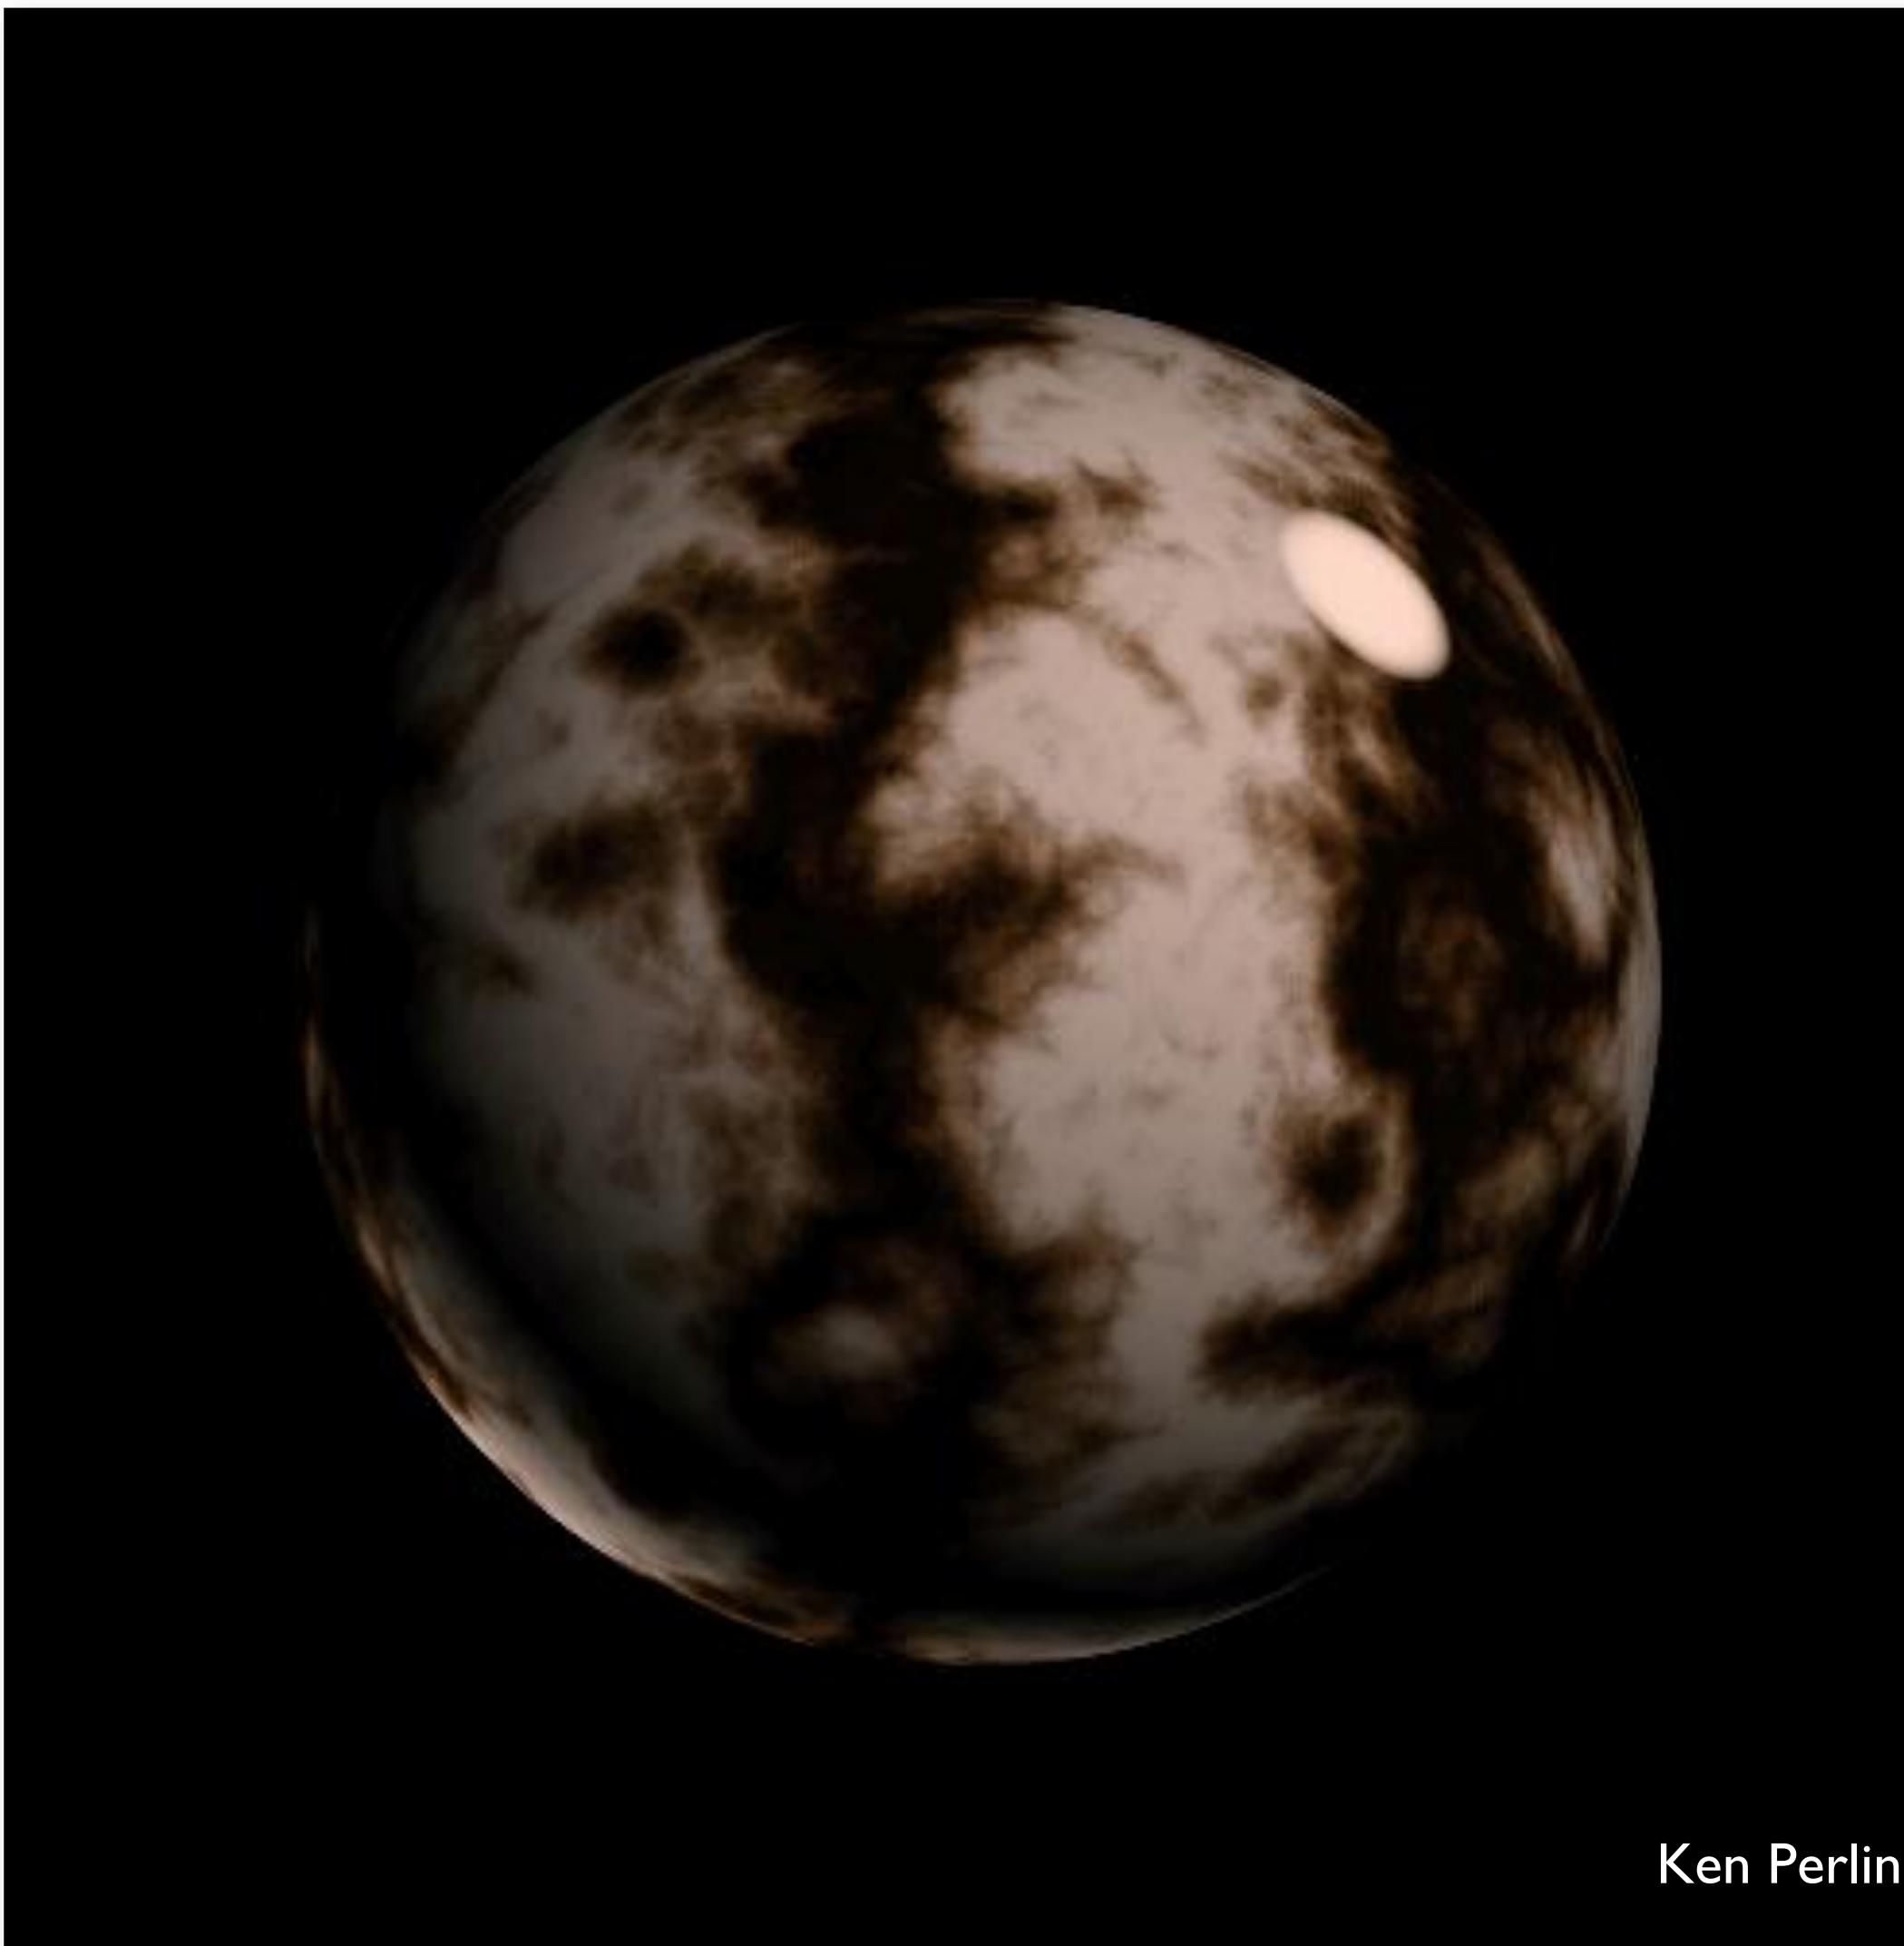
Another Wood

$(1 + \sin(\sqrt{p_x^2 + p_y^2}) + fBm(p))) / 2$



Marble

$(1 + \sin(k_1 p_x + \text{turbulence}(k_2 p))) / w) / 2$



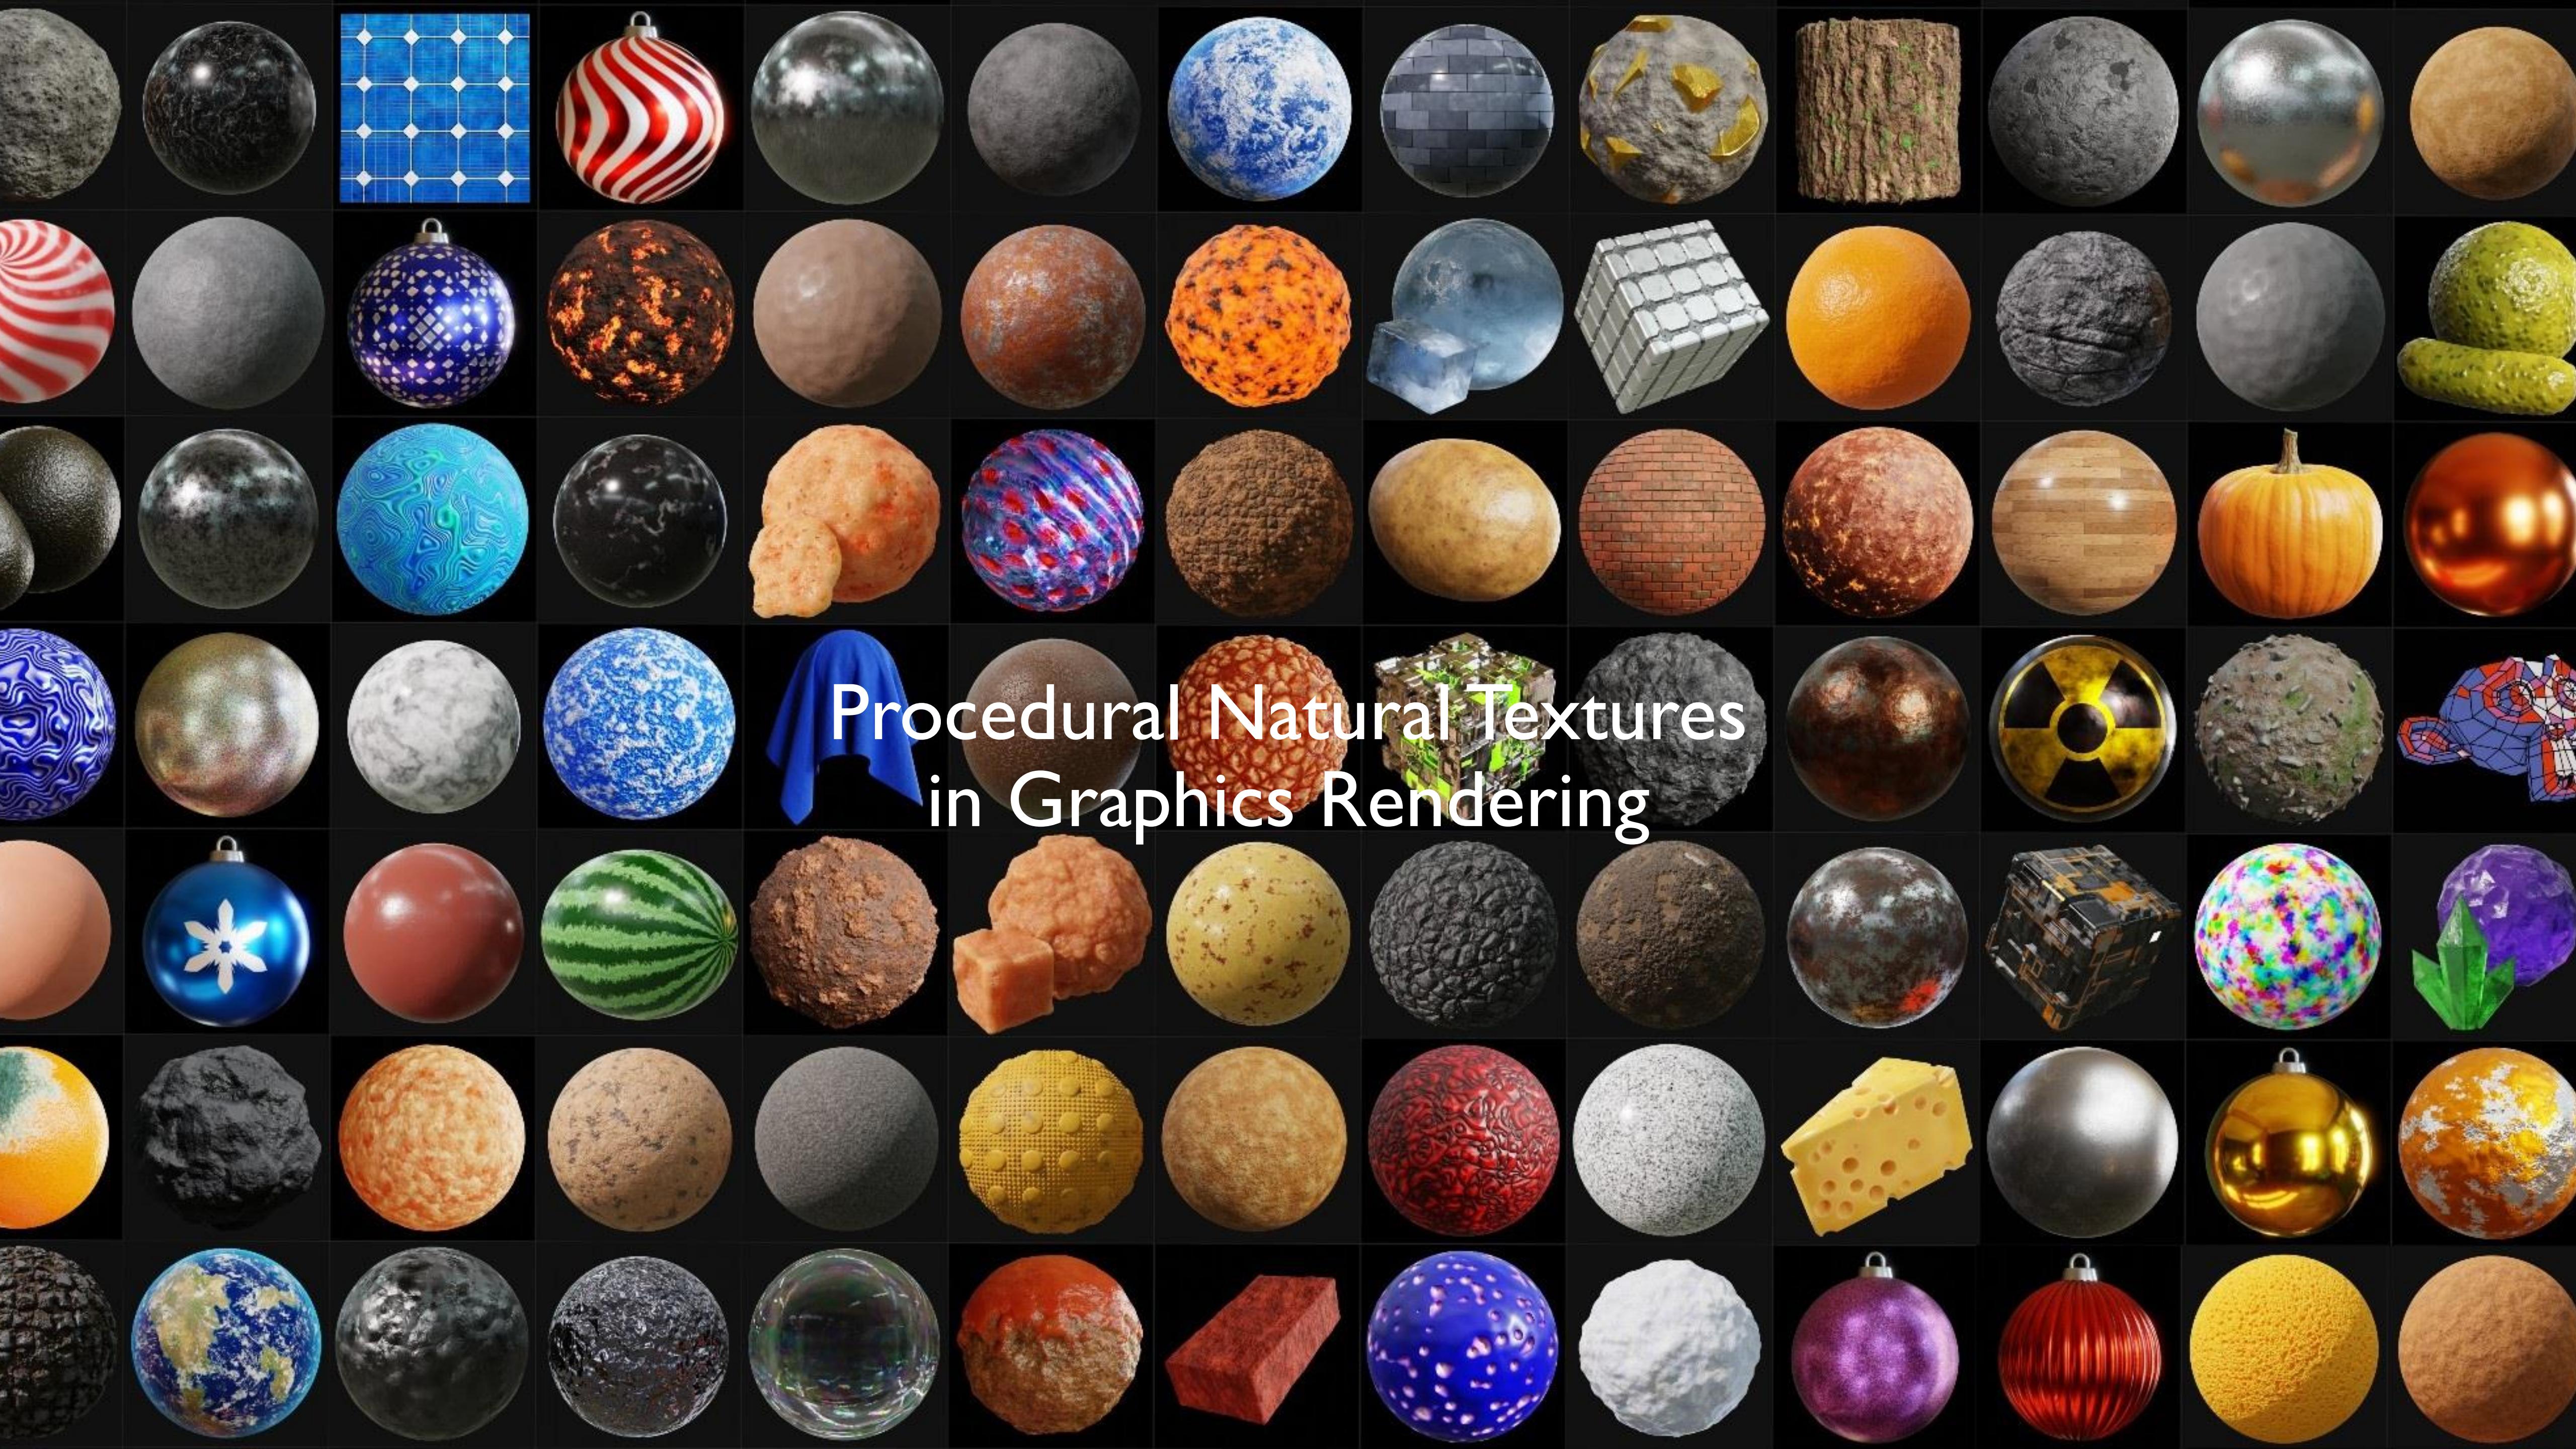
Ken Perlin



Ken Perlin



Procedural Natural Textures in Graphics Rendering



A wide-angle photograph of a majestic mountain range. In the foreground, a deep blue river flows through a valley, its surface reflecting the surrounding peaks. The mountains themselves are rugged, with steep slopes covered in patches of white snow and dark, rocky terrain. One prominent peak on the right side of the frame features a small, light-colored building with a distinctive green roof, possibly a weather station or a small temple. The sky is overcast with heavy clouds, creating a somber and dramatic atmosphere.

Procedural Natural Landscapes in Movies and Games

The background image shows a detailed 3D rendering of a city skyline at dusk or dawn. In the foreground, there's a large, dark building with a circular helipad on its roof, featuring a yellow 'H' symbol. A red flag-like marker is positioned on top of one of the buildings in the middle ground. The sky is filled with soft, scattered clouds.

Procedural City Scenes in Movies and Games

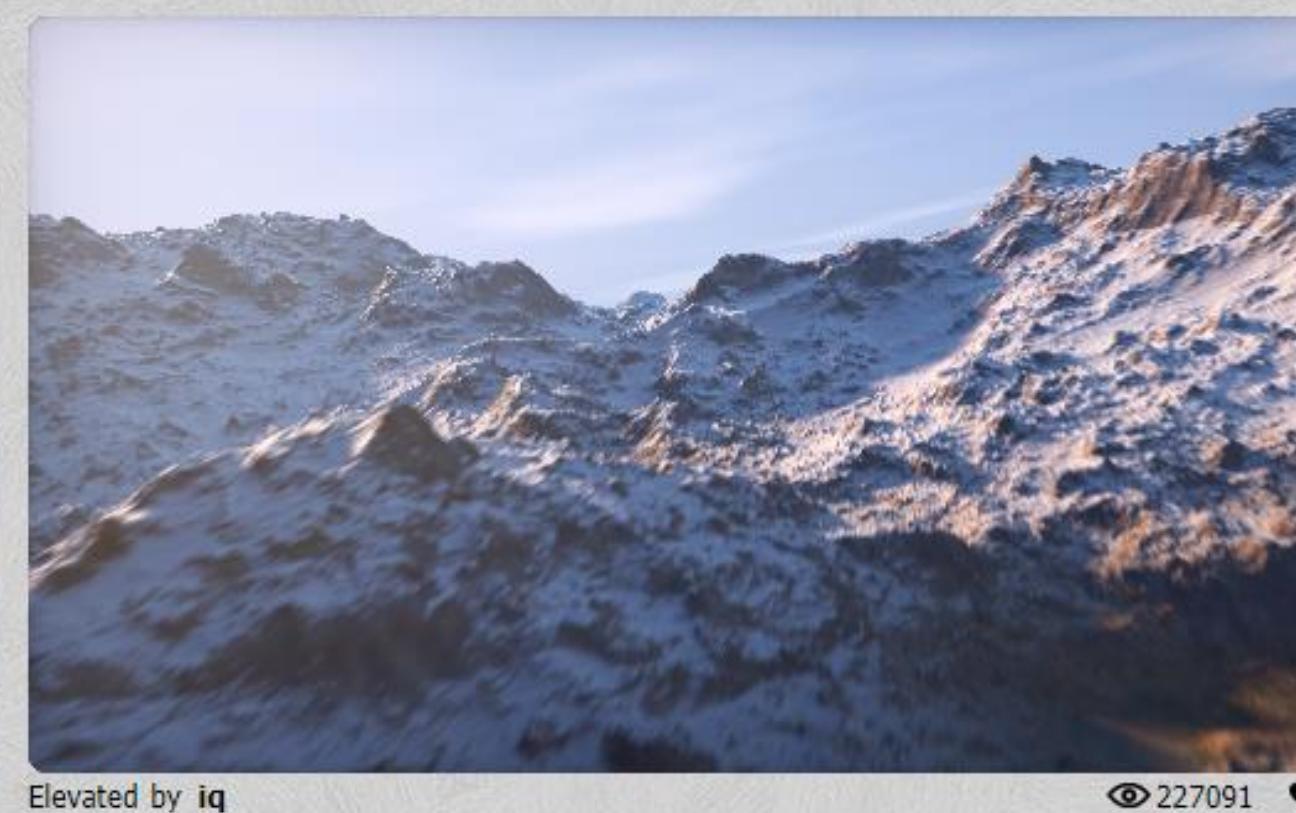
Filter: **Multipass** GPU Sound VR Microphone Soundcloud Webcam

View: **Slideshow**

Results (744): **1** **2** **3** ... **62**



30871 1383



Elevated by iq 227091 990



Planet Shadertoy by reinder 116419 481



Cloudy Terrain by iq 93949 534



68640 358



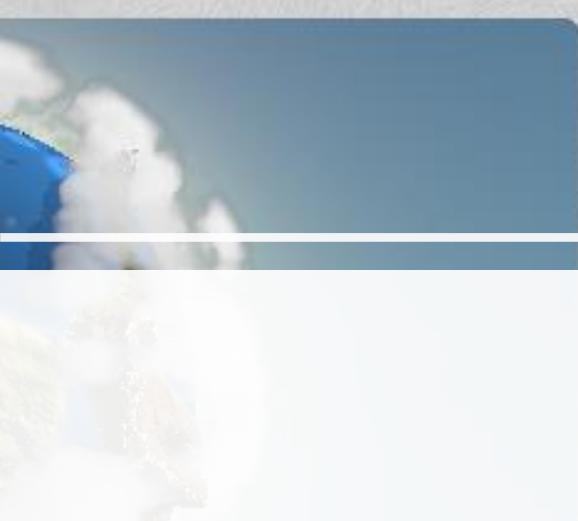
Canyon by iq 67283 360



Mountains by Dave_Hoskins 50417 329



Sirenia Dawn by nimitz 42871 711



38553 309



Hydraulic Erosion by davidar 37769 56



Plate Tectonics by davidar 35545 21



Paroi Jaune by XT95 31285 110

Generating World Scenes with Procedural Shaders

ShaderToy Demo: Procedural Terrain



<https://www.shadertoy.com/view/4ttSWf>



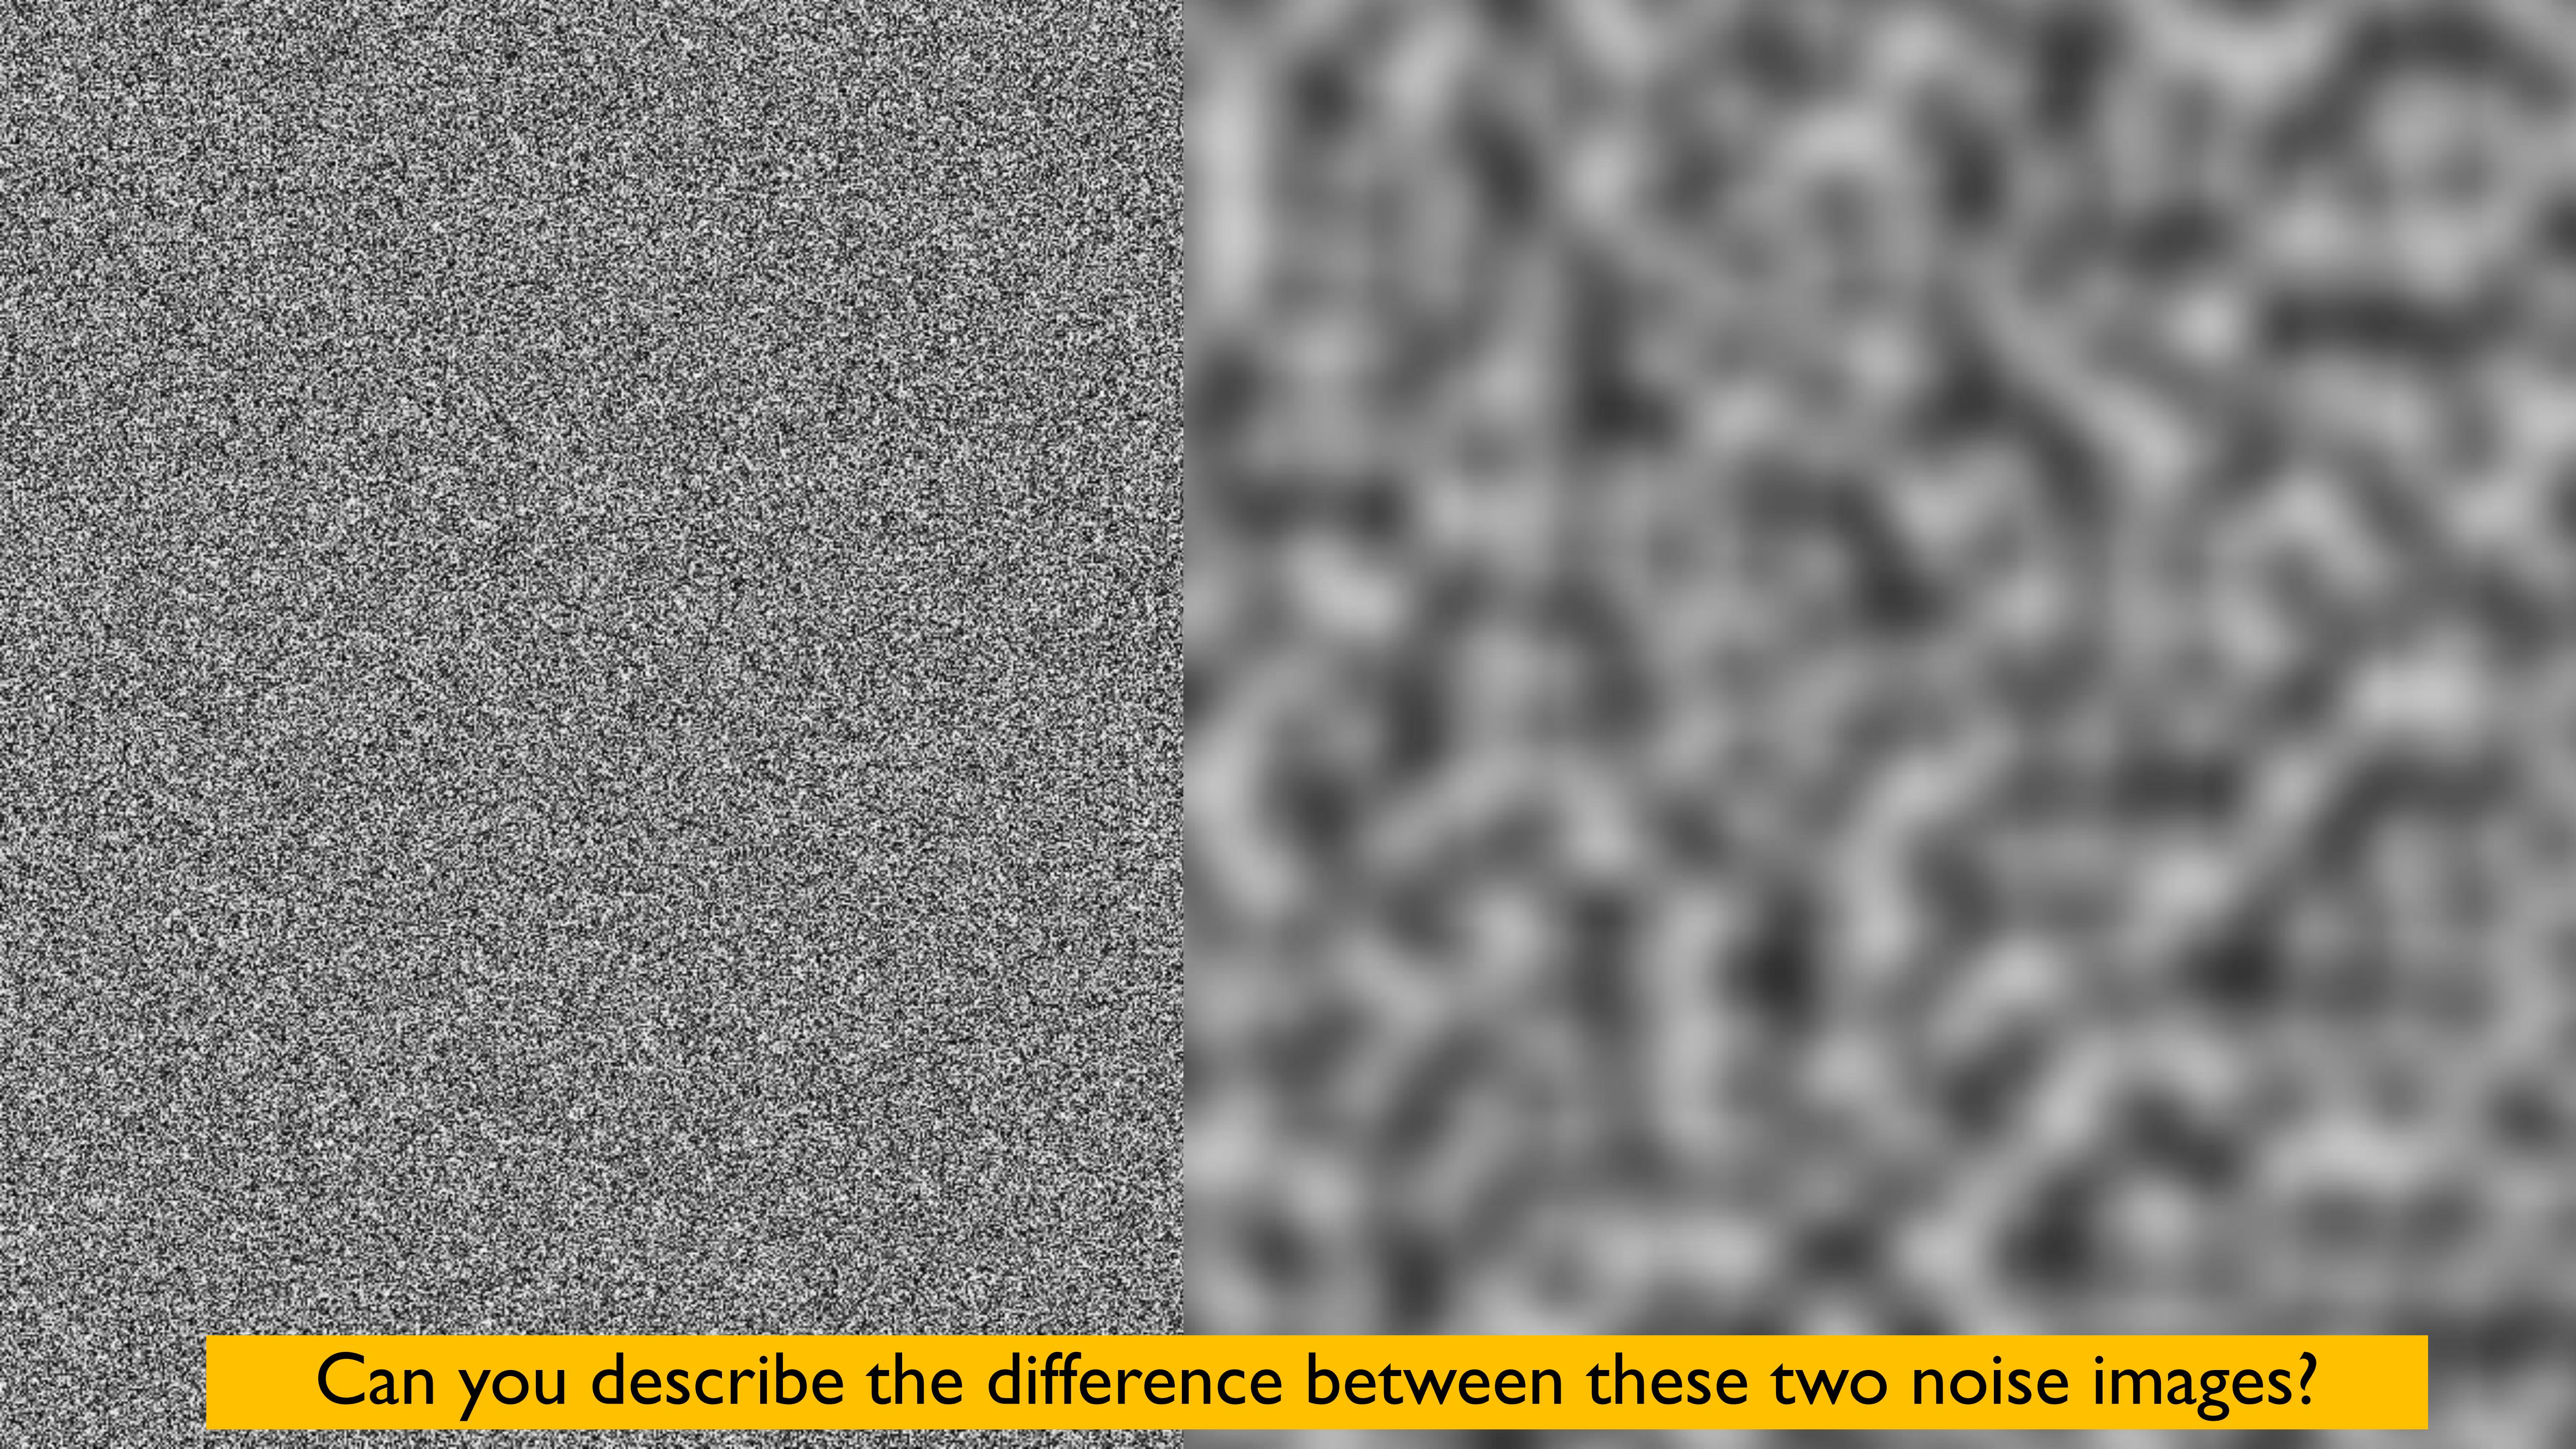
Study Plan

- Natural-looking Noise
- 1D Perlin Noise
- 2D Perlin Noise
- Terrain Generation
- Other Noise Functions

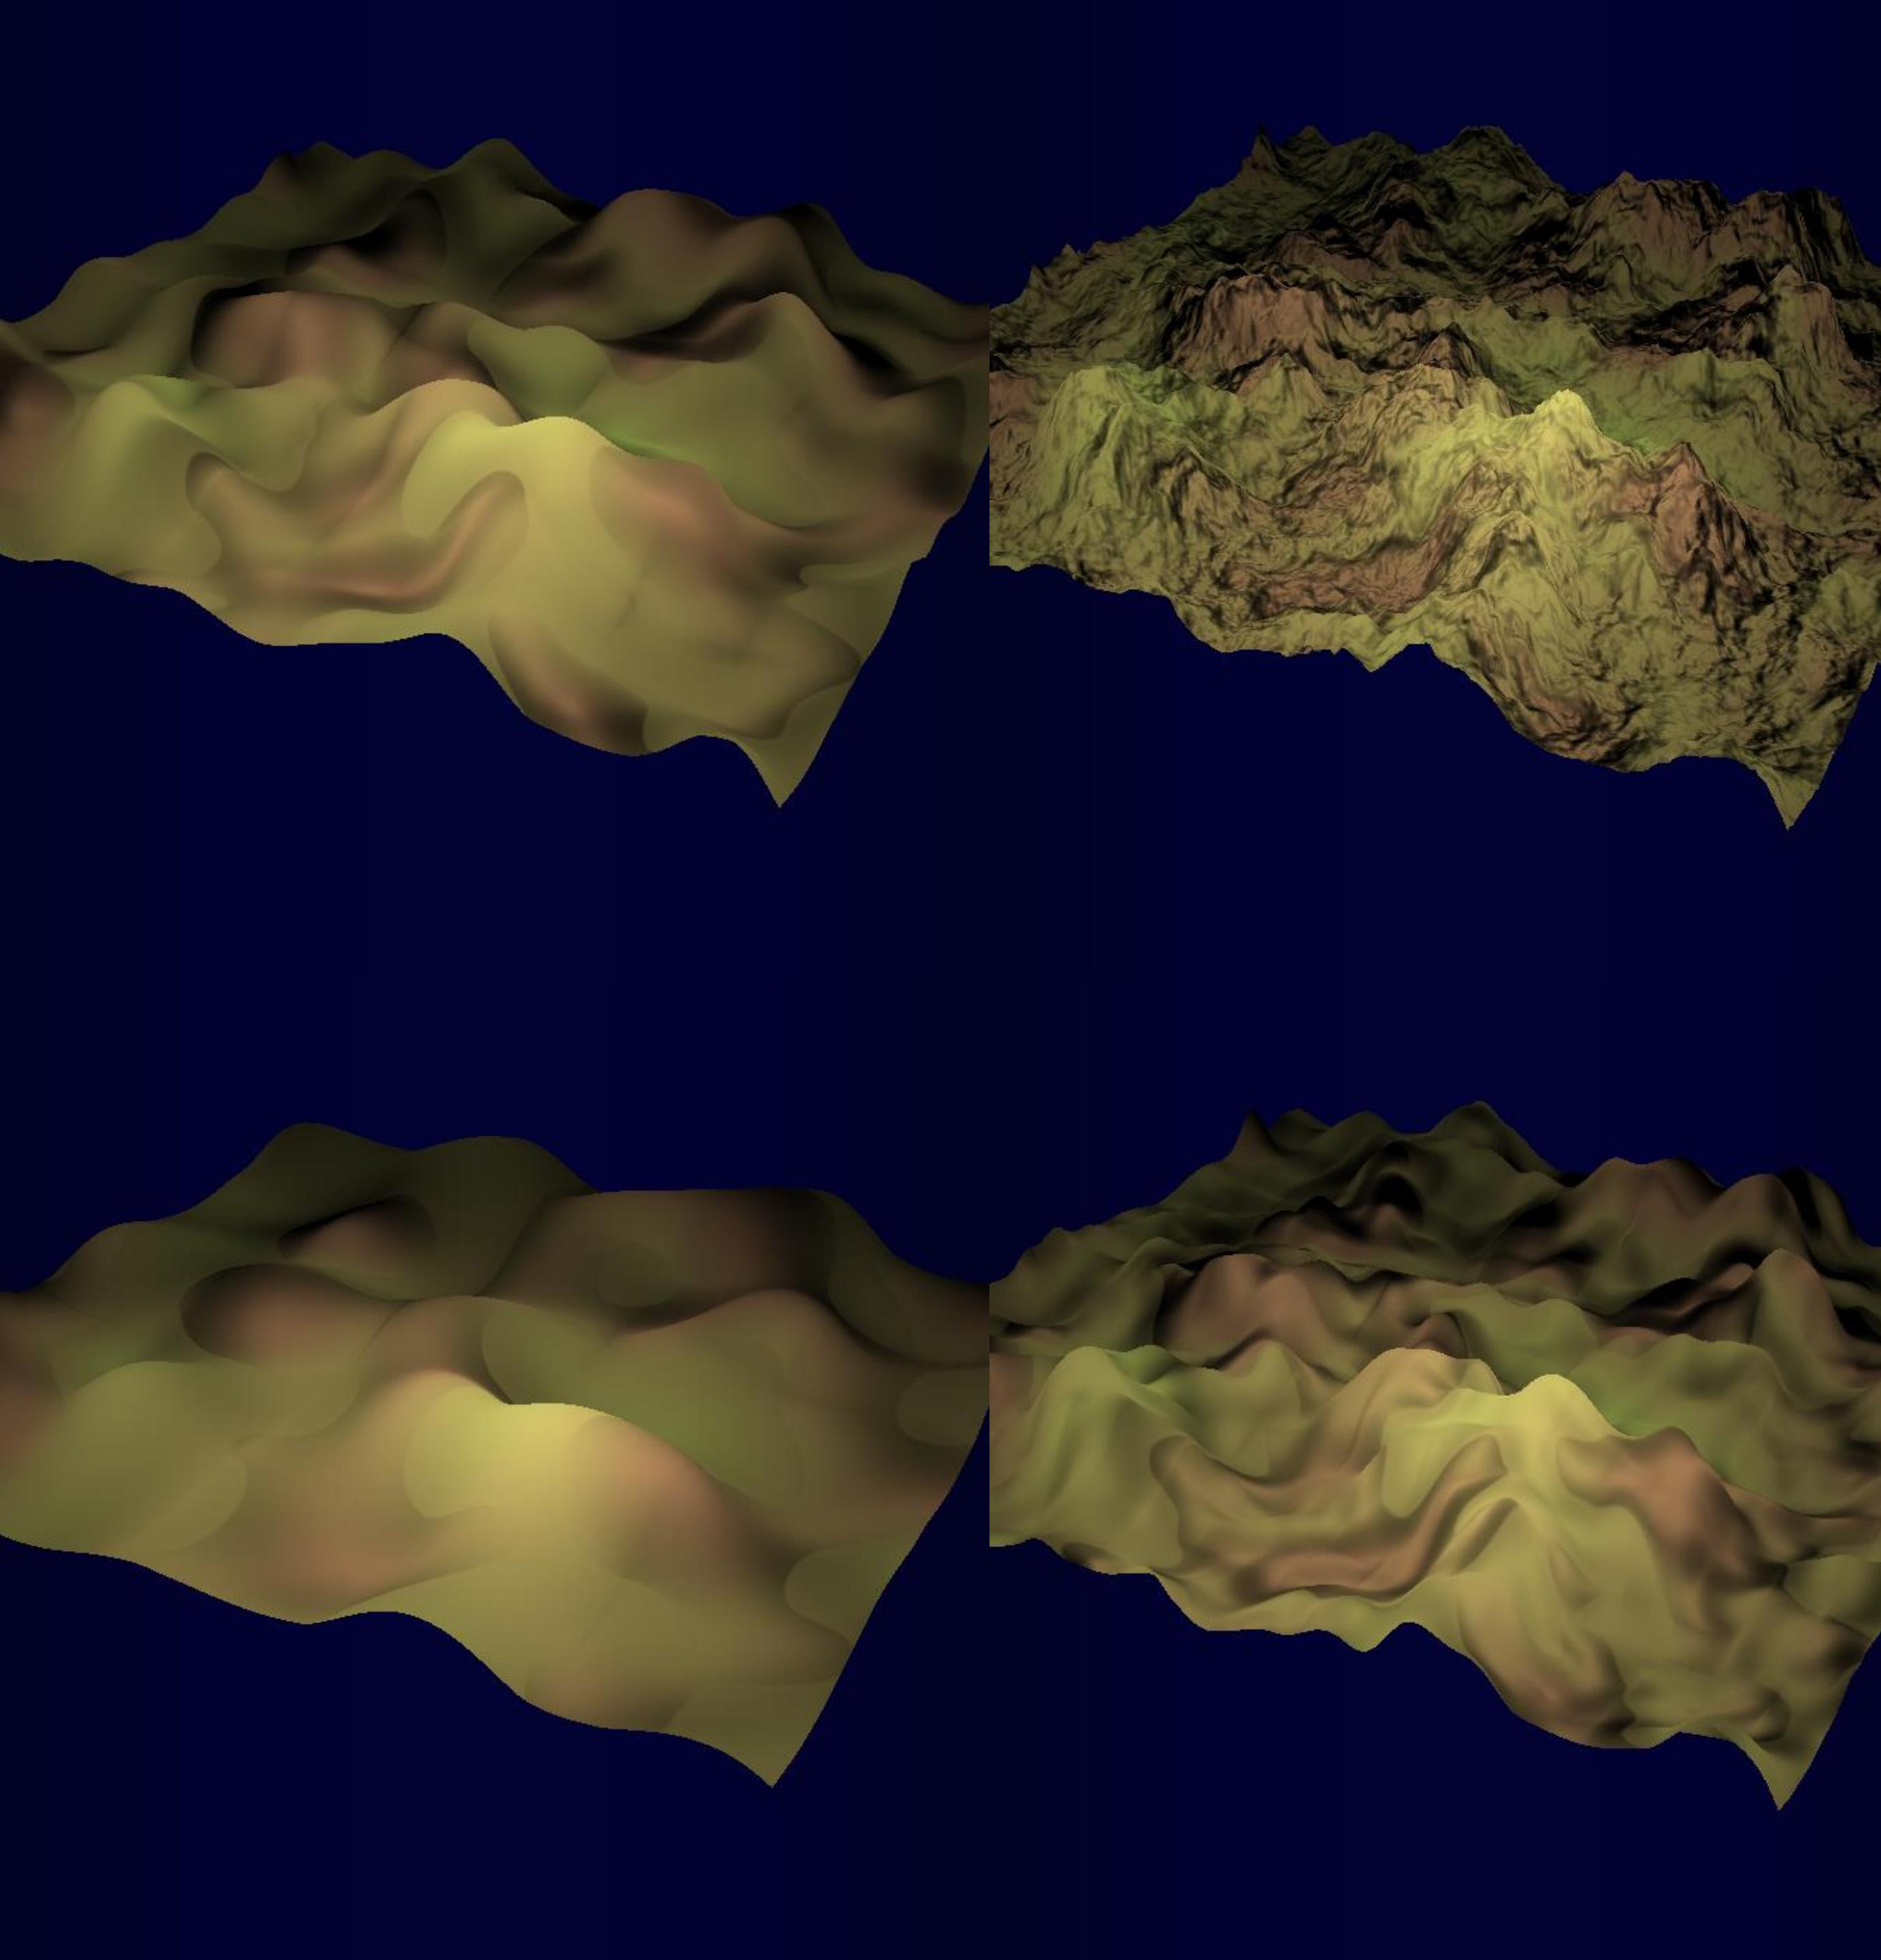


Natural-Looking Noise





Can you describe the difference between these two noise images?

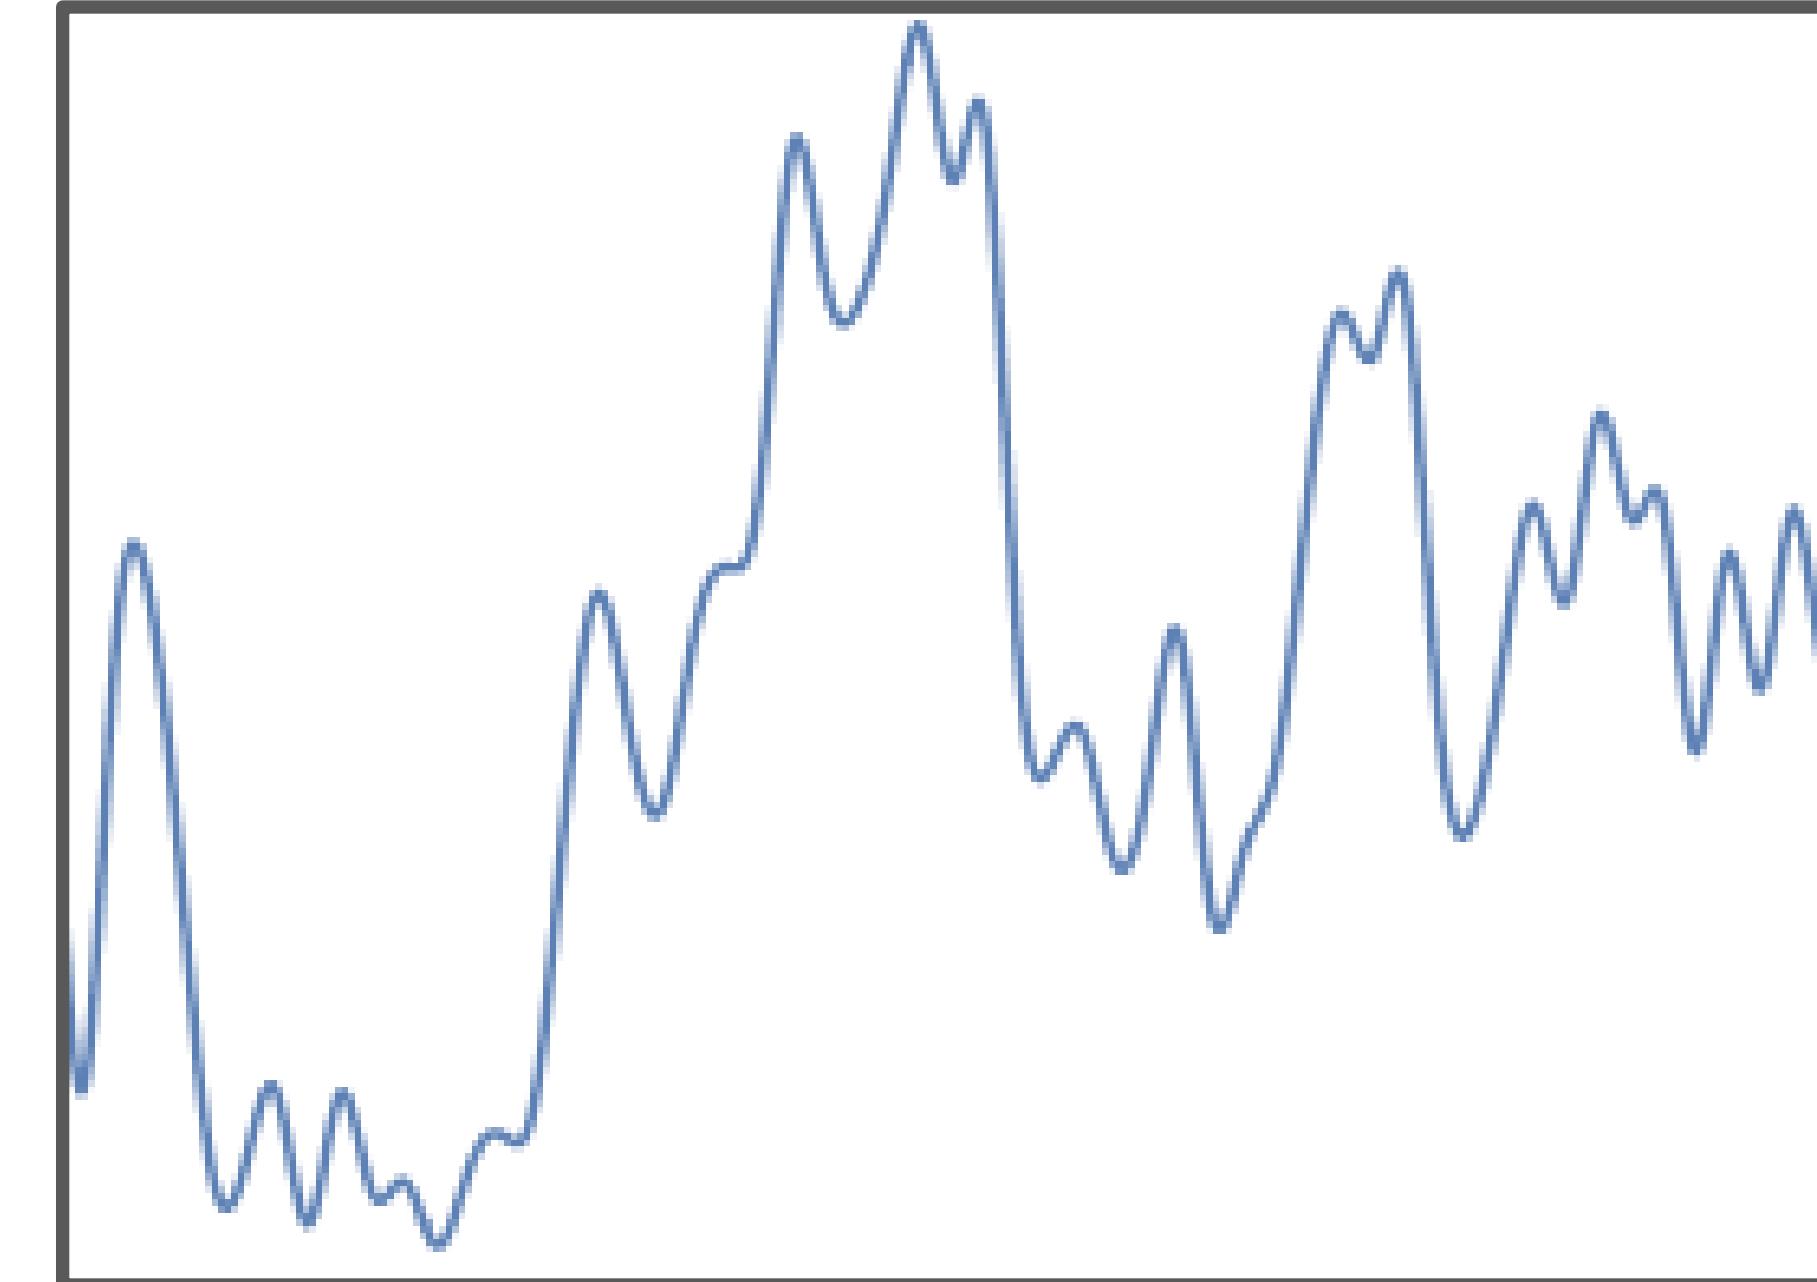
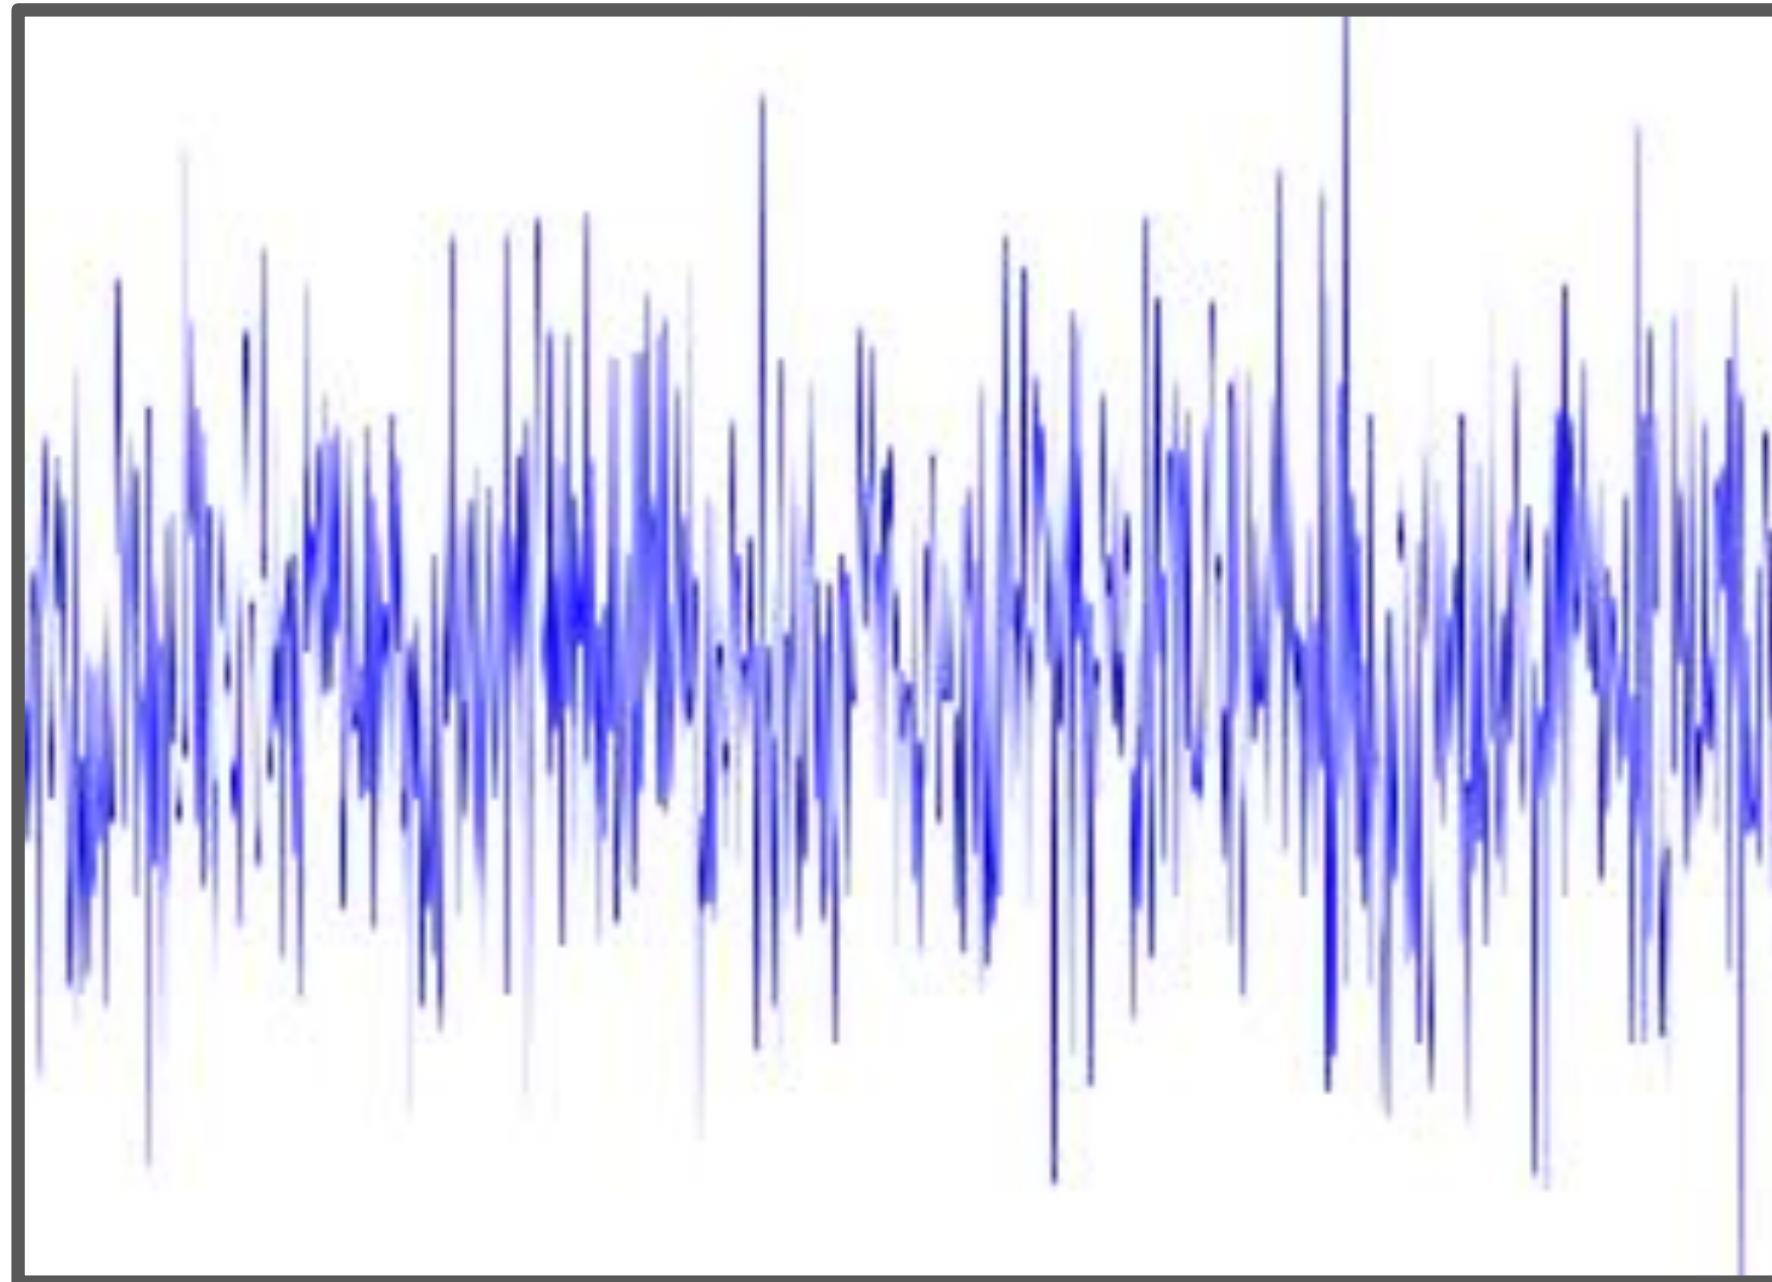


Can you tell the difference
among these terrain meshes?

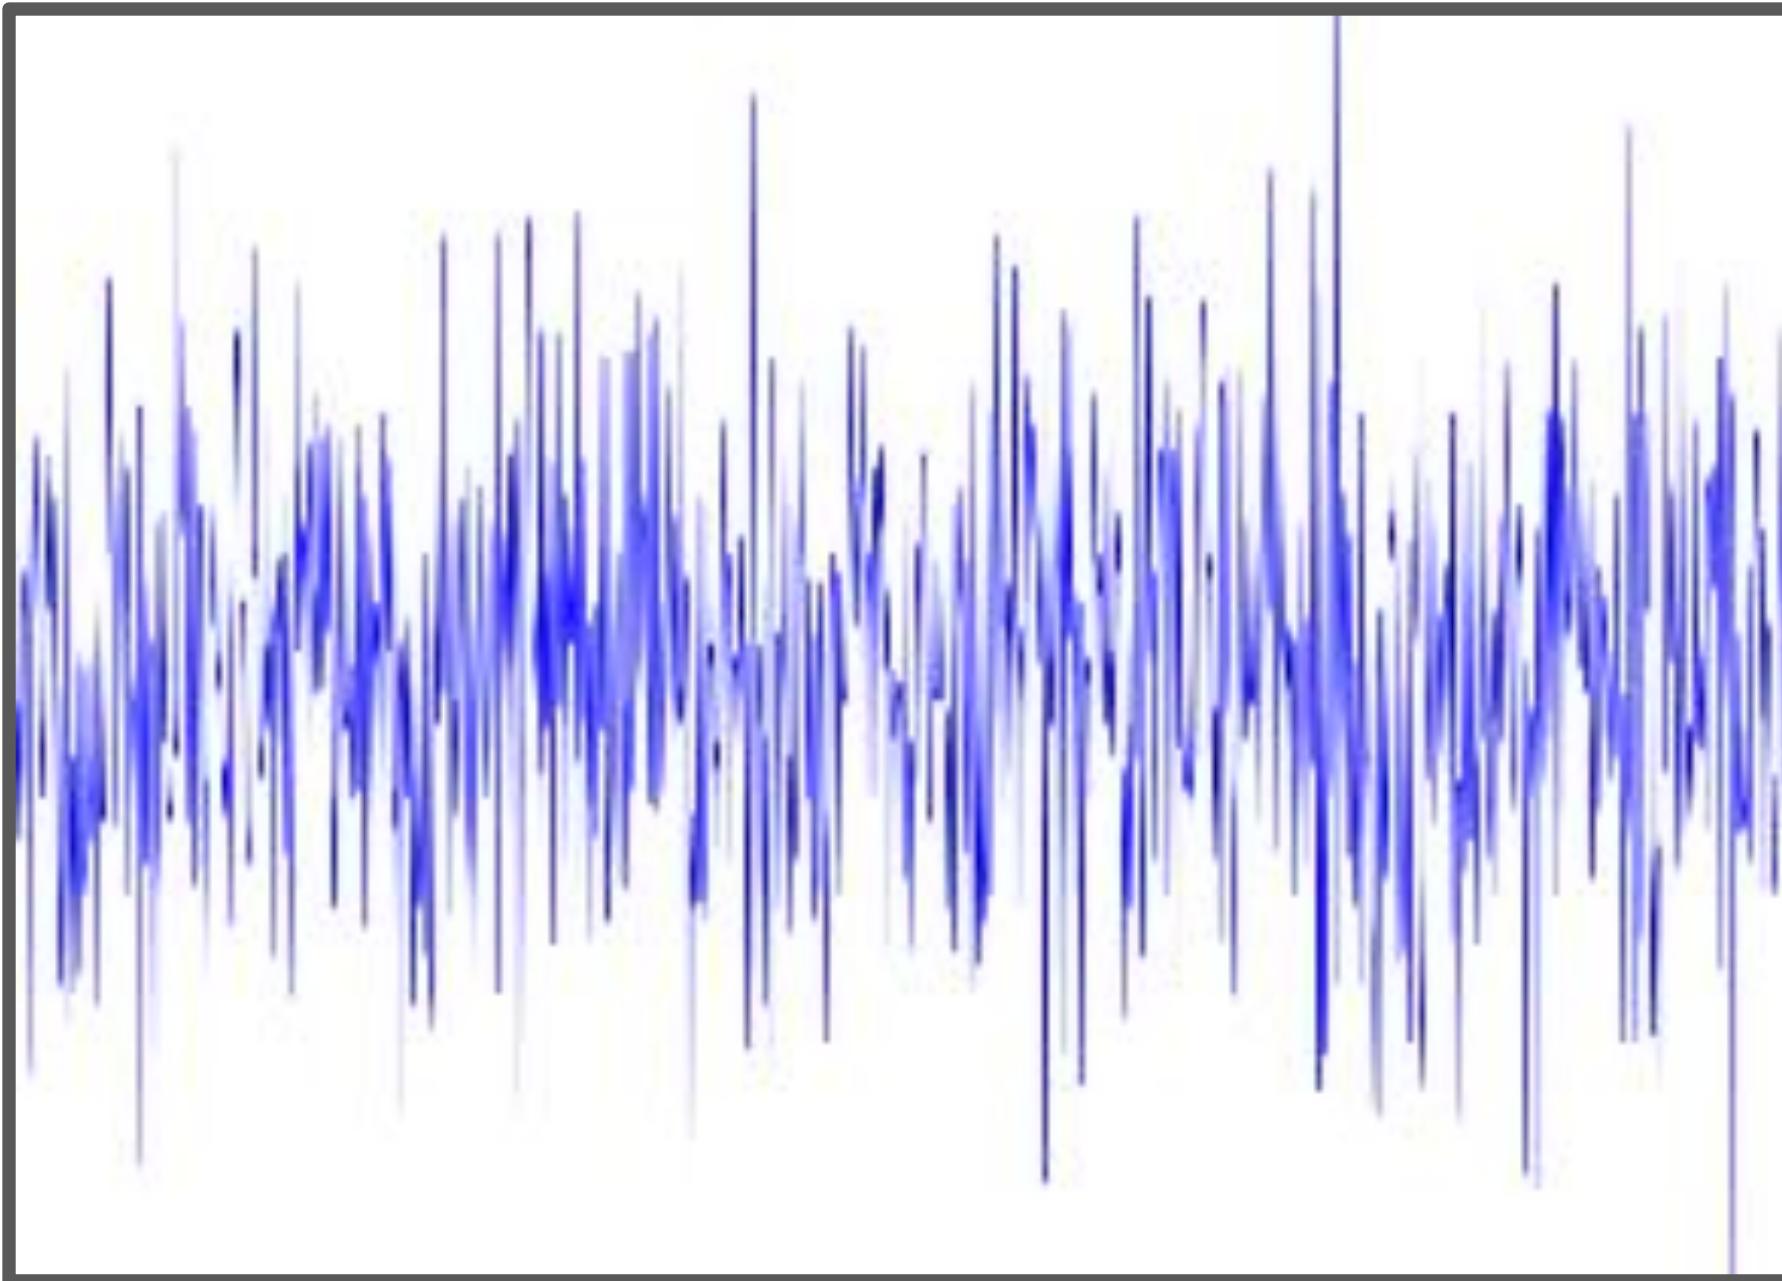
Which one looks more like a
real mountain?



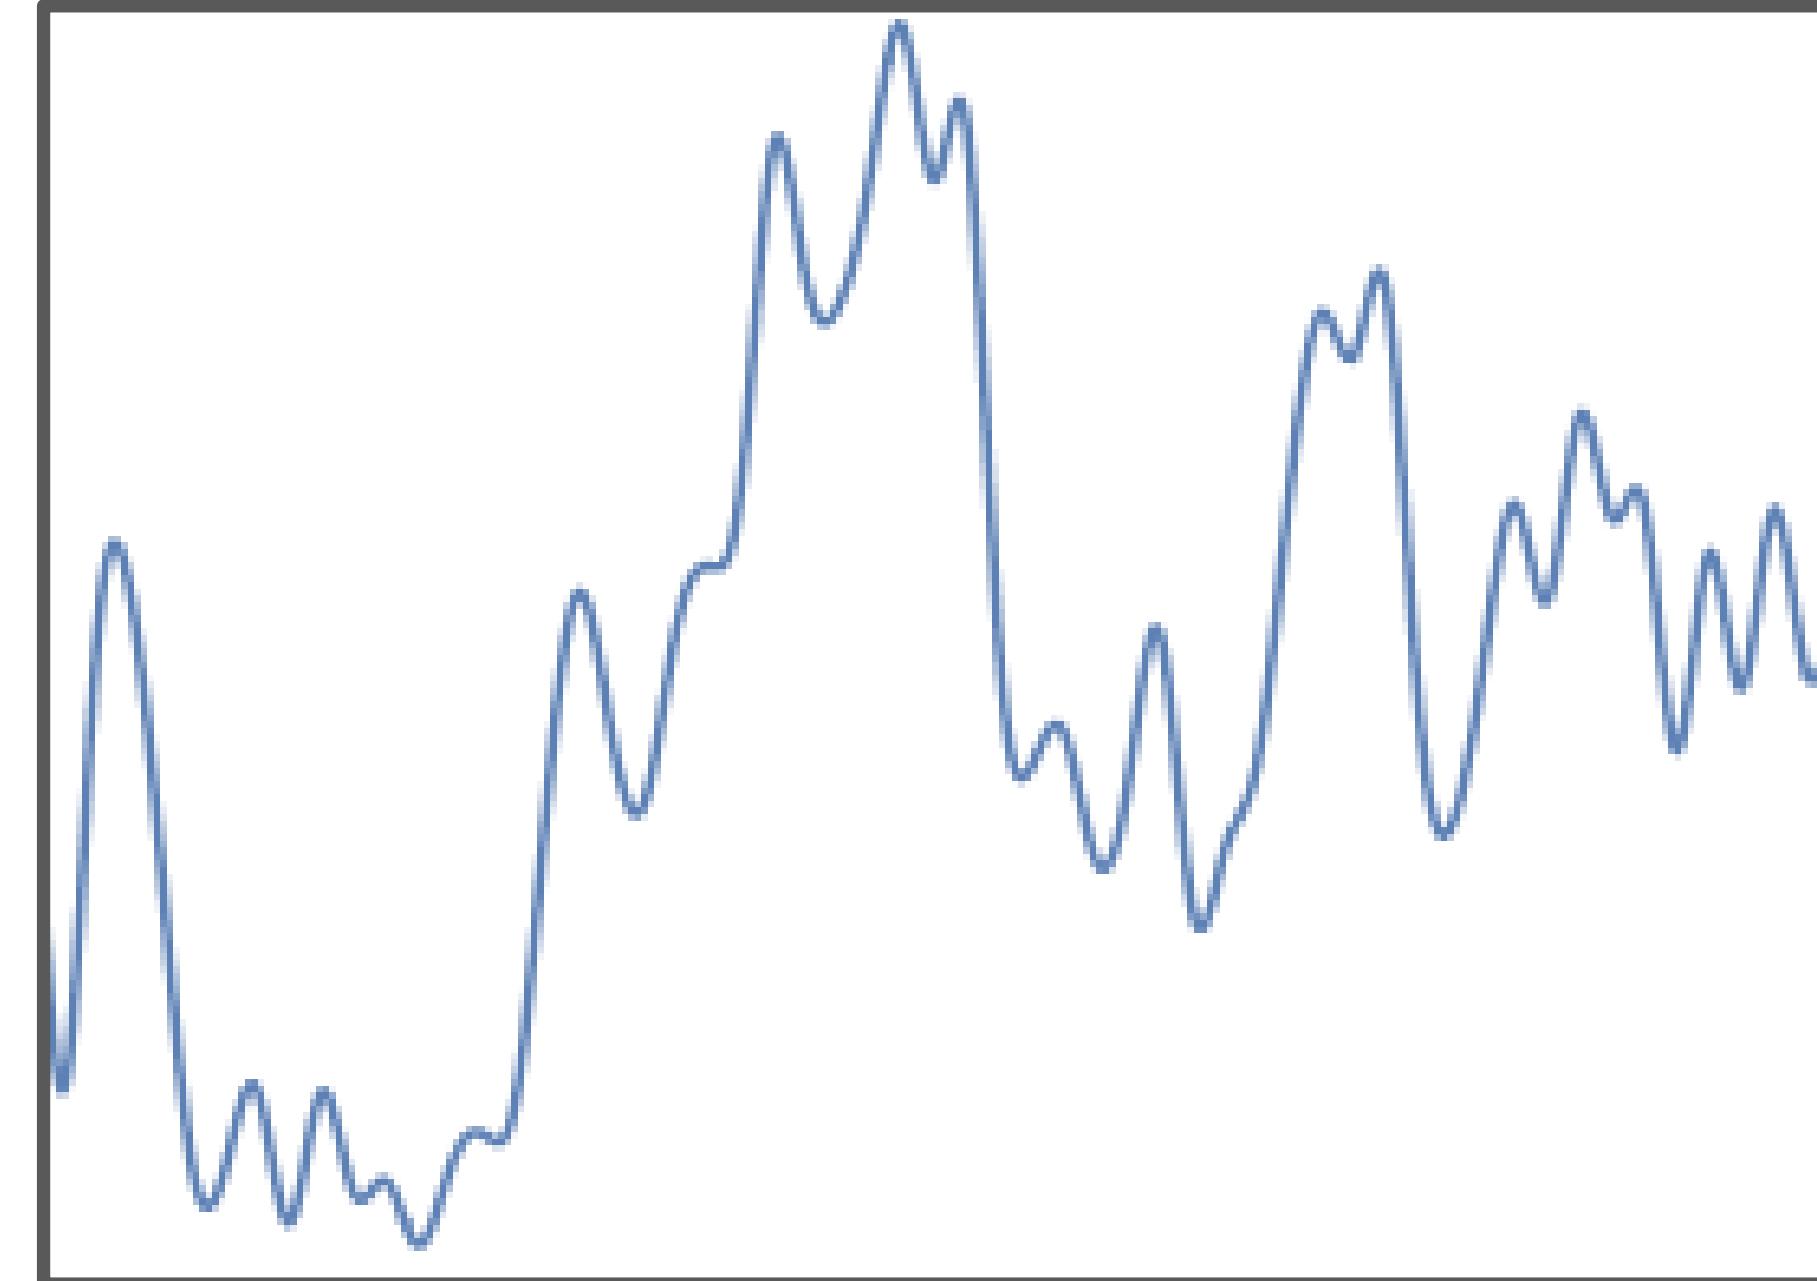
How do we describe the randomness of natural systems?



Maybe we can work on a simpler problem first:
Can we describe the difference between these two signals?



White Noise



Natural Noise

Natural-looking Noise Function

- A good noise is a noise that looks random, but **changes smoothly**
- The noise's features should have a **similar size** (band-limited)

Can we write down the mathematical function to describe a natural-looking noise?

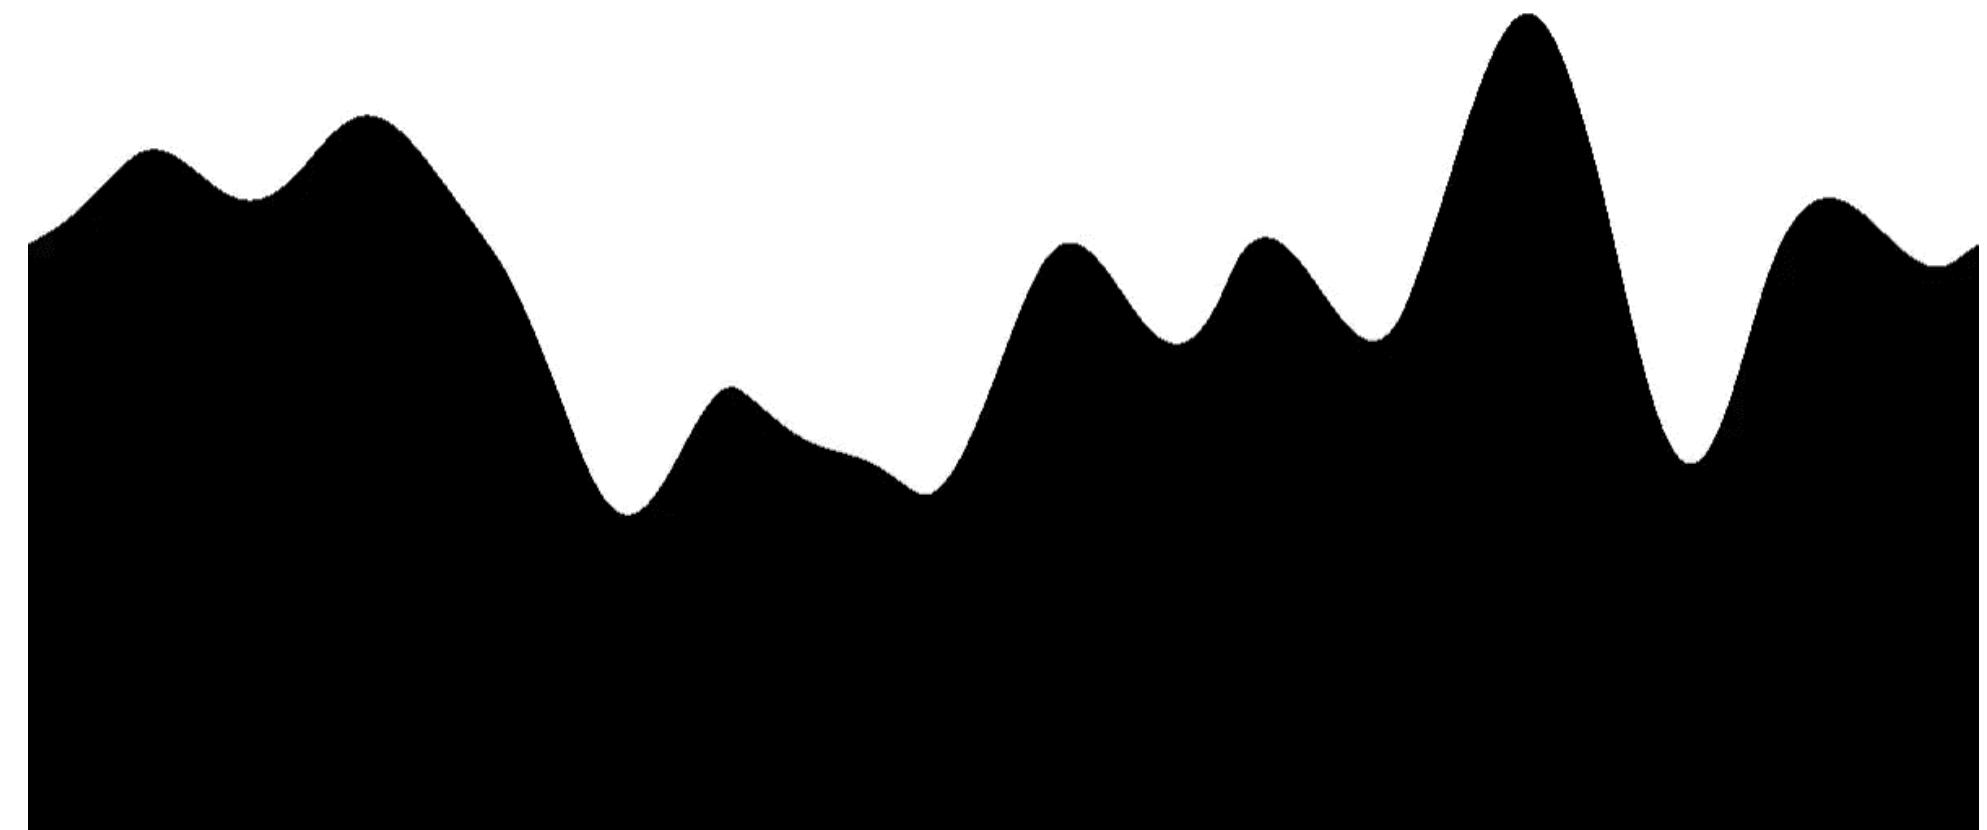


Natural Noise Functions

- Function: $\mathbf{R}^n \rightarrow [-1, 1]$, where $n = 1, 2, 3, \dots$
- Desirable properties:
 - **No obvious repetition**
 - **Frequency band-limited**
- Fundamental building block of most procedural textures

Let's figure out how to build these natural noise functions mathematically 😊

$$f(x) = ?$$



1D Natural Noise



2D Natural Noise

First Step: Let's create a **random** function in GLSL

- **Bad News:** There is no default `random()` function in GLSL. You need to implement it on your own
- **Good News:** You can implement it with one line of code
- The idea is: We want to return a “random number” with input coordinates x and y
- For the same x and y , we always return the same number
- This is called a **hash** function (or pseudorandom number, if you think of x and y as “seeds”)

$$r = h(x, y)$$

Random number

Seeds

How do we design a hash function?



Hash Function Examples

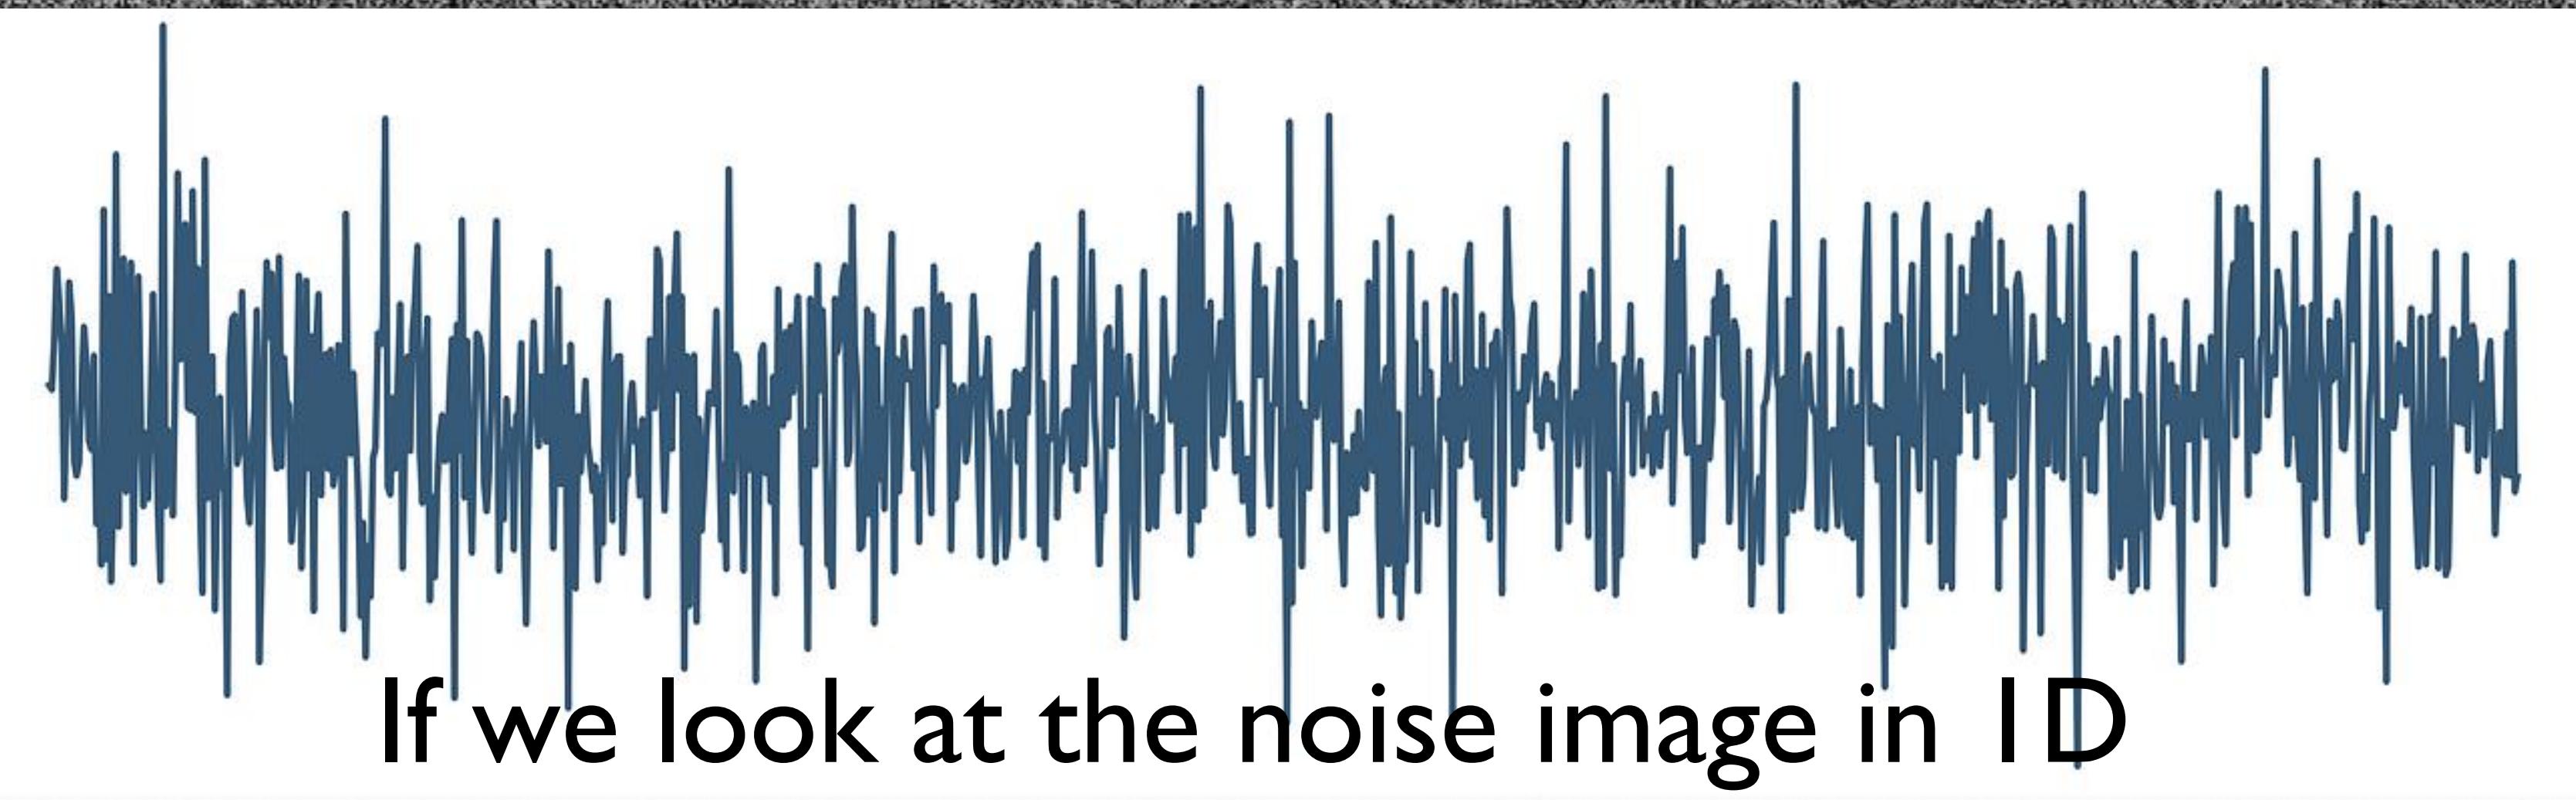
Example:

```
float hash(vec2 v)
{
    return fract(sin(dot(v.xy ,vec2(12.9898,78.233))) * 43758.5453);
}
```

Implementation Tip:

- You may use the combination of GLSL built-in functions like sin, cos, mod, fract, dot, etc. to design your own hash (noise) function
- Search “noise” or “hash” in ShaderToy and you will find many different versions of implementation



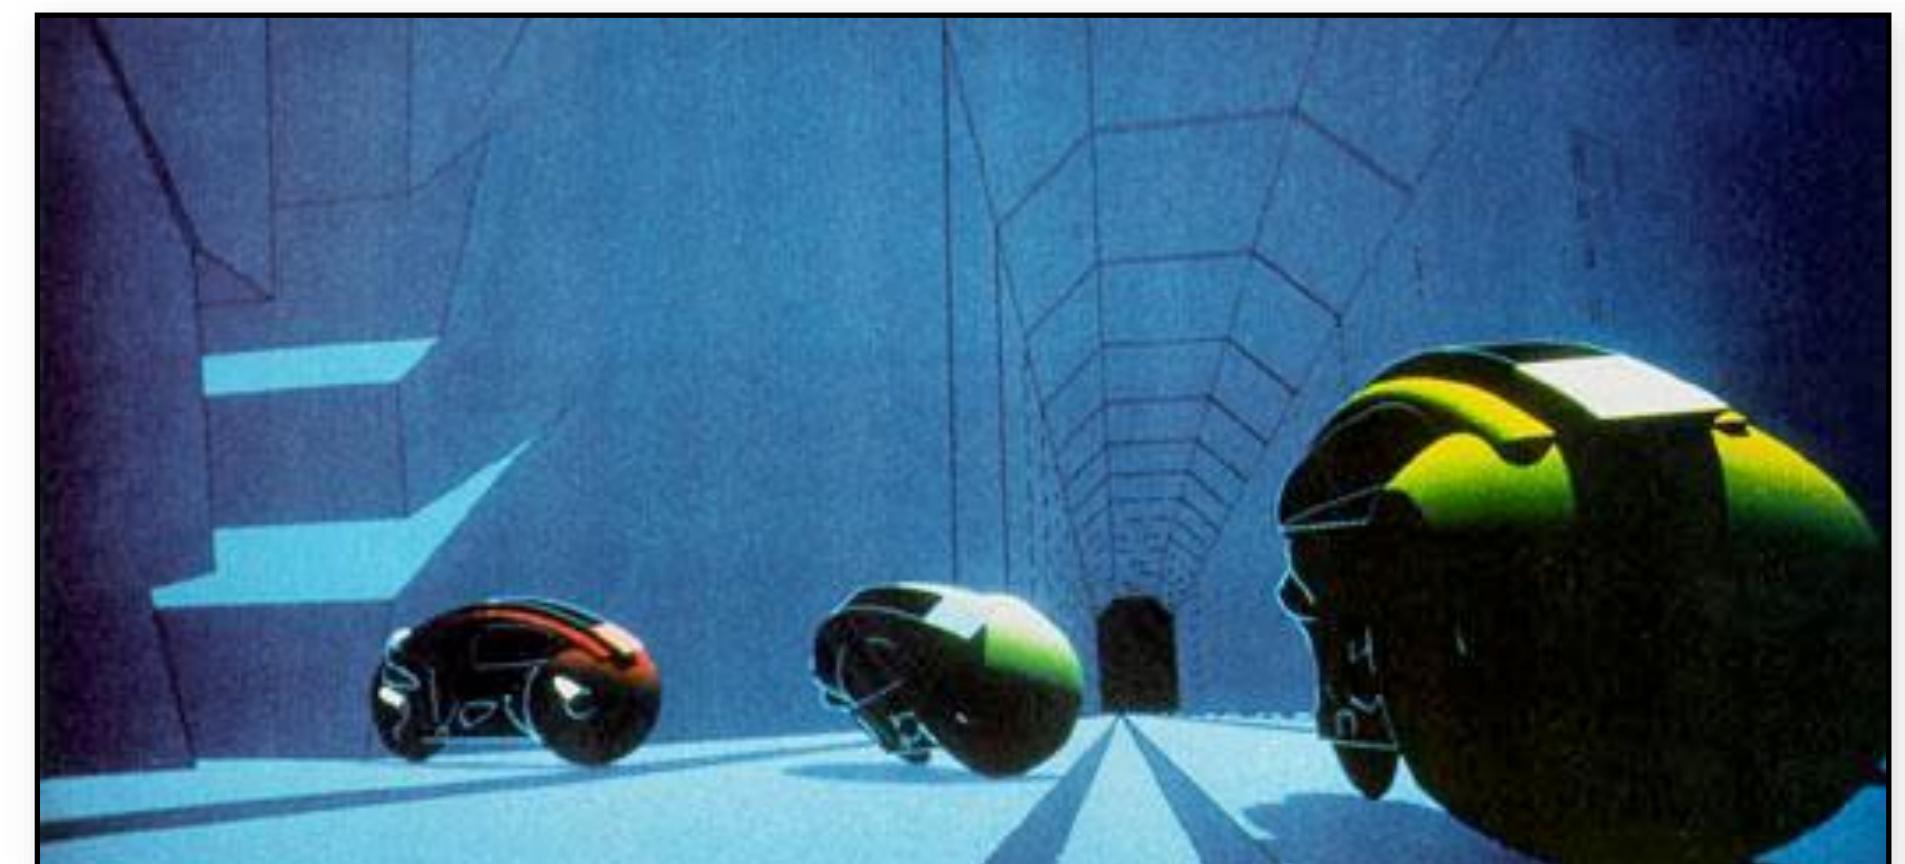
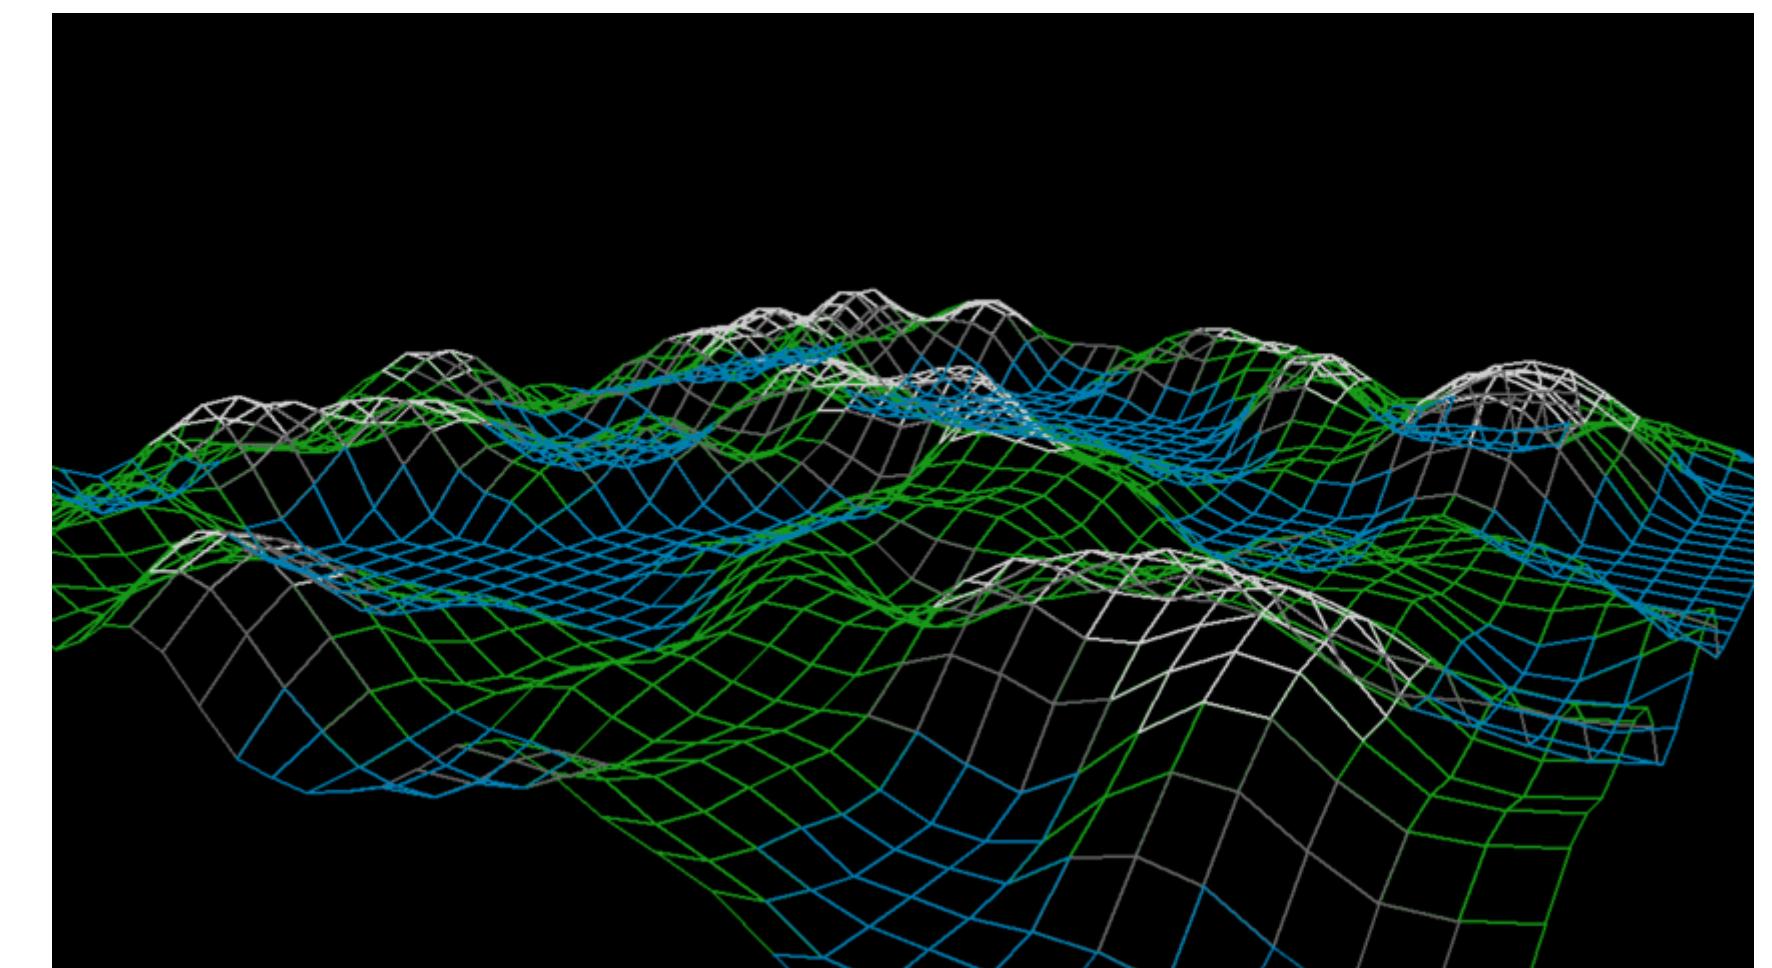
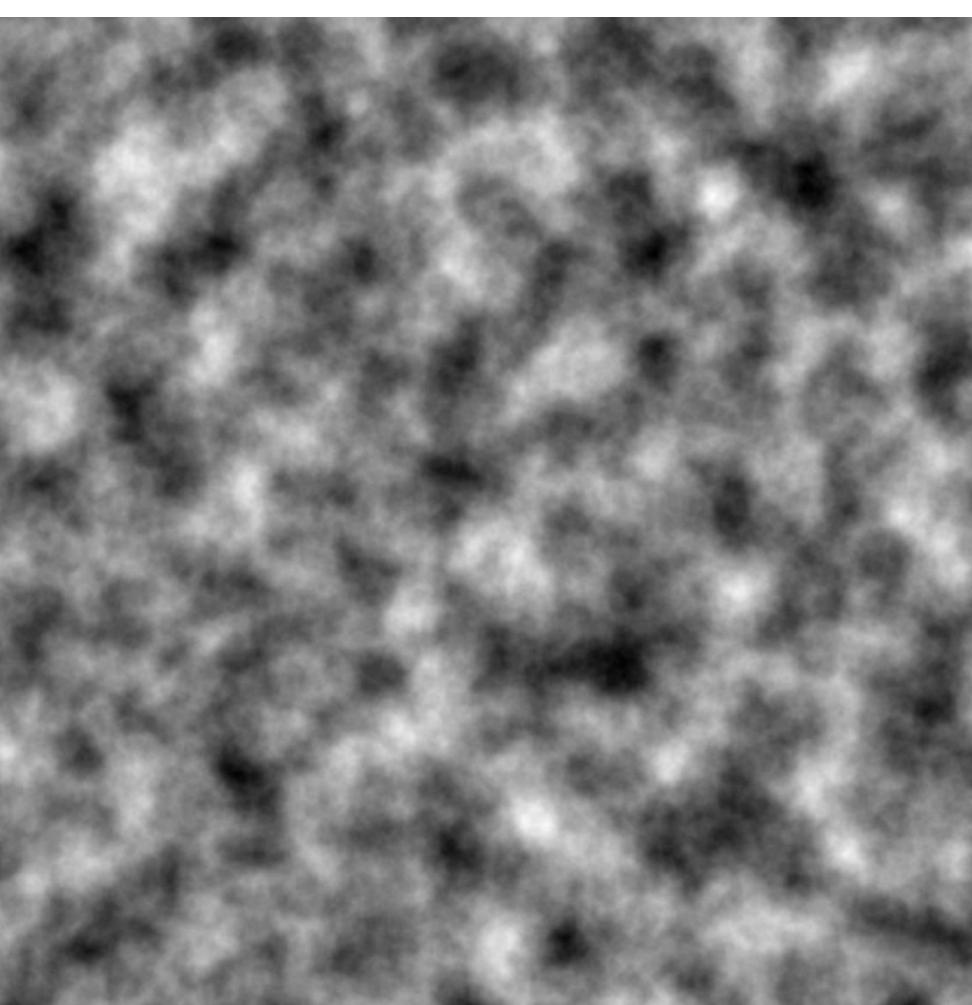


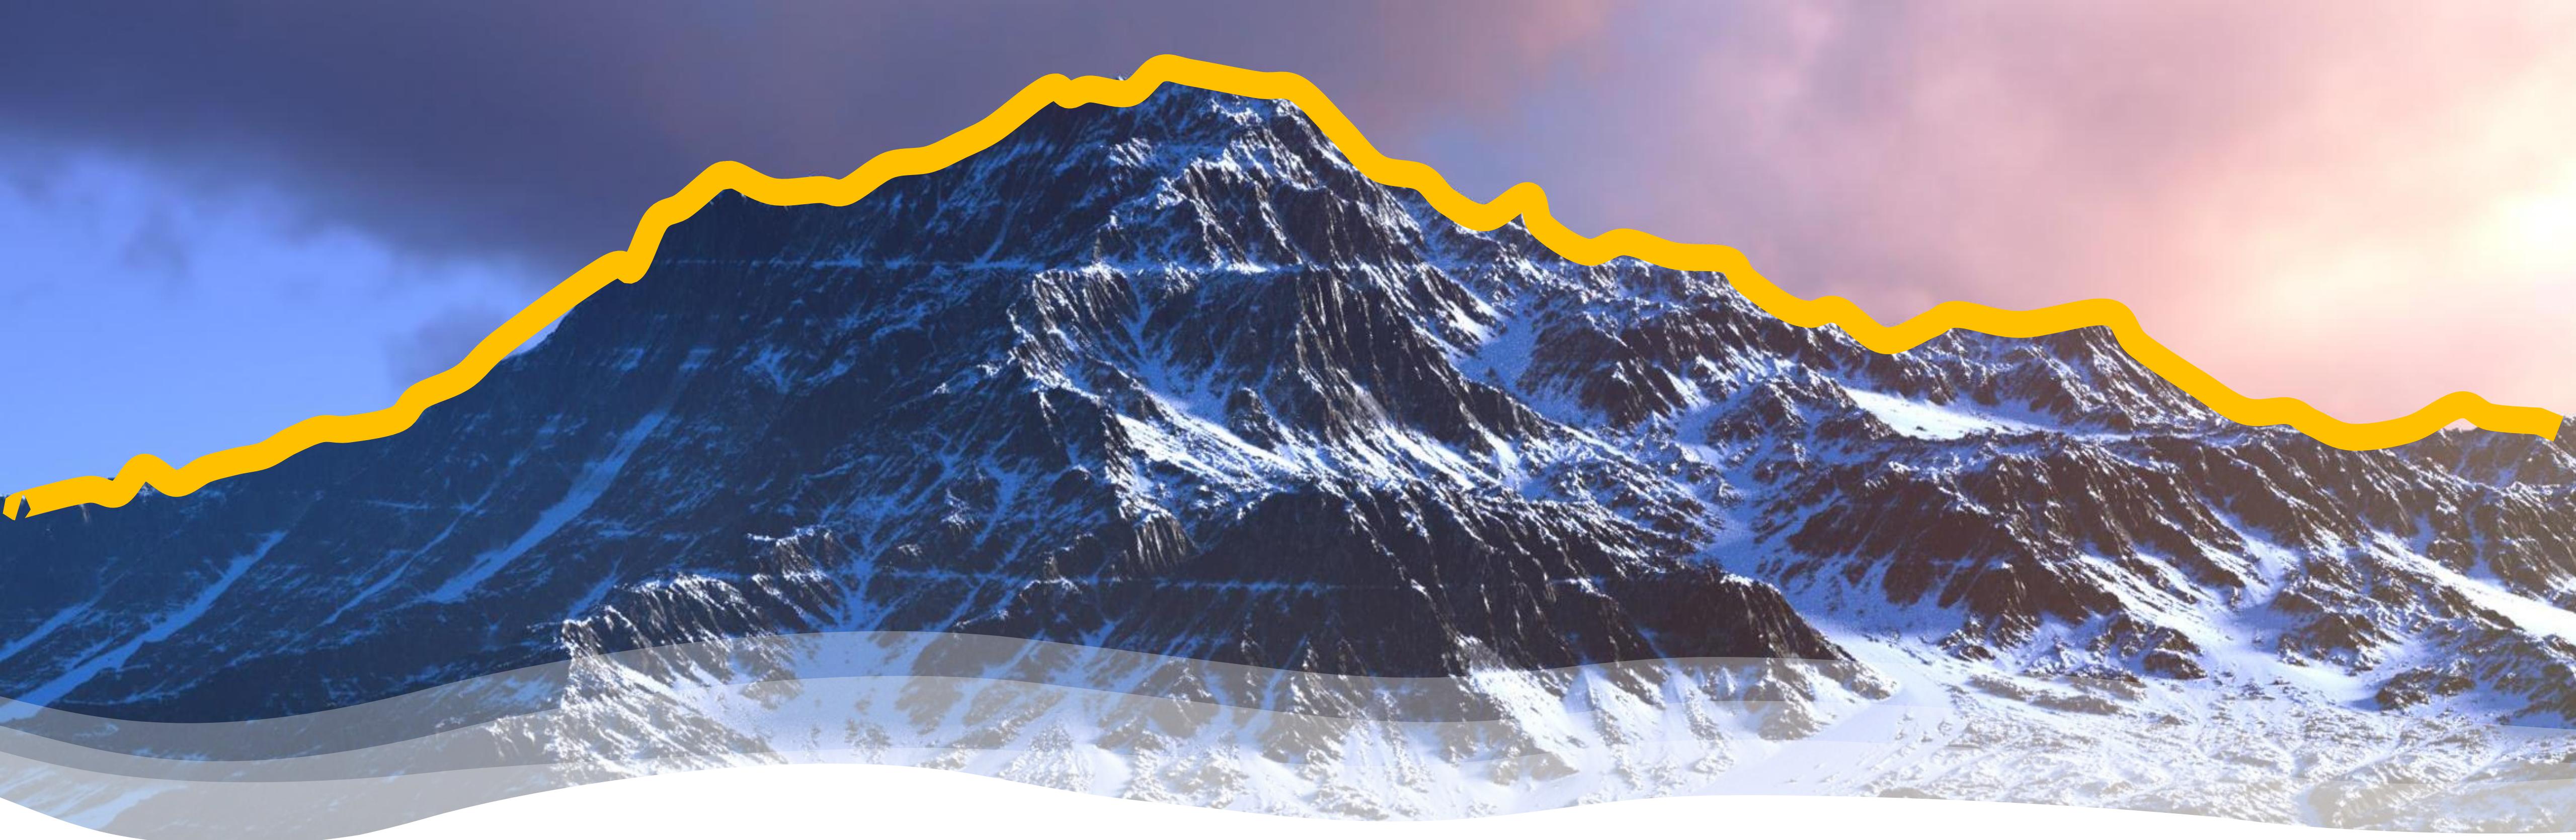
If you call the hash function for each fragment in the fragment shader:
You will get noise images that look like an untuned TV screen

Question: How do we convert such a noise image to a natural-looking texture?

Perlin Noise

- Perlin noise, invented by Ken Perlin in 1982
 - First used in the movie Tron!
- In 1996, Ken Perlin received an Academy Award for Technical Achievement for the development of Perlin noise.





Perlin noise is one of the greatest invention in the history of computer graphics

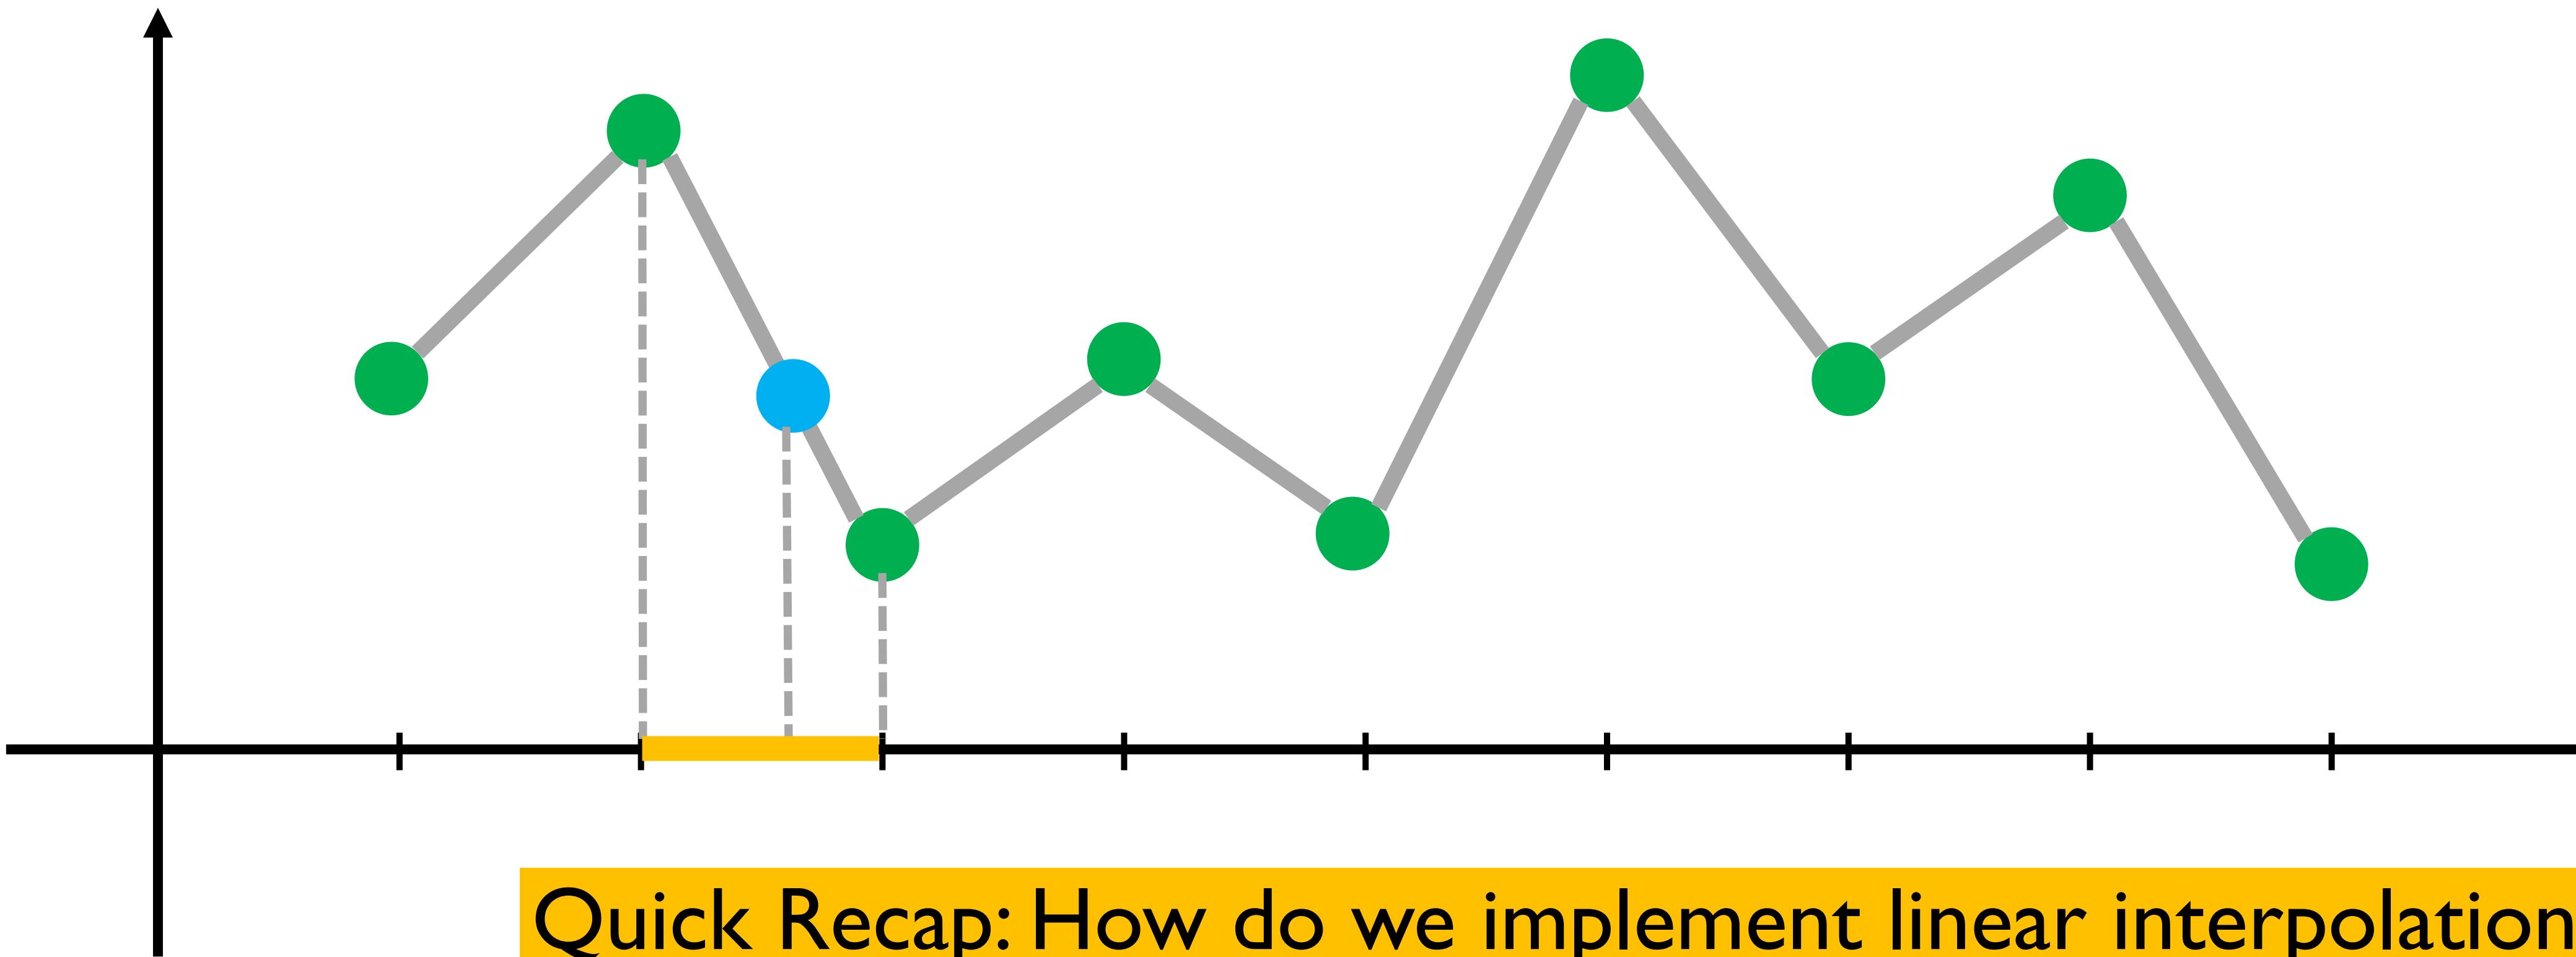
- It serves as the mathematical foundation for all kinds of procedural model generation --- ranging from texture colors to terrain meshes

1D Perlin Noise



A Naïve Idea: Noise with Linear Interpolation

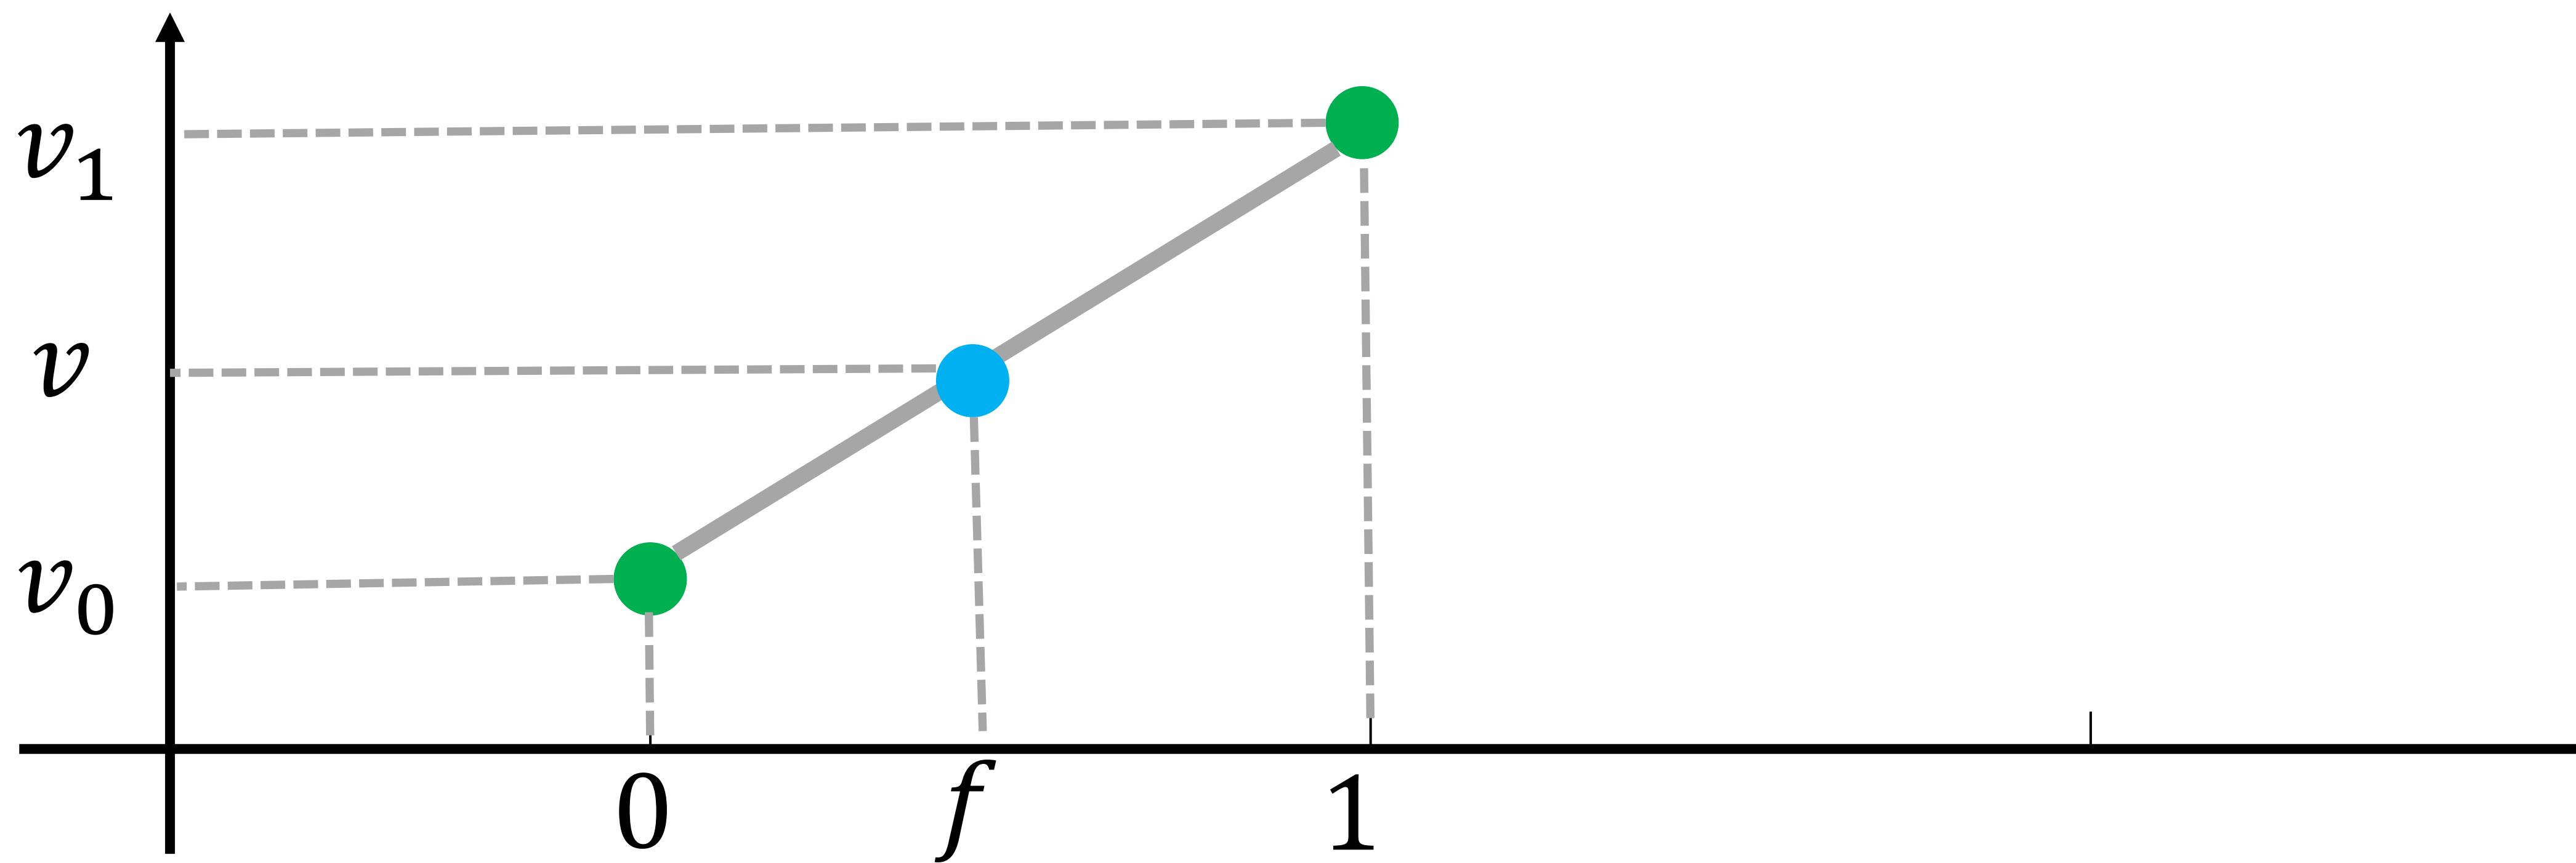
- Discretize the domain with a **uniform one-dimensional (1D) grid**
- Generate random values for each grid node using a **1D Hash() function**
- For a point, we first determine which grid cell it is in, and then interpolate its value using the neighboring grid nodes with **linear interpolation**



GLSL Function: **mix()**

float mix(float v0, float v1, float f)

- Linearly interpolate between two values
- The returned value is calculated as $v = (1 - f) * v_0 + f * v_1$



Pseudocode for 1D Linearly Interpolated Noise

Input: x

Output: noise

```
int i = floor(x);
```

```
float f = fract(x);
```

```
float v0=hash(i);
```

```
float v1=hash(i+1);
```

```
float noise = (1-f)*v0+f*v1;
```

```
return noise;
```

GLSL built-in `floor()`:

- find the nearest integer less than or equal to the parameter
- e.g., `floor(3.5)` will return 3

GLSL built-in `fract()`:

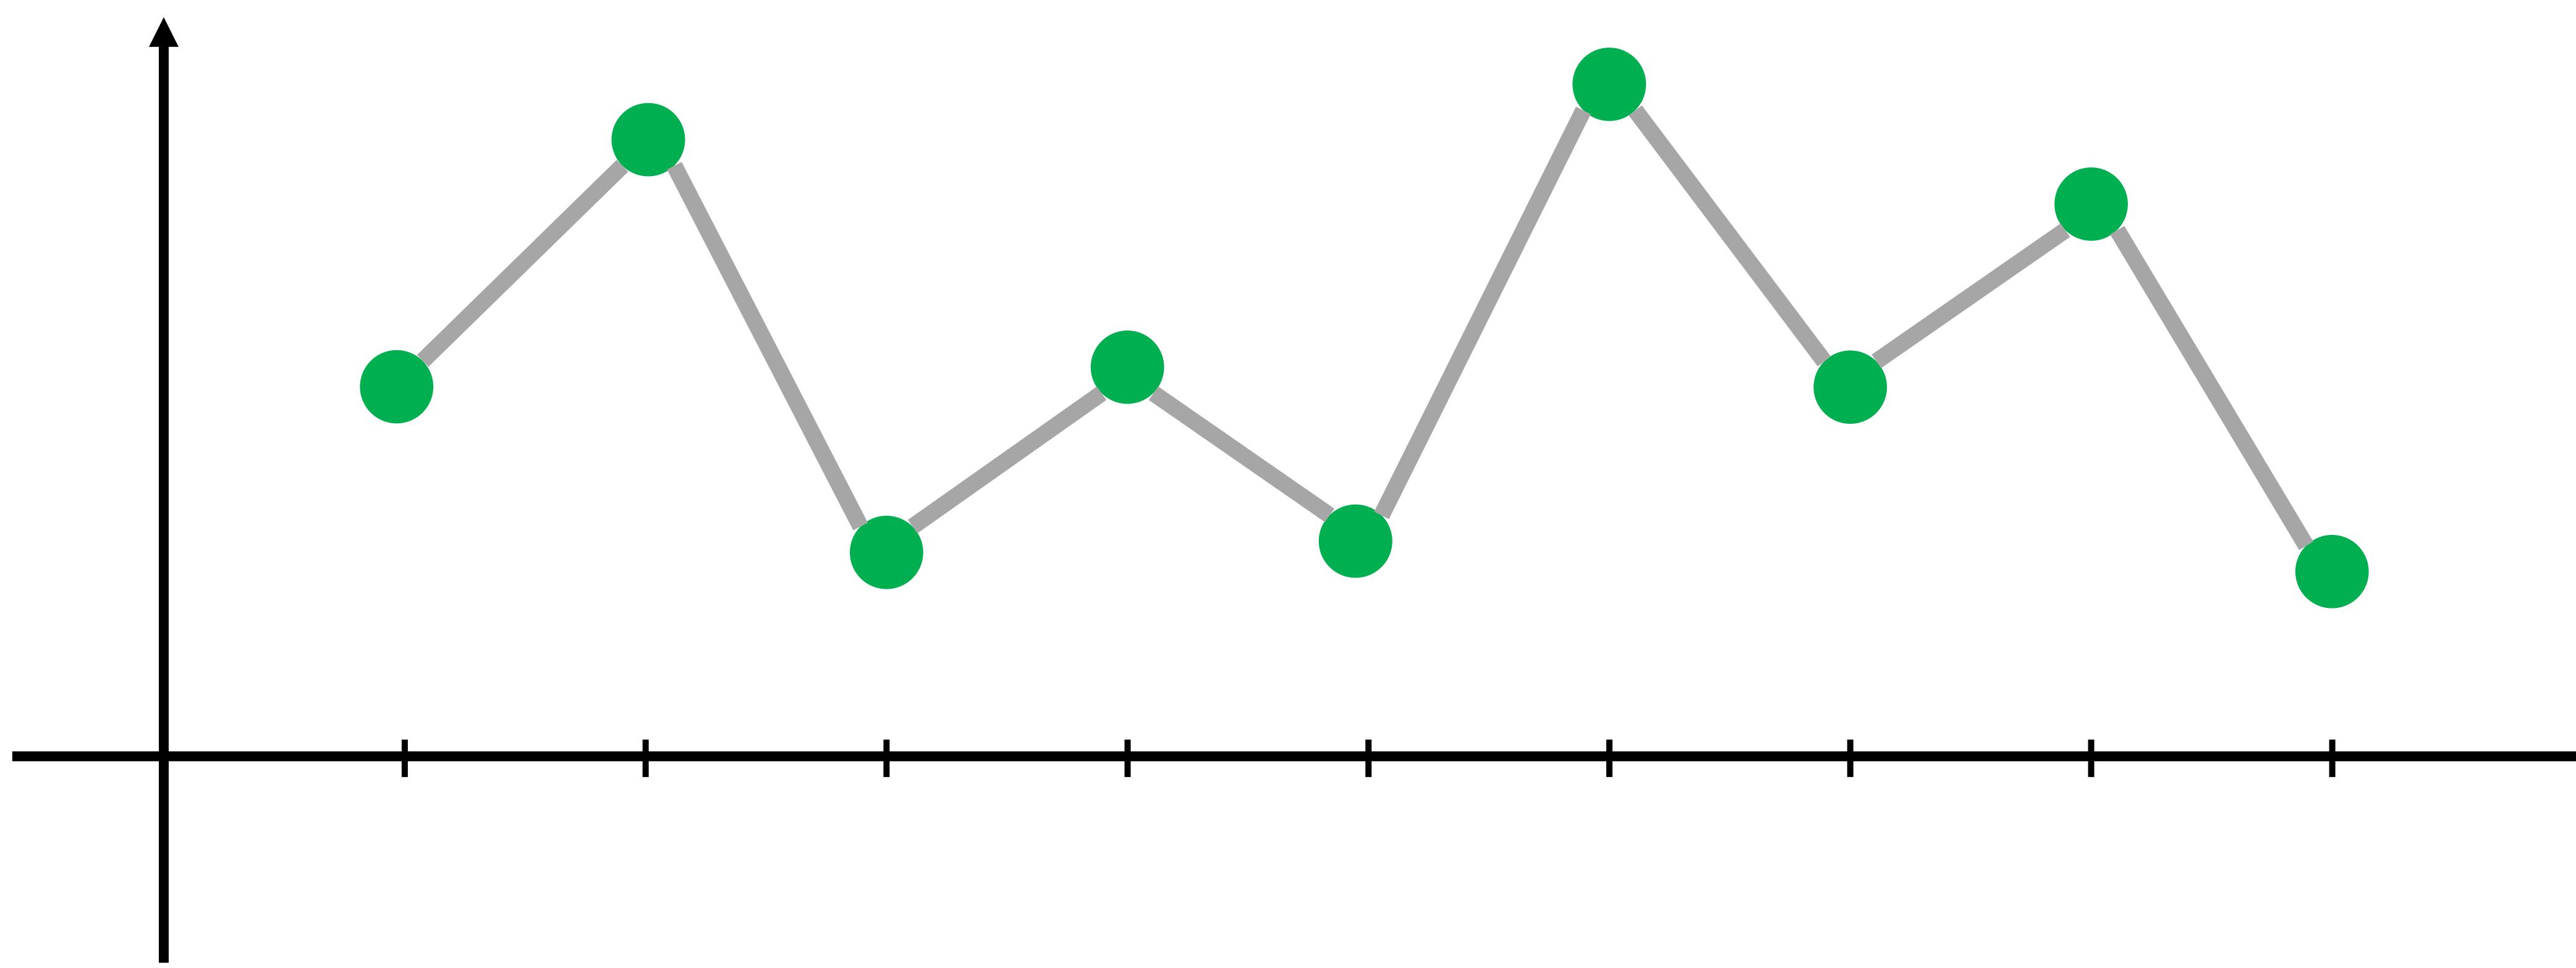
- compute the fractional part of the argument
- e.g., `fract(1.52)` will return 0.52

Or, we can call `mix()` in GLSL:
 $\text{noise} = \text{mix}(v0, v1, s);$



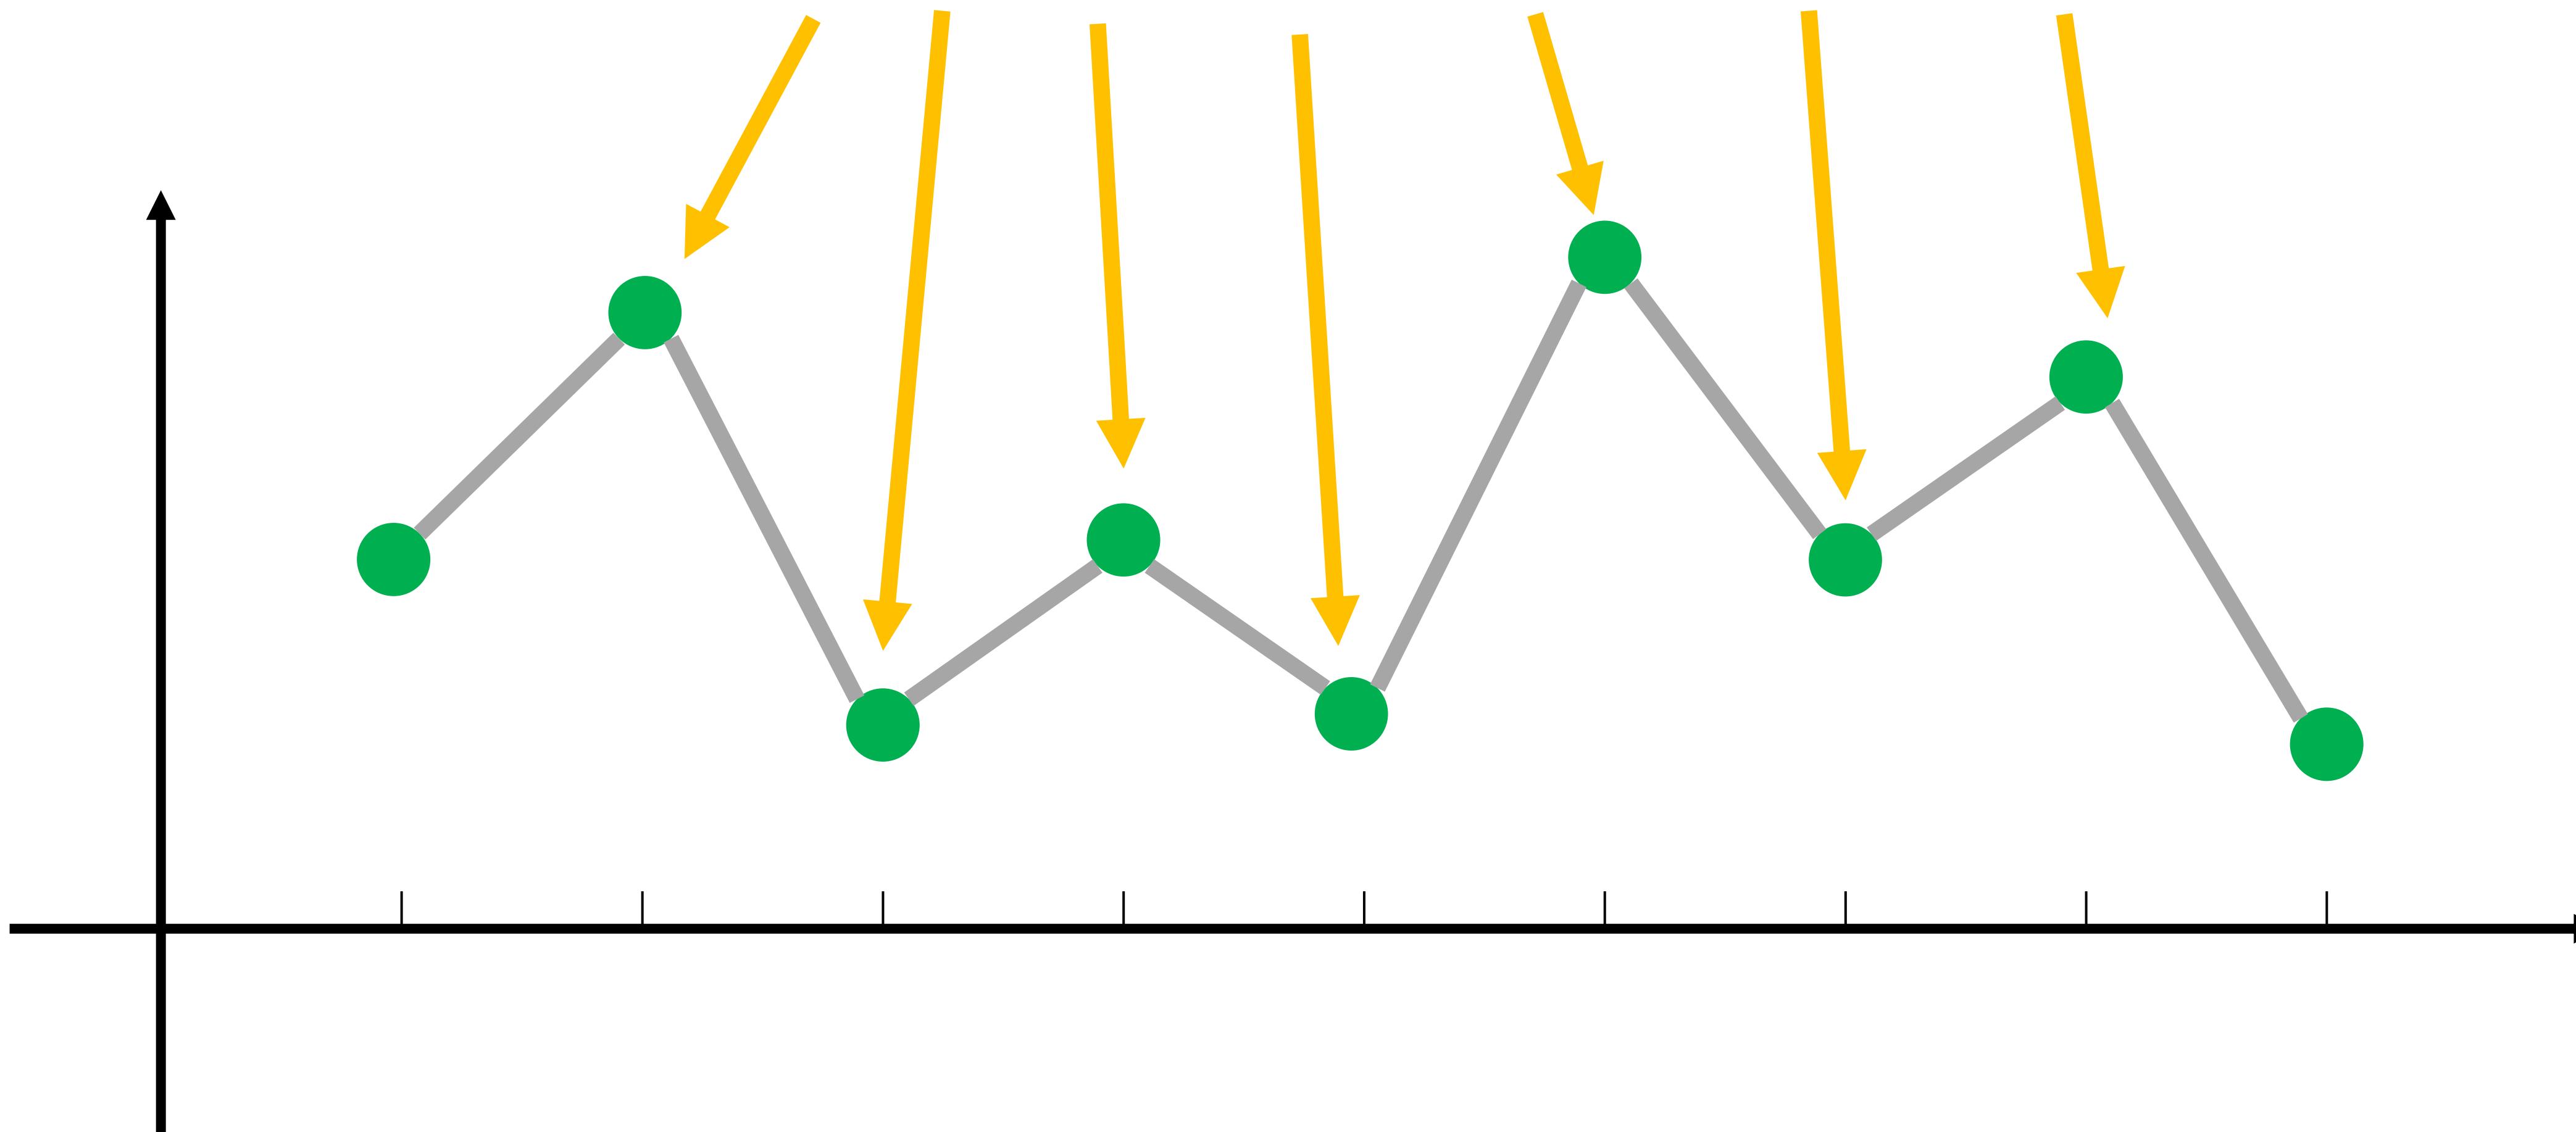
Problems with Linear Interpolation

- **Problem I:** The noise is not smooth
- **Problem II:** The noise is not band-limited
- **Problem III:** The noise does not have multi-scale features



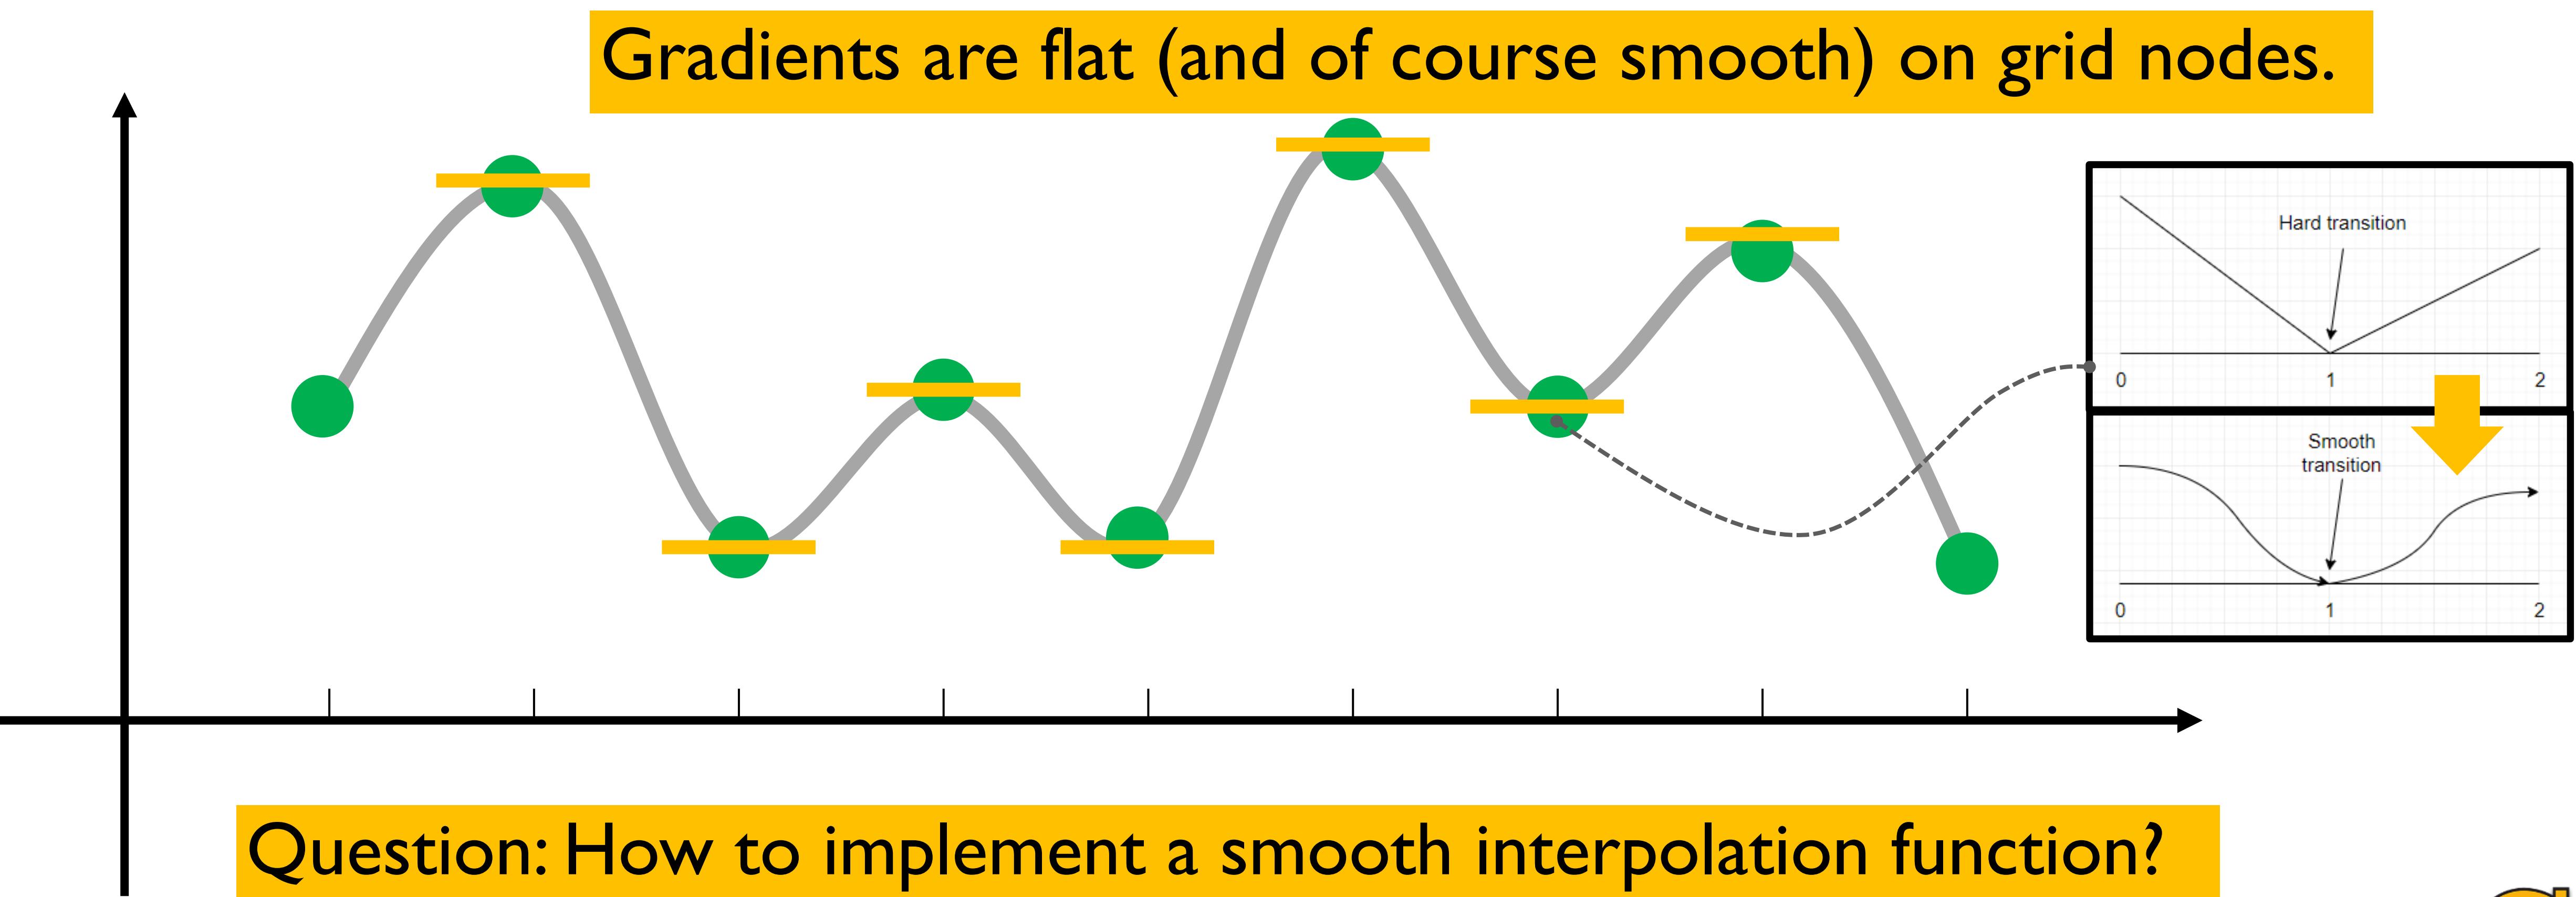
Problem I: Values are not smooth

Gradients change sharply on grid notes.

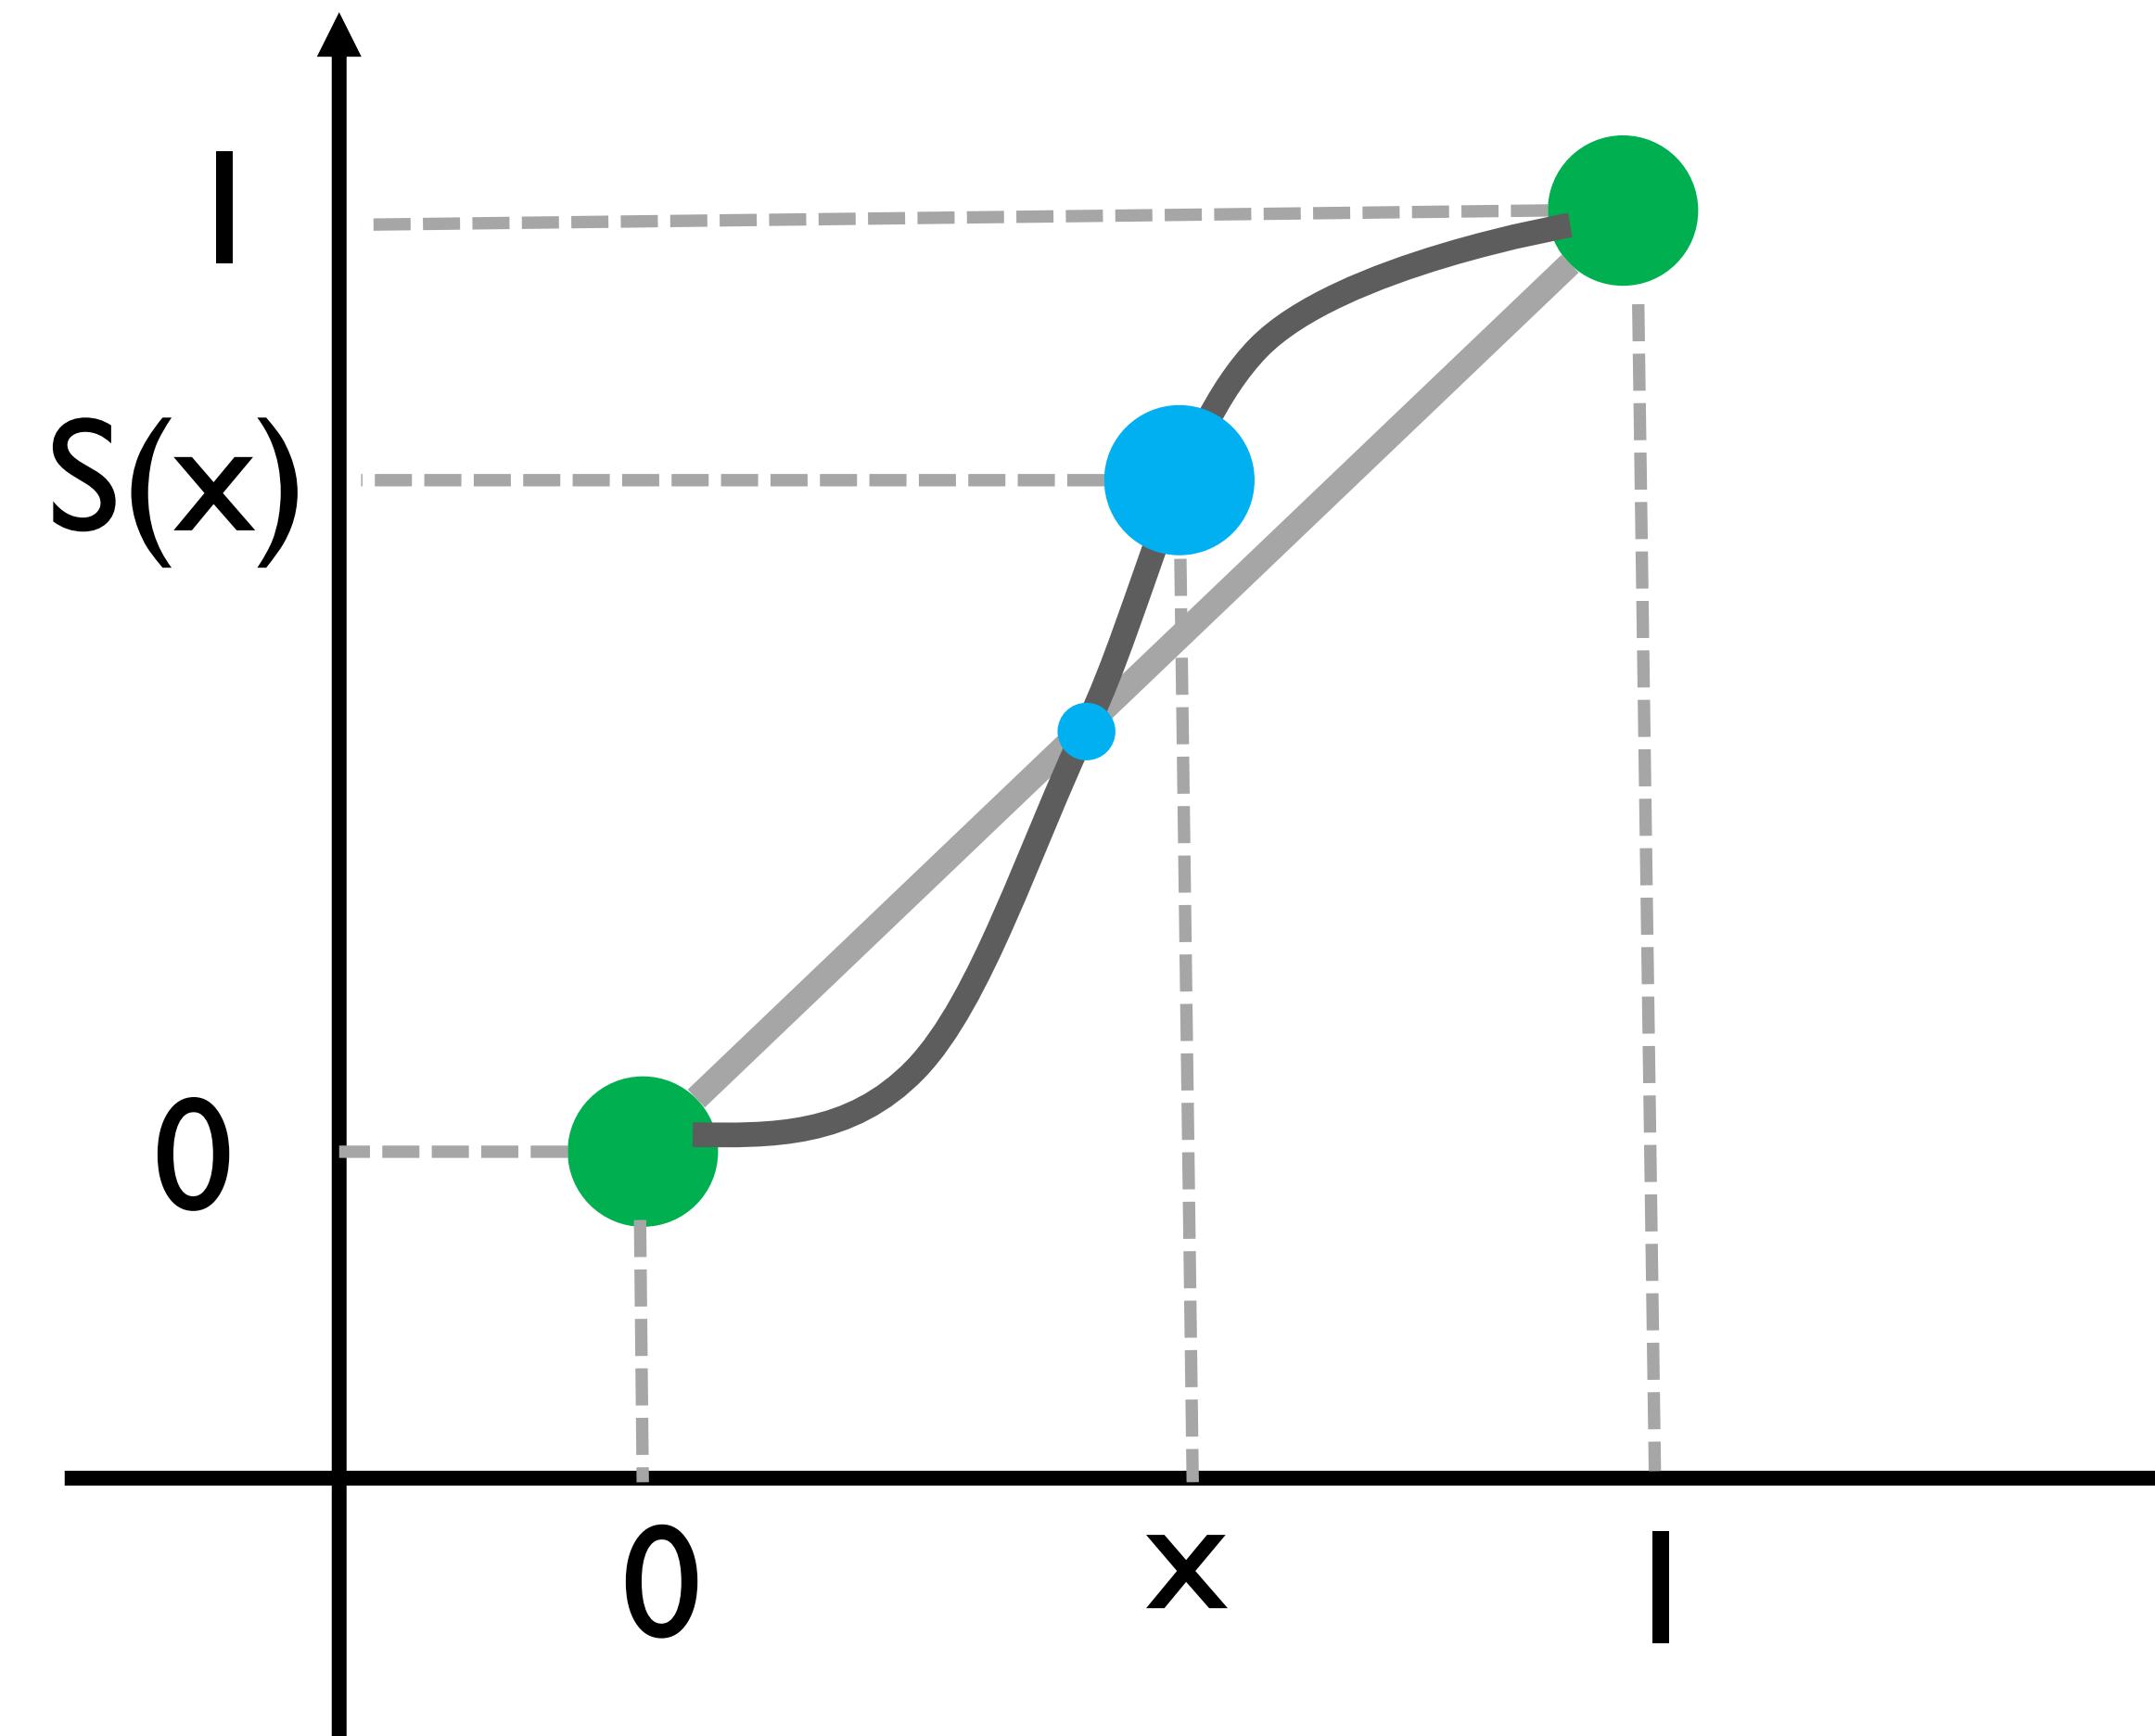


Solution: Smooth Interpolation

- Interpolate value from neighboring nodes using **smooth interpolation** instead of linear interpolation



Mathematical Definition of a Smooth Interpolation



$$S(x) = 3x^2 - 2x^3 \quad x \in [0,1]$$

Intuition: $S(x)$ will make an x that is close to 0 (or 1) even closer to 0 (or 1)

The derivative of $S(x)$ on $x=0$ and $x=1$ are both 0

Implementation Tip:
You may either implement it with one line of code or call smoothstep in GLSL

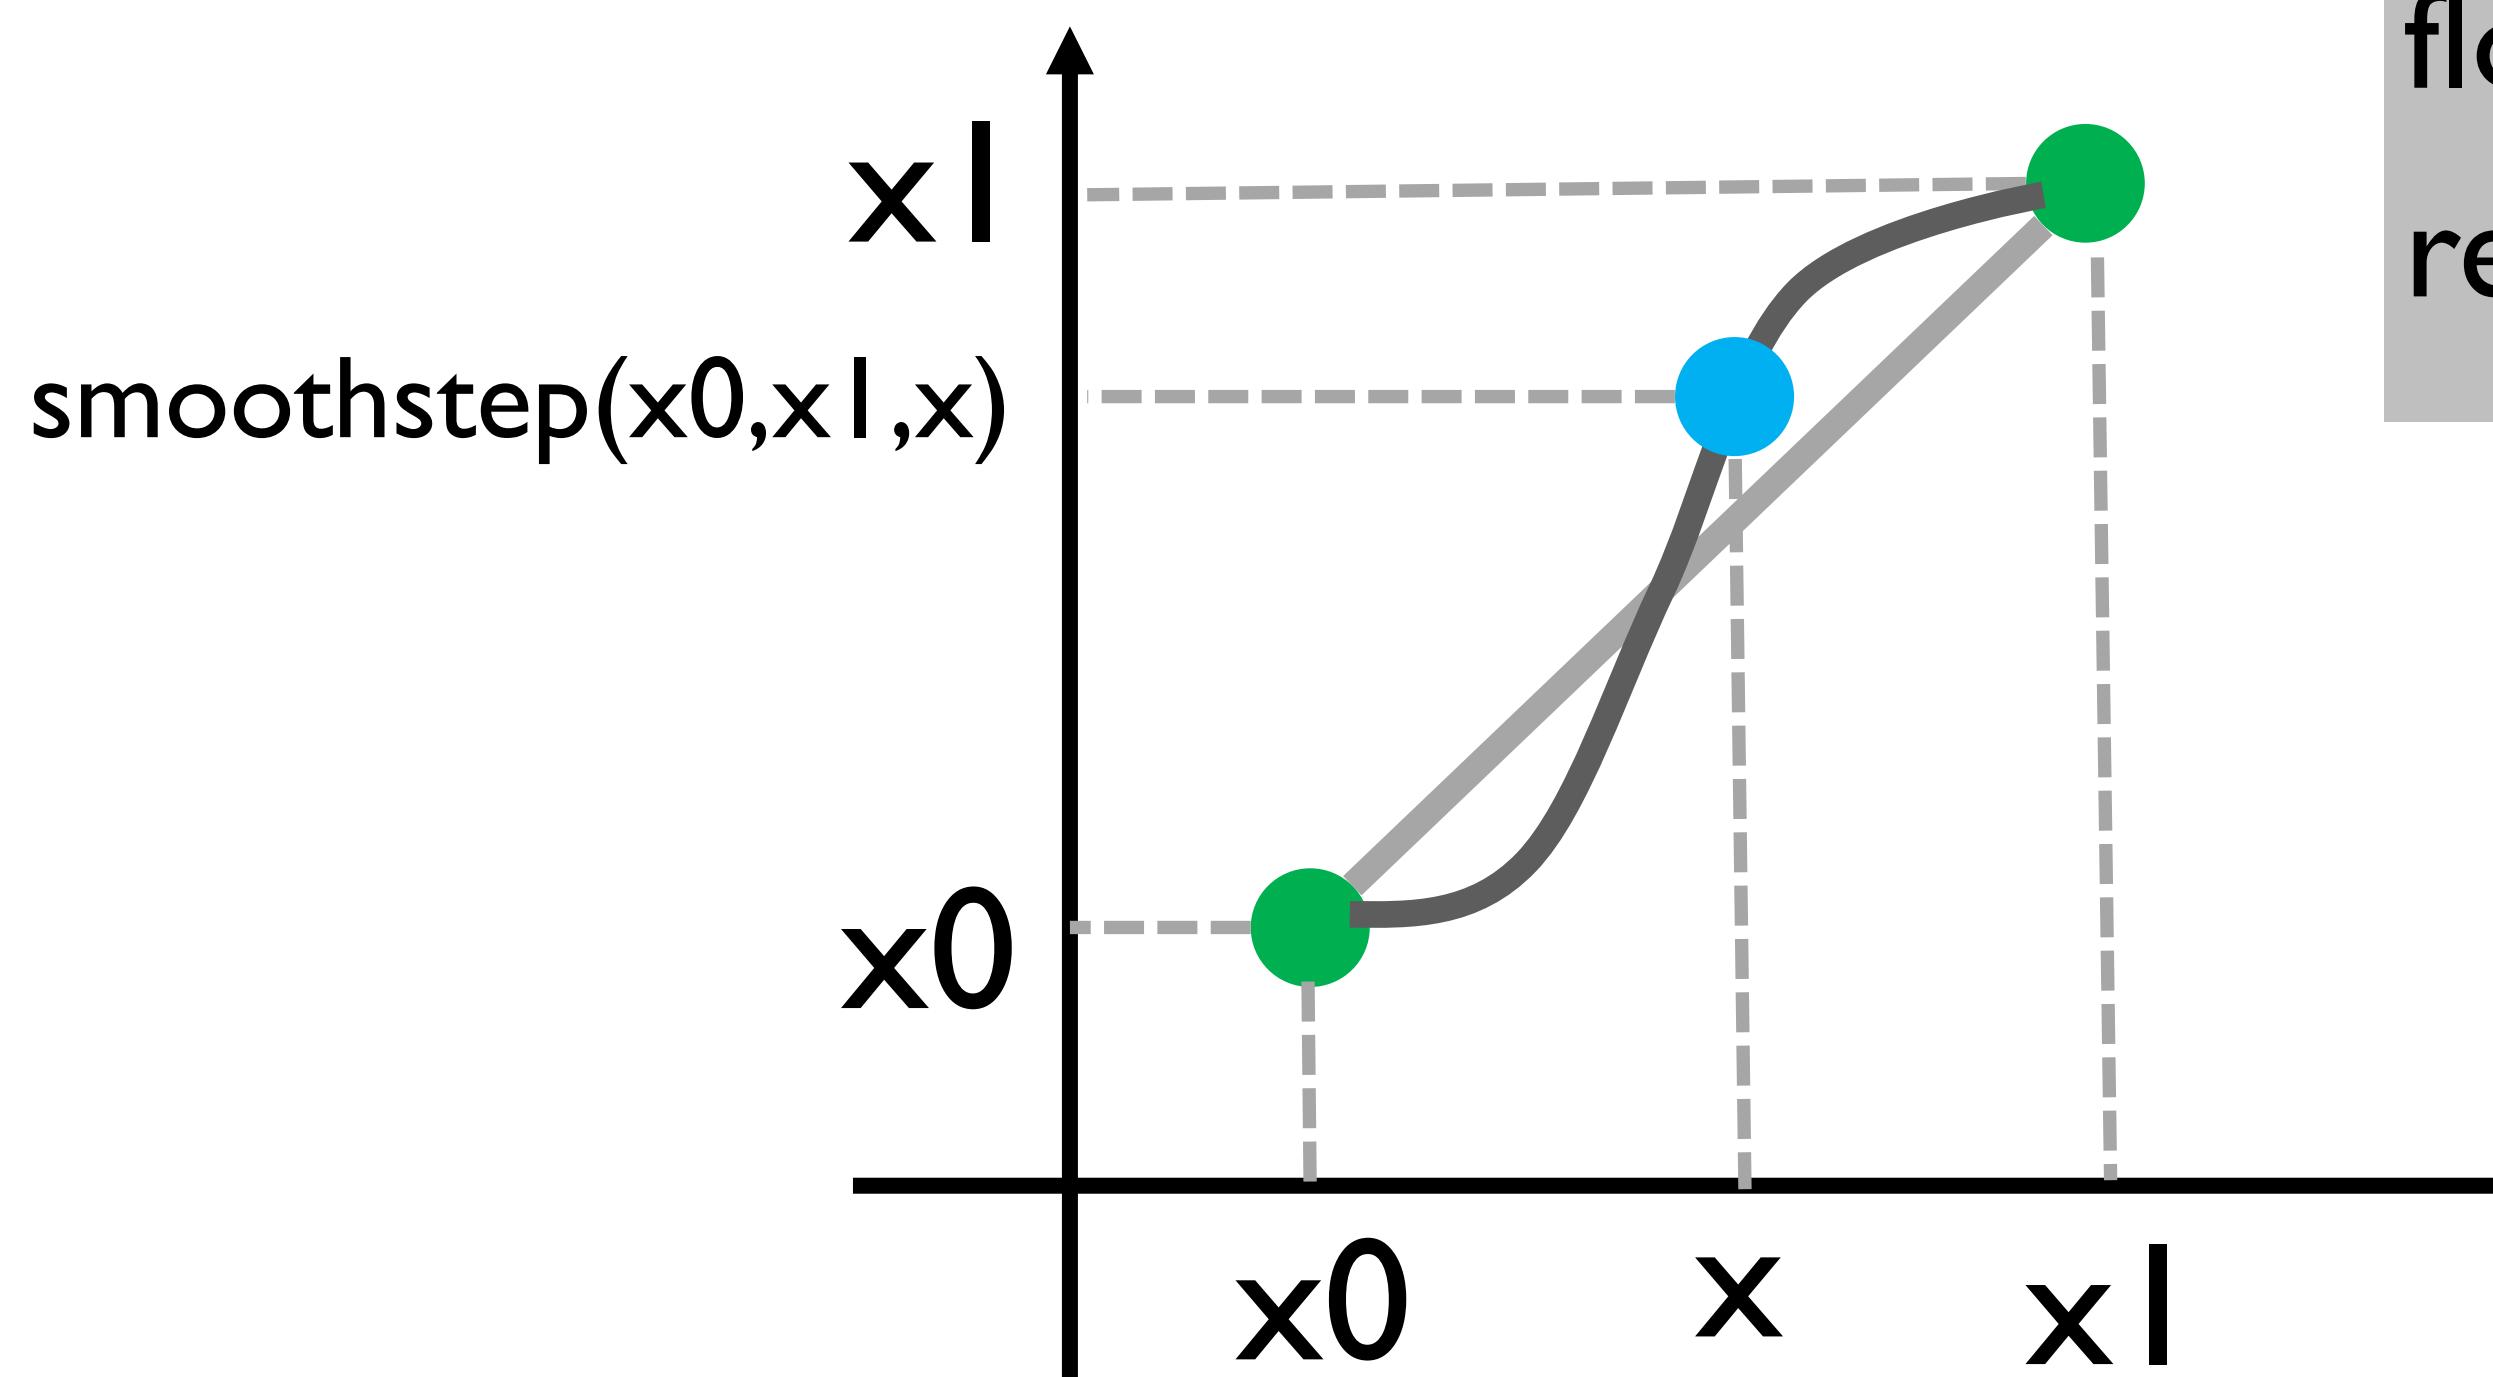
Such smooth function is also called *Hermite interpolation*, *fade function*, *ease curve*, or *smooth step* (in GLSL)



GLSL Function: smoothstep()

float smoothstep(float x0, float x1, float x)

- Create a smooth transition between x_0 and x_1
- The returned value is calculated as follows:



```
float t = clamp( (x-x0)/(x1-x0),0.0,1.0);  
return t*t*(3.0-2.0*t);
```

We typically use it to calculate a smoothed interpolation weight as:

float s=smoothstep(0.0,1.0,t);

by setting $x_0=0.0$ and $x_1=1.0$, and then use s in mix function as:

float v=mix(v0,v1,s);



Pseudocode for 1D Smoothstep Noise

Input: x

Output: noise

```
int i = floor(x);
float f = fract(x);
float s = f*f*(3.0-2.0*f);
float v0=hash(i);
float v1=hash(i+1);
float noise = (1-s)*v0+s*v1;
return noise;
```

Or, we can call **smoothstep** in GLSL:
`float s = smoothstep(0.0,1.0,f);`

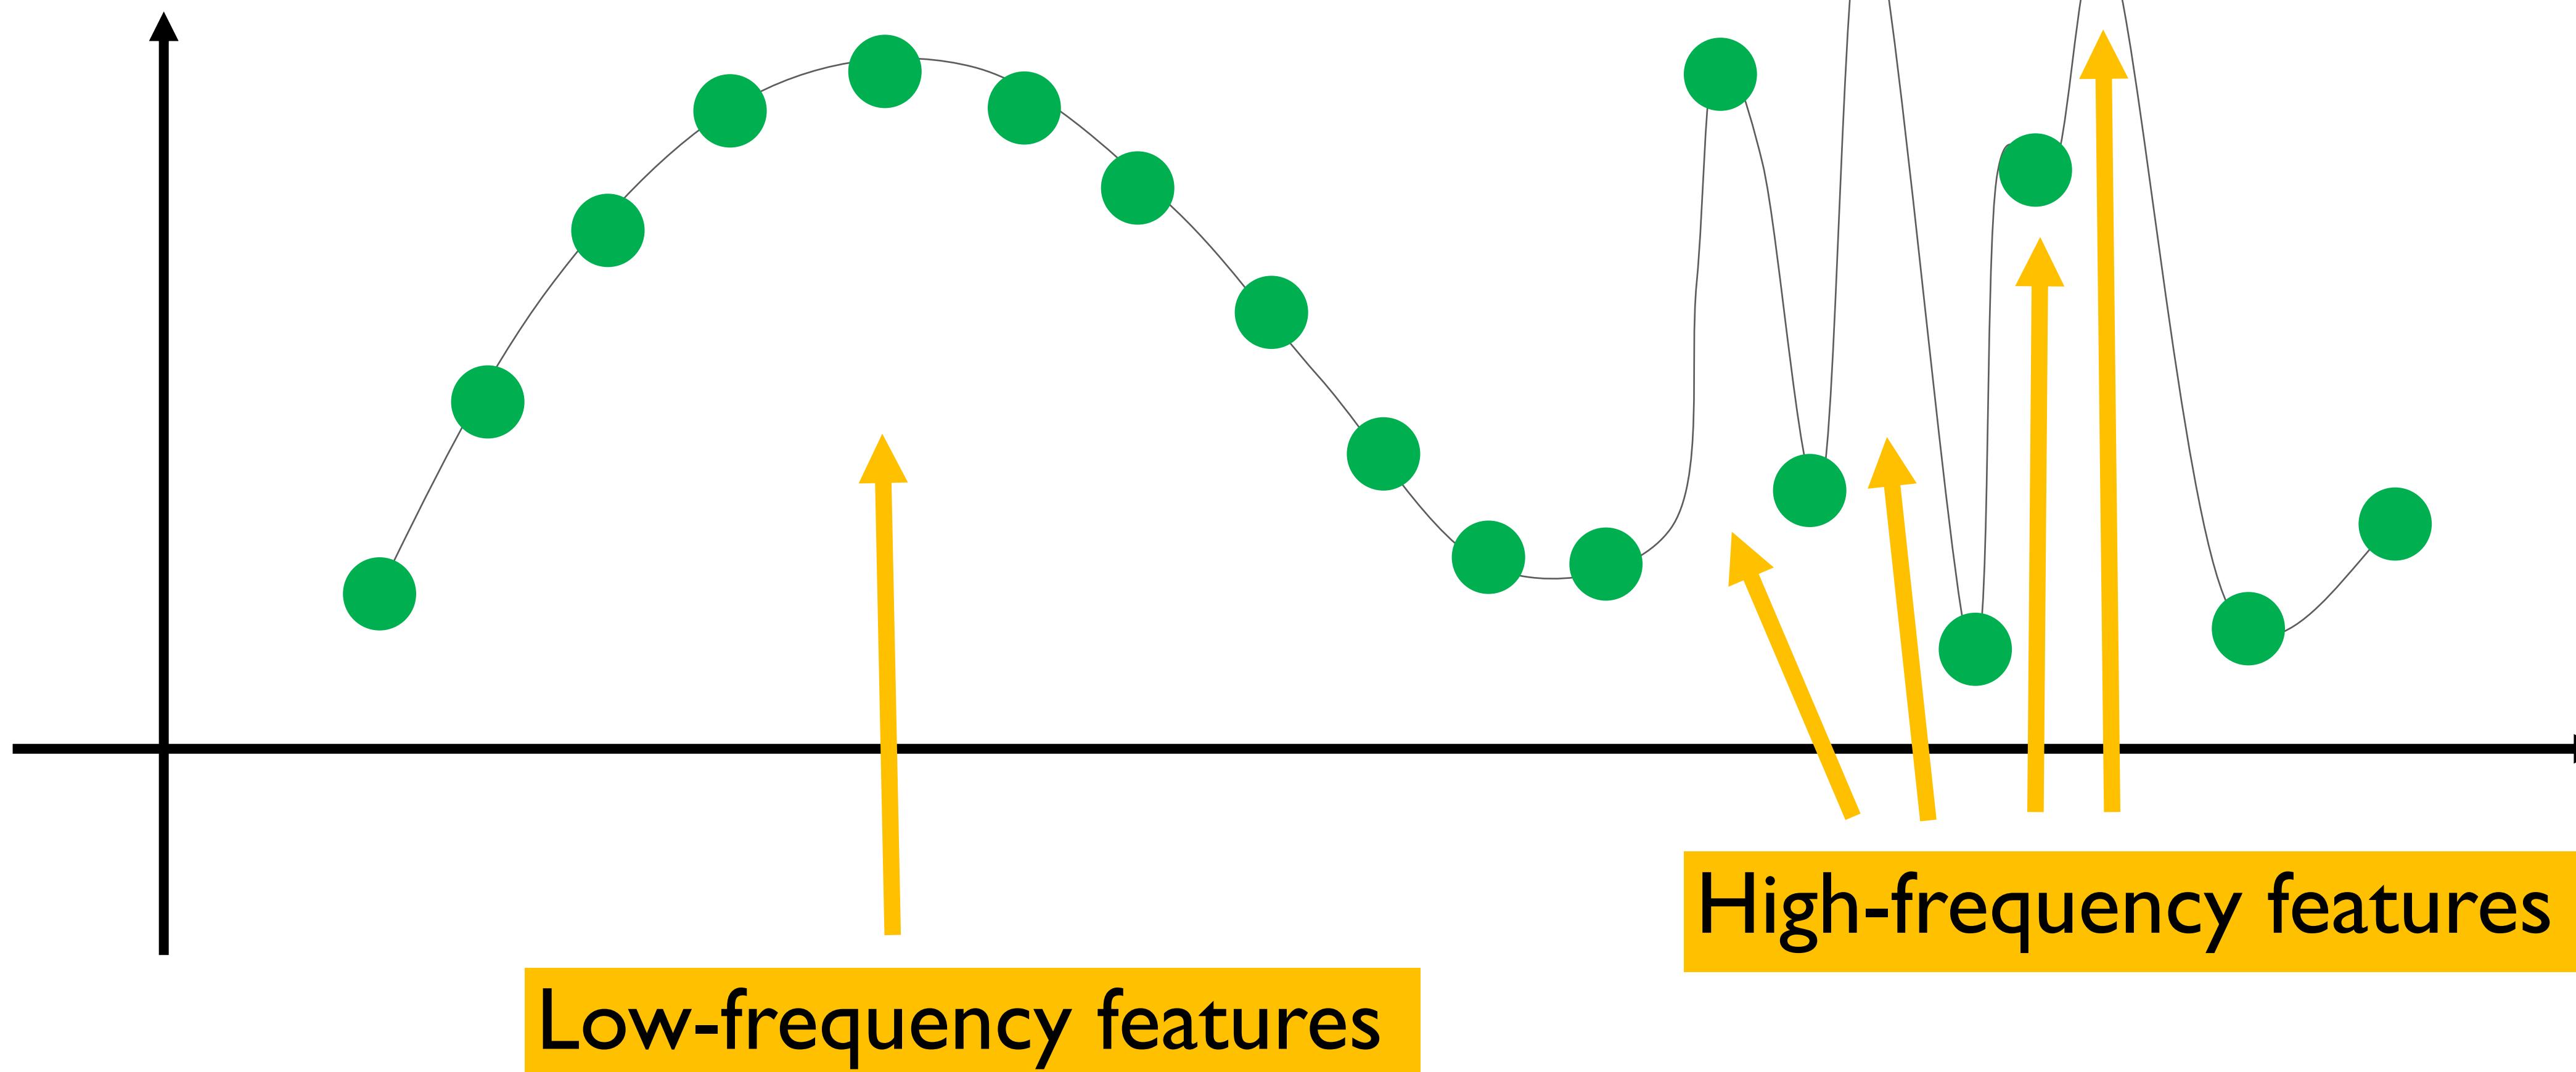
Or, we can call **mix** in GLSL:
`noise = mix(v0, v1, s);`



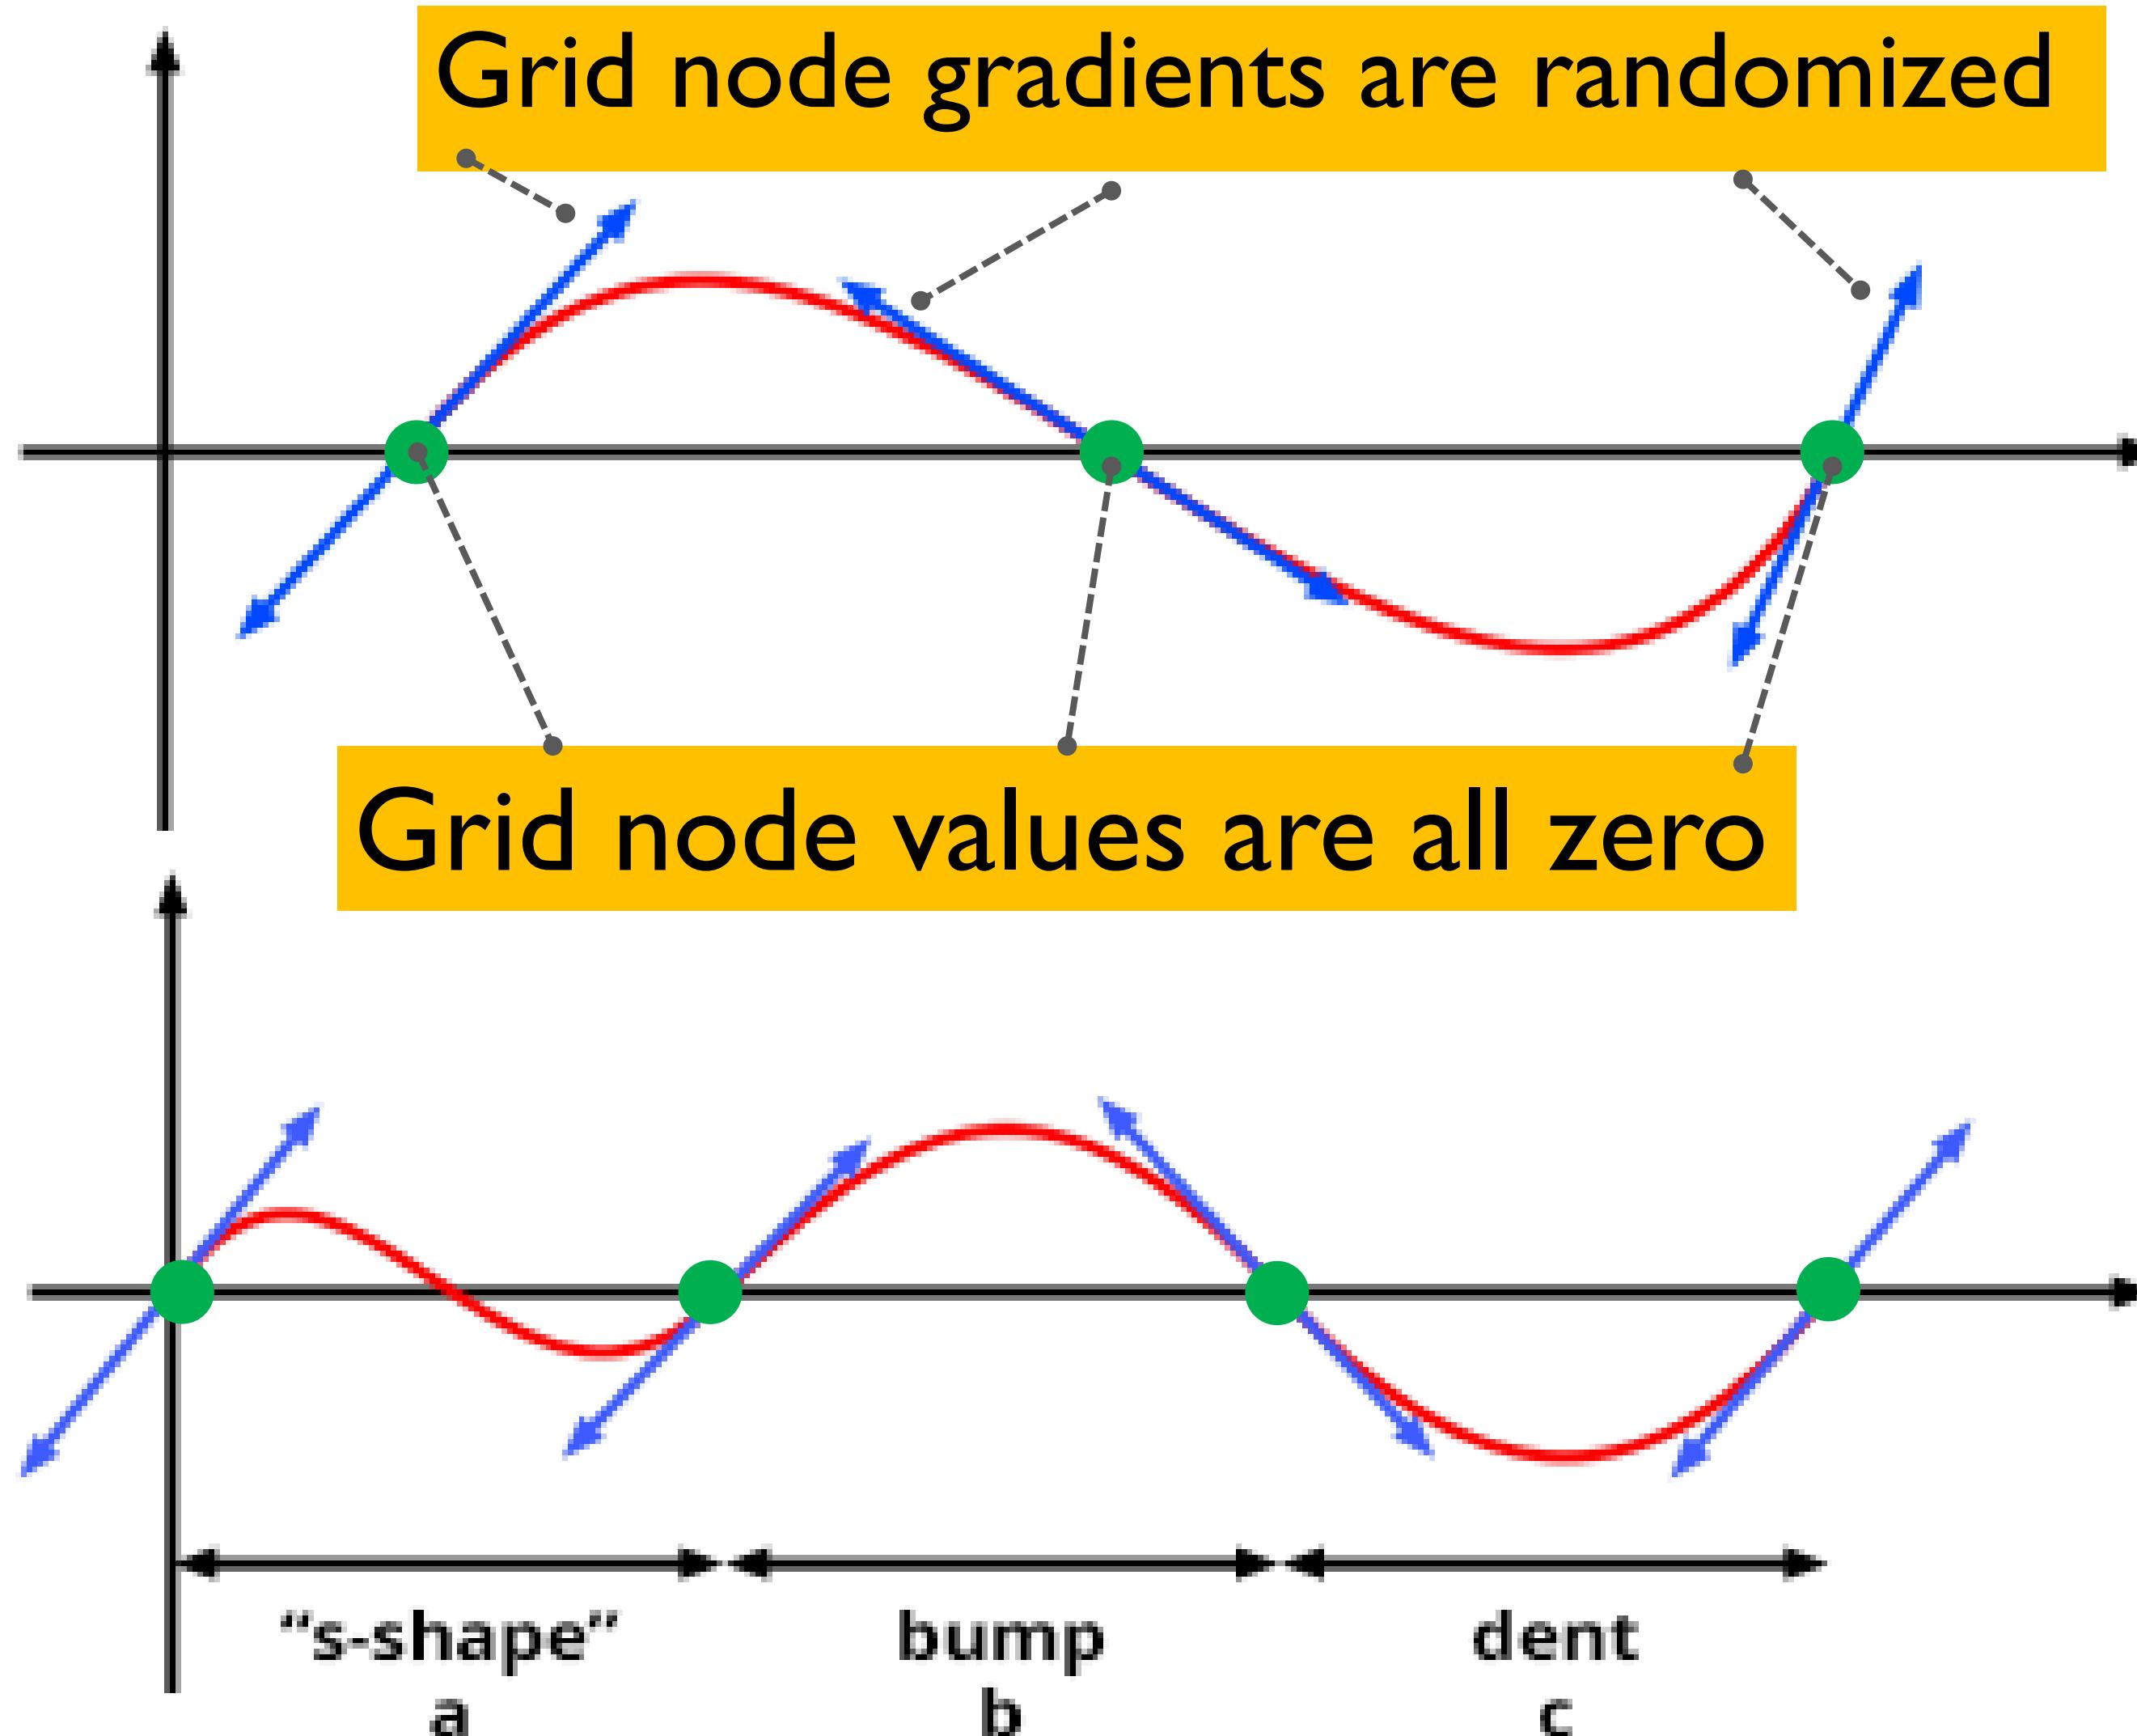
Problem II: Frequencies are not band-limited

- We might have features exhibiting dramatically different frequencies in different local regions --- some very slow slopes, and some very steep cliffs!

How do we produce a function that is band limited?



Solution: Use Gradients in Interpolation



Key Idea: We (1) randomize the **gradient** instead of the value of a grid node, and (2) assume node values are always **zero**, and then (3) **interpolate** the function using both gradients and values on grid nodes.

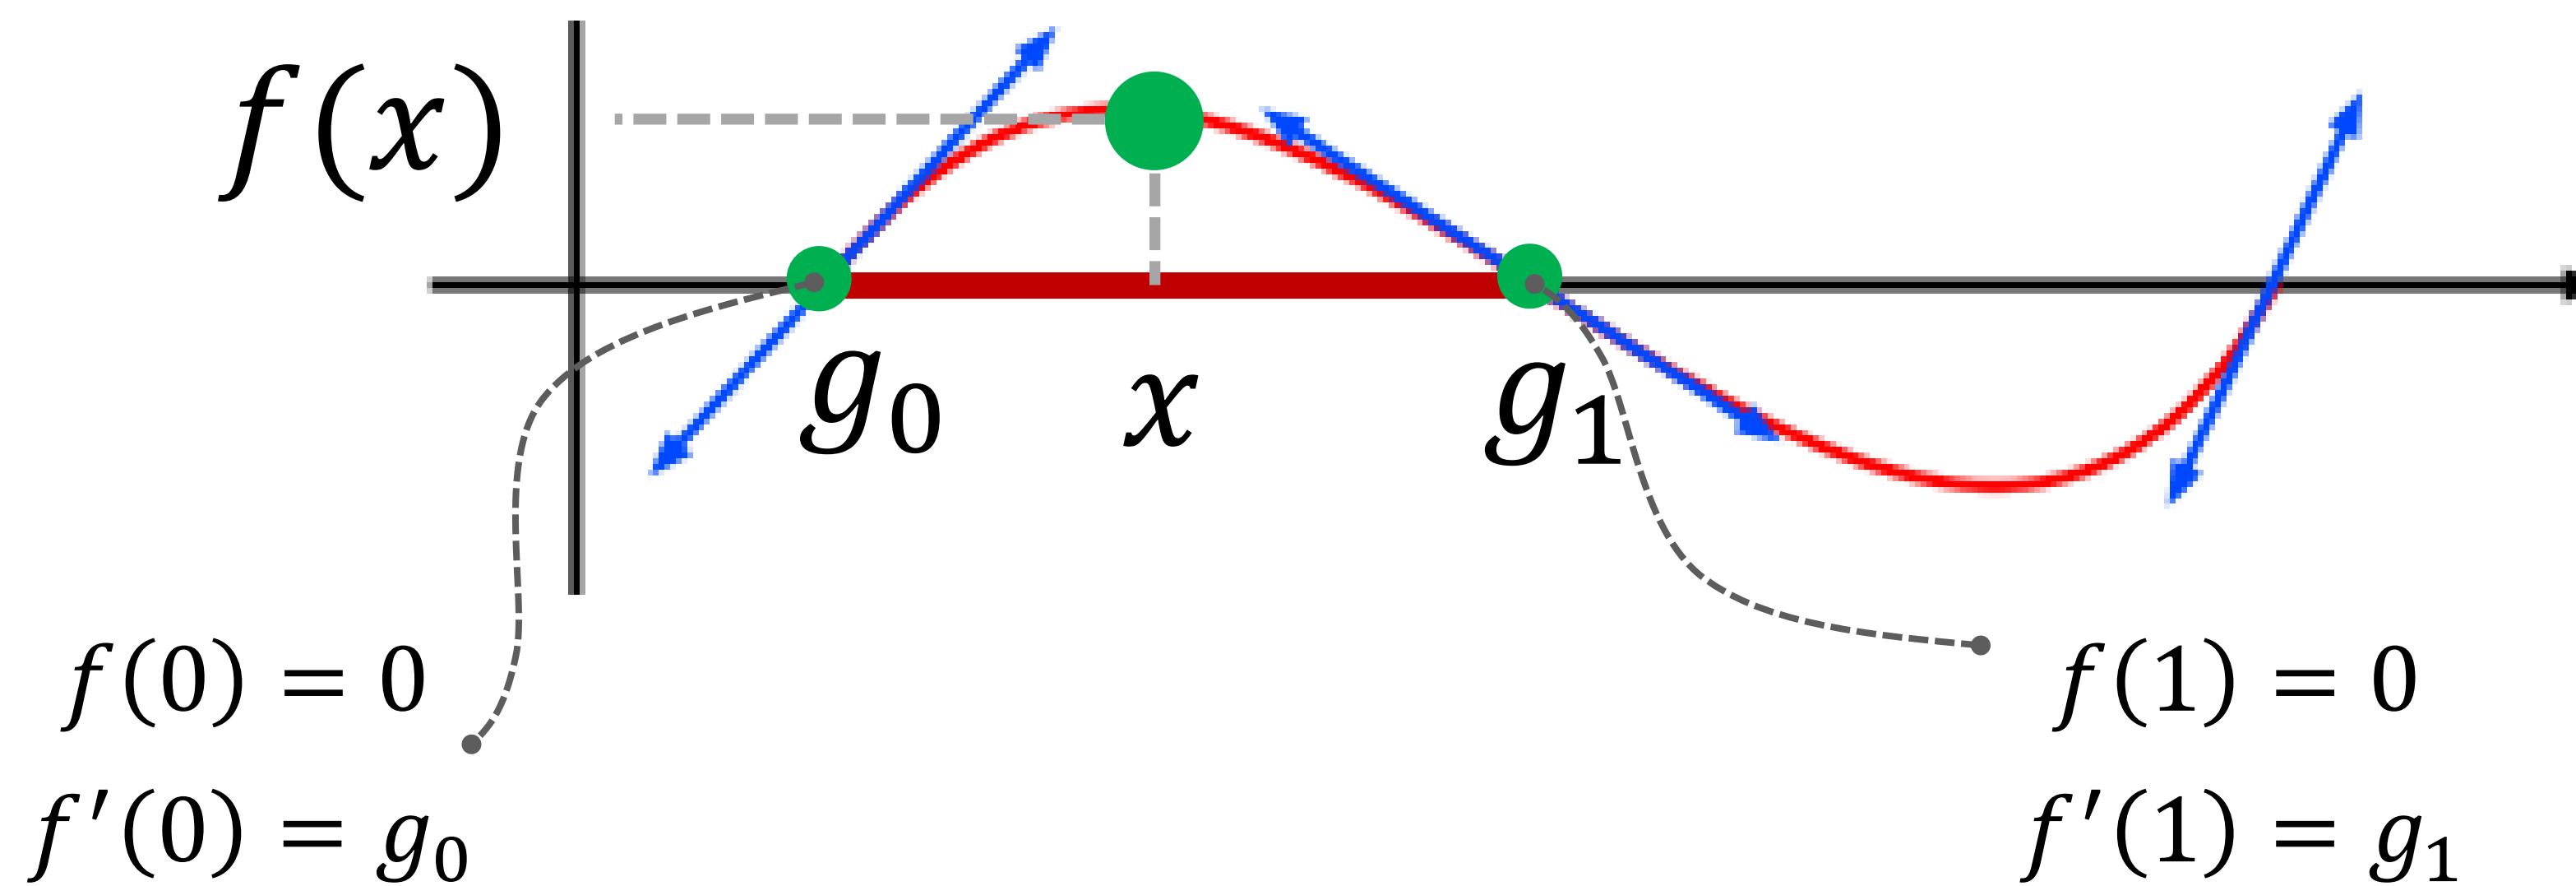
All features have more or less the same size, either a “S” or a bump/dent.

Mathematical Definition of Gradient Interpolation

- We can **design** a function satisfying both value constraints and gradient constraints on the grid nodes:

$$f(x) = (1 - s(x))xg_0 + s(x)(x - 1)g_1$$

with $s(x) = 3x^2 - 2x^3$, $x \in [0,1]$



Pseudocode for 1D Gradient Noise

Input: x

Output: noise

```
int i = floor(x);
float f = fract(x);
float s = f*f*(3.0-2.0*f);
float g0=hash(i);
float g1=hash(i+1);
float noise = (1-s)*f*g0+s*(f-1)*g1;
return noise;
```

This version is the 1D Perlin noise with a single frequency.

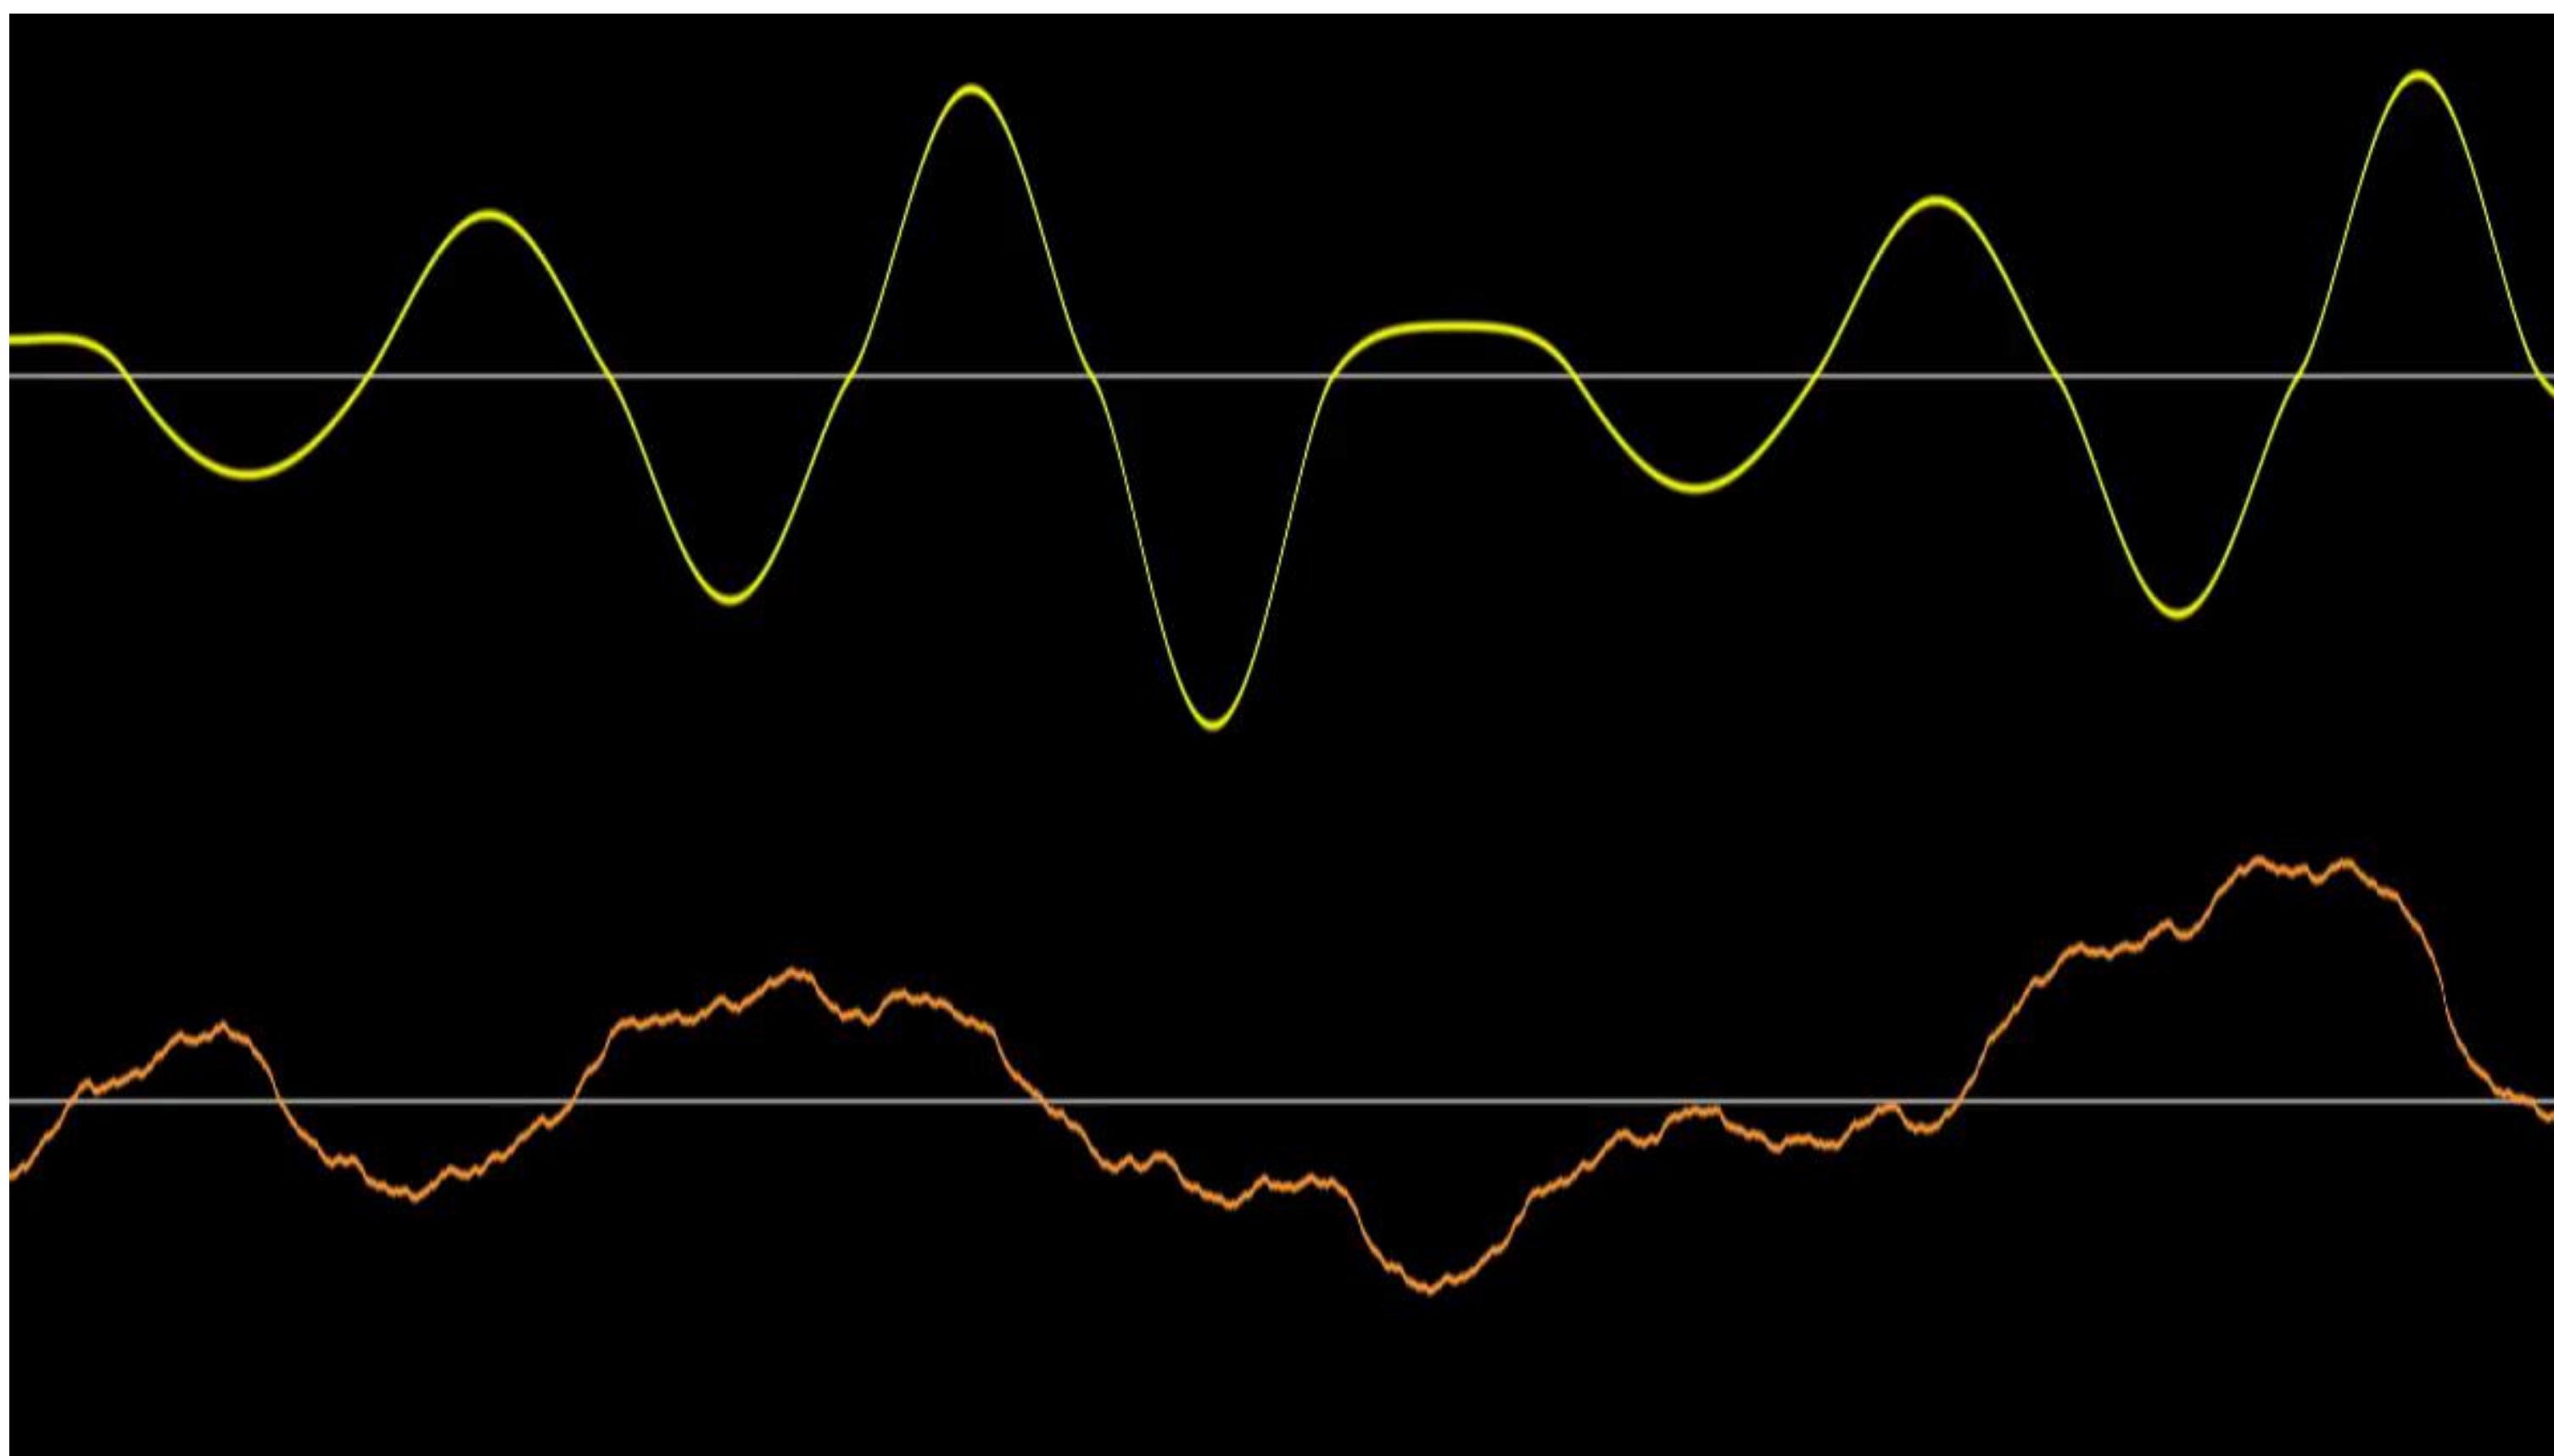
Or, we call **smoothstep()** in GLSL:
float s = smoothstep(0.0, 1.0, f);

Or, we call **mix()** in GLSL:
noise = mix(f*g0, (f-1)*g1, s);



Problem III: Synthesizing Multiple Frequencies

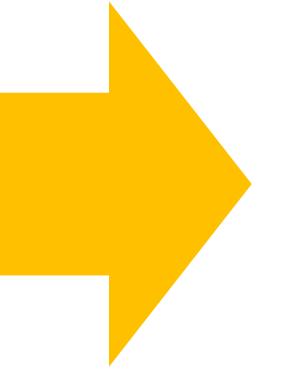
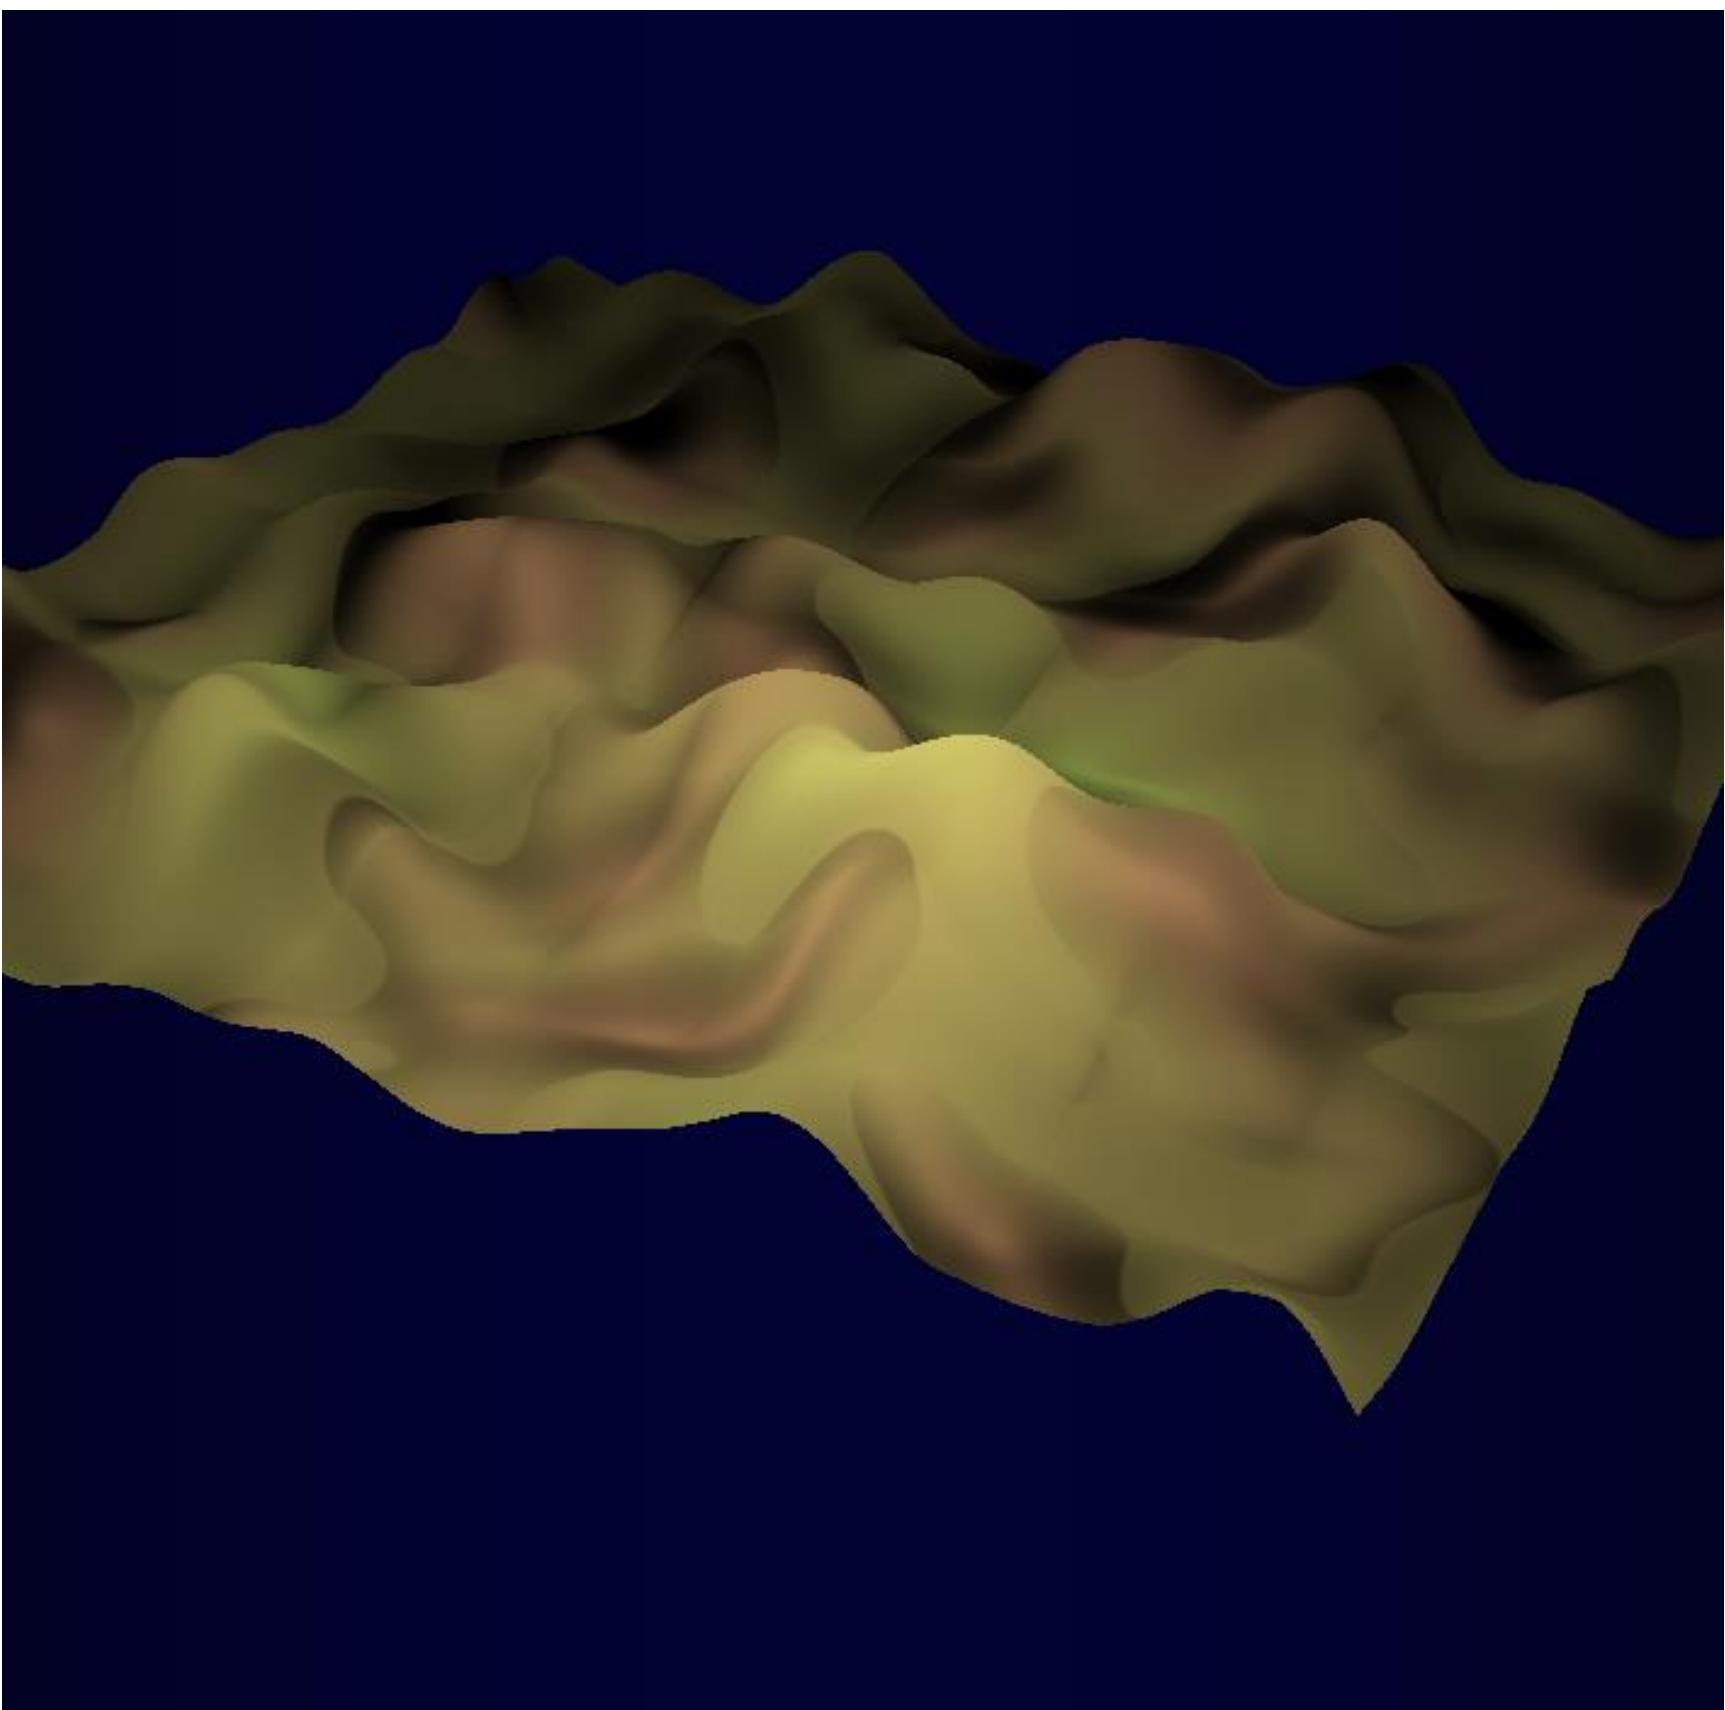
- Real-world textures/shapes/terrains are not always smooth: they exhibit features by mixing different frequencies together



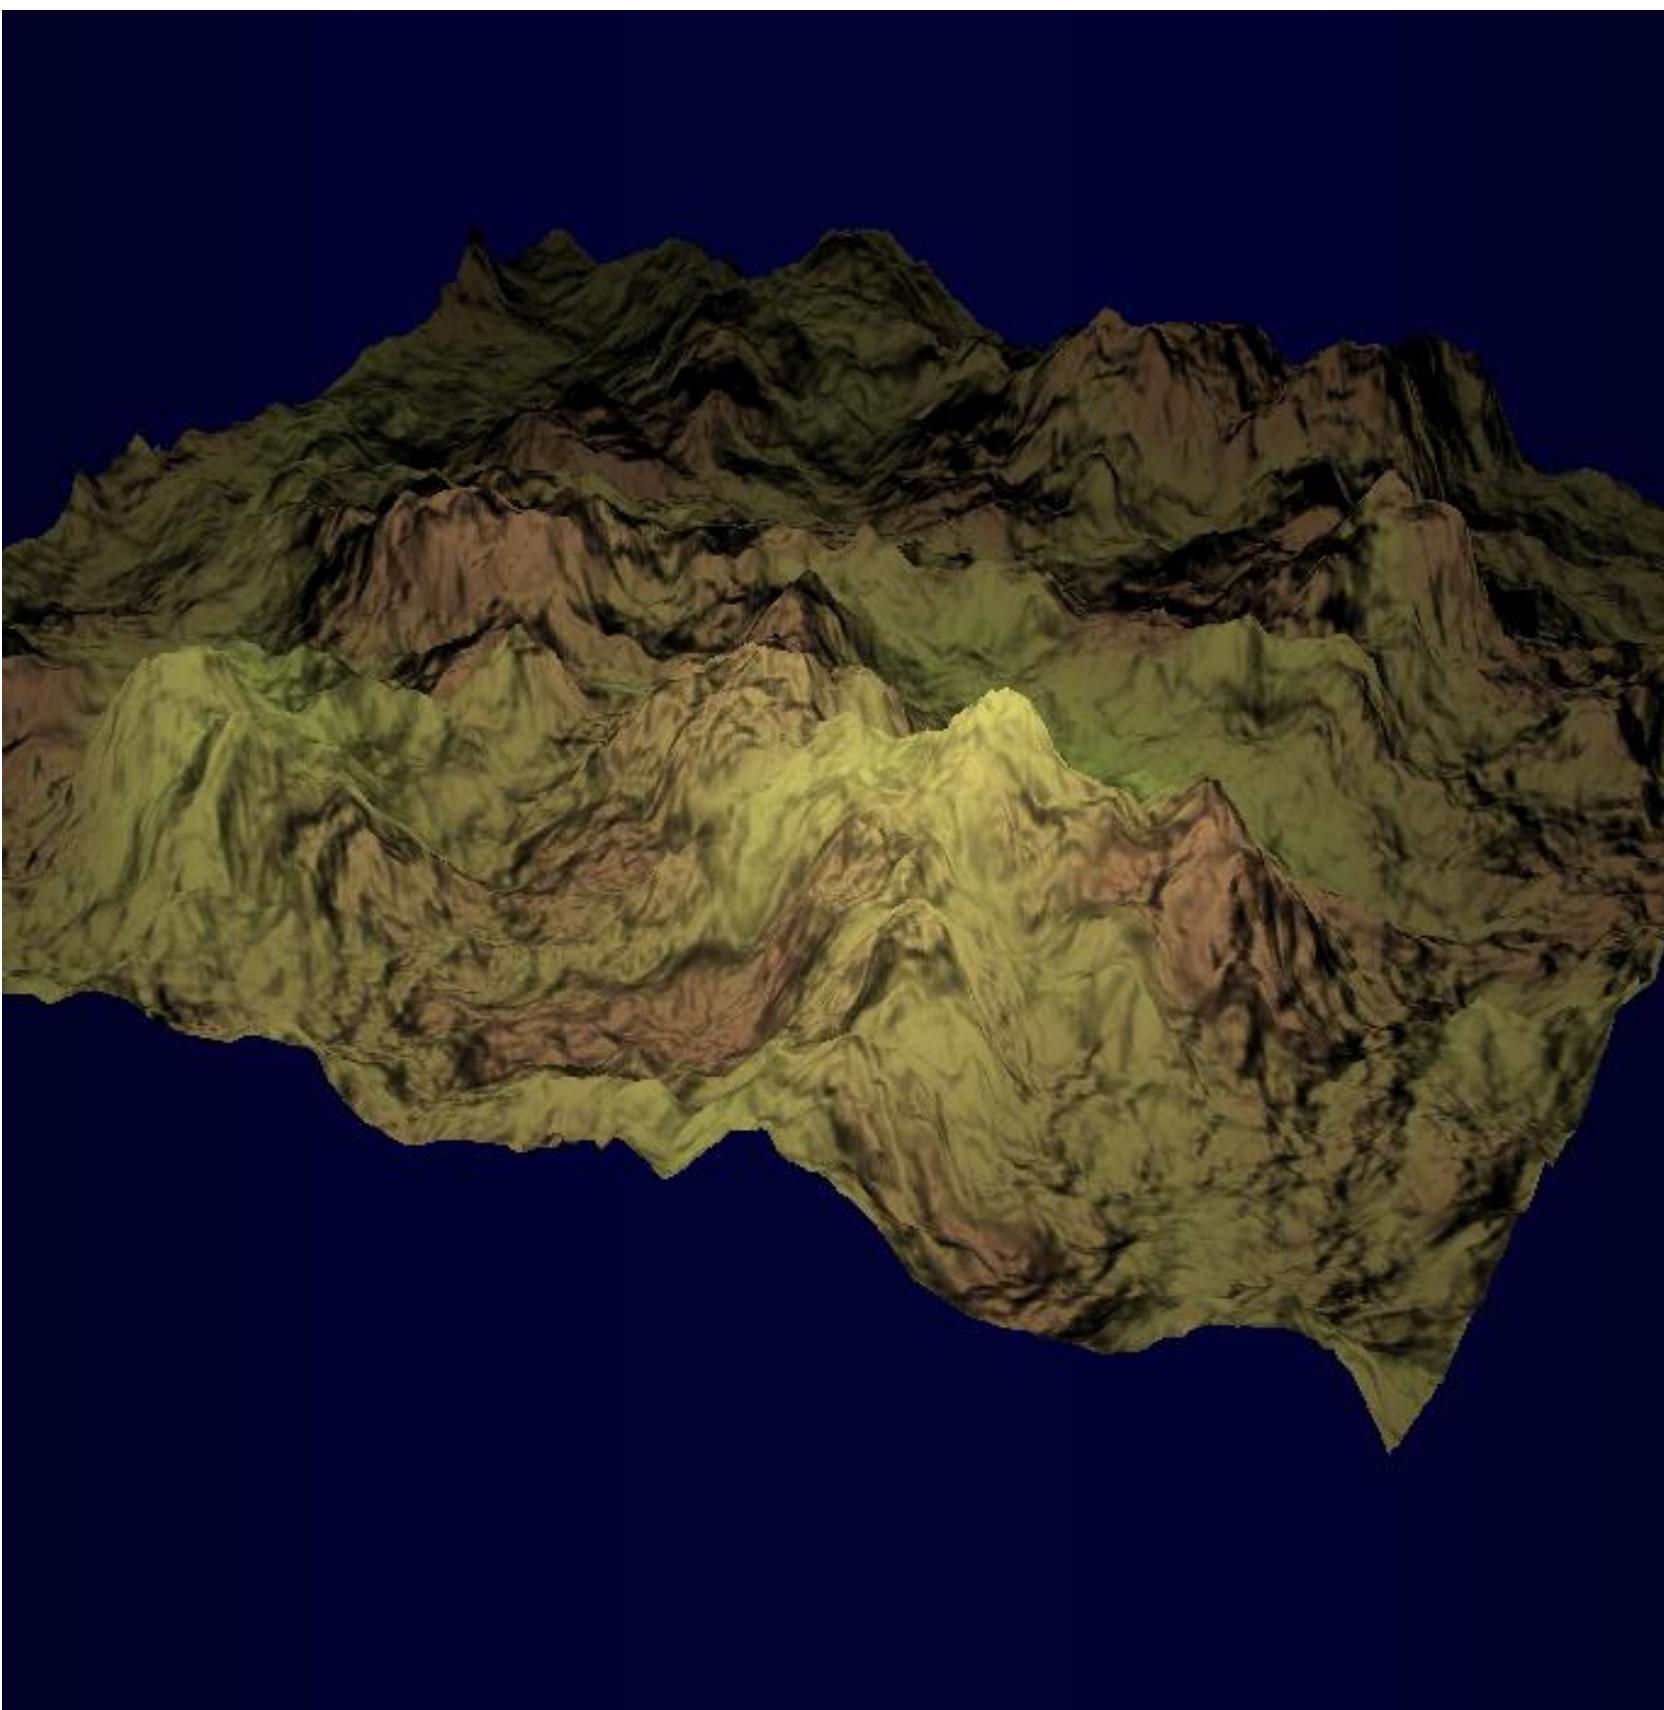
Noise with a
single frequency

Noise with multiple
frequencies mixed

Intuition: Multi-scale Features



?



Solution: Octave Synthesis

- Representing a complex function $f_s(x)$ by a sum of weighted contributions from a scaled function $f(x)$:

Key Idea:
*double the frequency,
half the amplitude.*

$$f_s(x) = \sum_i \omega_i f(s_i x)$$

with $\omega_i = \frac{1}{2^i}$, $s_i = 2^i$

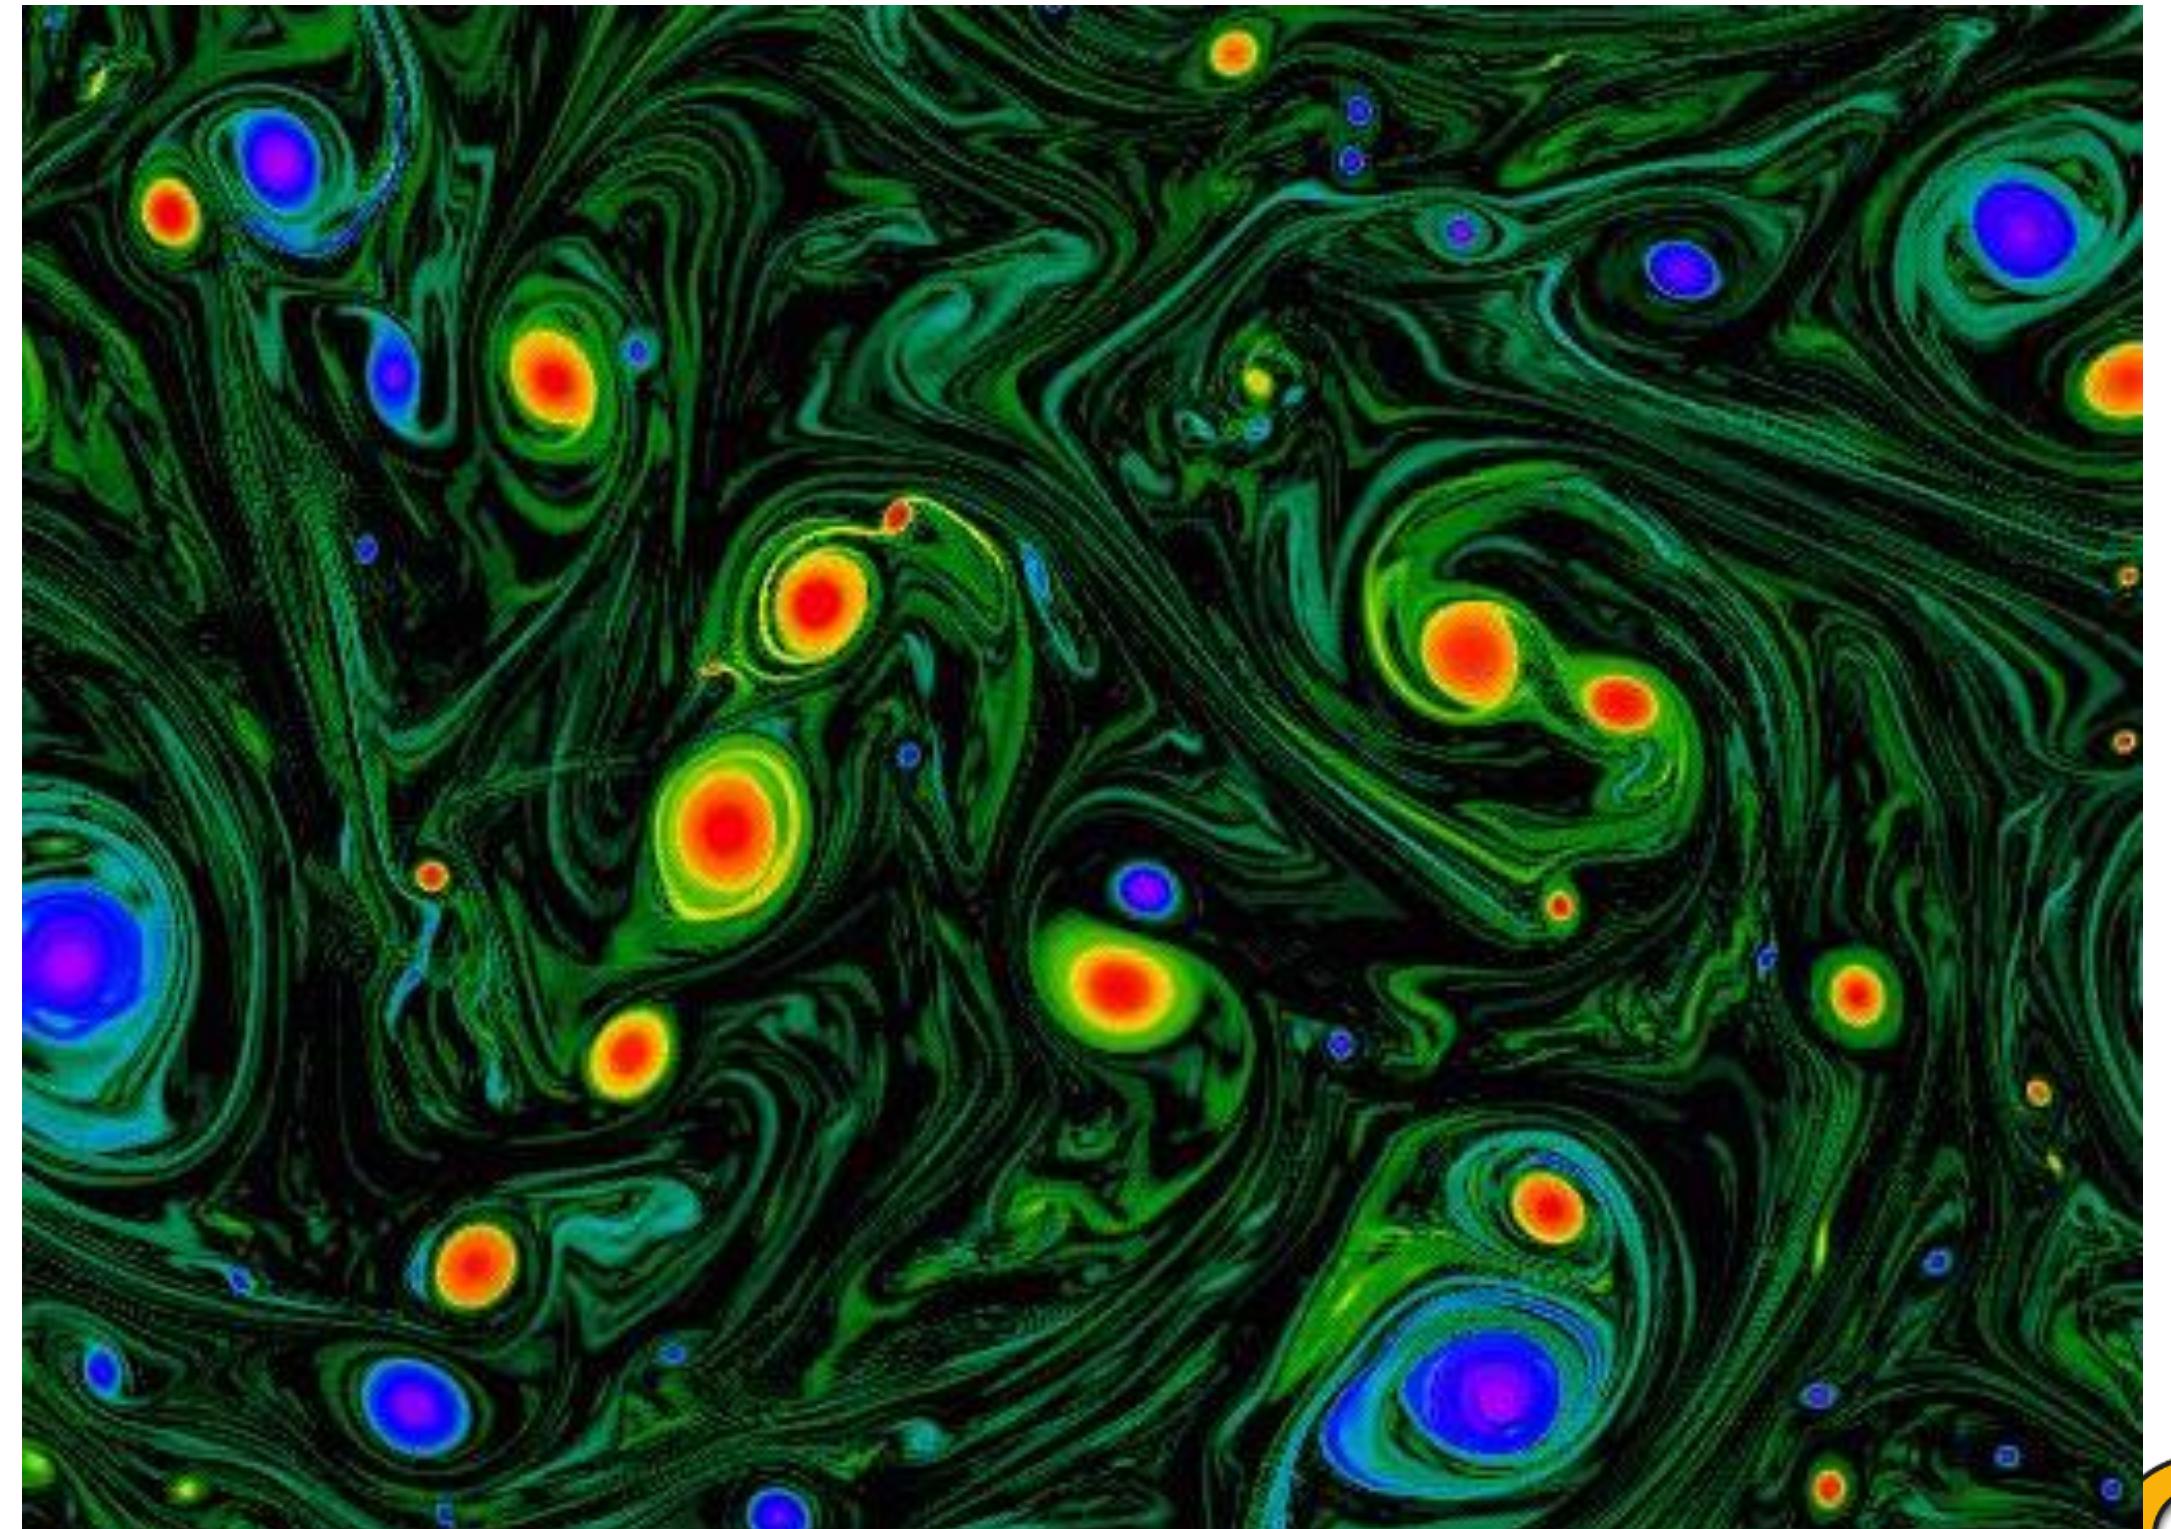
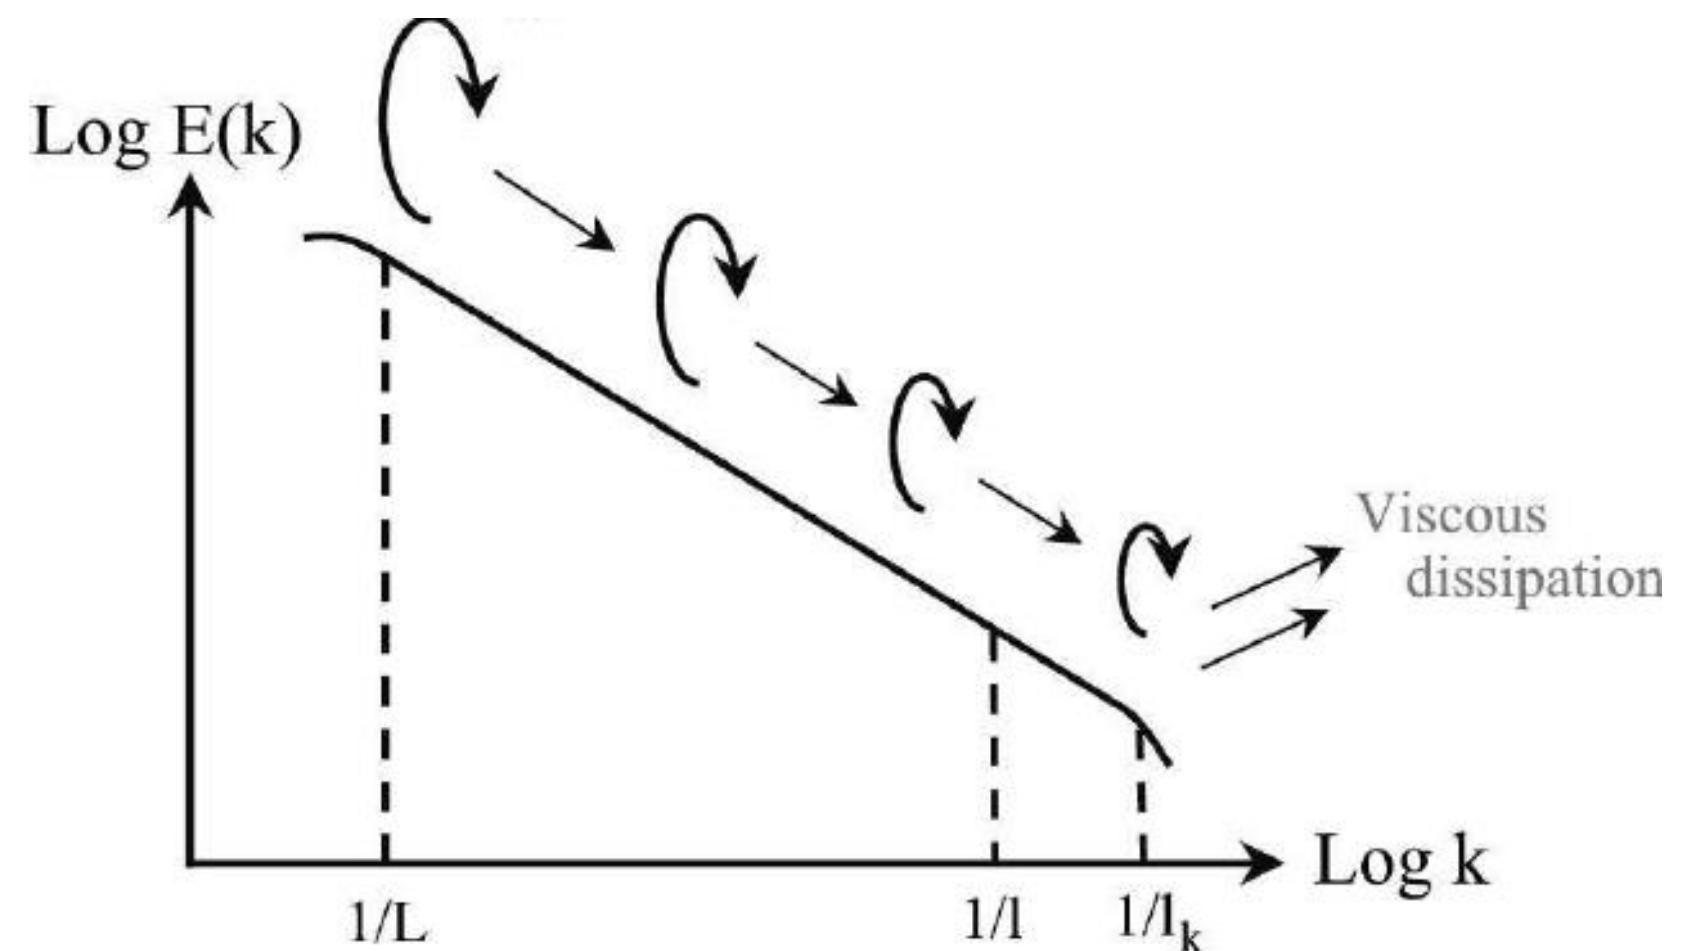
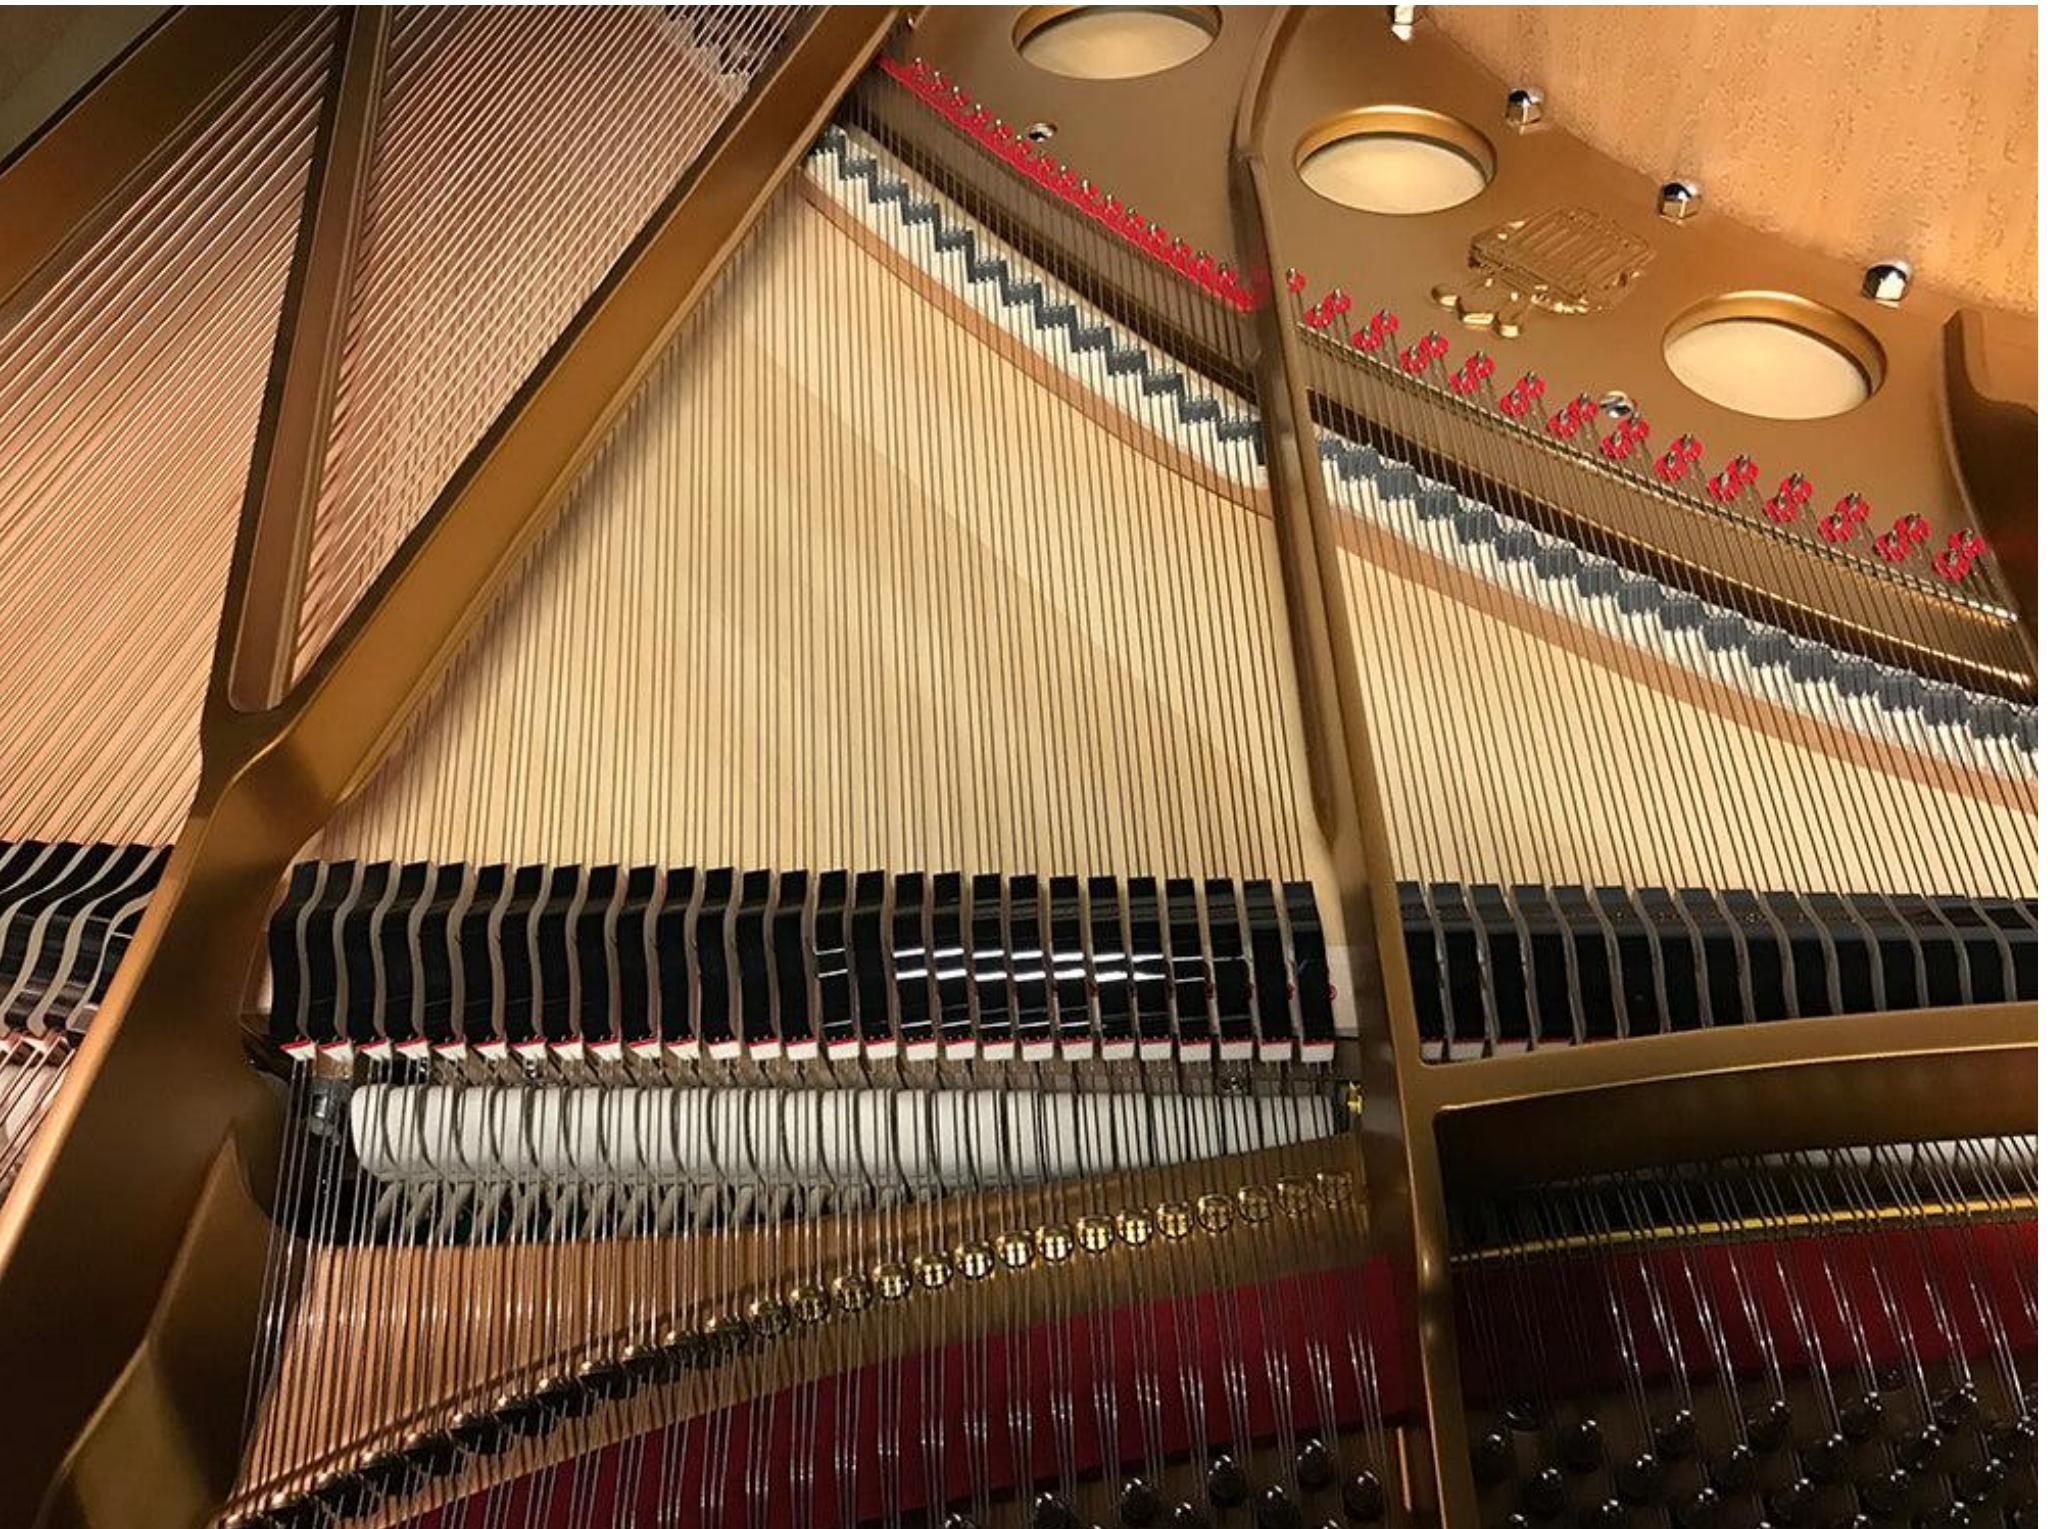
Each term in the summation is called an “**octave**”

$$\text{E.g., } f_s(x) = f(x) + \frac{1}{2}f(2x) + \frac{1}{4}f(4x) + \frac{1}{8}f(8x) + \dots$$



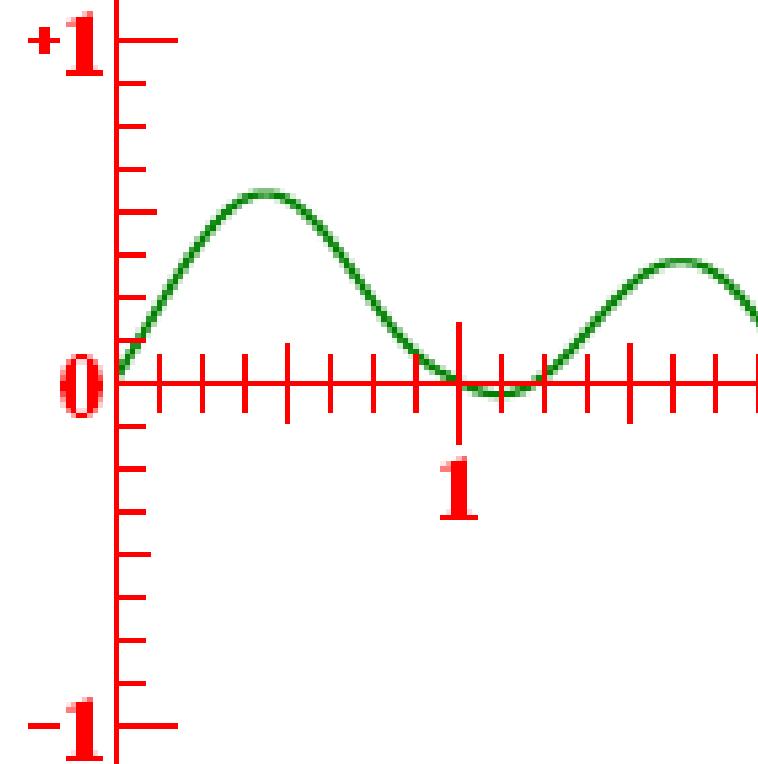
Octave

Beethoven Violin Concerto in D major

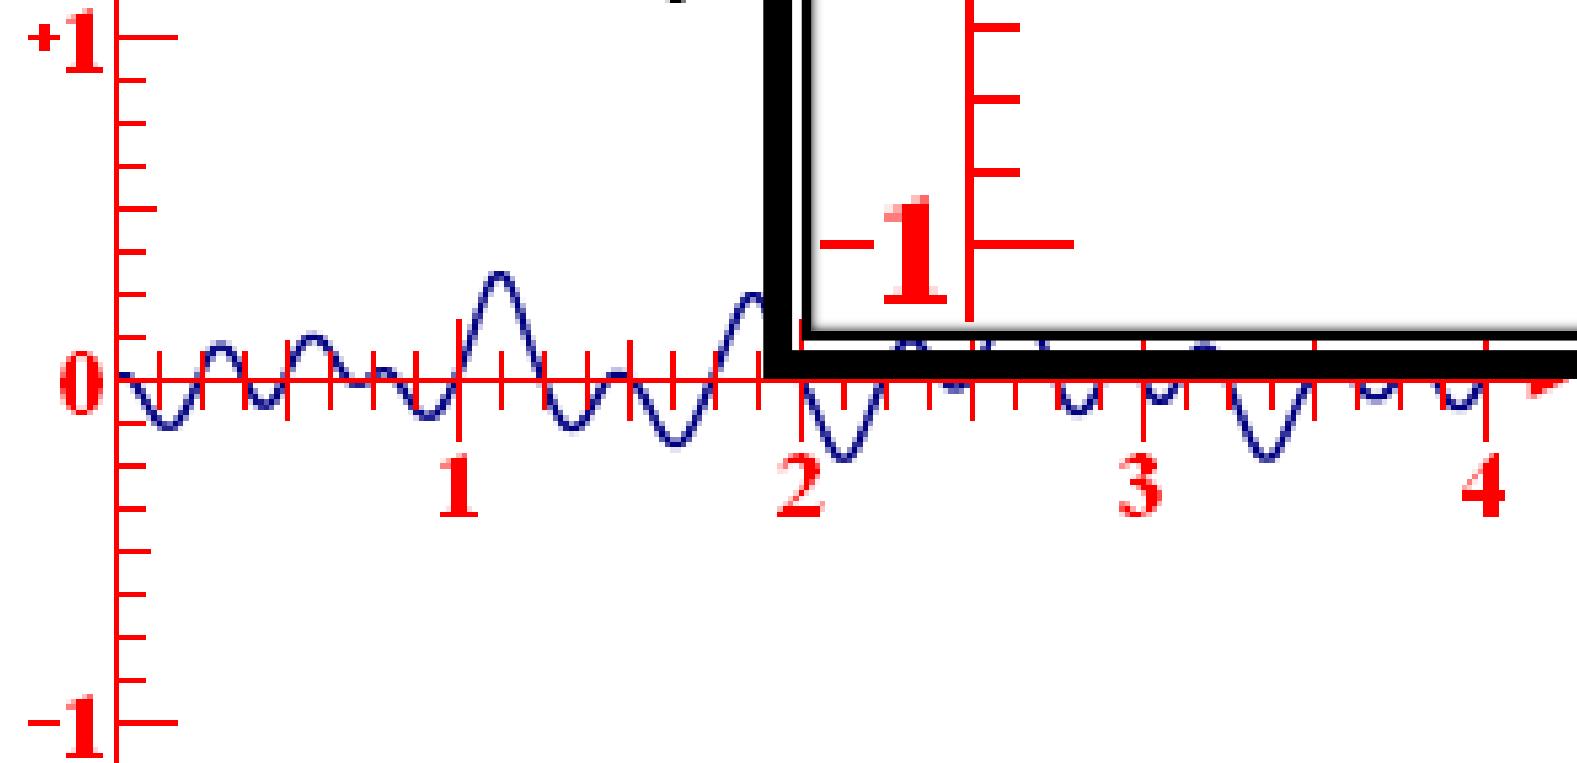


Example: Noise Synthesis with Four Octaves

$n(x)$ $f=1, a=1$



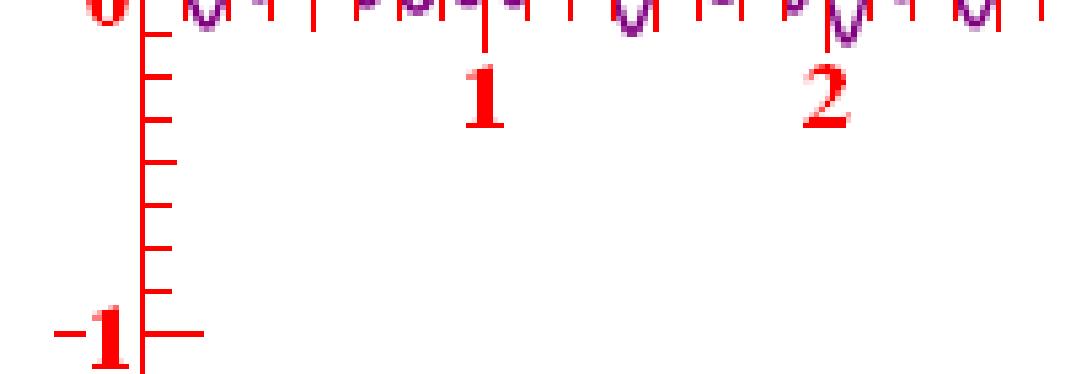
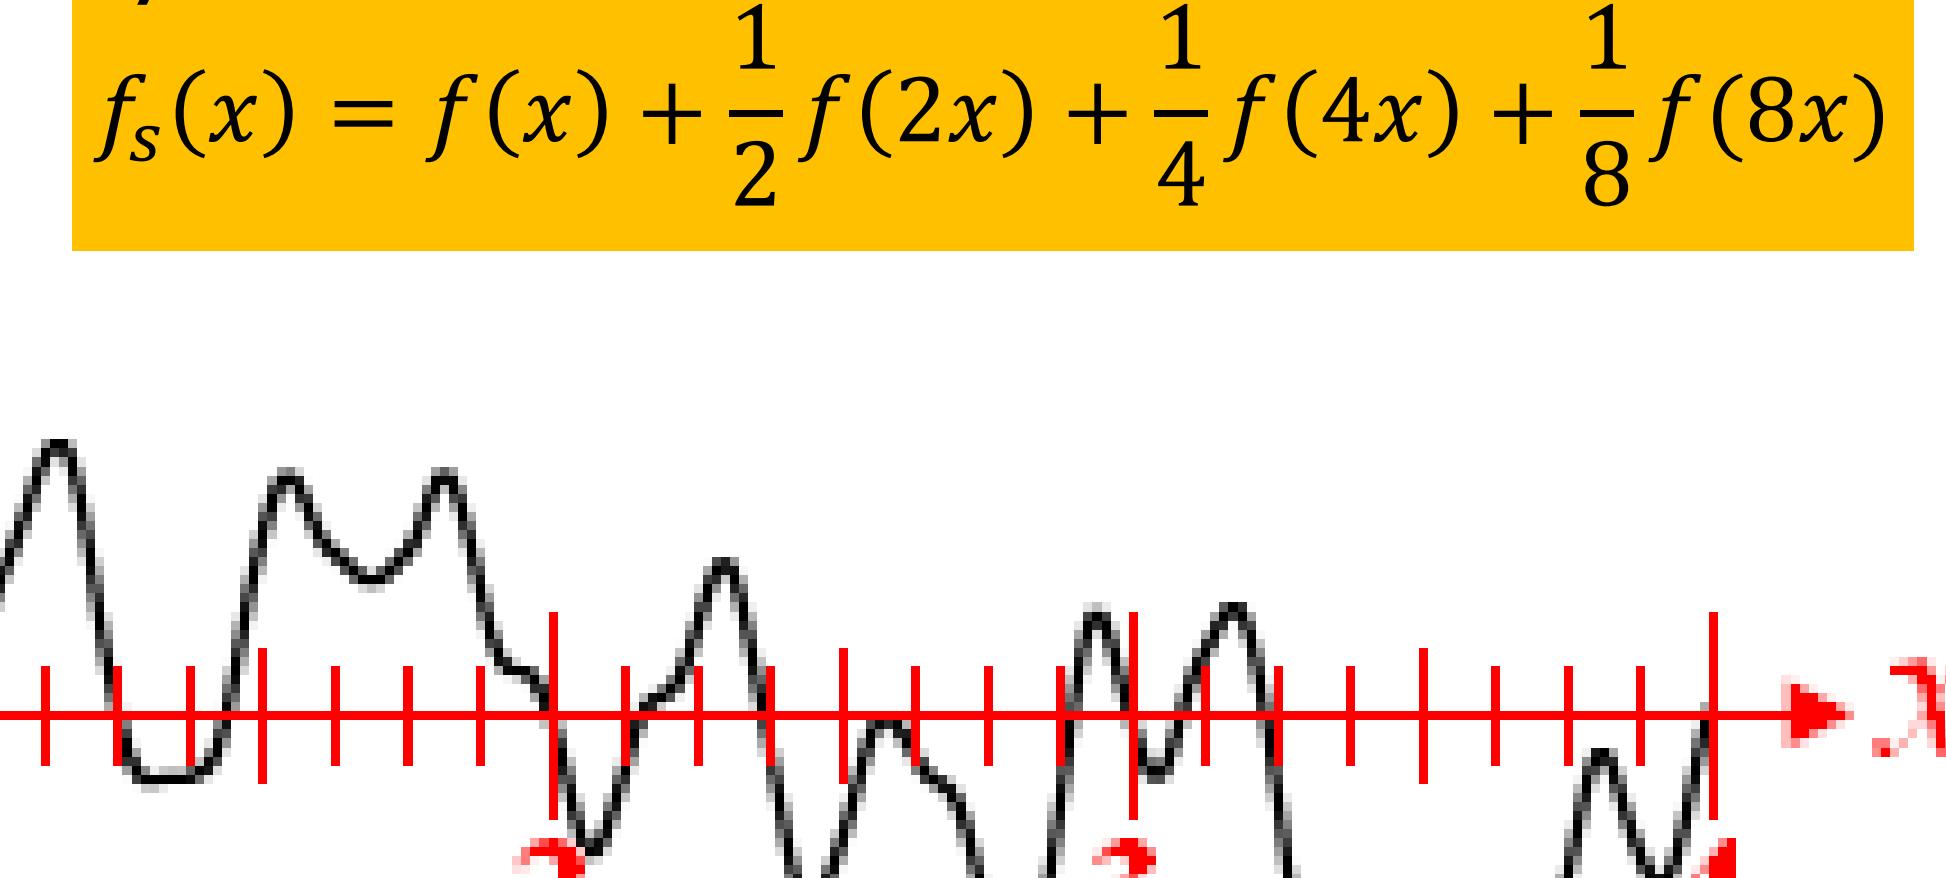
$n(x)$ $f=4, a=\frac{1}{4}$



$n(x)$ $f=2, a=\frac{1}{2}$

Synthesized noise function:
$$f_s(x) = f(x) + \frac{1}{2}f(2x) + \frac{1}{4}f(4x) + \frac{1}{8}f(8x)$$

$n(x)$



Pseudocode for 1D Octave Noise

Input: x

Output: noise

```
float noise=0.0;  
for i = 0 to n-1:  
    noise += 2^(-i) * perlin(2^i * x);  
return noise;
```

Here **perlin()** is a pre-implemented
1D gradient noise function



That's it! We have developed the 1D Perlin noise!

- If you put together these ideas: **grid discretization**, **hash function**, **smooth step**, **gradient interpolation**, and **octave synthesis**, we will get the full model of a **1D Perlin noise function**:

$$s(x) = 3x^2 - 2x^3$$

Smooth Step

$$f(x) = (1 - s(x))xg_0 + s(x)(x - 1)g_1$$

Hash Functions

Gradient Interpolation

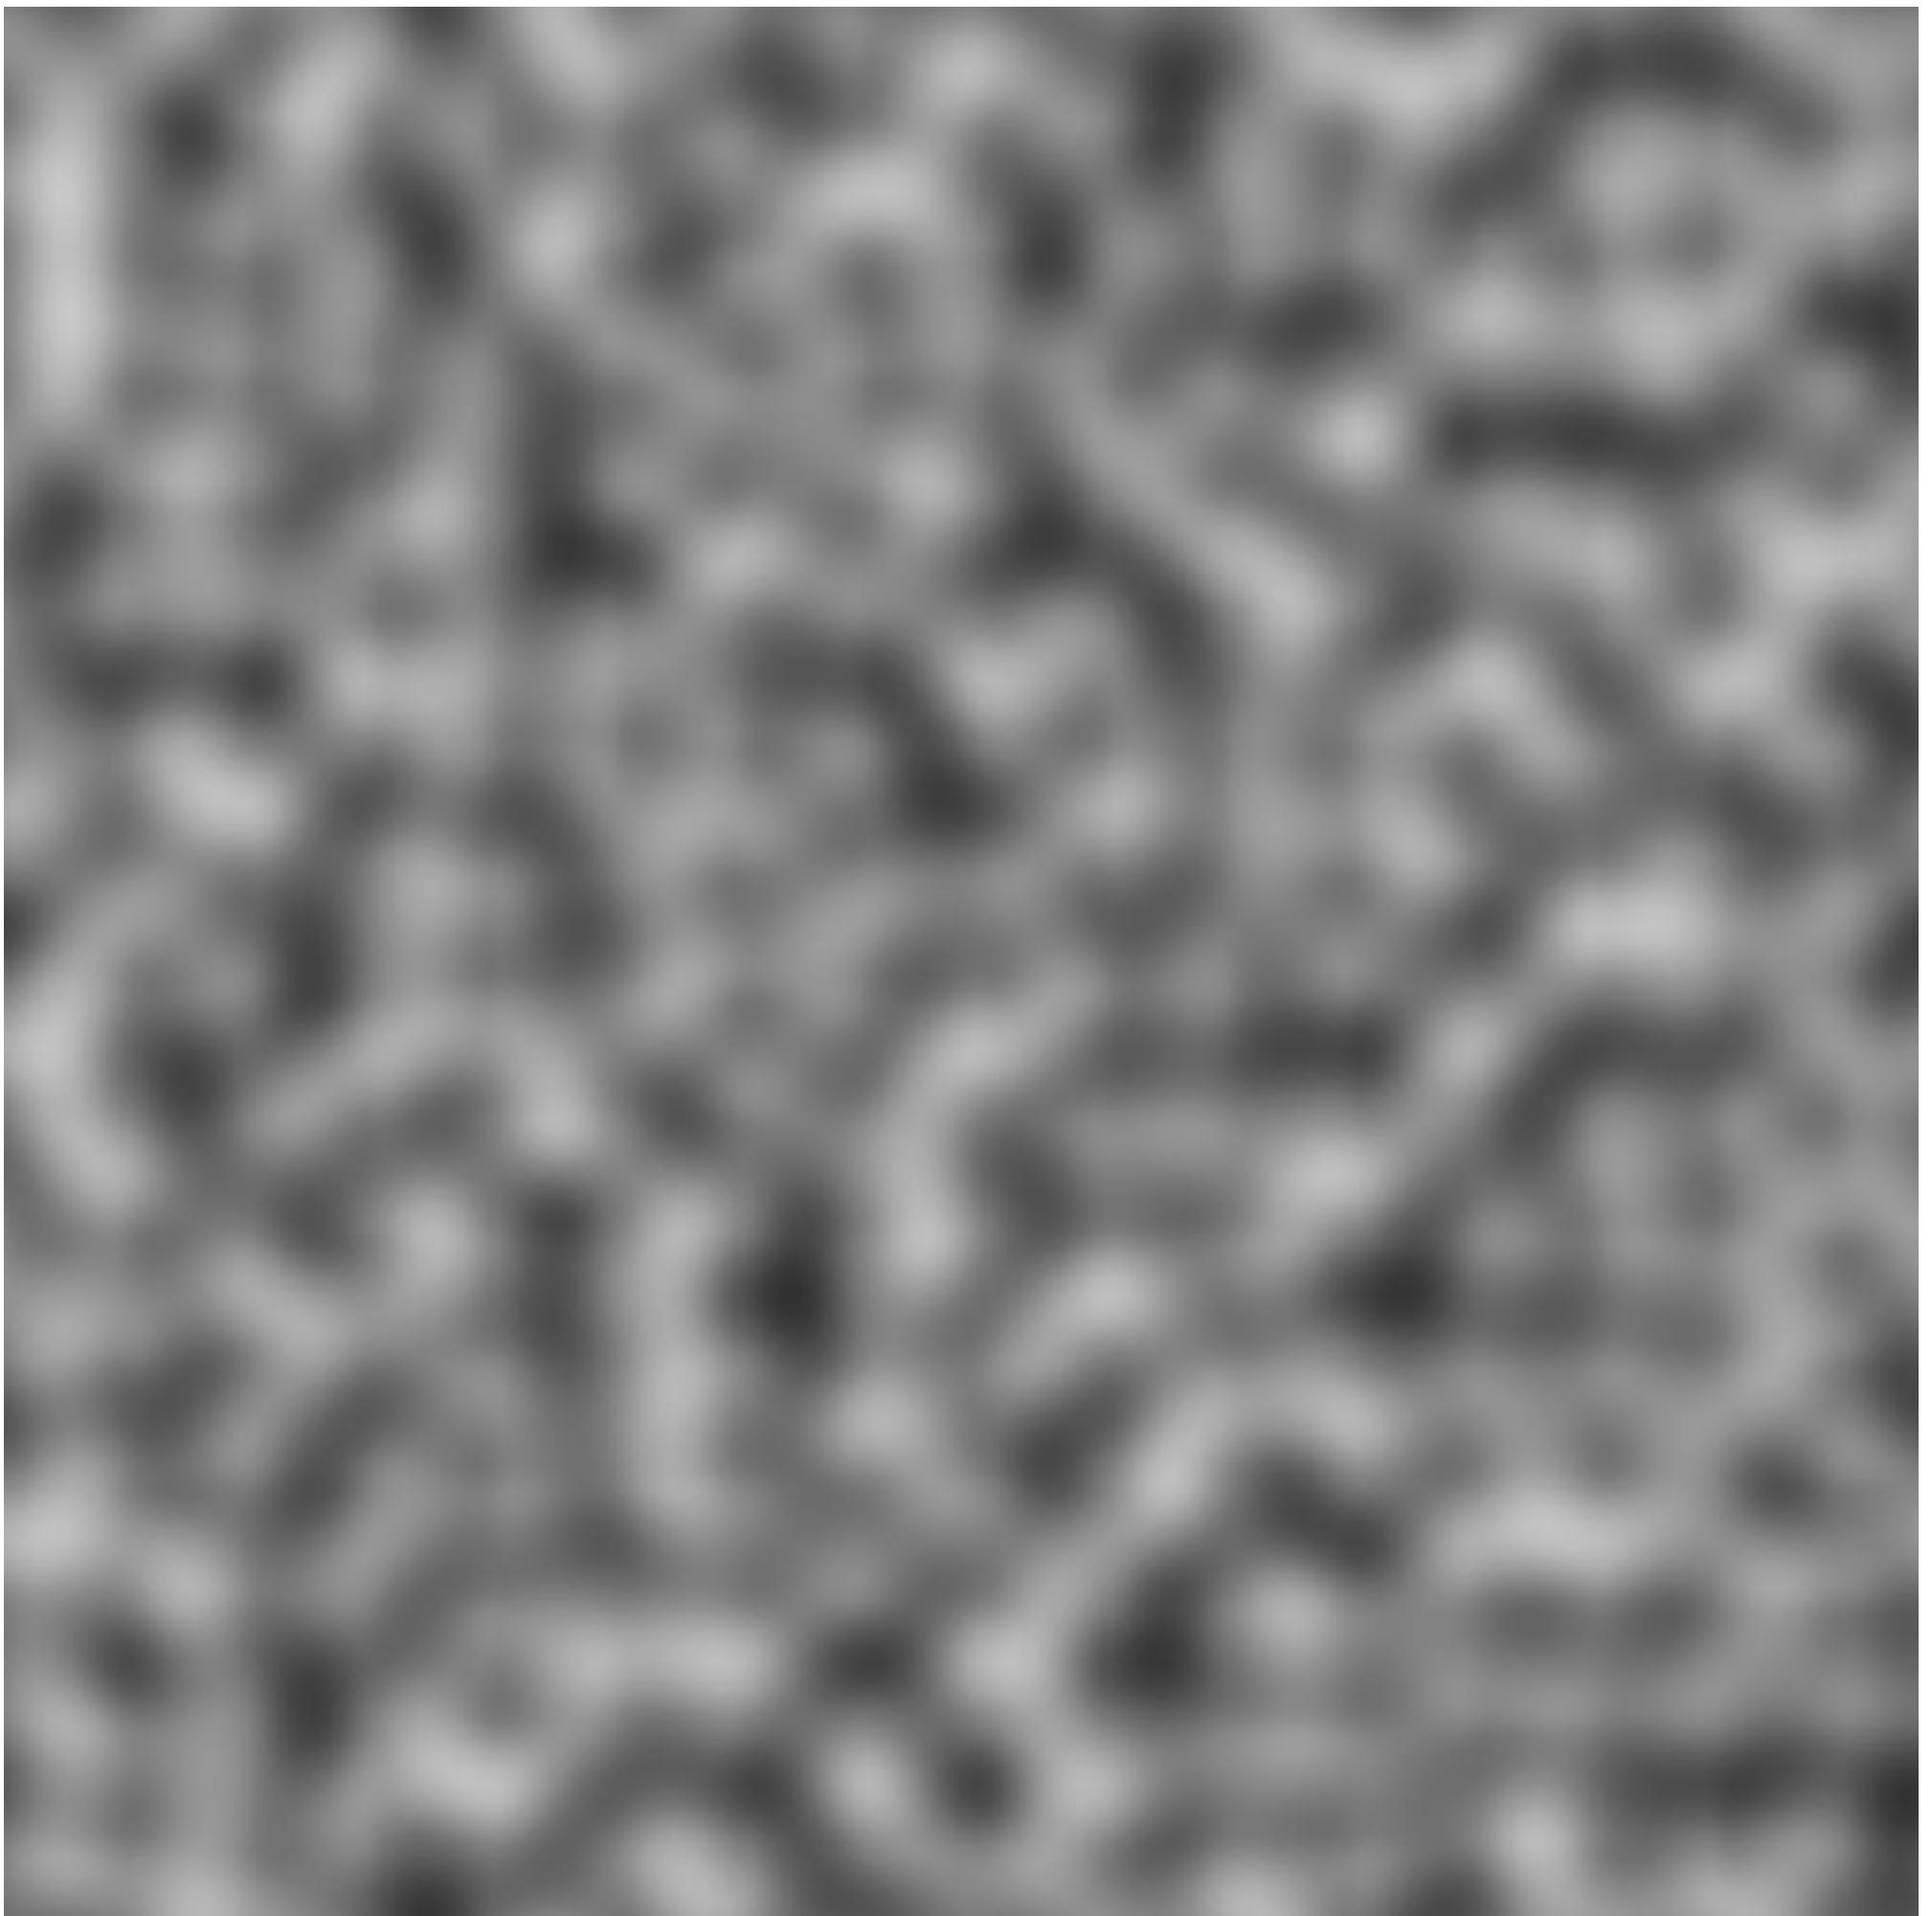
$$f_s(x) = f(x) + \frac{1}{2}f(2x) + \frac{1}{4}f(4x) + \frac{1}{8}f(8x) + \dots$$

Octave Synthesis

How to go from 1D to 2D?



2D Perlin Noise

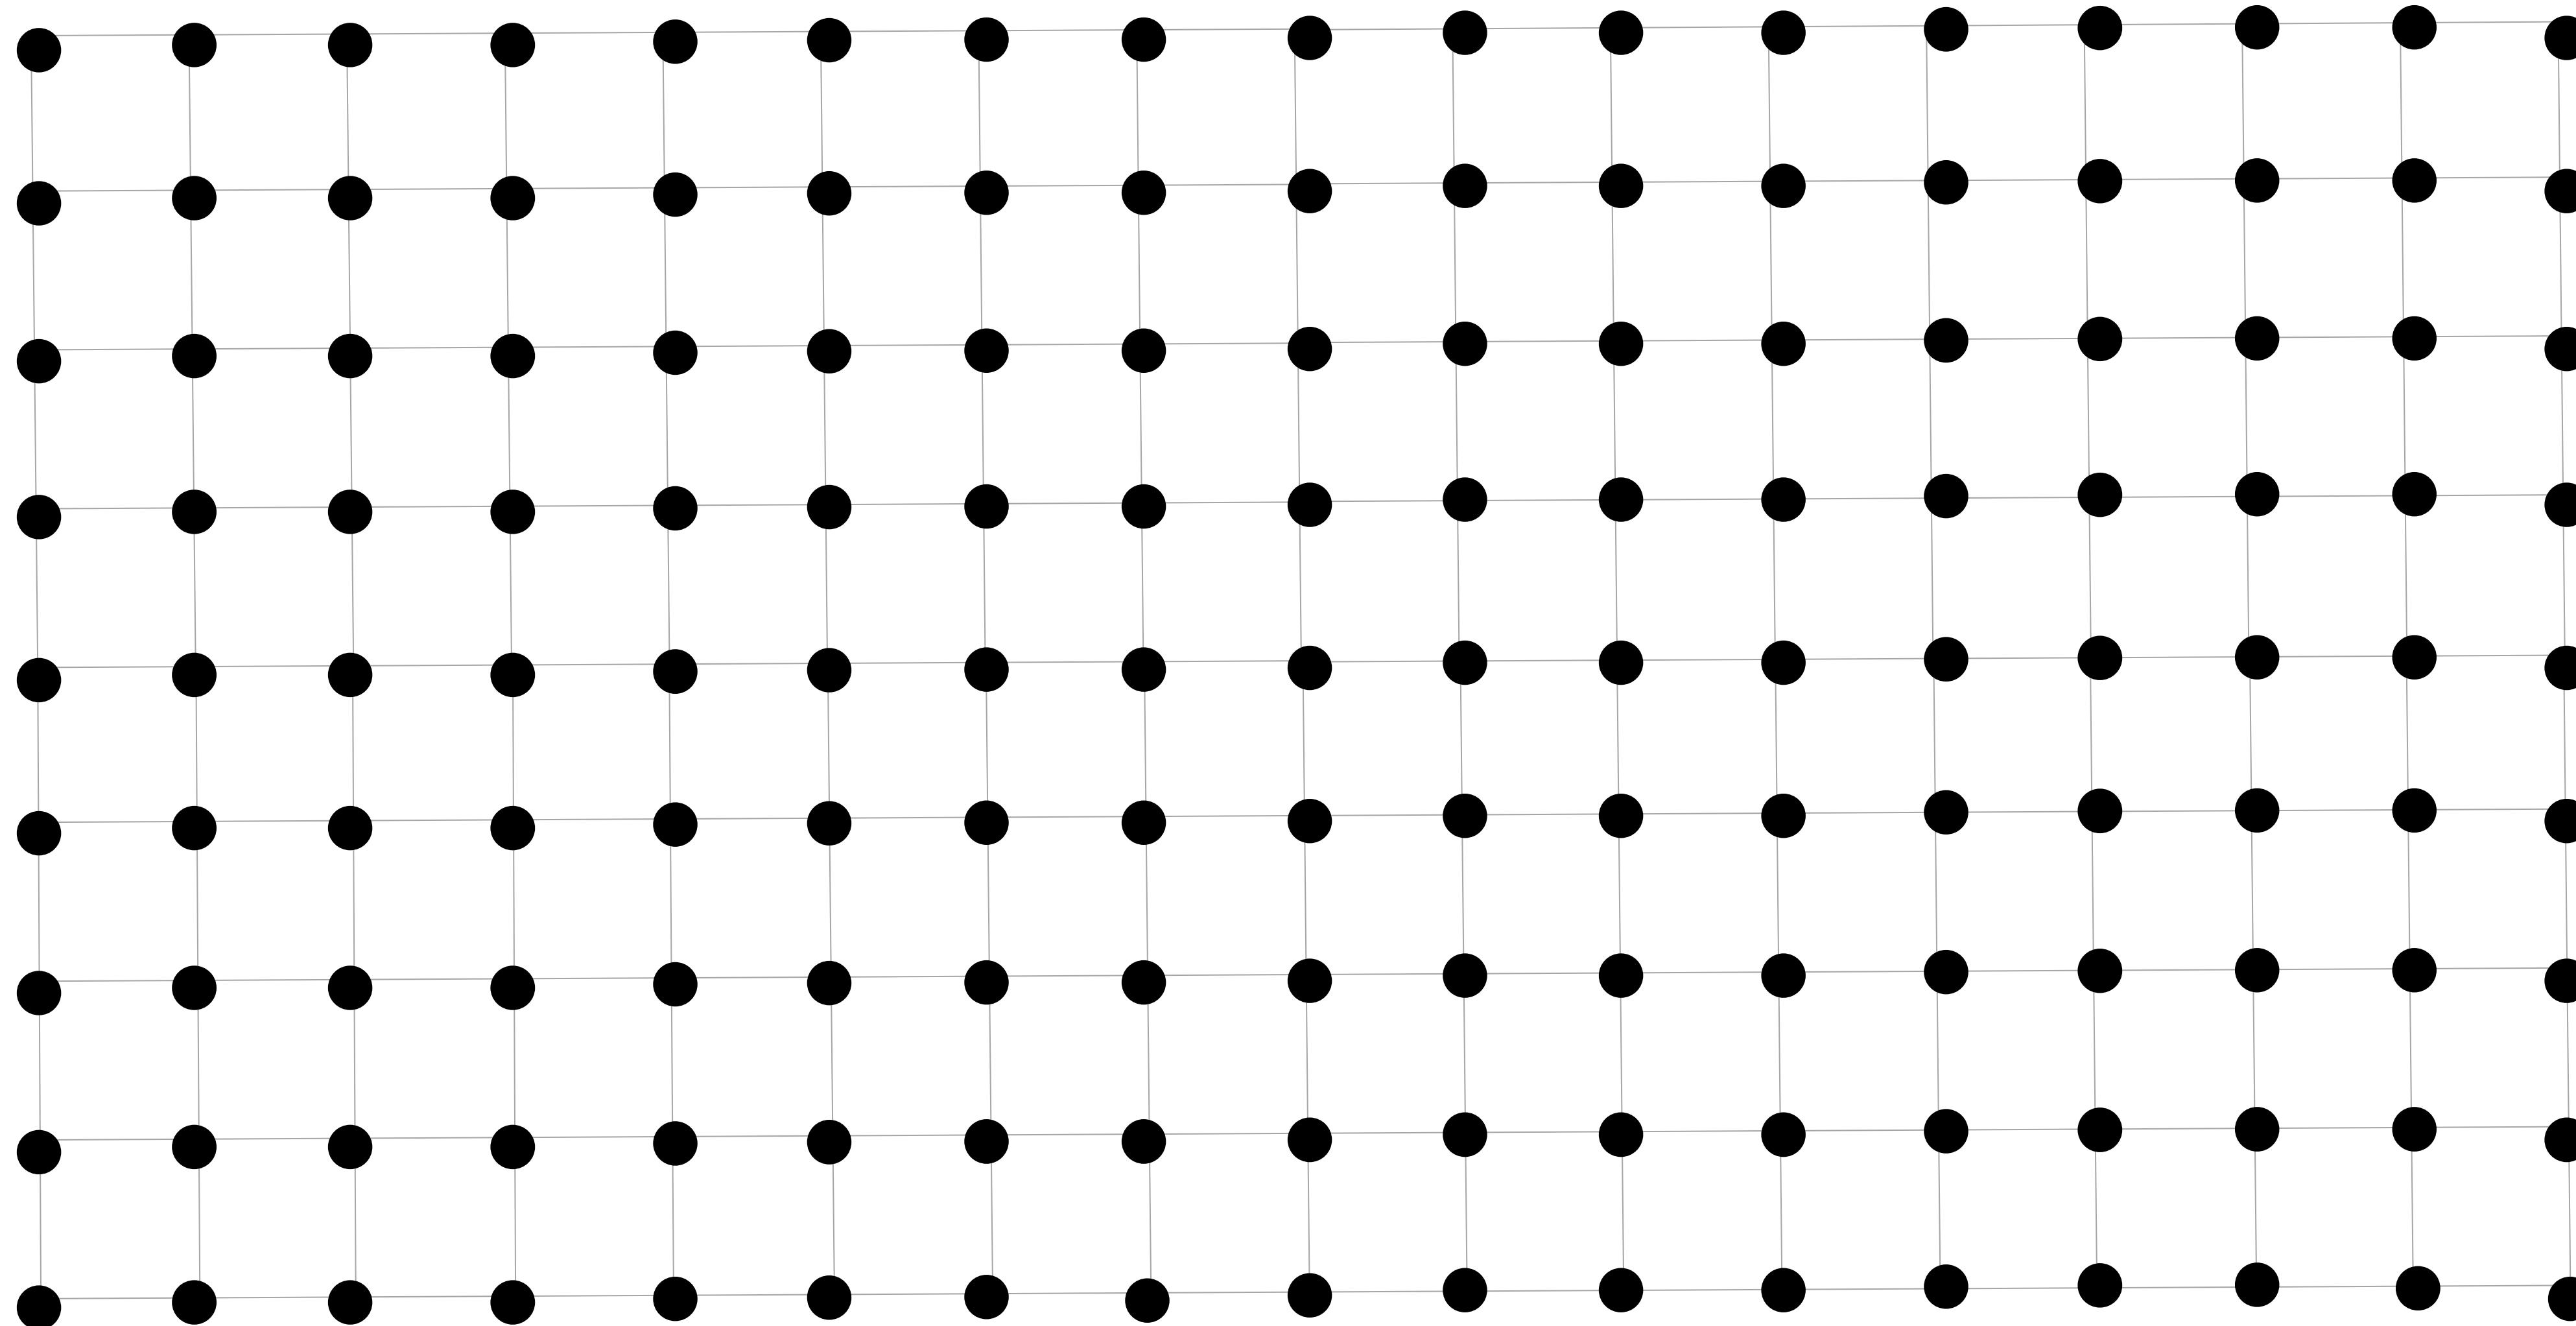


We can extend the concept of the 1D Perlin noise to calculate the 2D Perlin noise

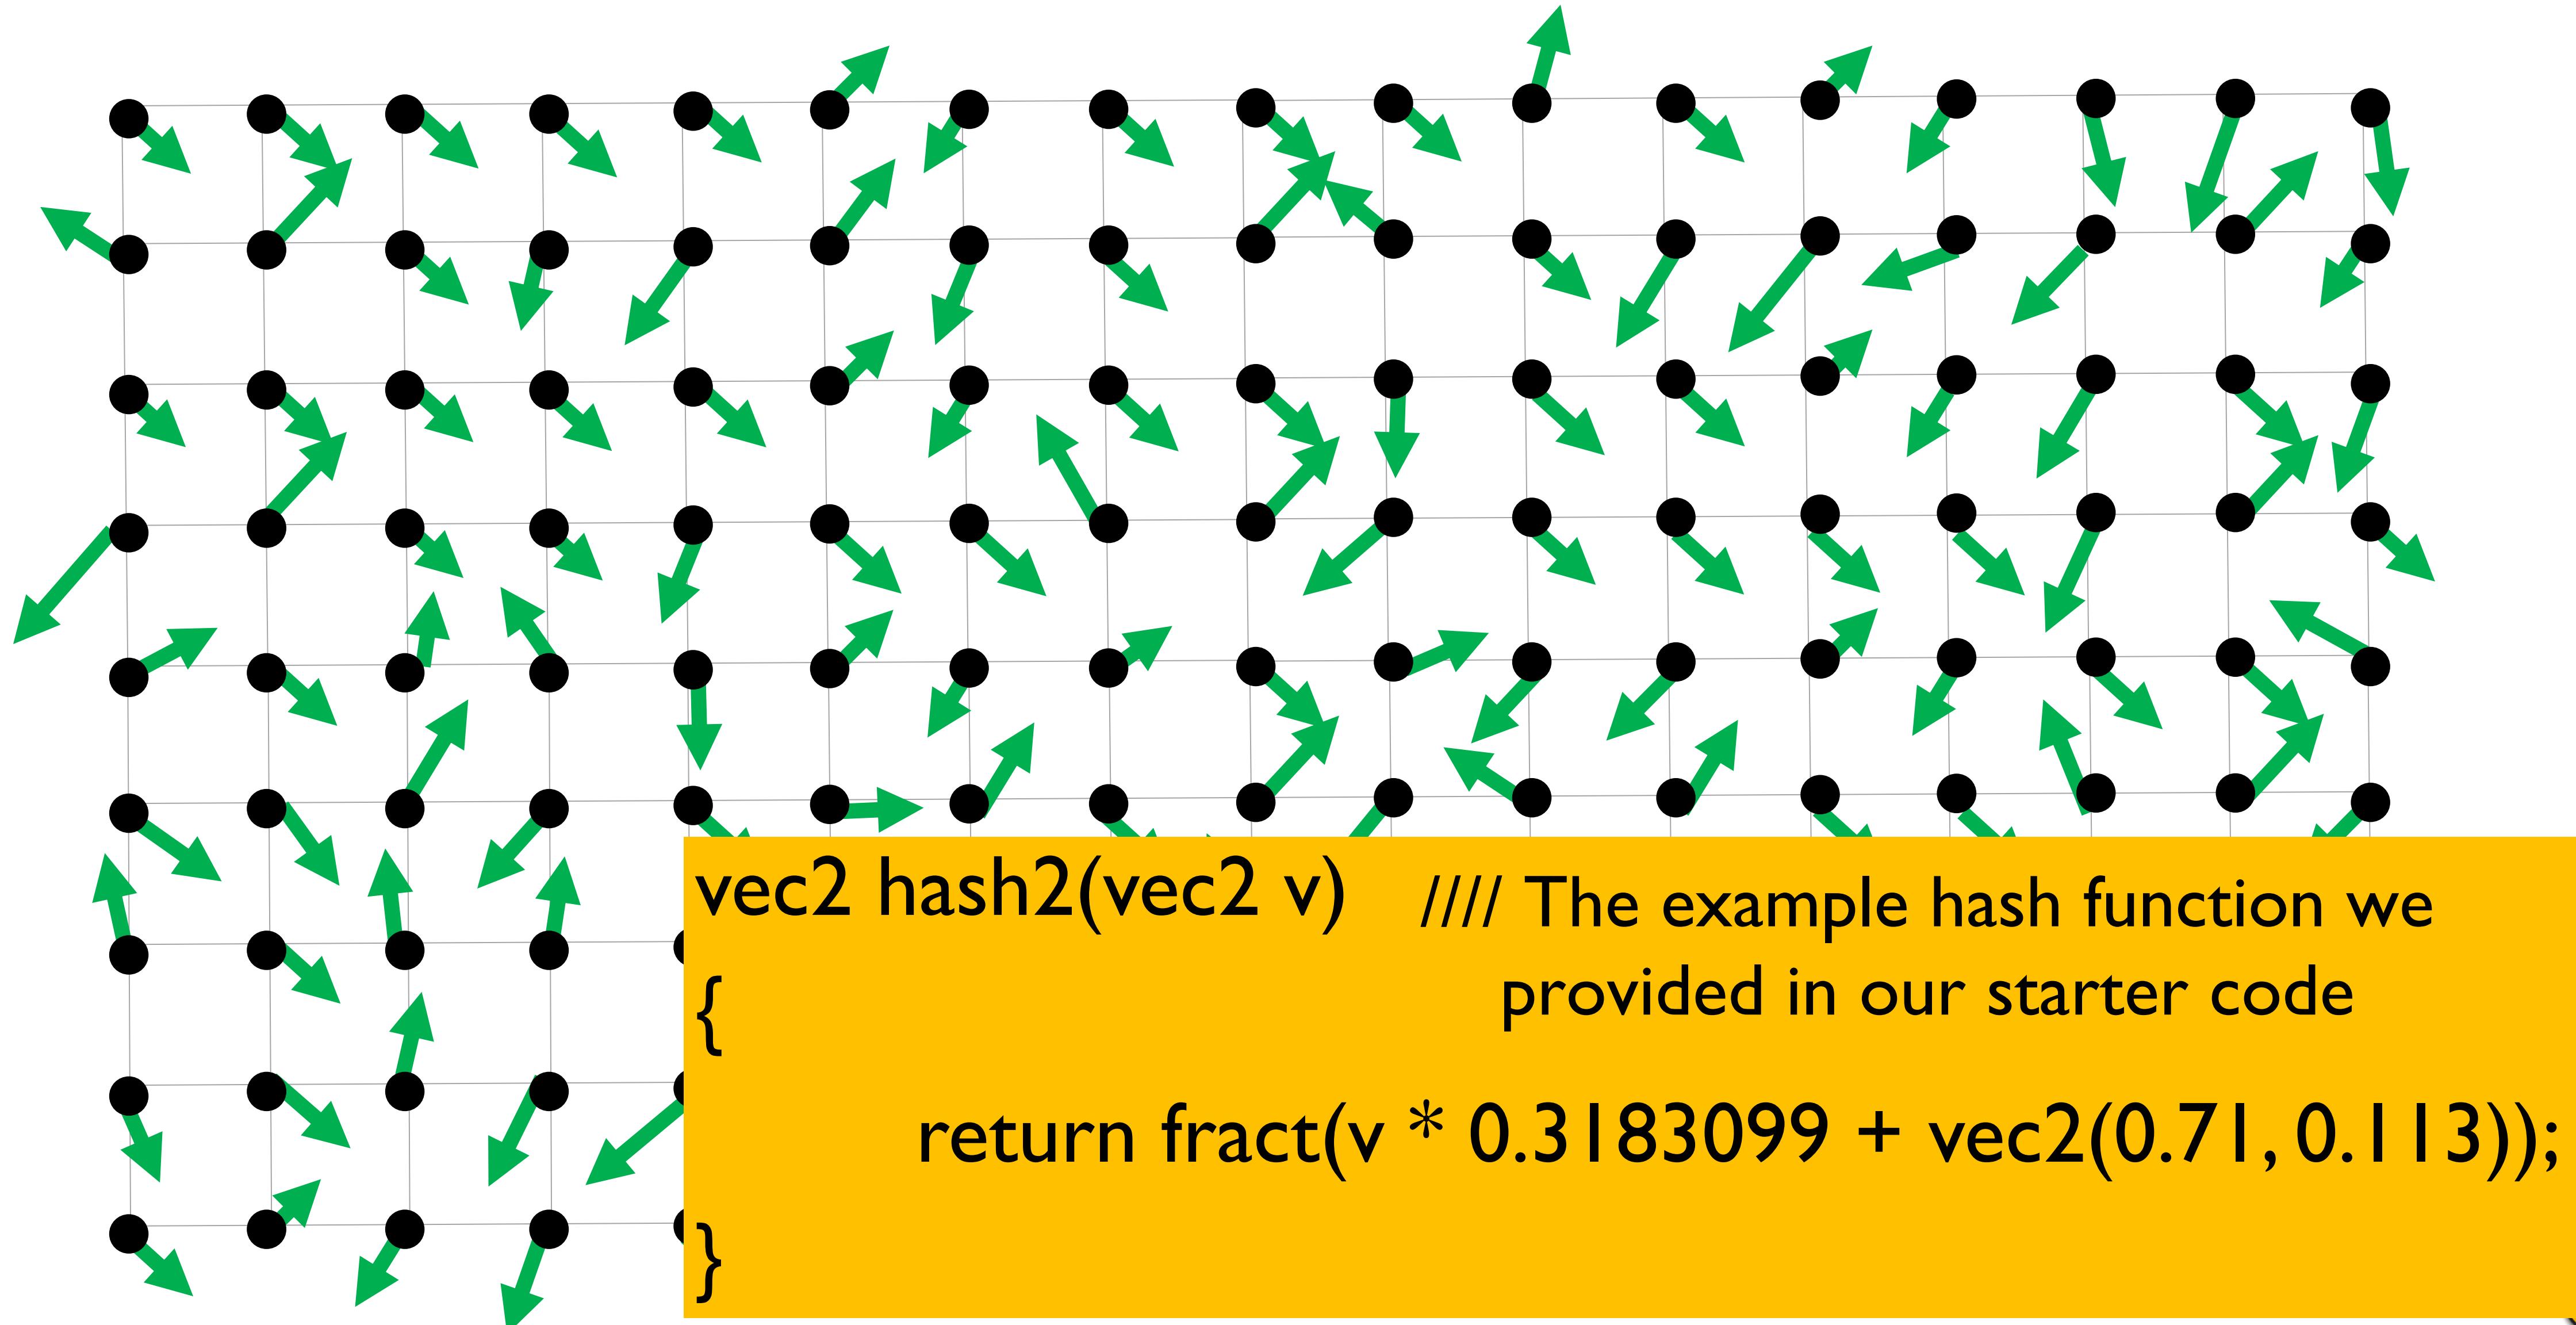
1. Discretize the 2D domain with a 2D grid
2. Generate a randomized gradient vector for each grid nodes using a 2D hash function
3. Interpolate gradient vectors within each 2D grid cell with the smooth step function
4. Octave synthesis on 2D grids



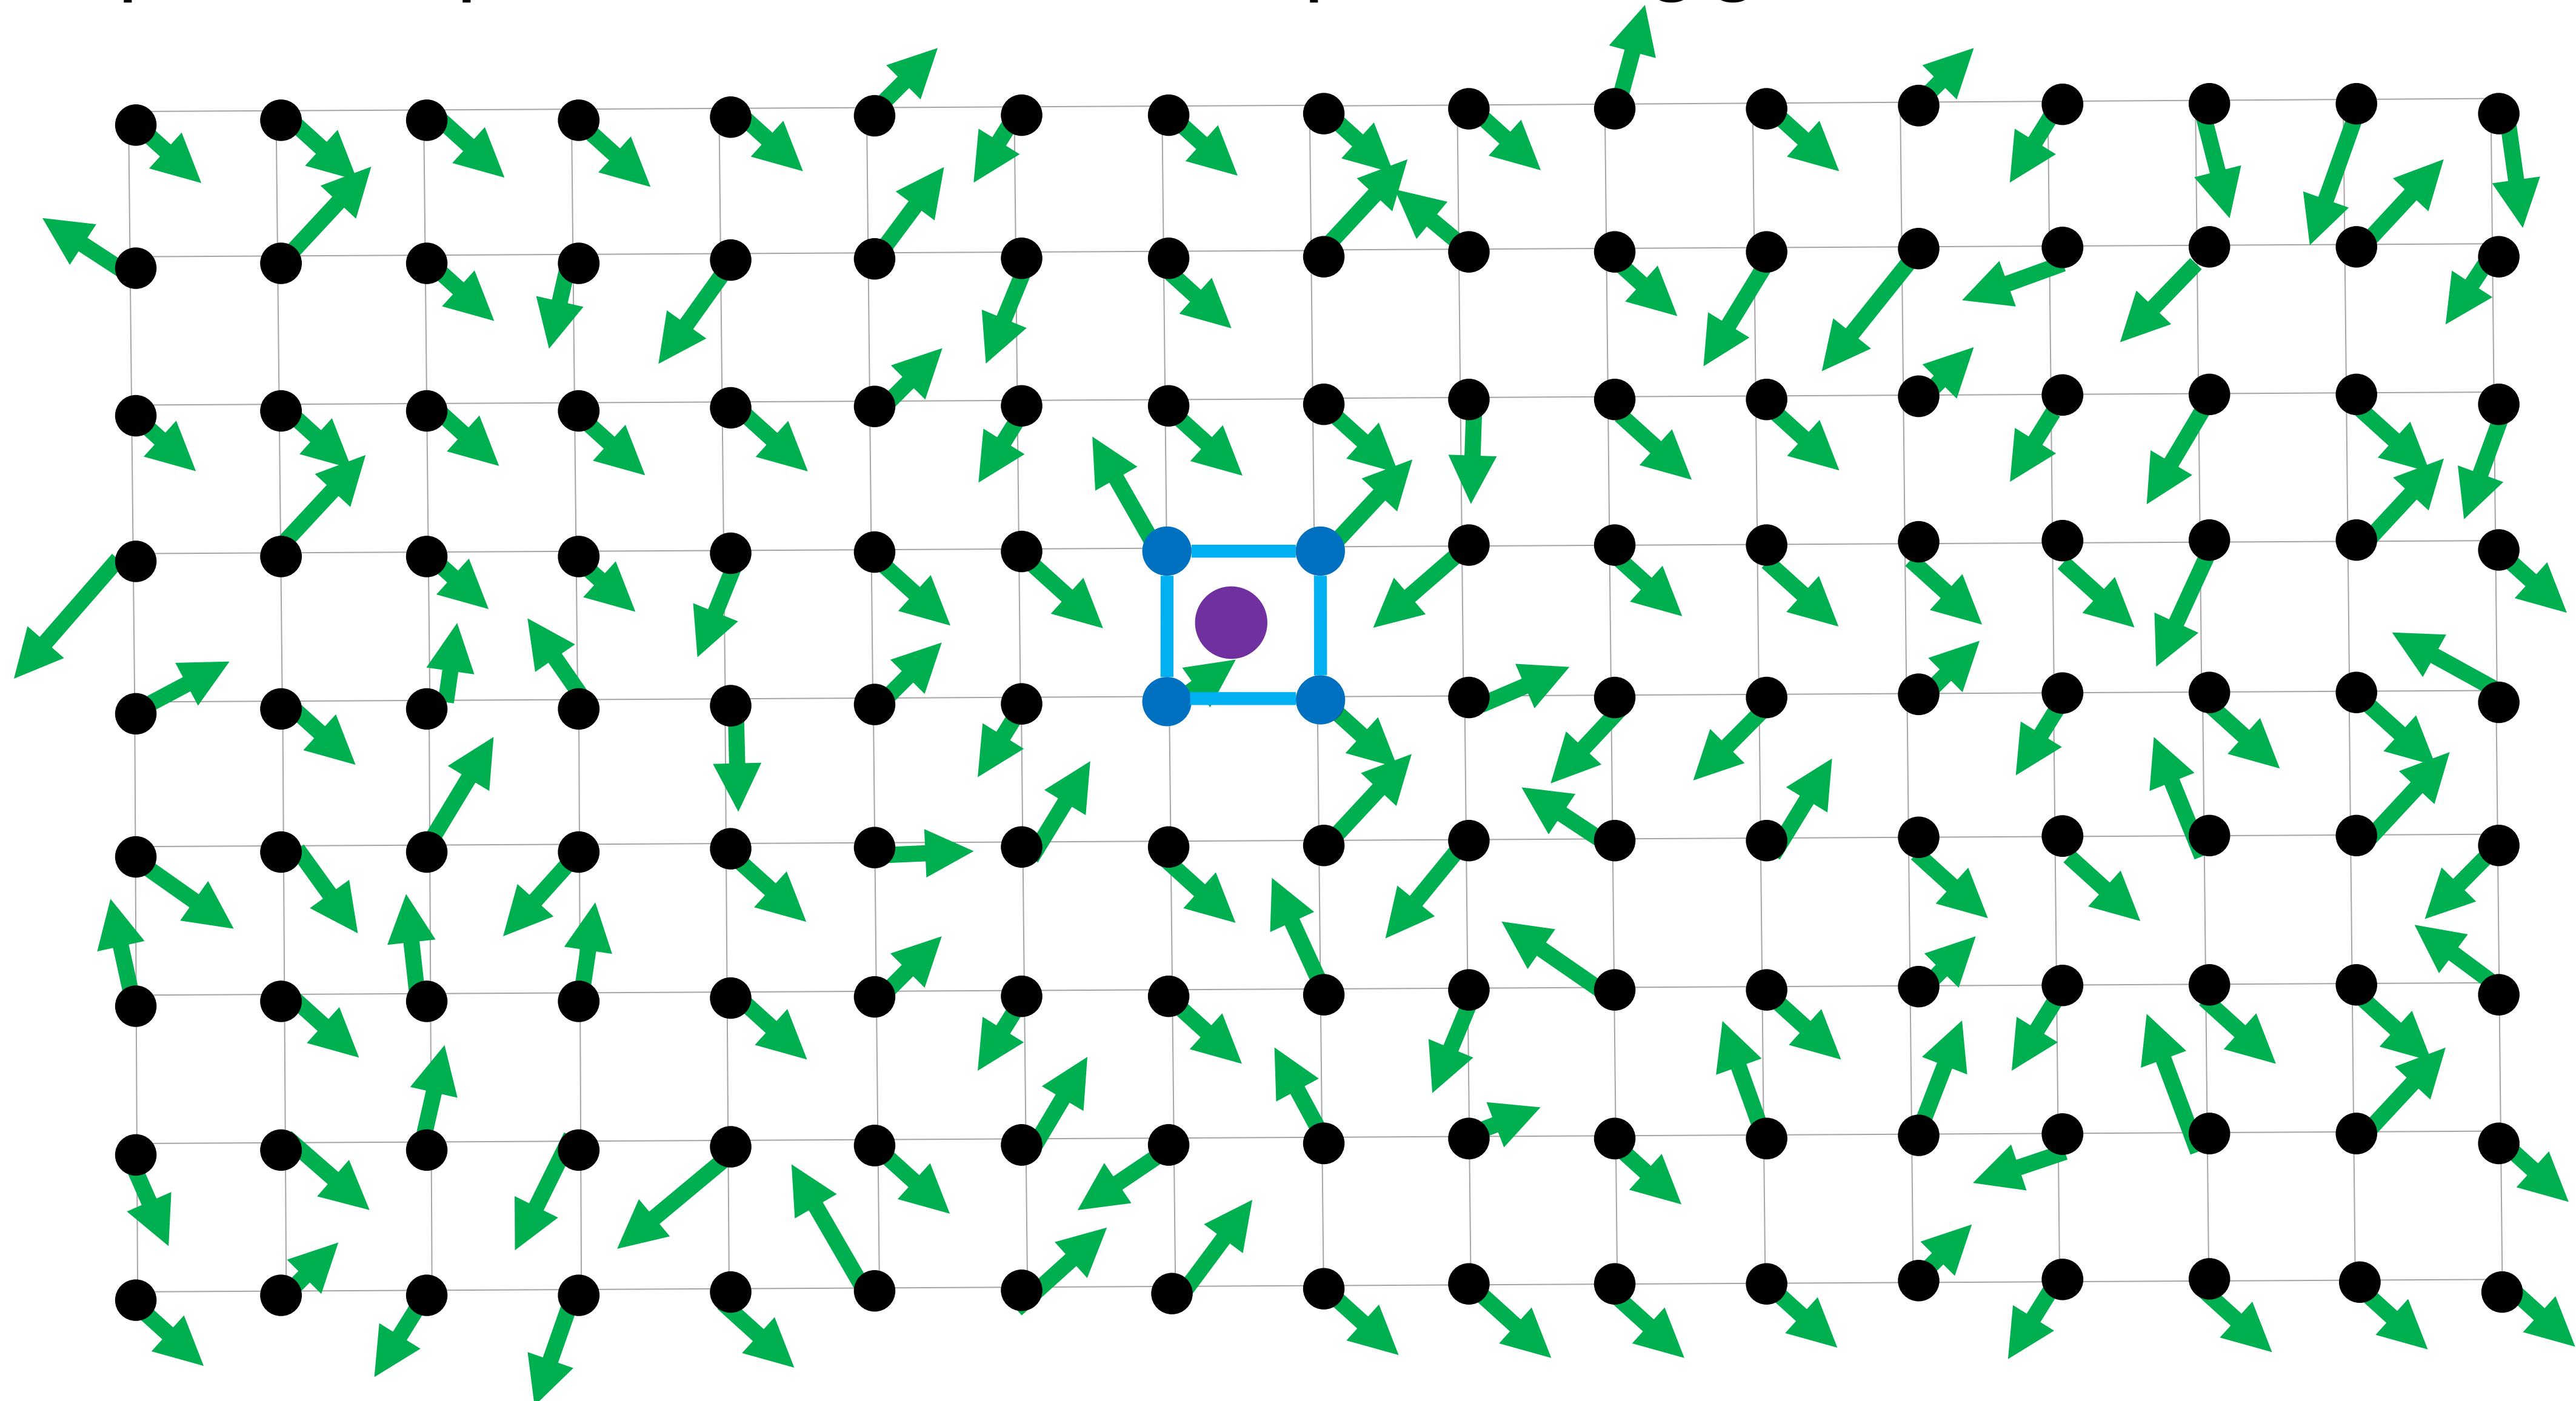
Step 1: Discretize the domain with a grid



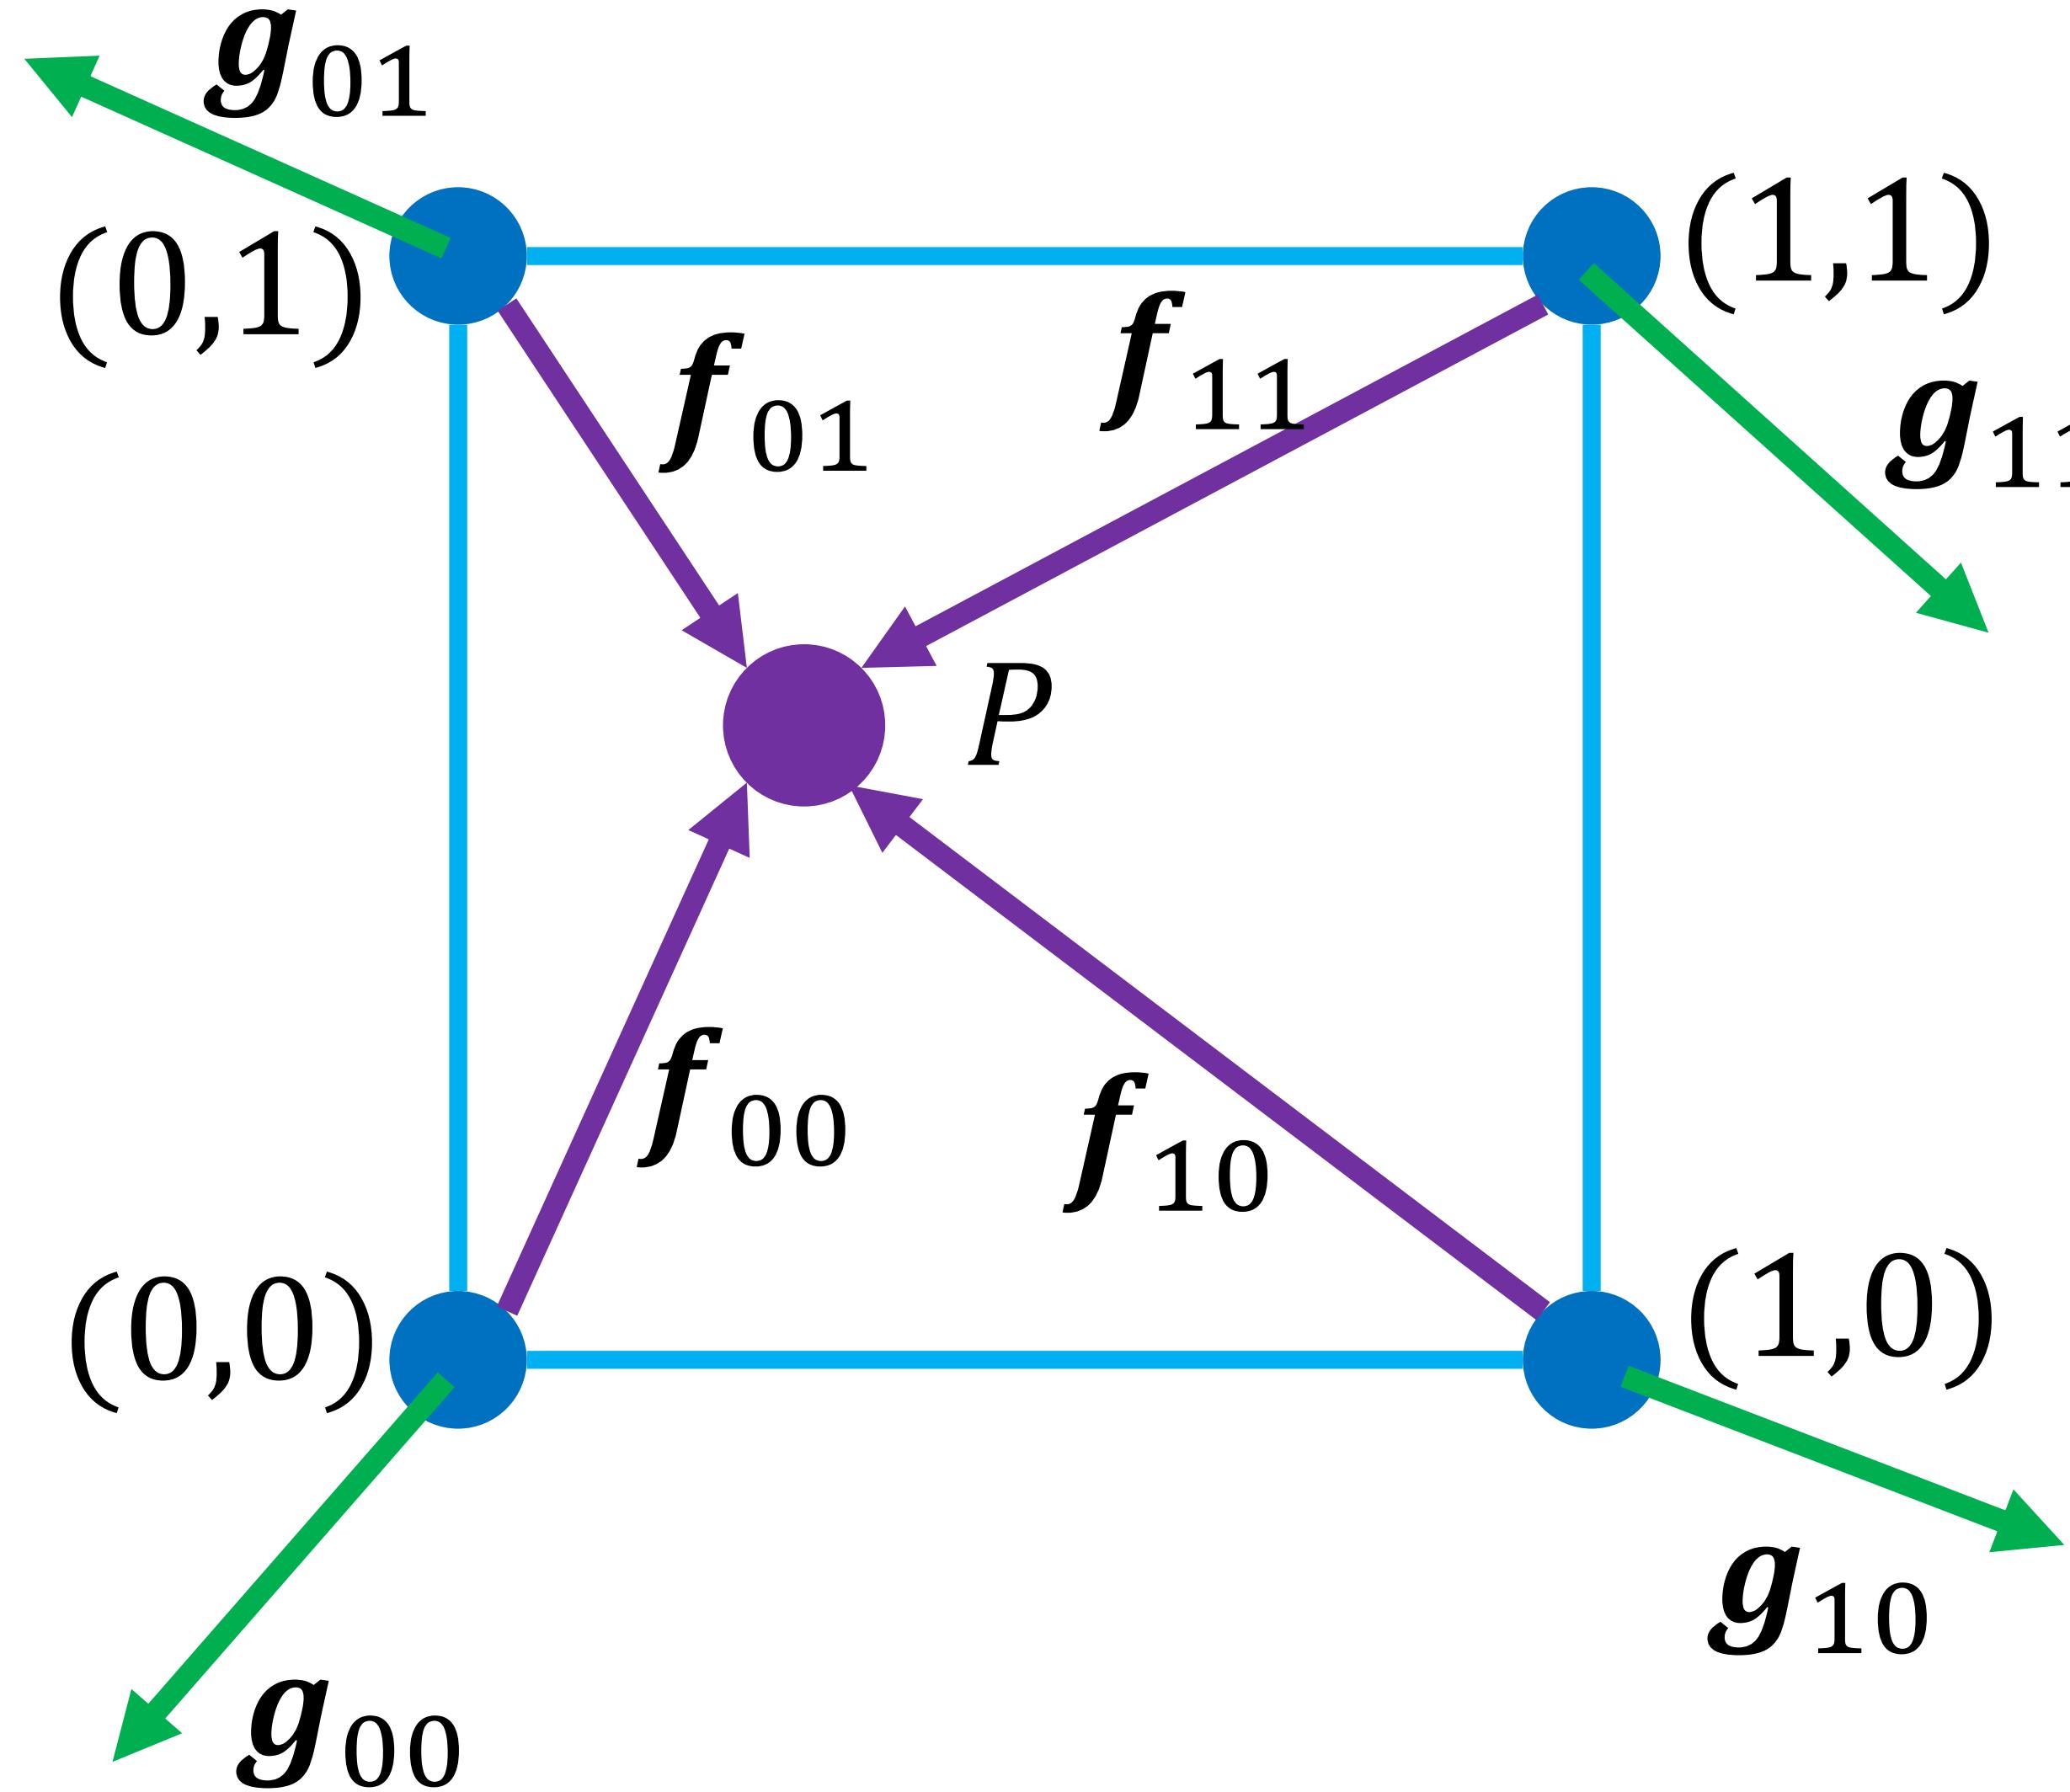
Step II: Generate random gradient vectors on grid nodes with hash functions



Step III: Interpolate noise value for a point using grid nodes



Step III.I: Calculate the **offset vector** for each grid node



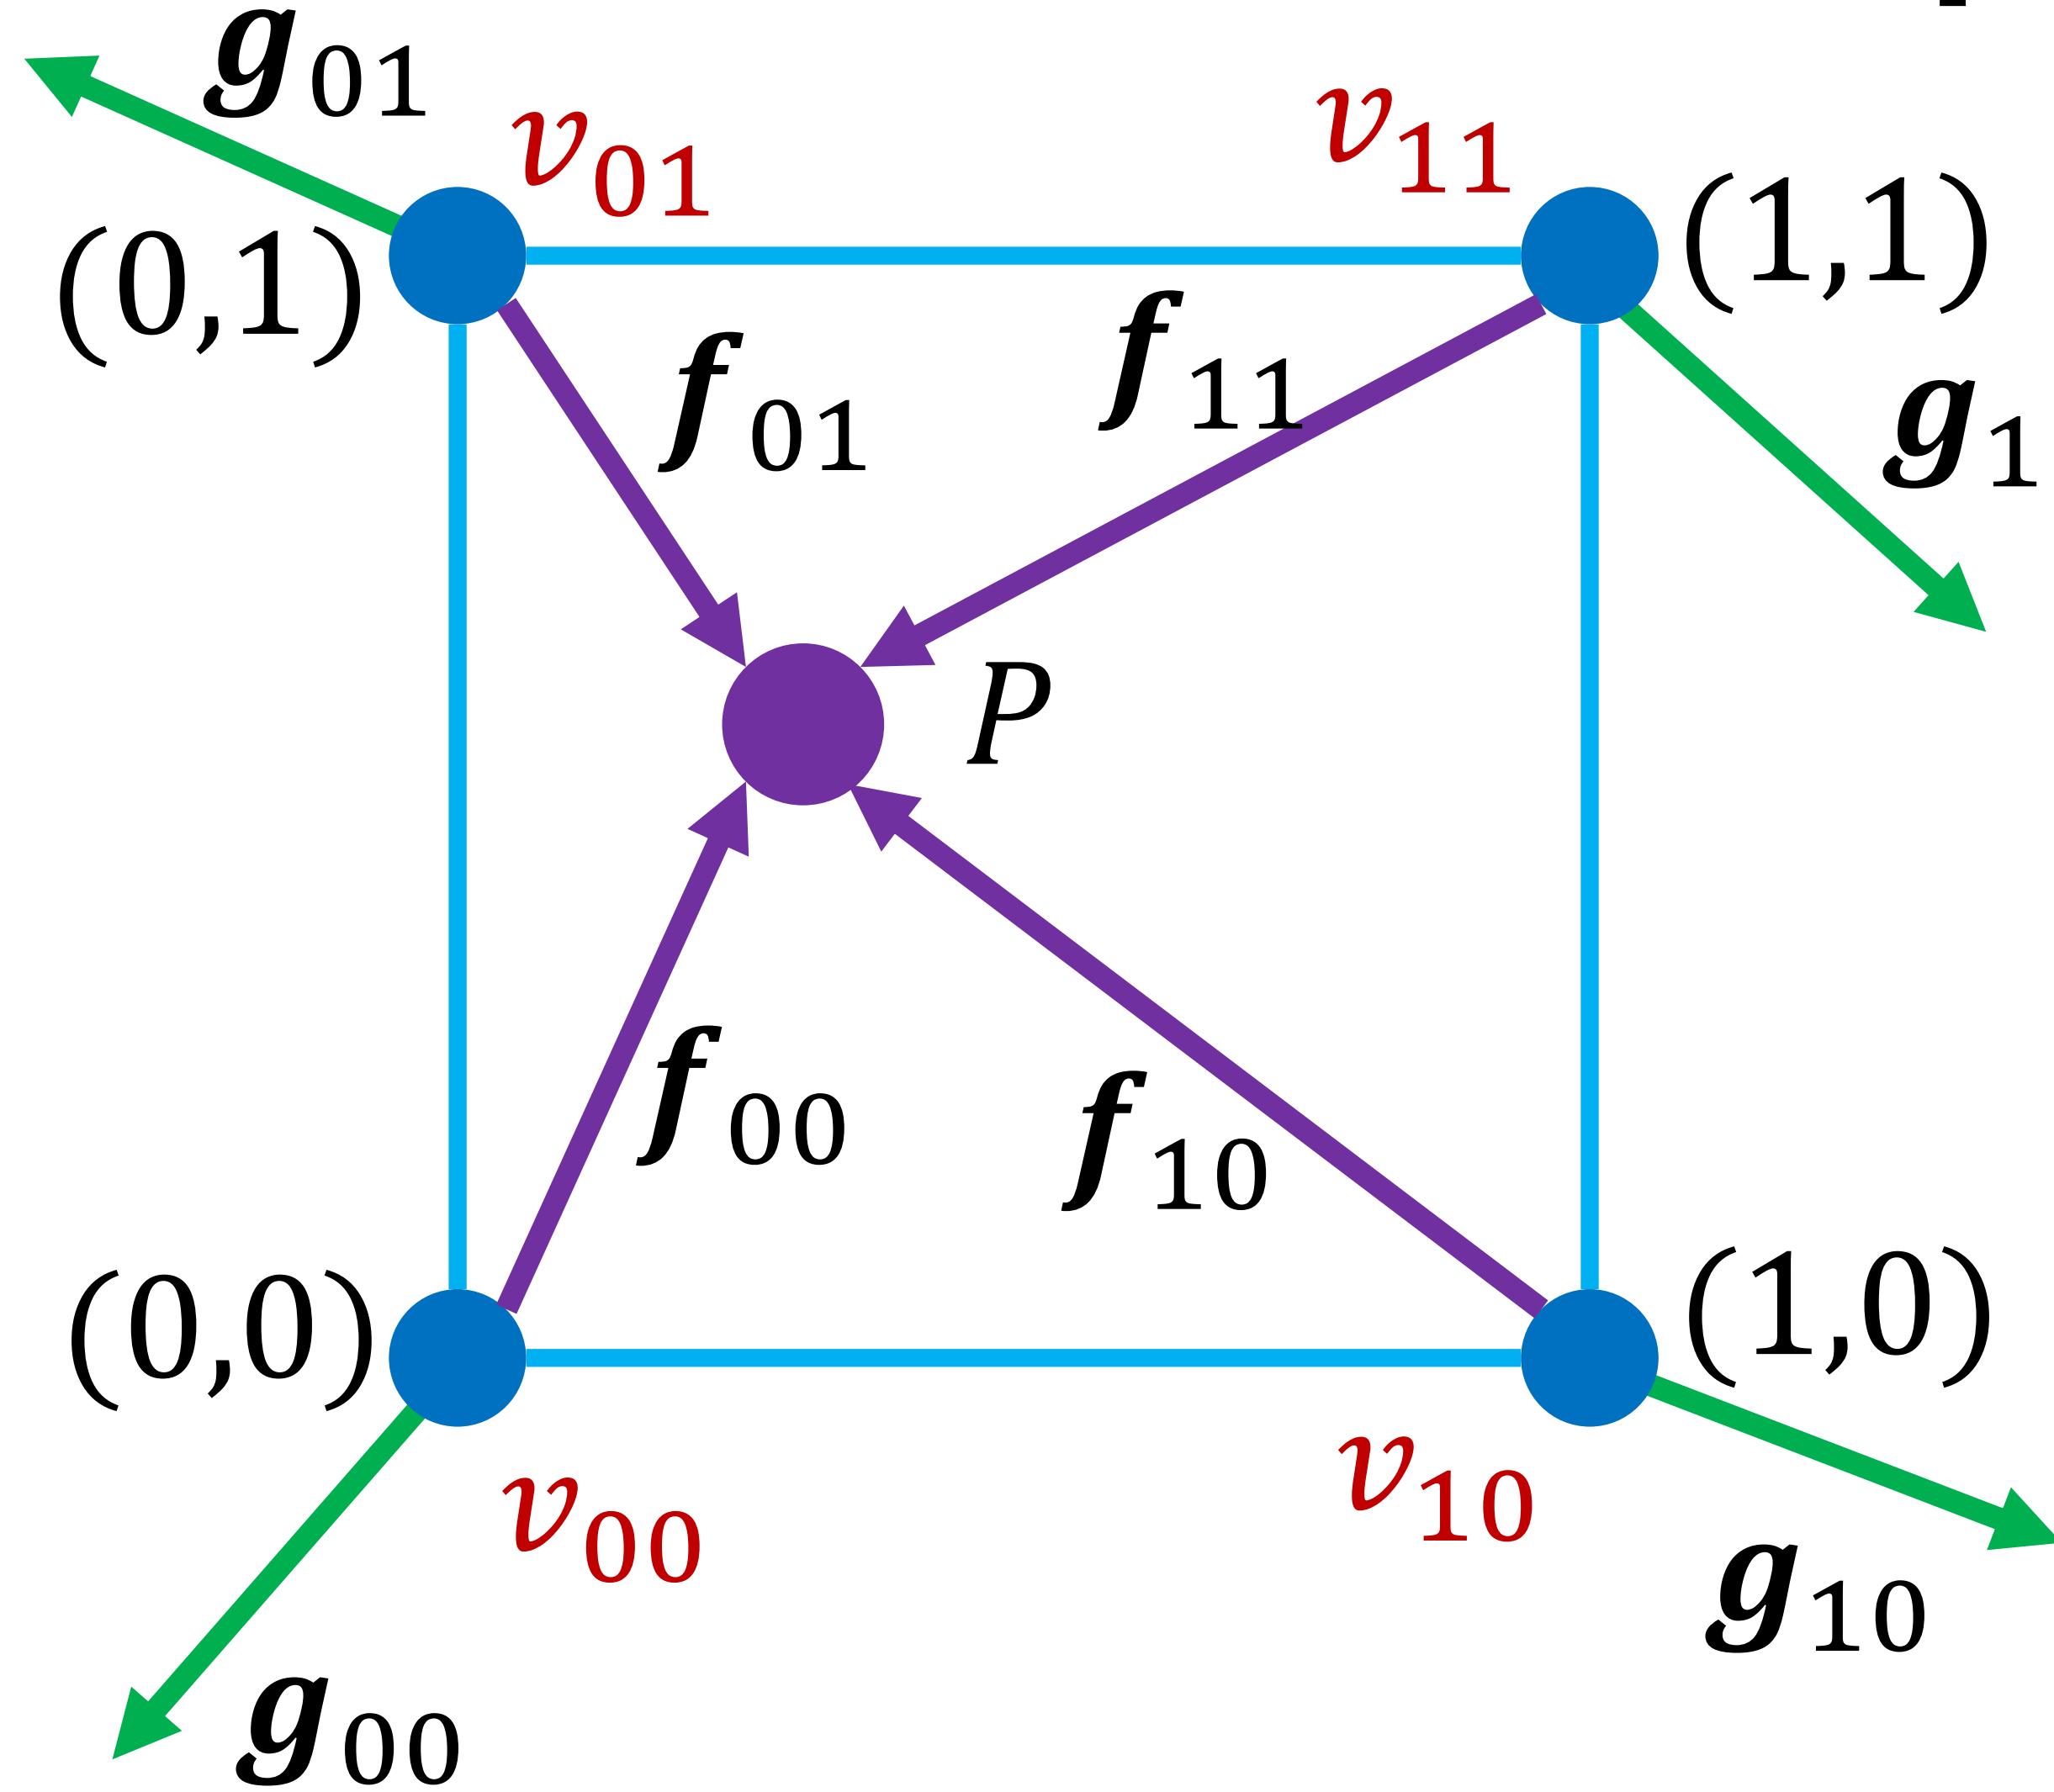
$$f_{00} = P - (0,0)$$

$$f_{10} = P - (1,0)$$

$$f_{01} = P - (0,1)$$

$$f_{11} = P - (1,1)$$

Step III.2: Calculate the **dot product** for each grid node



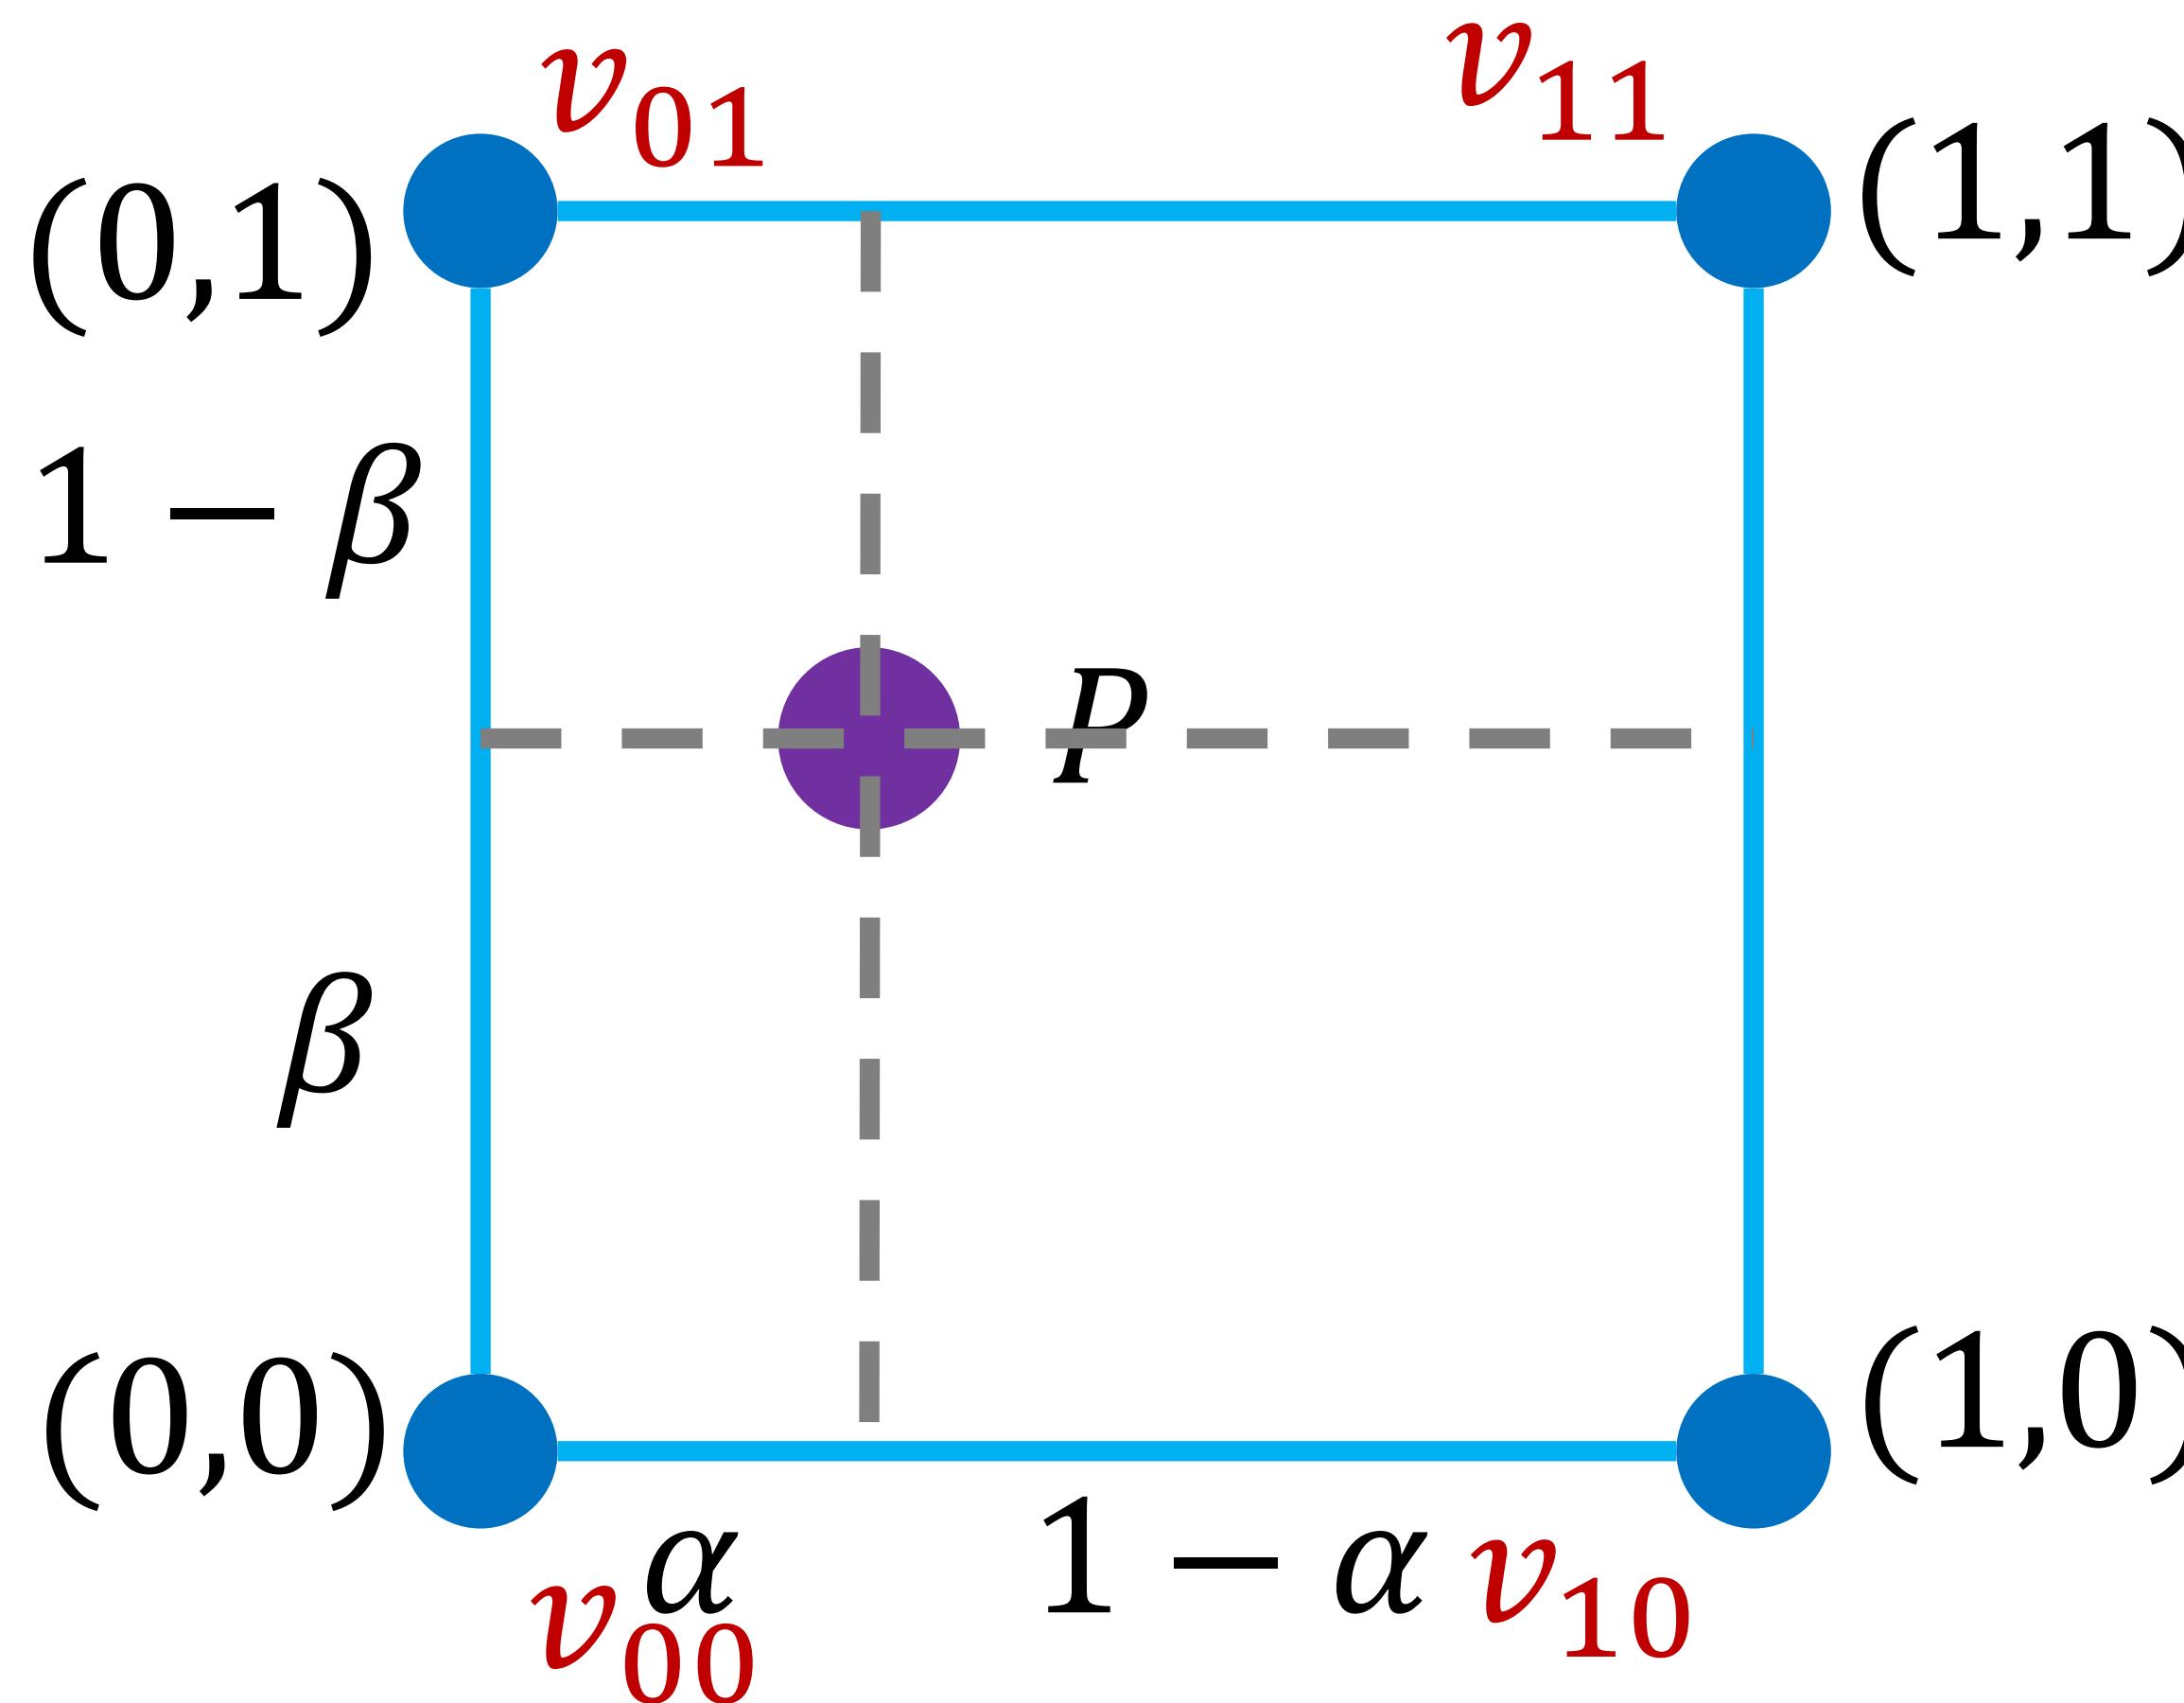
$$v_{00} = f_{00} \cdot g_{00}$$

$$v_{10} = f_{10} \cdot g_{10}$$

$$v_{01} = f_{01} \cdot g_{01}$$

$$v_{11} = f_{11} \cdot g_{11}$$

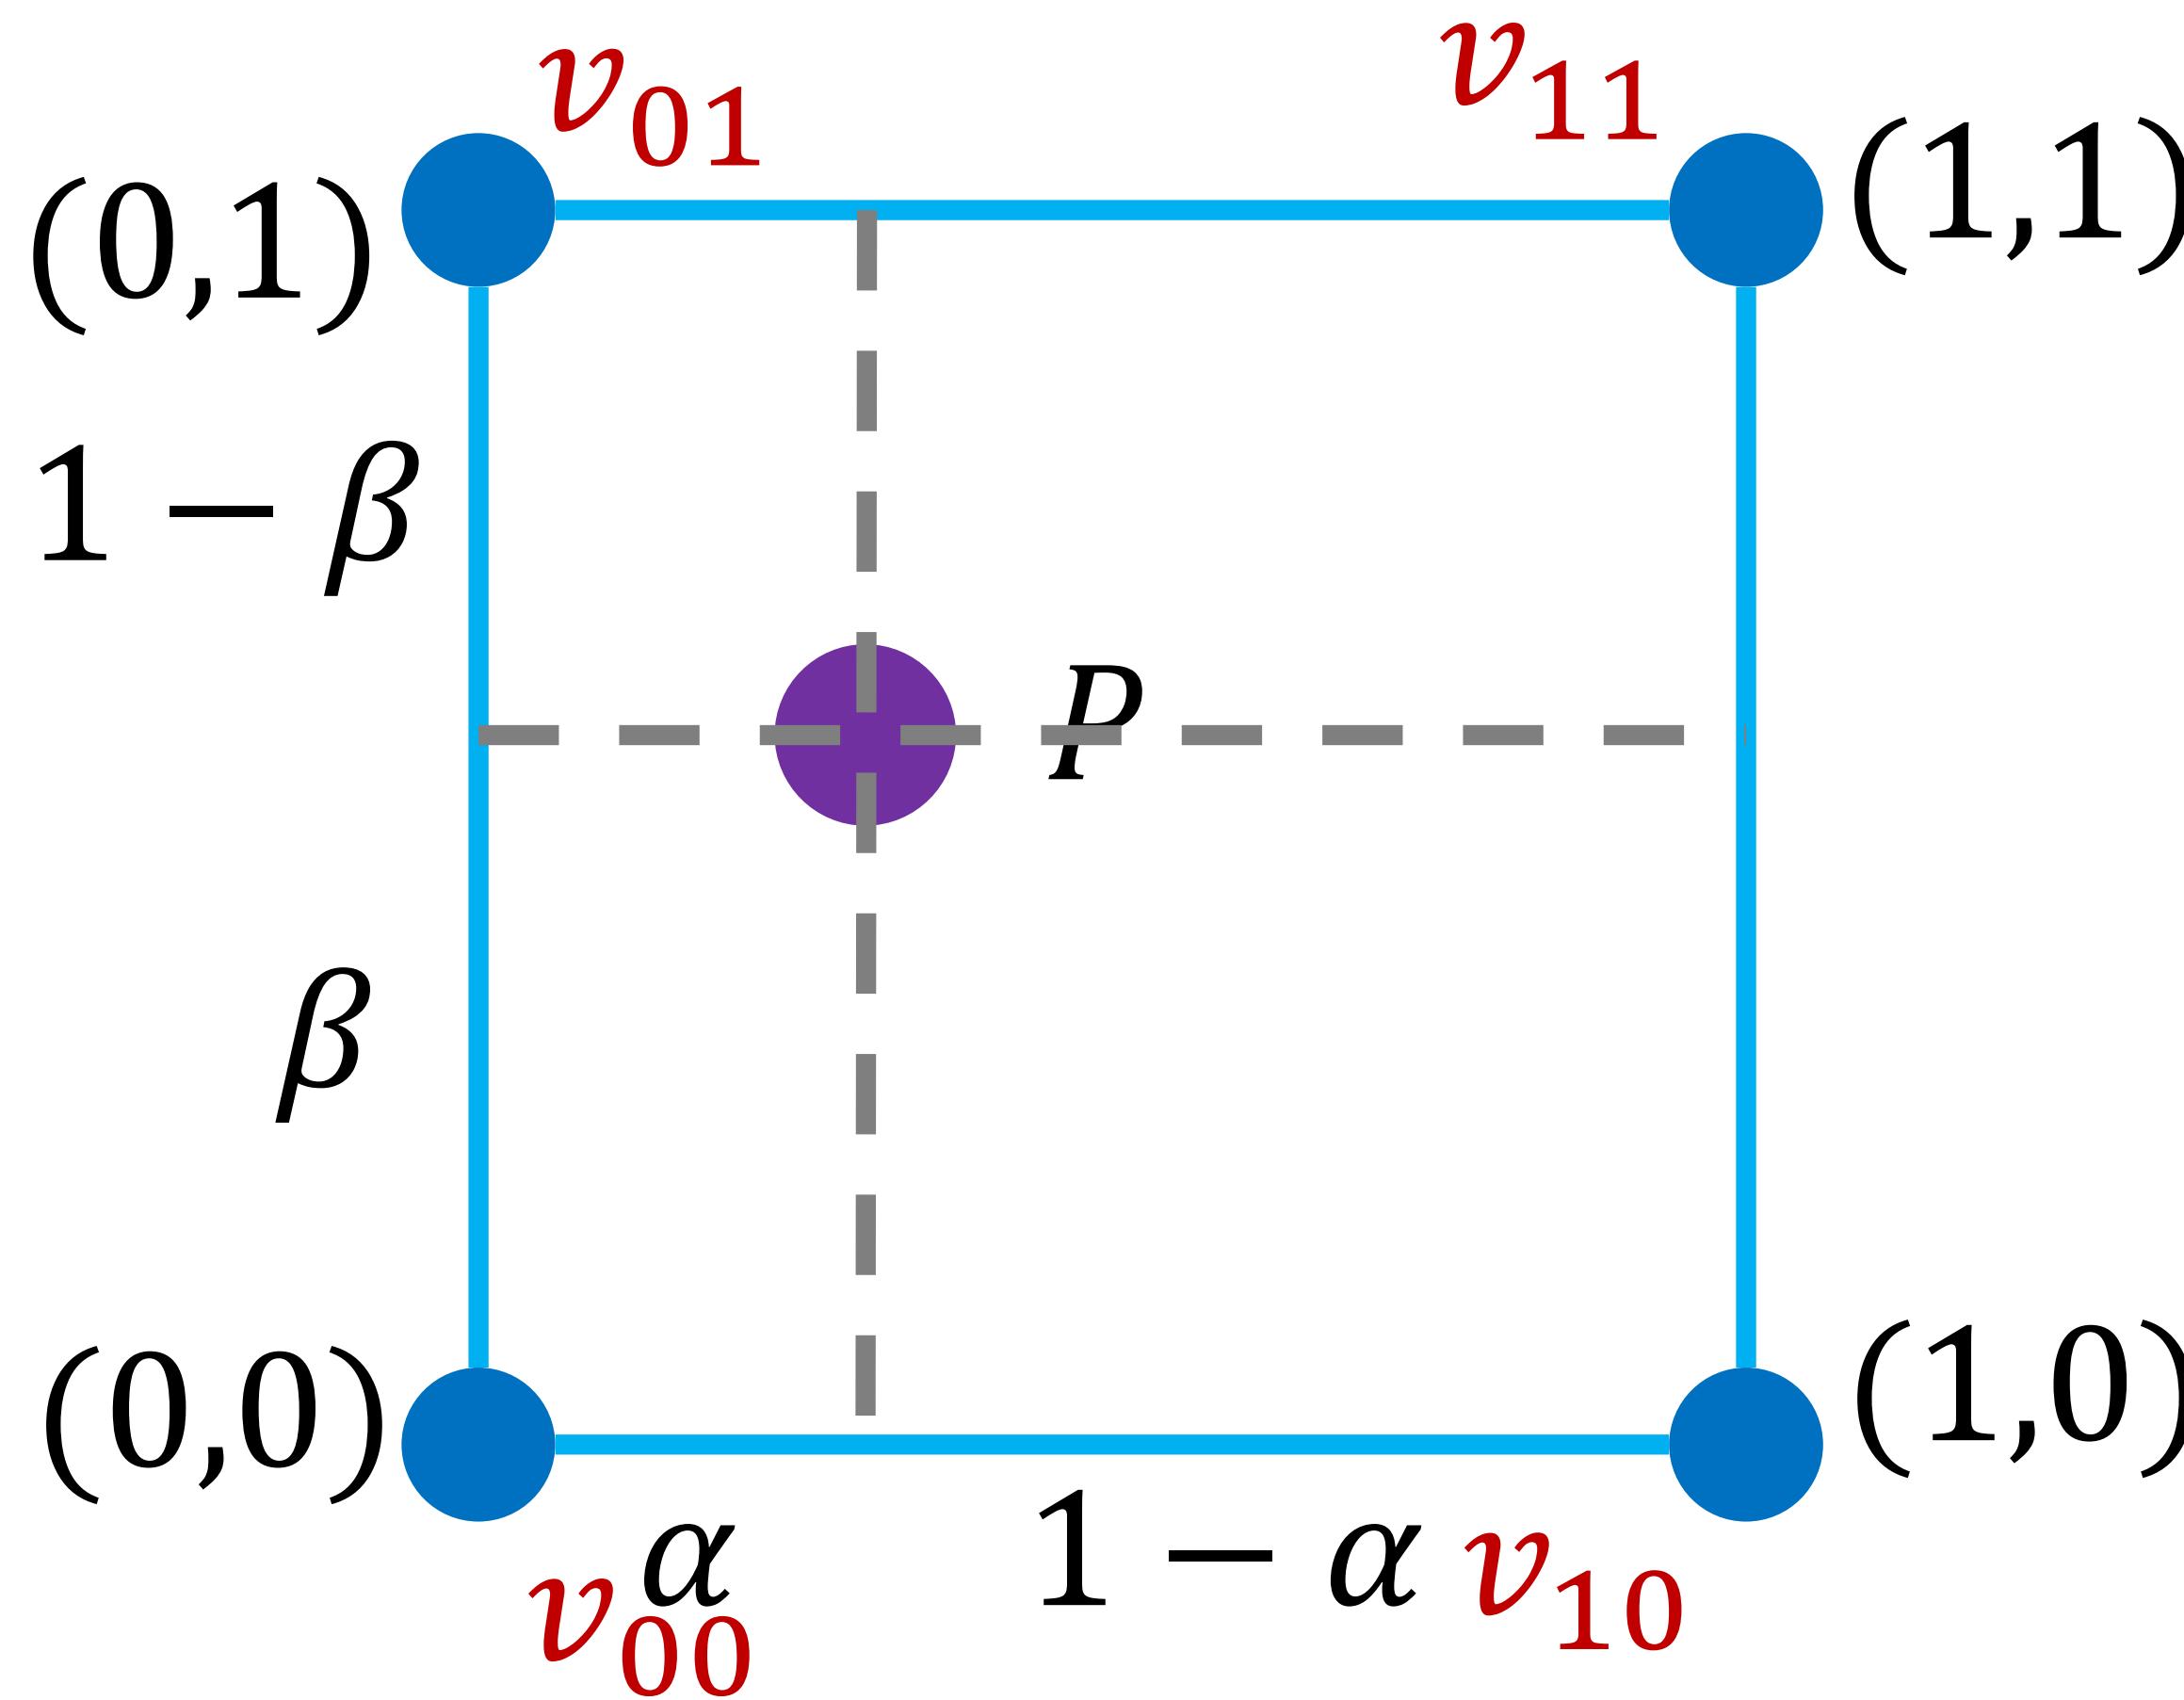
Step III.3: Use bi-linear interpolation to interpolate node values



$$\begin{aligned} \text{bilinear_intp}(P) &= (1 - \alpha)(1 - \beta)v_{11} \\ &\quad + \alpha(1 - \beta)v_{21} \\ &\quad + (1 - \alpha)\beta v_{12} + \alpha\beta v_{22} \end{aligned}$$

One missing point here: We need to use $S(\alpha), S(\beta)$ instead of α and β to enable a smoothed interpolation!

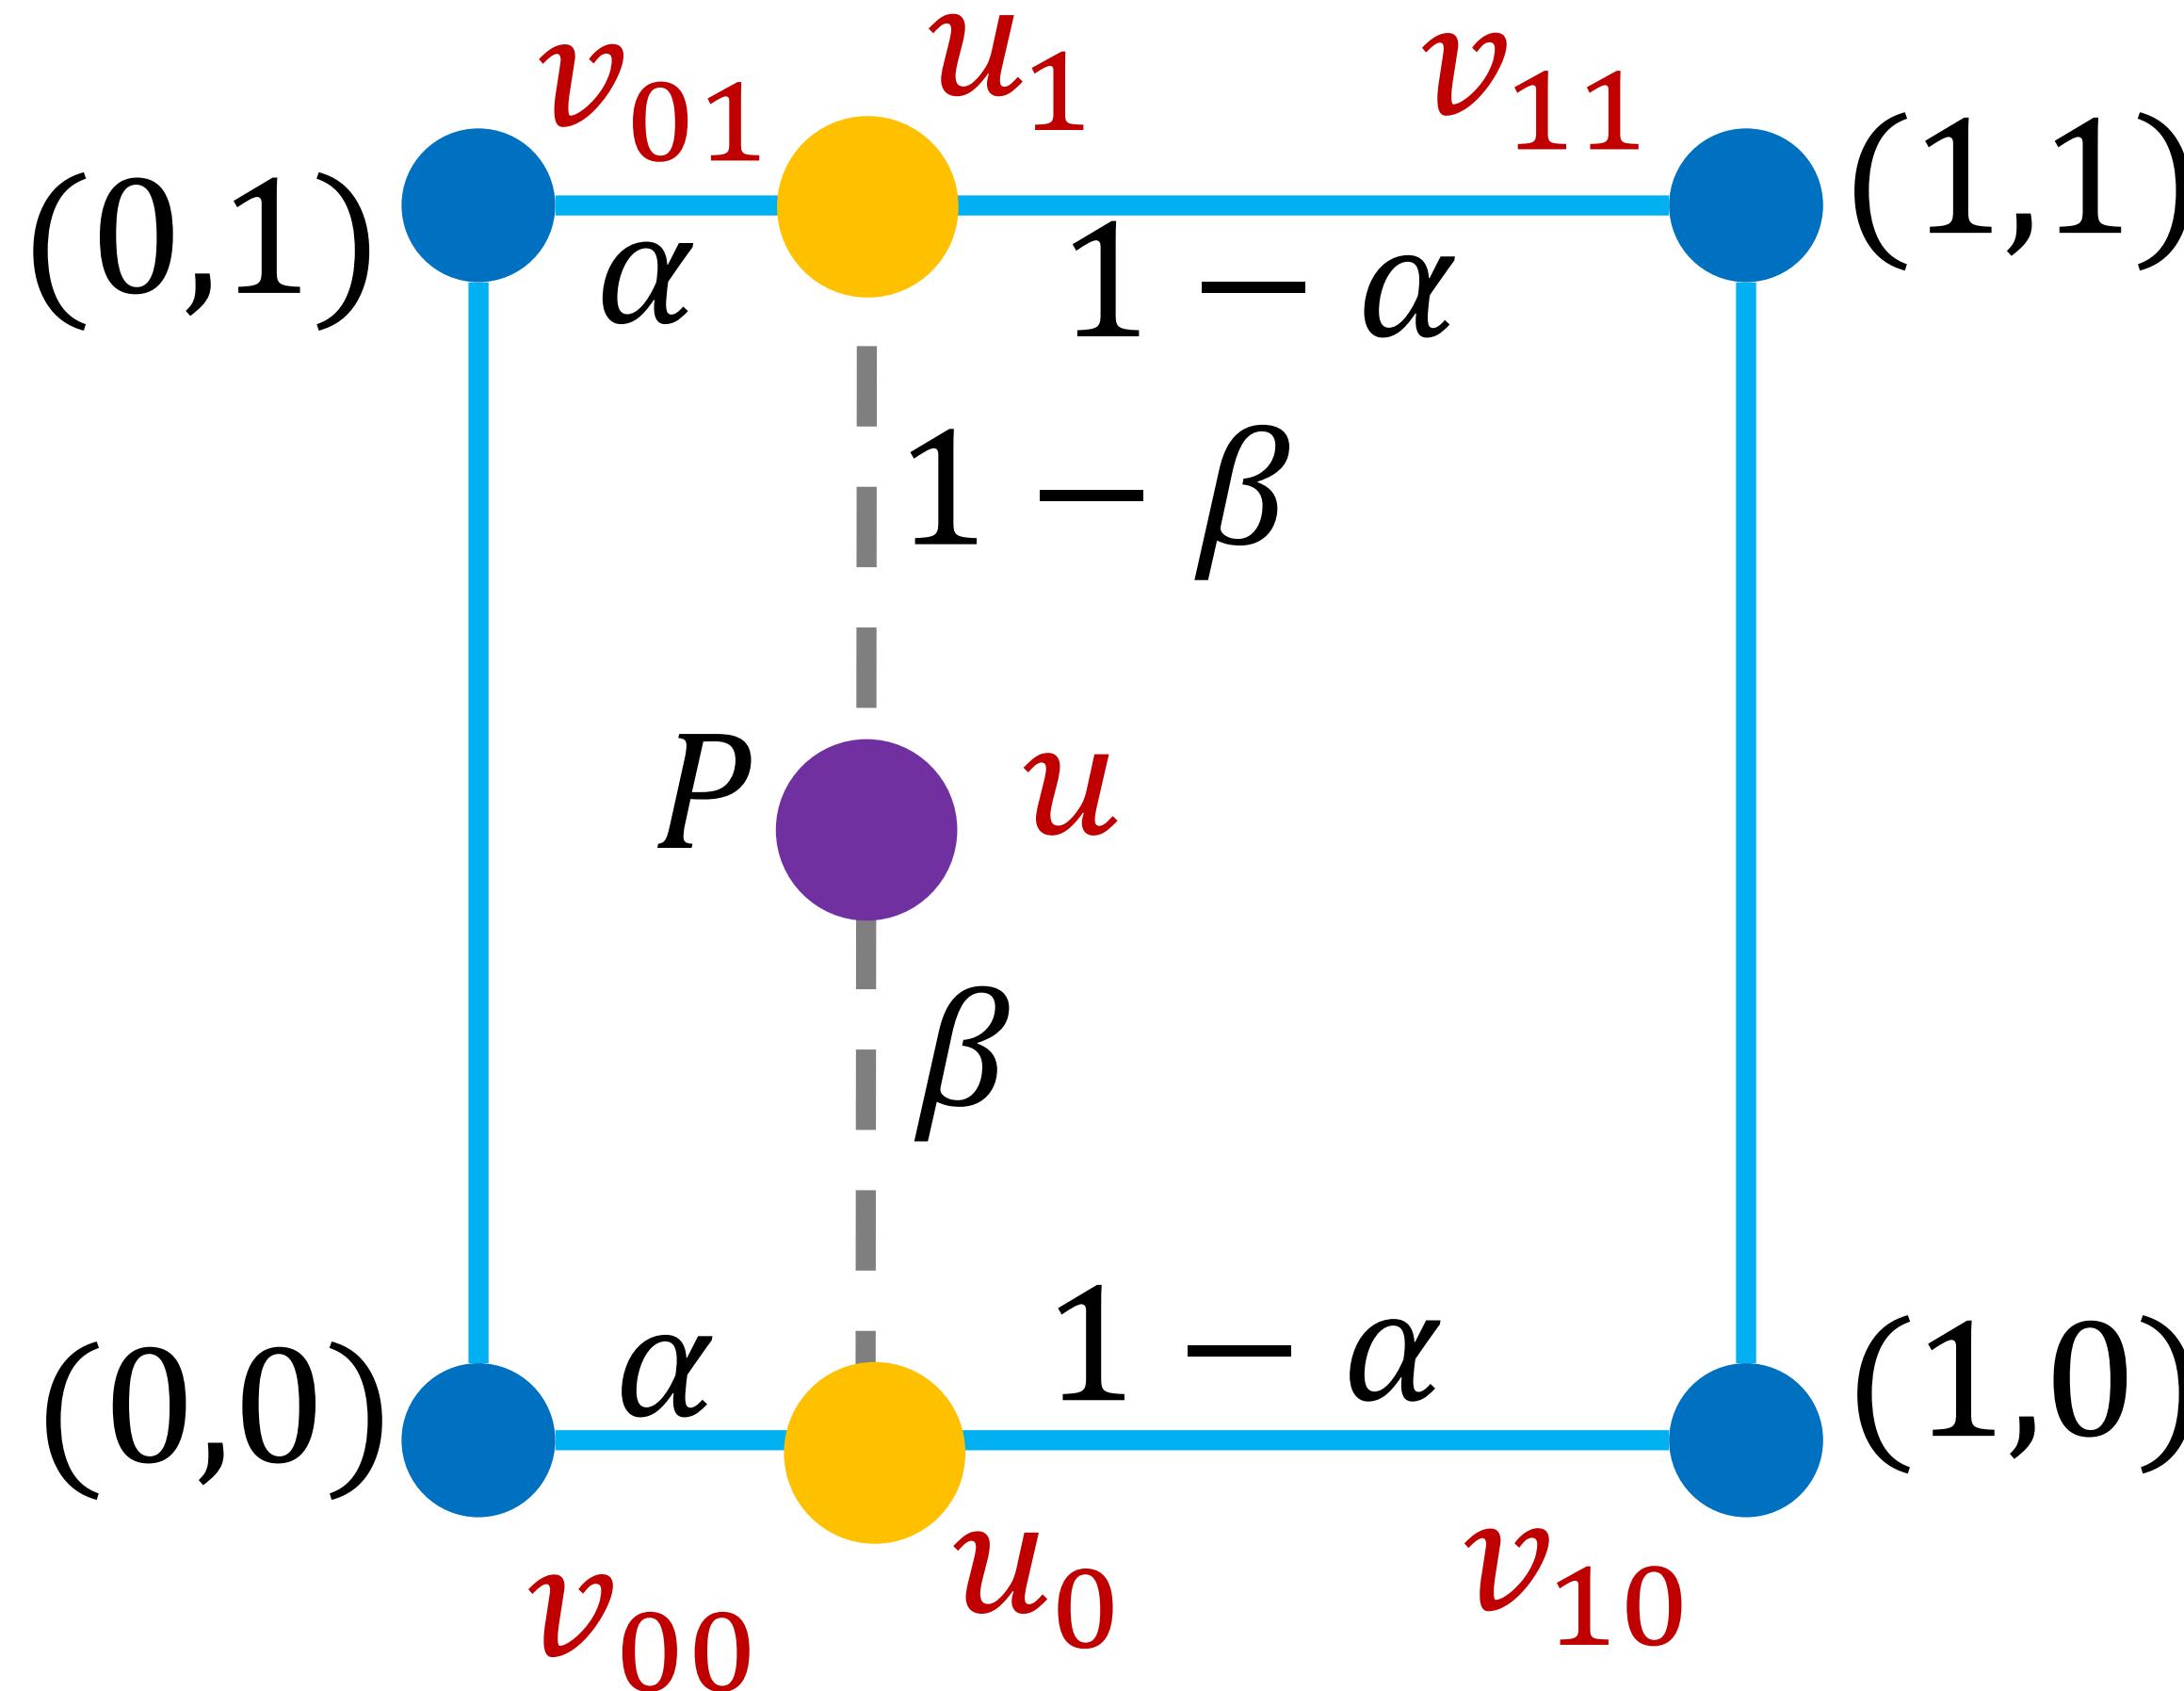
Step III.3: Use bi-smoothed interpolation to calculate interpolation



$$\begin{aligned}
 & \text{smoothed_intp}(P) \\
 &= (1 - S(\alpha))(1 - S(\beta))v_{11} \\
 &+ \alpha(1 - S(\beta))v_{21} \\
 &+ (1 - S(\alpha))\beta v_{12} \\
 &+ S(\alpha)S(\beta)v_{22}
 \end{aligned}$$

with $S(x) = 3x^2 - 2x^3, x \in [0,1]$

For **implementation**, this interpolation can be realized with three linear interpolations



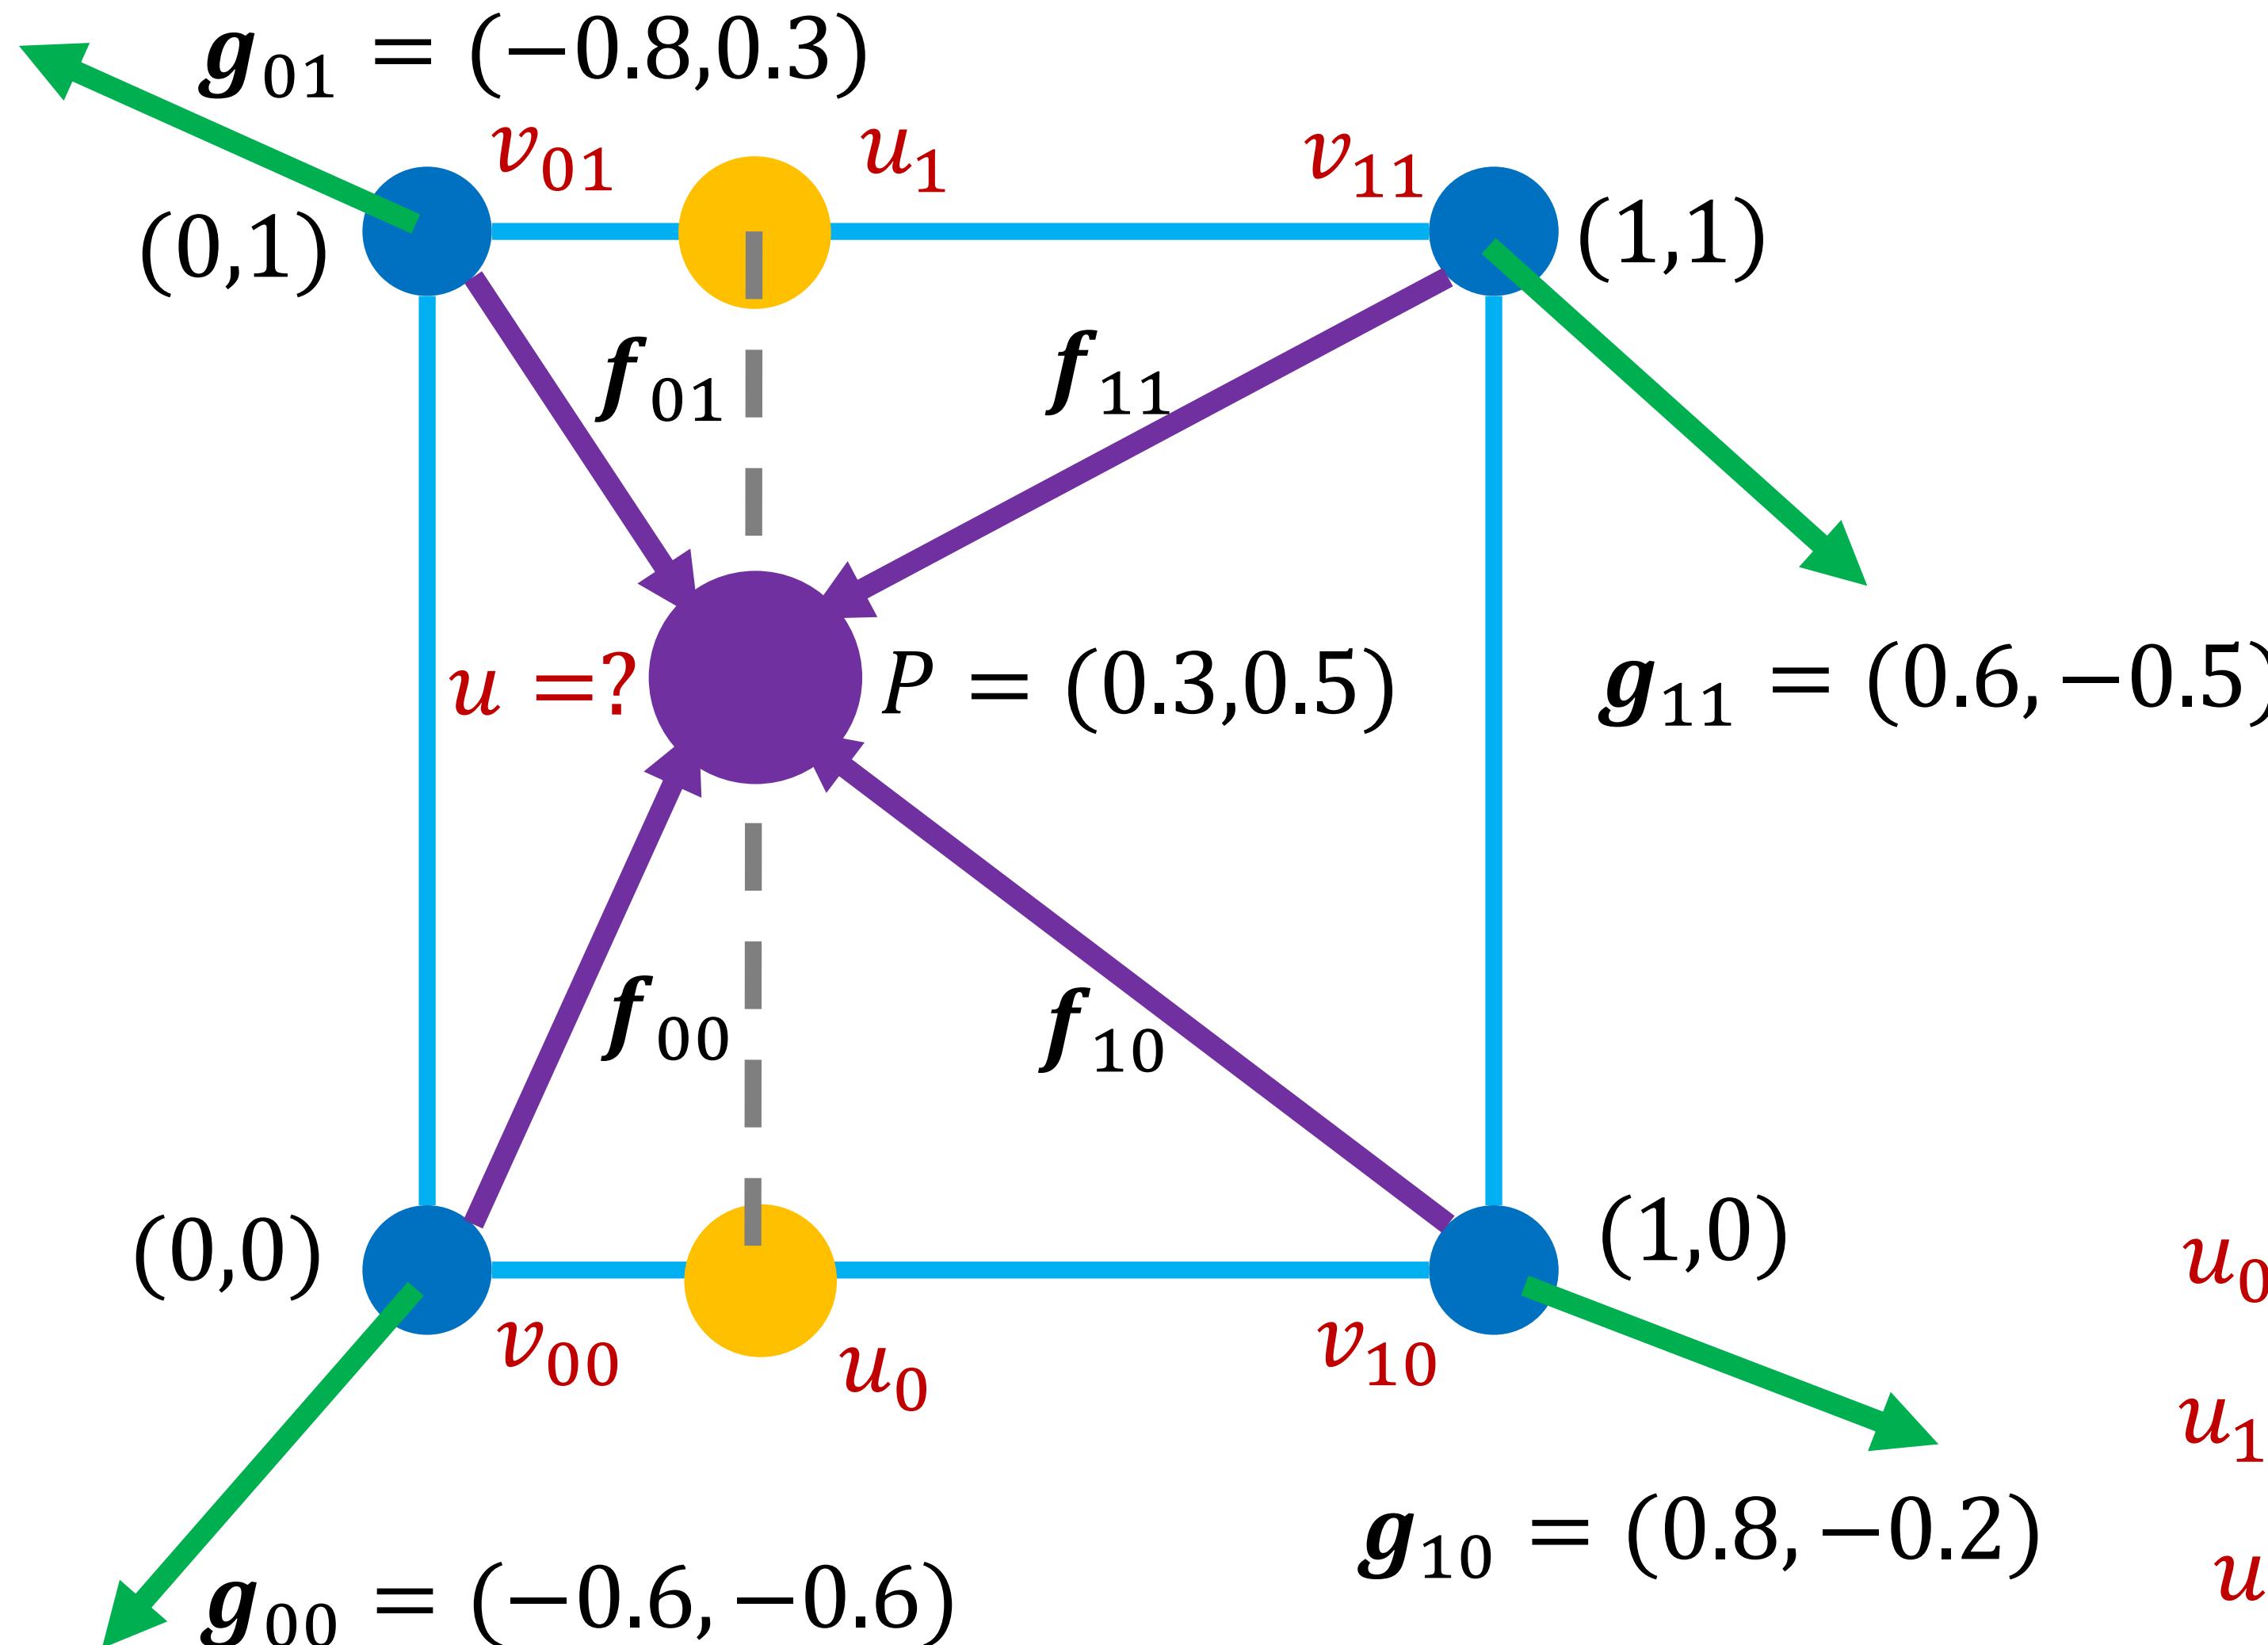
$$u_0 = (1 - S(\alpha))v_{00} + S(\alpha)v_{10}$$
$$u_1 = (1 - S(\alpha))v_{01} + S(\alpha)v_{11}$$
$$u = (1 - S(\beta))u_0 + S(\beta)u_1$$

Or, with GLSL implementation:
 $u0=\text{mix}(v00,v10,S(alpha))$
 $u1=\text{mix}(v01,v11,S(alpha))$
 $u=\text{mix}(u0,u1,S(beta))$

Or, even simpler:
 $u=\text{mix}(\text{mix}(v00,v10,S(alpha)), \text{mix}(v01,v11,S(alpha)), S(beta))$



A Numerical Example: Calculate Perlin Noise for Point P



$$\mathbf{f}_{00} = \mathbf{P} - (0,0) = (0.3, 0.5)$$

$$\mathbf{f}_{10} = \mathbf{P} - (1,0) = (-0.7, 0.5)$$

$$\mathbf{f}_{01} = \mathbf{P} - (0,1) = (0.3, -0.5)$$

$$\mathbf{f}_{11} = \mathbf{P} - (1,1) = (-0.7, -0.5)$$

$$v_{00} = \mathbf{f}_{00} \cdot \mathbf{g}_{00} = -0.48$$

$$v_{10} = \mathbf{f}_{10} \cdot \mathbf{g}_{10} = -0.66$$

$$v_{01} = \mathbf{f}_{01} \cdot \mathbf{g}_{01} = -0.39$$

$$v_{11} = \mathbf{f}_{11} \cdot \mathbf{g}_{11} = -0.17$$

$$u_0 = (1 - S(0.3))v_{00} + S(0.3)v_{10} = -0.52$$

$$u_1 = (1 - S(0.3))v_{01} + S(0.3)v_{11} = -0.34$$

$$u = (1 - S(0.5))u_0 + S(0.5)u_1 = -0.43$$



Pseudocode for 2D Perlin Noise with a Single Octave

Input: x

Output: noise

```
vec2 i = floor(x);
vec2 f = fract(x);
vec2 s = f*f*(3.0-2.0*f);
vec2 g00 = hash2(i);
vec2 g10 = hash2(i+vec2(1,0));
vec2 g01 = hash2(i+vec2(0,1));
vec2 g11 = hash2(i+vec2(1,1));
vec2 f00 = f-vec2(0,0);
vec2 f10 = f-vec2(1,0);
vec2 f01 = f-vec2(0,1);
vec2 f11 = f-vec2(1,1);
float noise=(calculate with g00,g10,g01,g11, f00, f10, f01, f11, and s)
return noise;
```

This step can be implemented with
three **mix()** function calls



Step IV: Octave Synthesis

Input: x, n */// x – position, n – number of octaves*

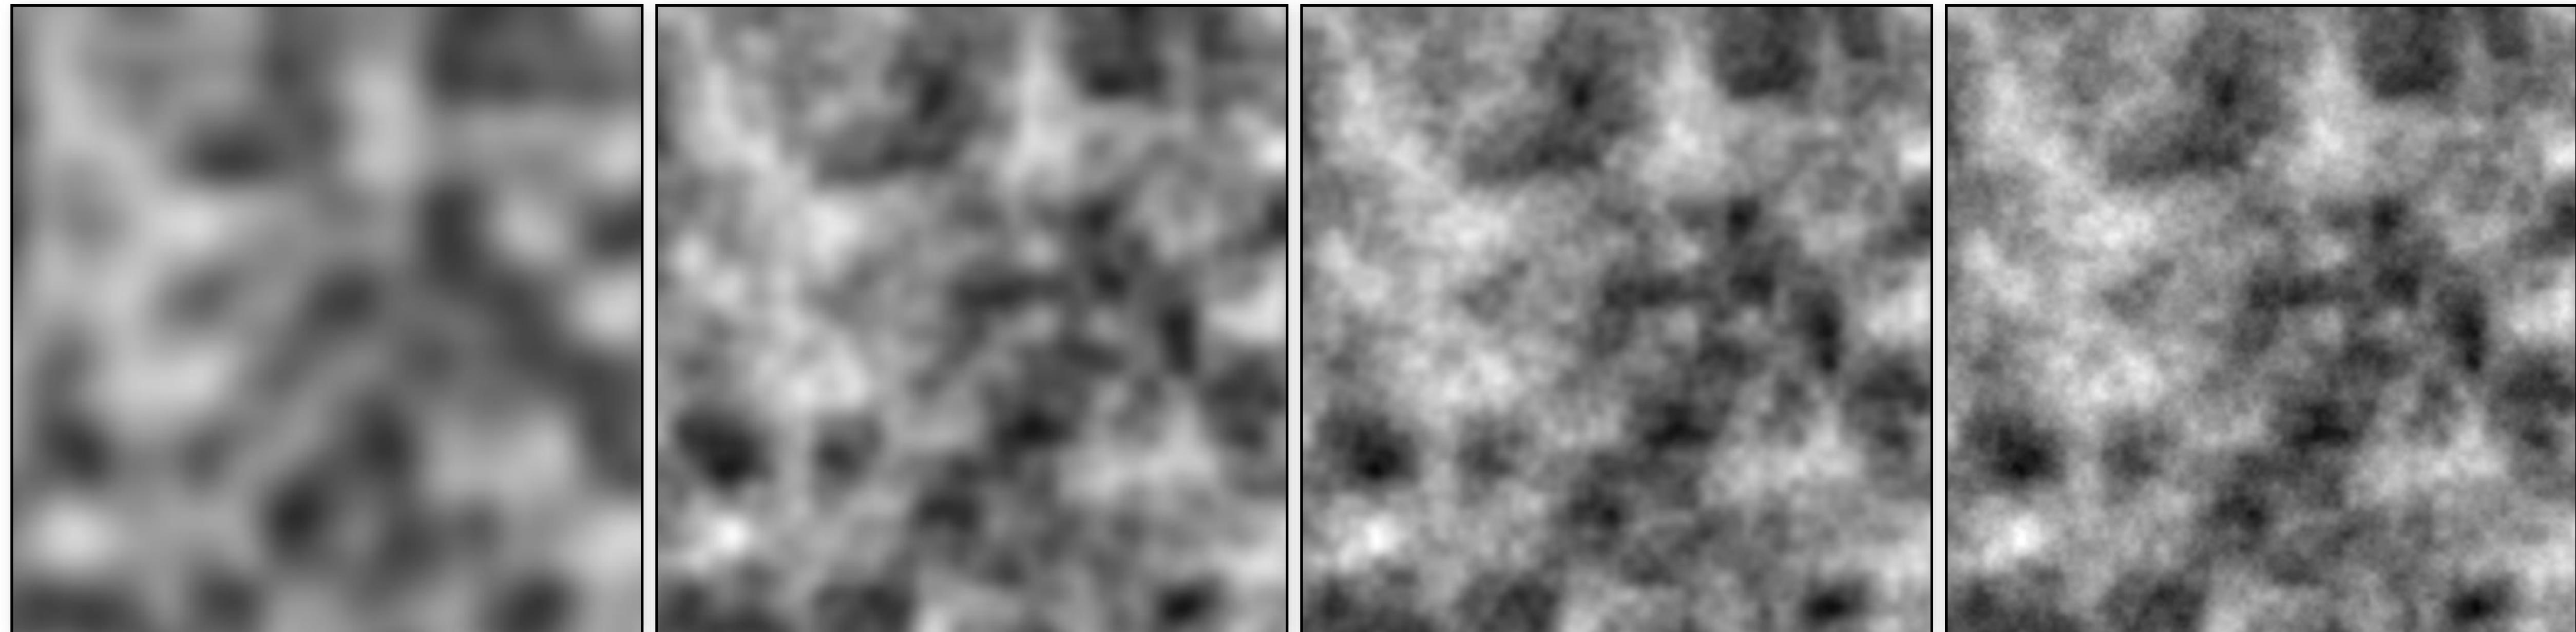
Output: noise

```
float noise=0.0;  
for i = 0 to n-1:  
    noise += 2^-i * perlin(2^i * x);  
return noise;
```

Here **perlin()** is the 2D Perlin noise function implemented previously



Octave Synthesis on a Texture Image



1 octaves

2 octaves

3 octaves

4 octaves

This technique is also called "fractal Brownian Motion" (fBM), or simply "fractal noise"

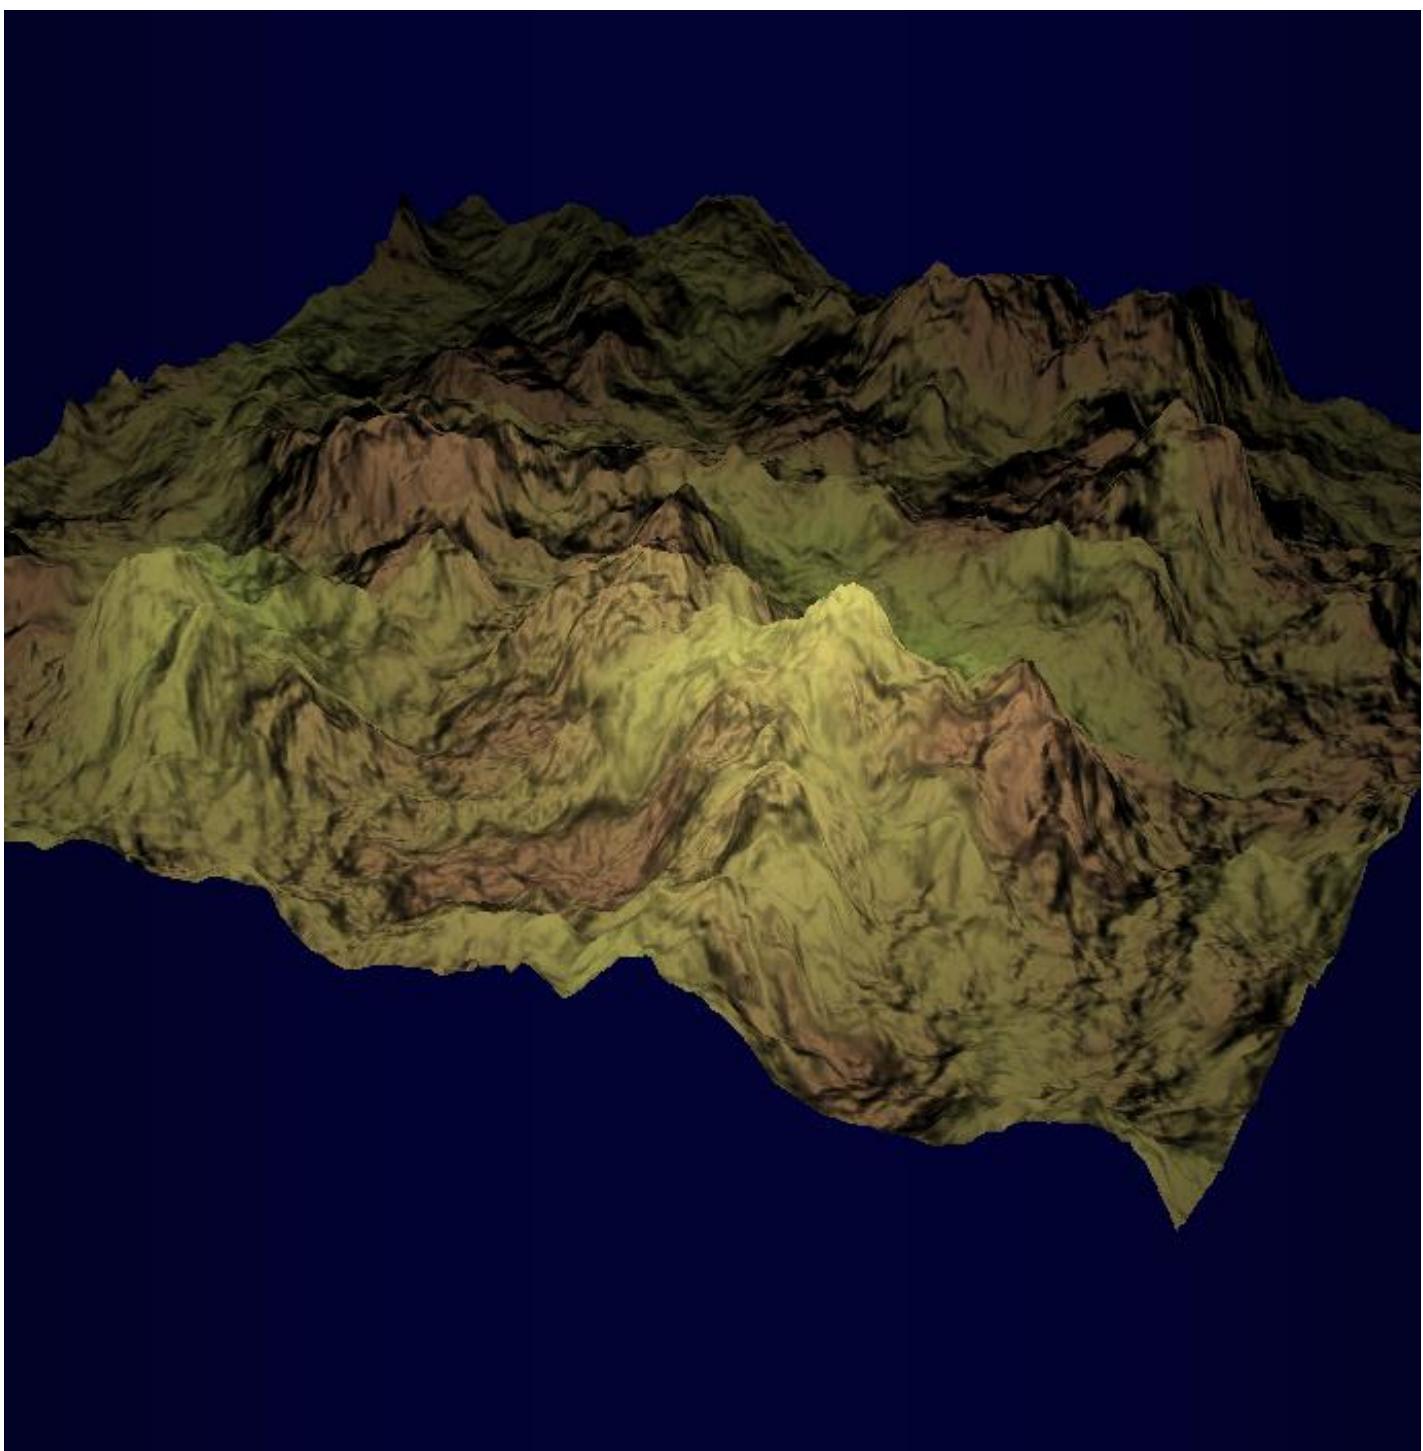


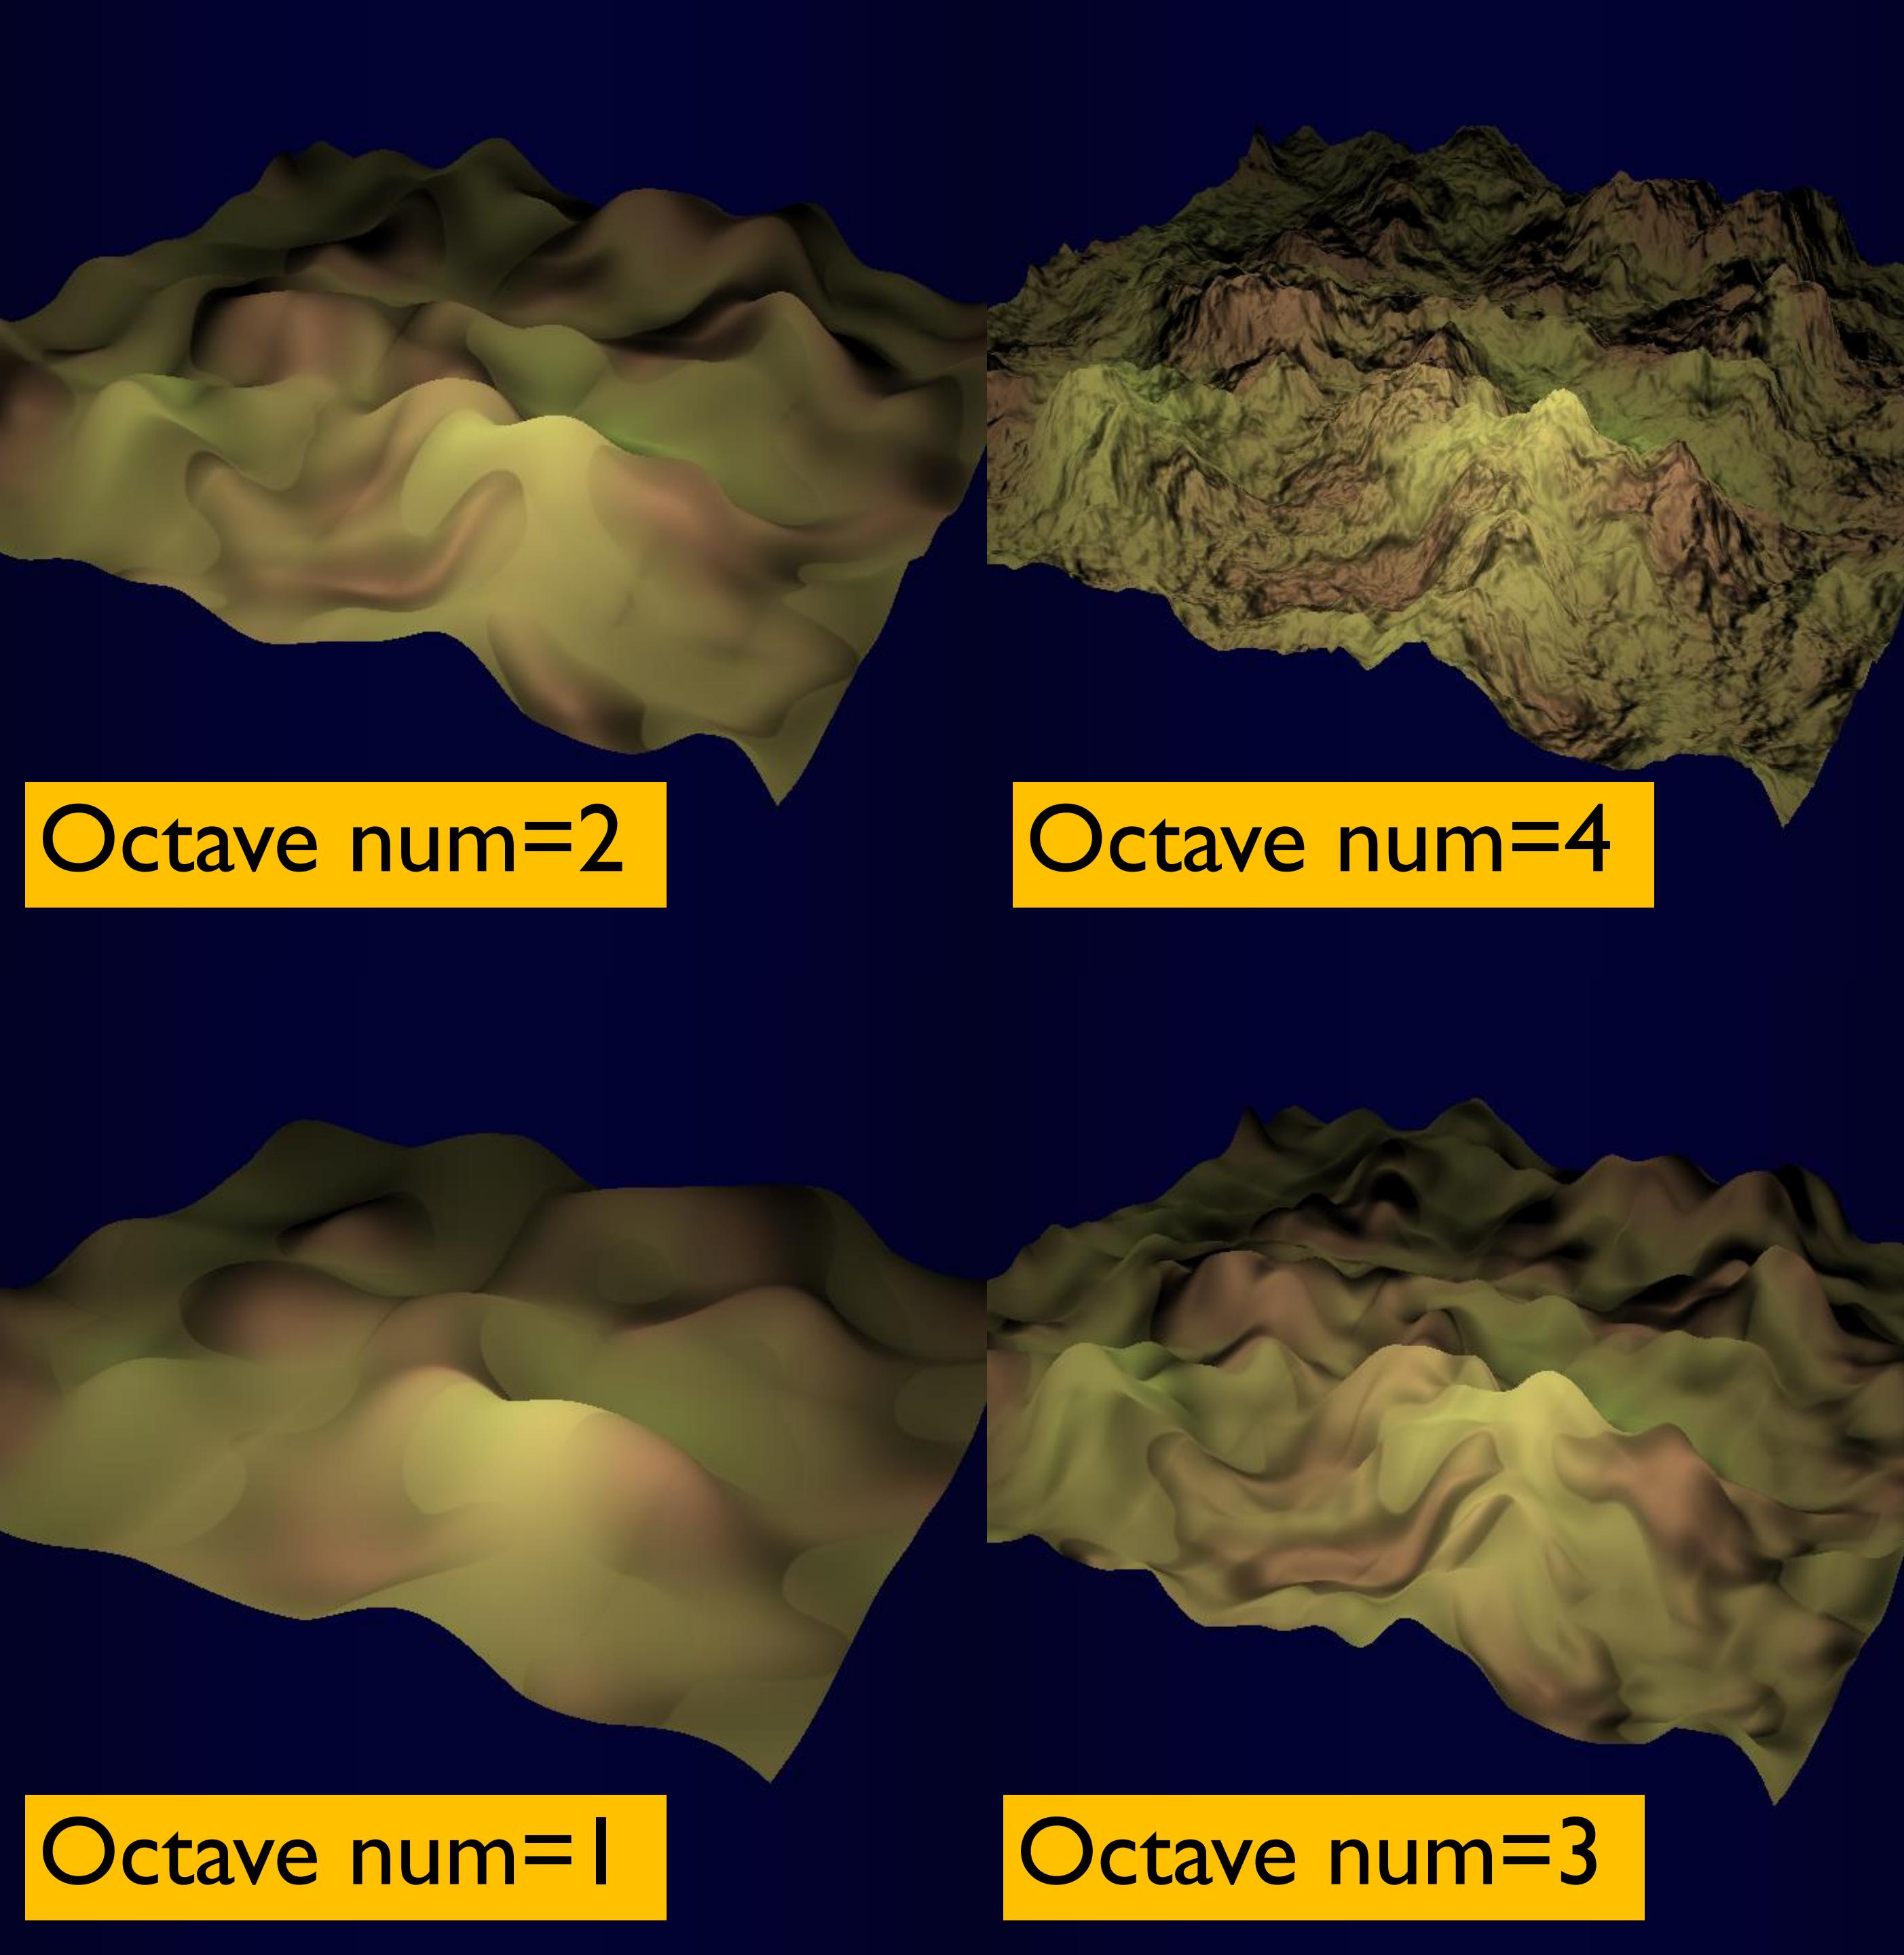
Key Takeaways for Perlin Noise

- 1. Grid Discretization:** Generate a random gradient vector on each grid node using Hash() function
- 2. Random Gradients:** For each grid node, dot product between gradient vectors and offset vectors
- 3. Smooth Interpolation:** For an arbitrary point with a grid cell, interpolate the dot products on the four grid nodes using smooth step function and return it as the noise value
- 4. Octave Synthesis:** Synthesize the noise values on different octaves



Terrain Generation

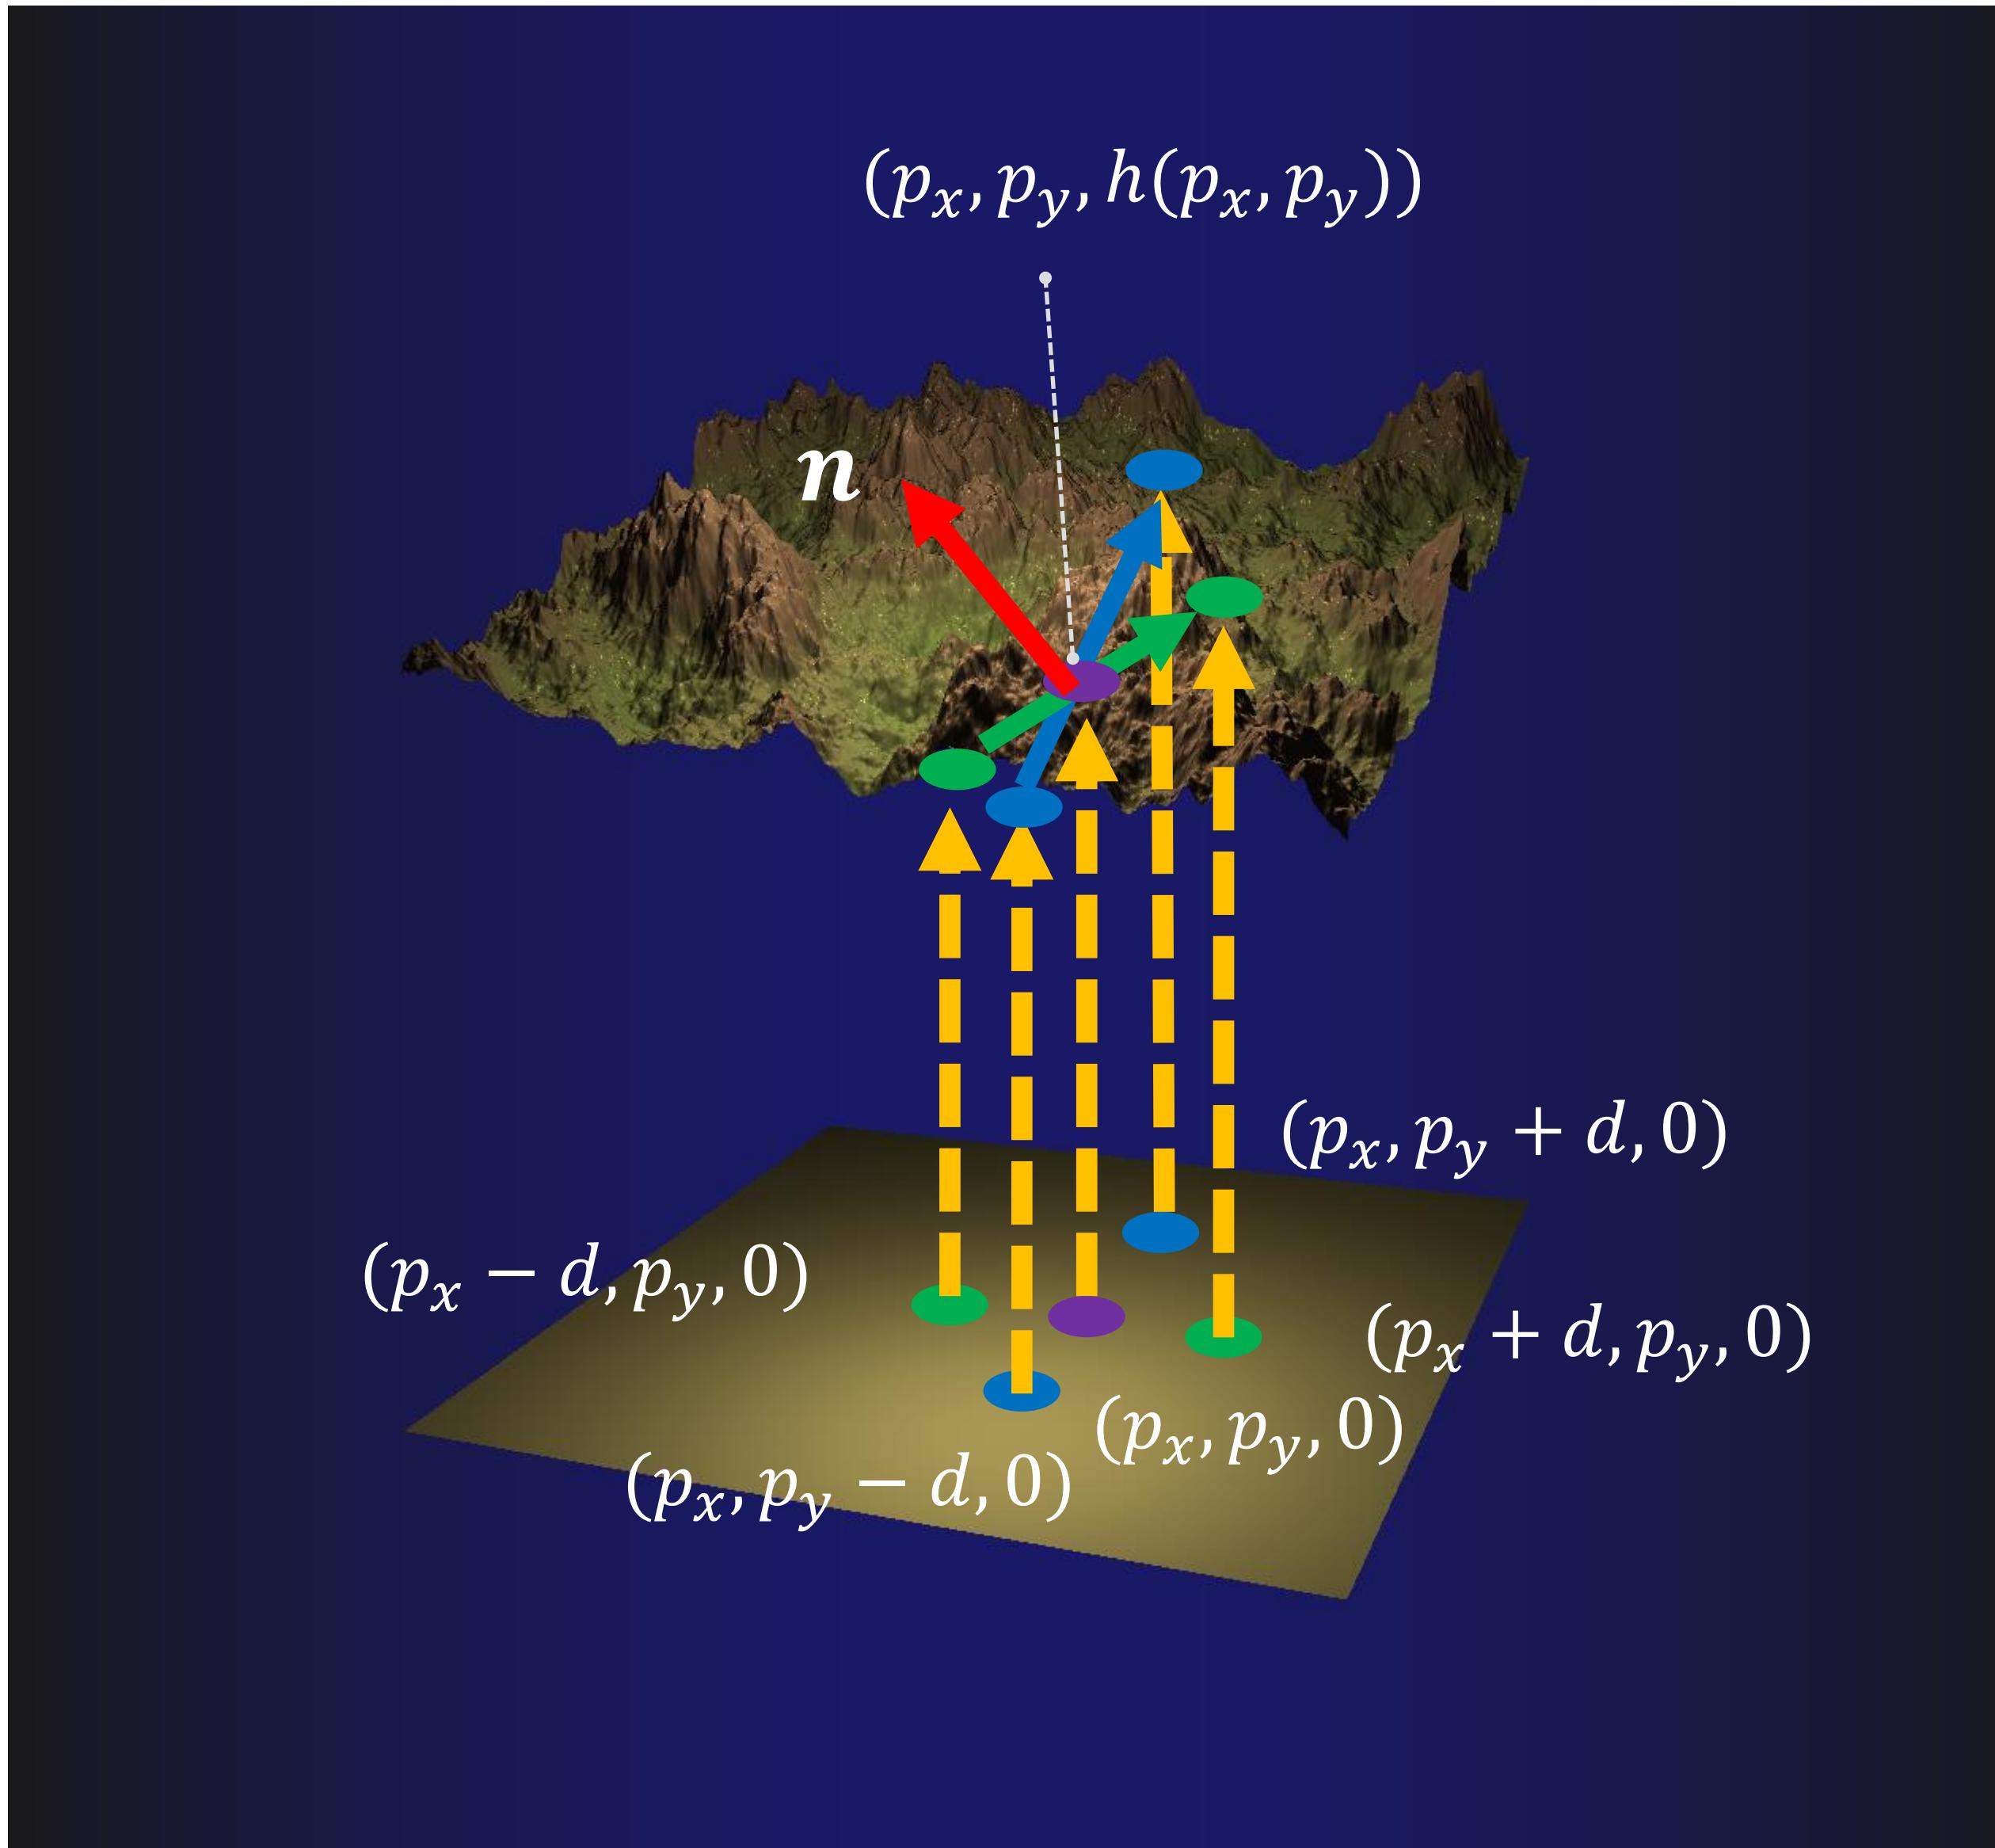




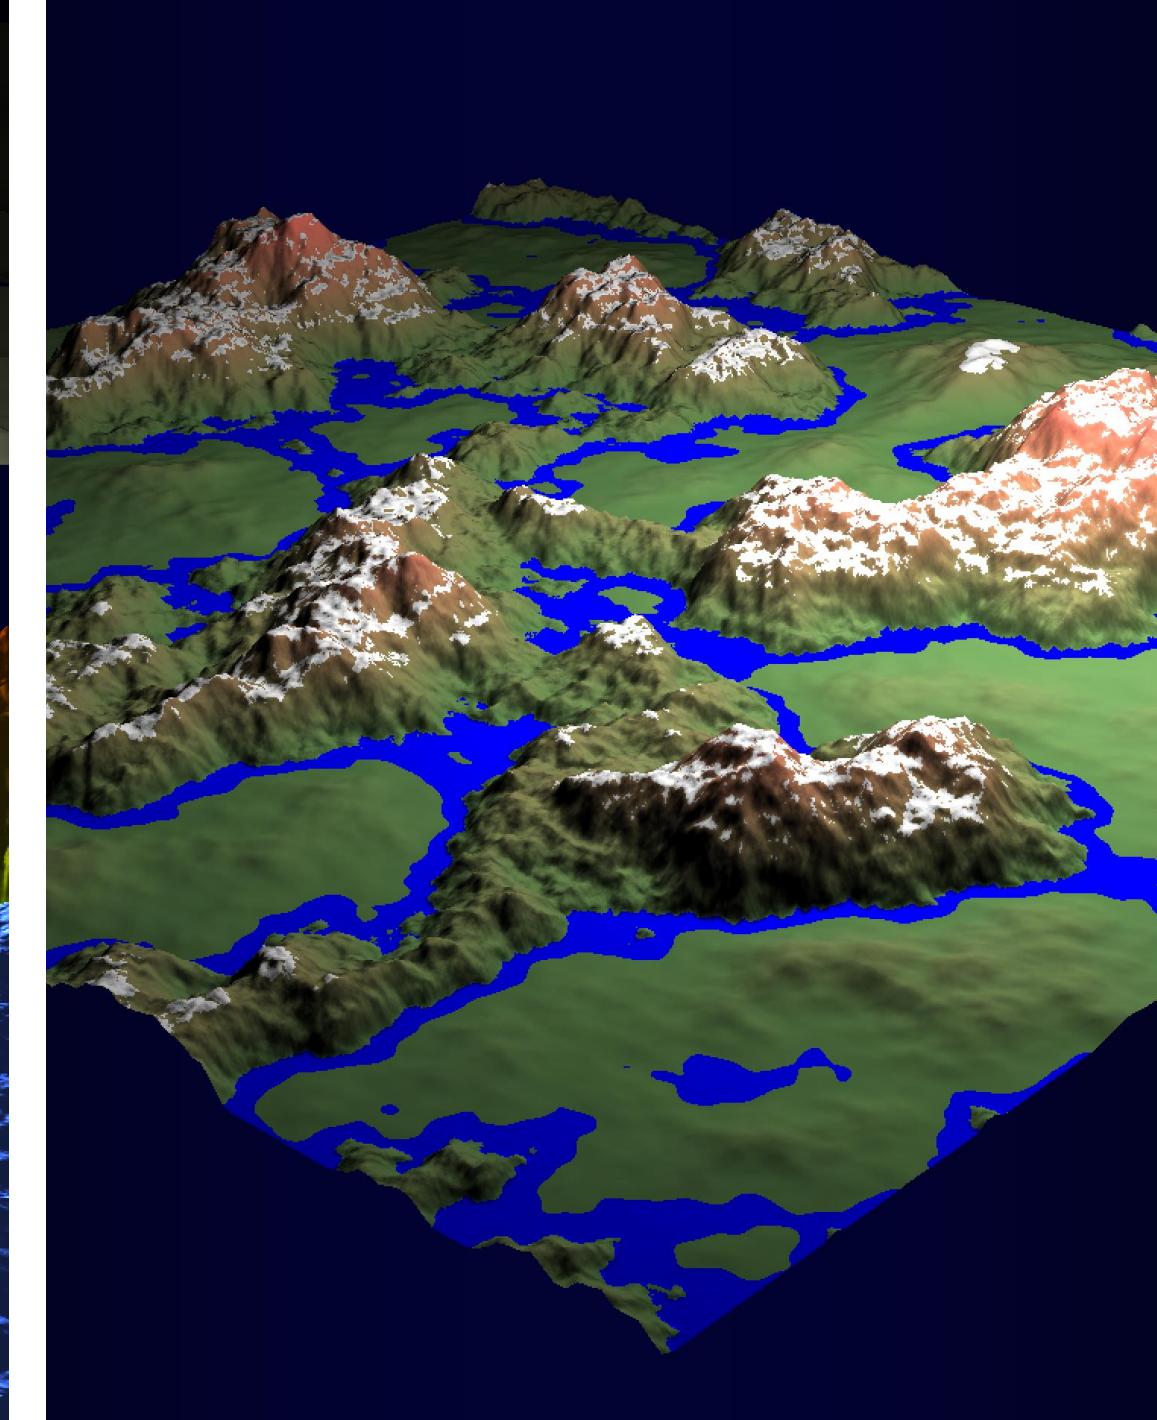
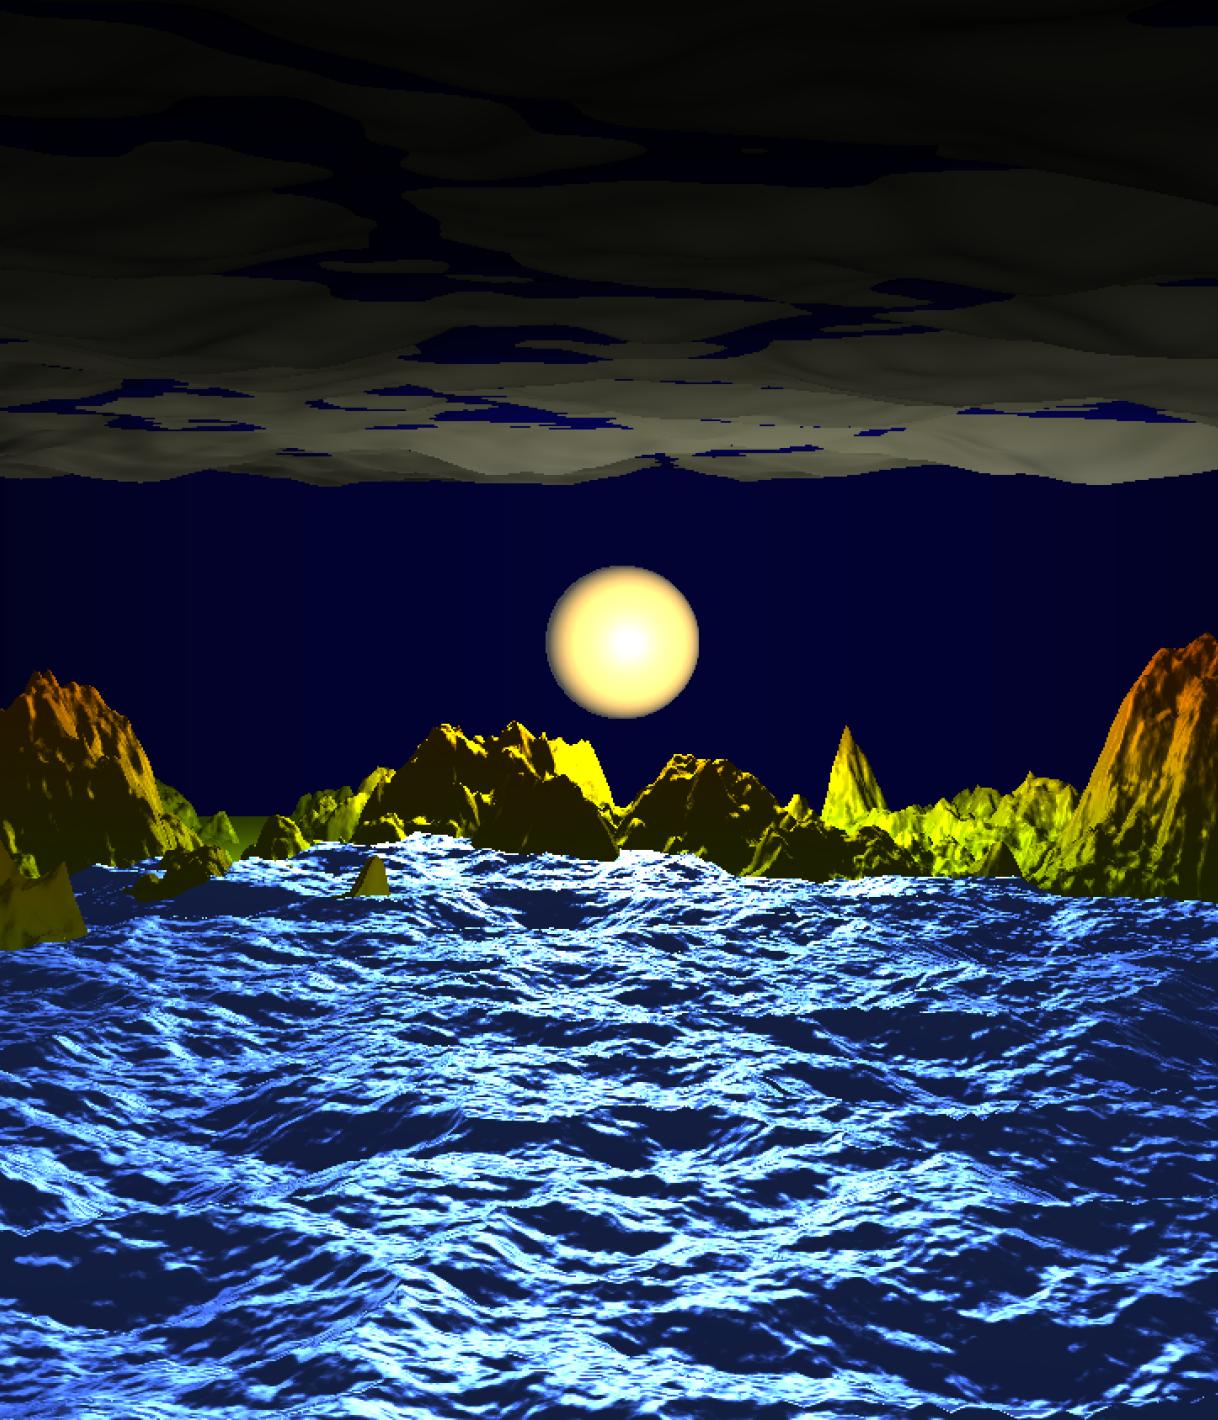
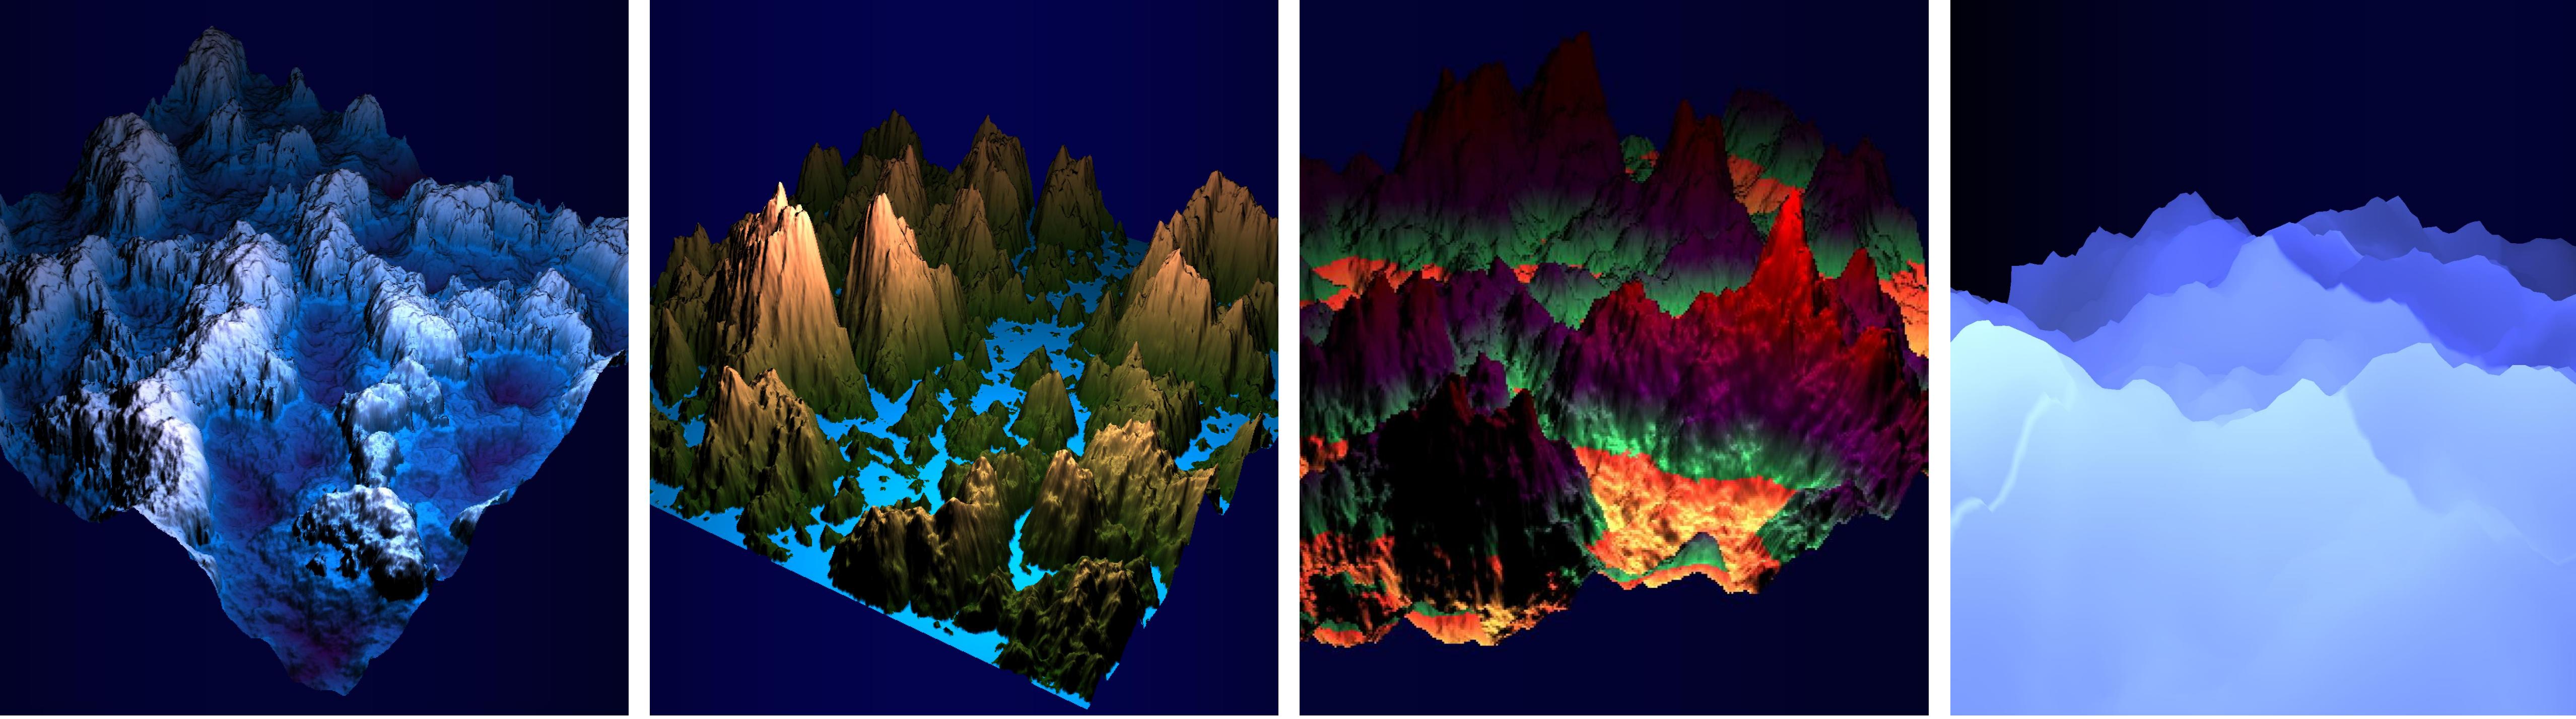
Use Perlin Noise to Generate Terrain in GLSL shaders

- We can use Perlin noise to in both vertex shader and fragment shader:
- In vertex shader, we use it to calculate vertex height
- In fragment shader, we use it to calculate fragment color

Calculate Terrain's Normals and Colors



- We start from a flat triangle mesh with zero height for each vertex
- [Vertex shader – Height]: For each vertex of the mesh, we calculate its **height** using a shader function, and update vertex position
- [Fragment shader – Normal]: We calculate the **normal** of a point by looking at the height of its neighboring points and calculate its normal (red) as the **cross product** of the two axis vectors on the mesh surface (blue and green)
- [Fragment shader – Emission]: We calculate an **emissive color** of each fragment based on its height (we use this color for the color of snow, trees, water, and soil!)
- The final color of a fragment is calculated as the multiplication of the lighting color and the emissive color



Mountains and Rivers
Depict the grandeur in your heart
with Perlin noise functions

Additional Reading Materials for Perlin Noise

Ken Perlin, Making Noise:

<https://web.archive.org/web/20071011035810/http://noisemachine.com/talk1/>

The book of shaders, noise: <https://thebookofshaders.com/11/>

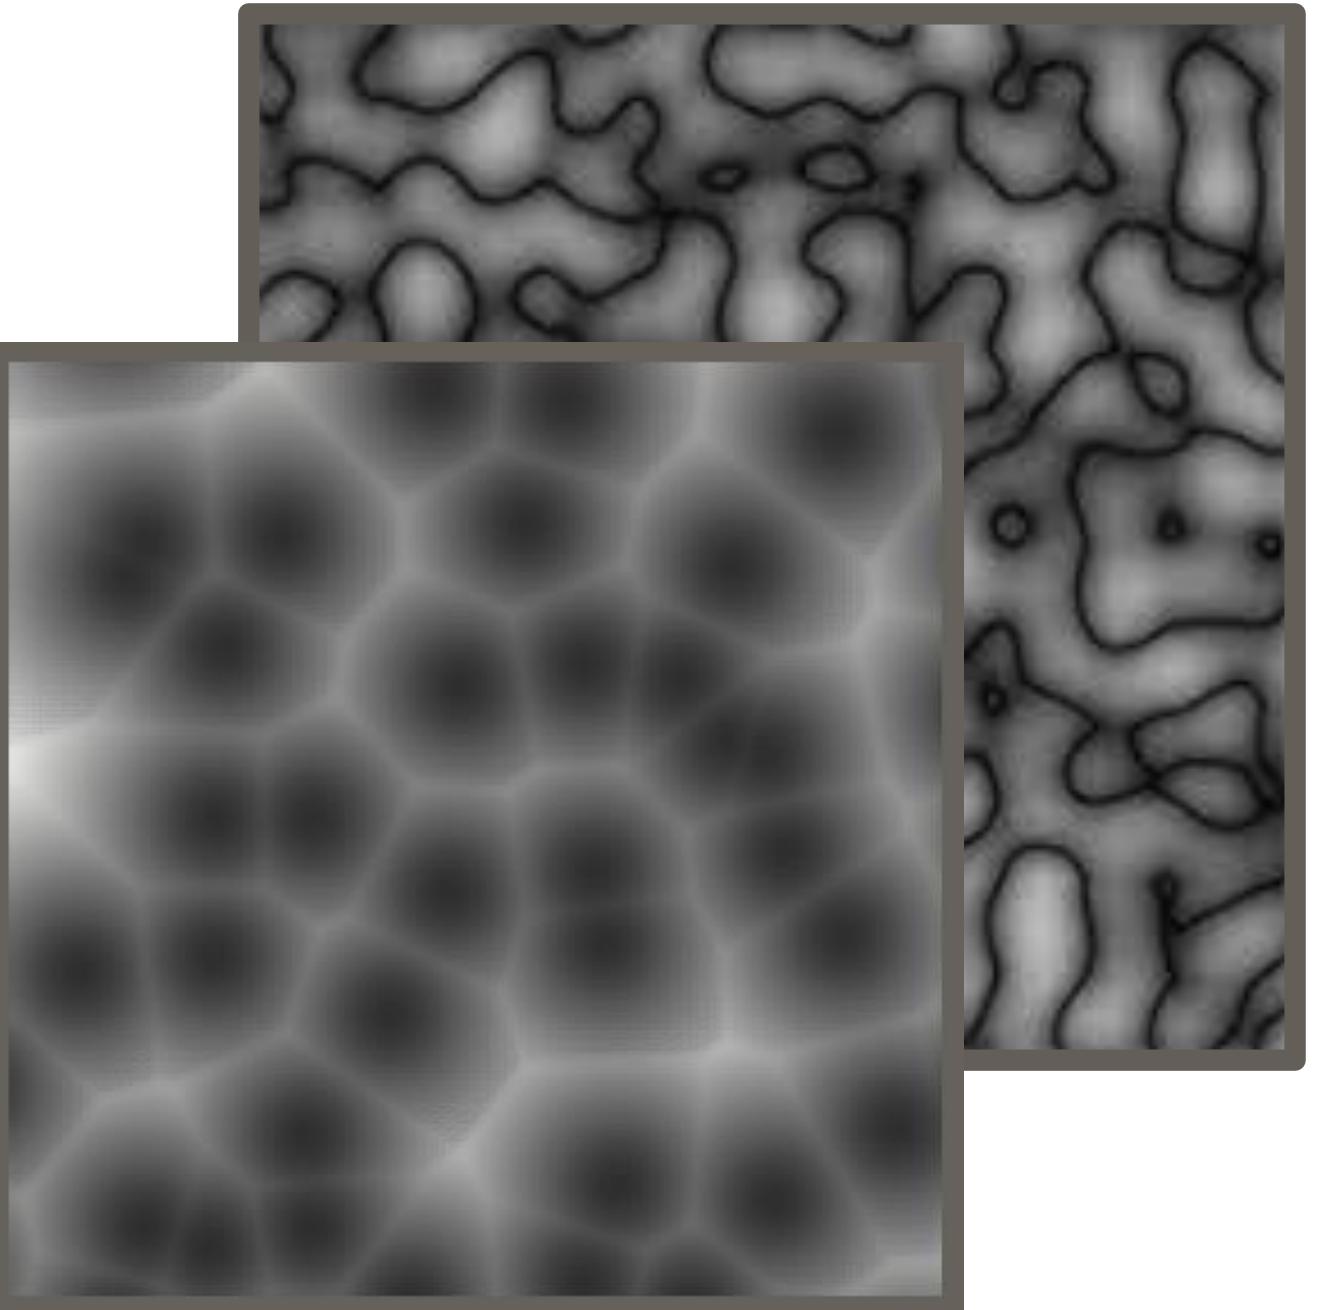
Fragment shader code: <https://thebookofshaders.com/edit.php#11/2d-gnoise.frag>

The Perlin noise math FAQ:

<https://mzucker.github.io/html/perlin-noise-math-faq.html>

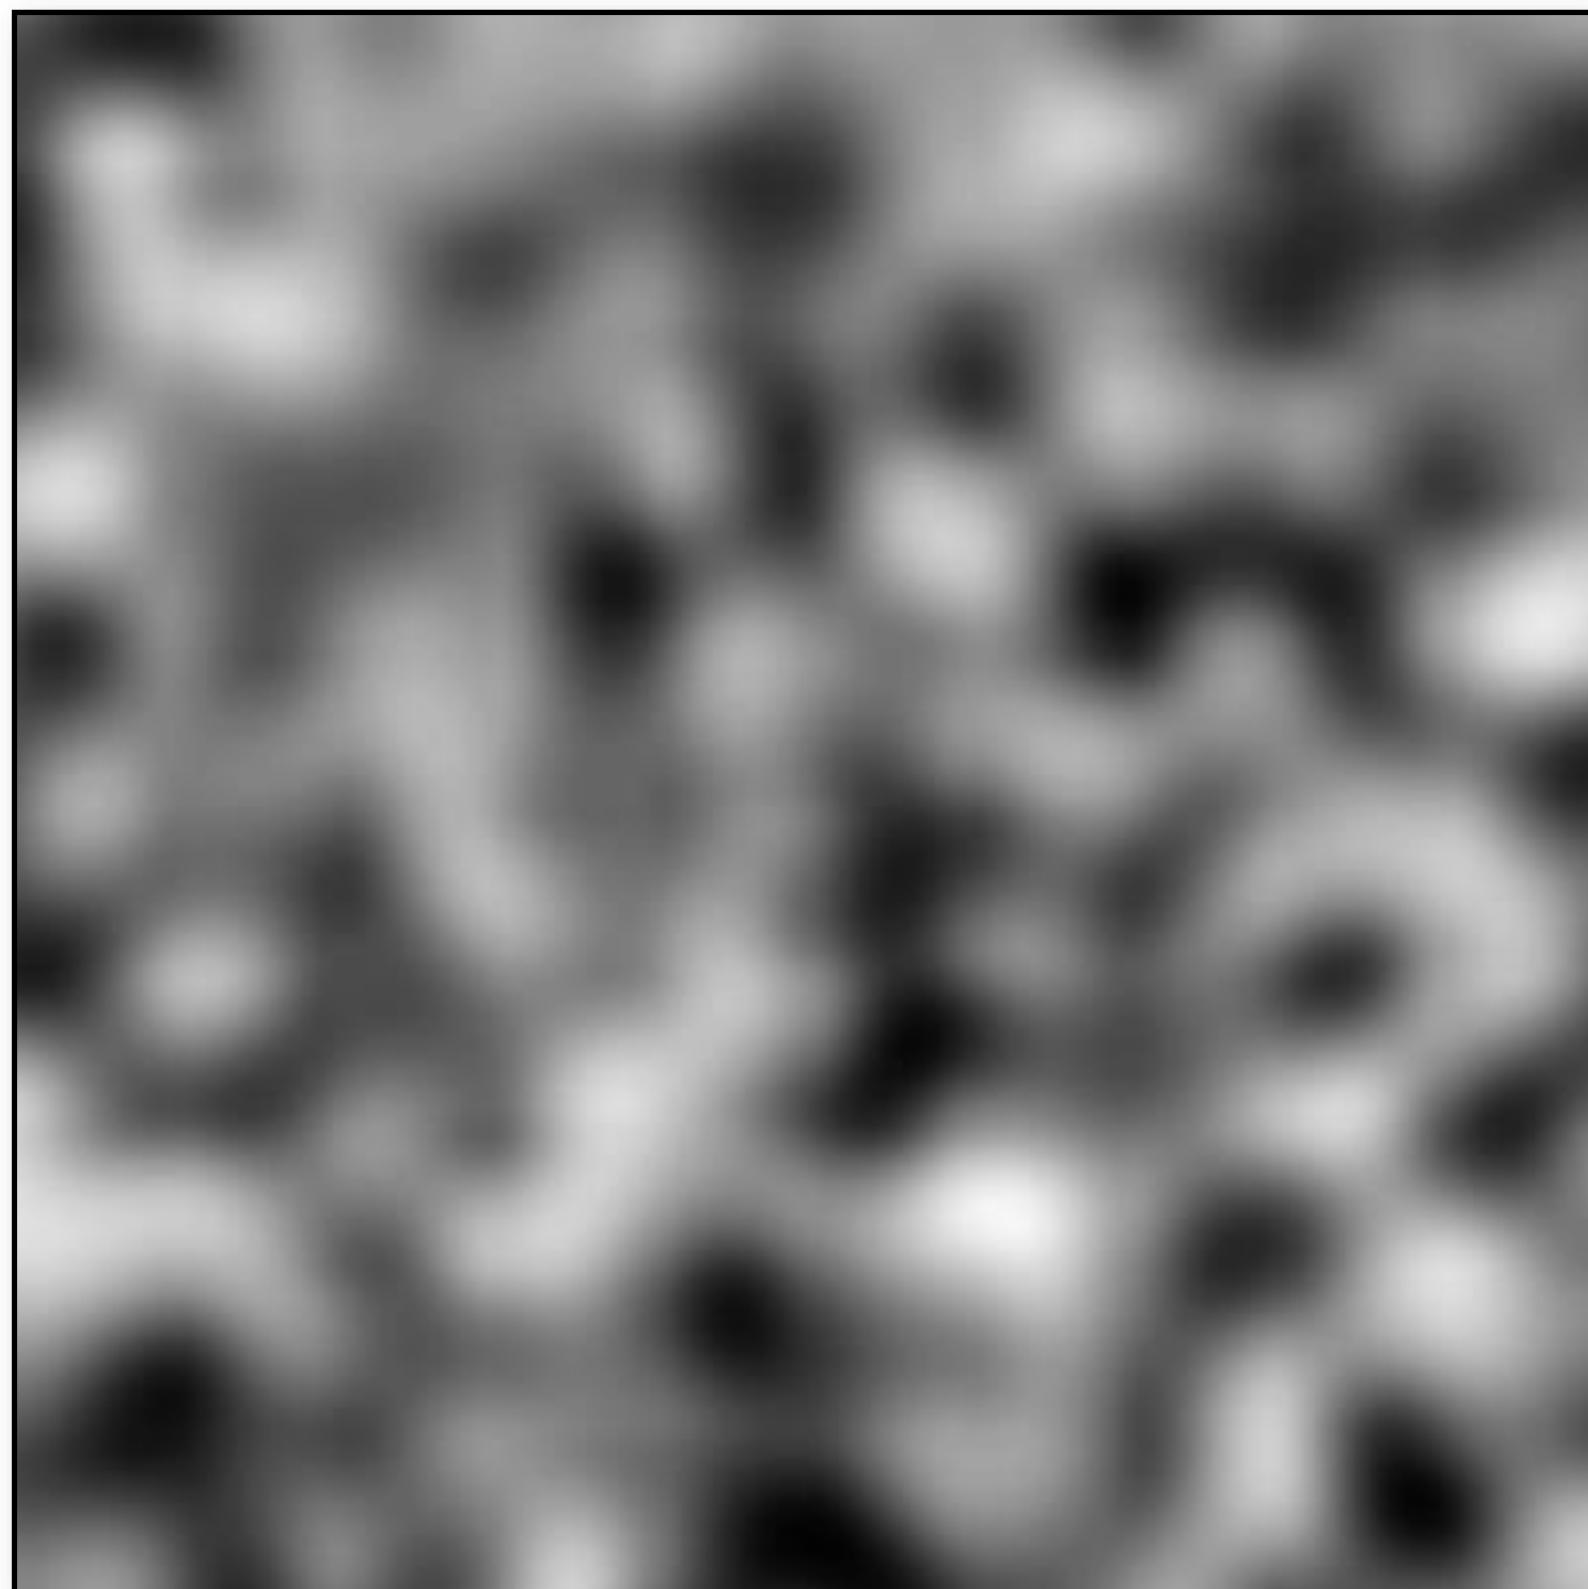


Other Noise Functions

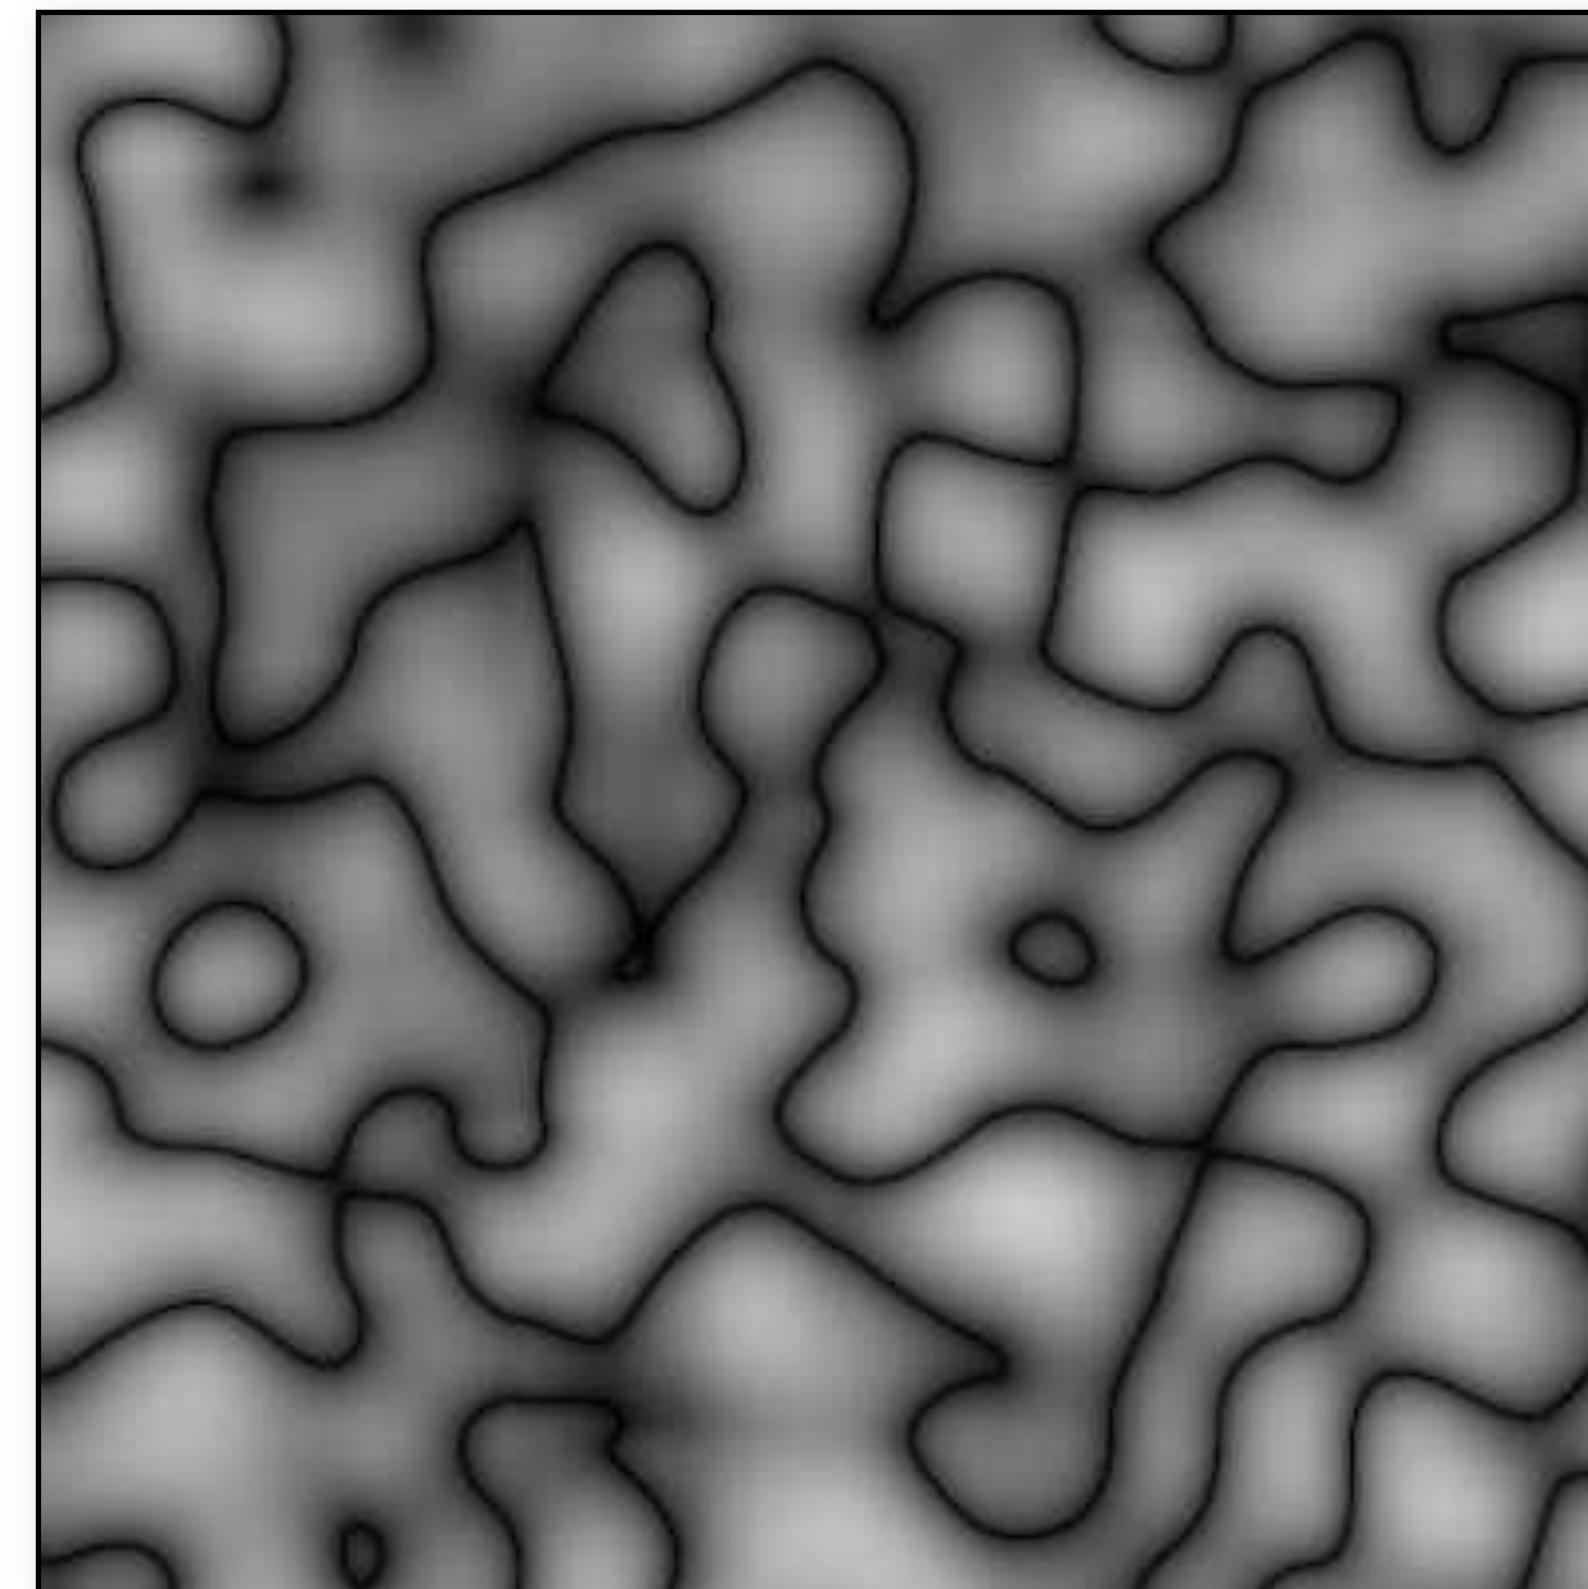


Perlin Noise with Absolute Value

- Typically signed by default, ~in $[-1, 1]$ with a mean of 0
offset/scale to put into $[0,1]$ range take absolute value



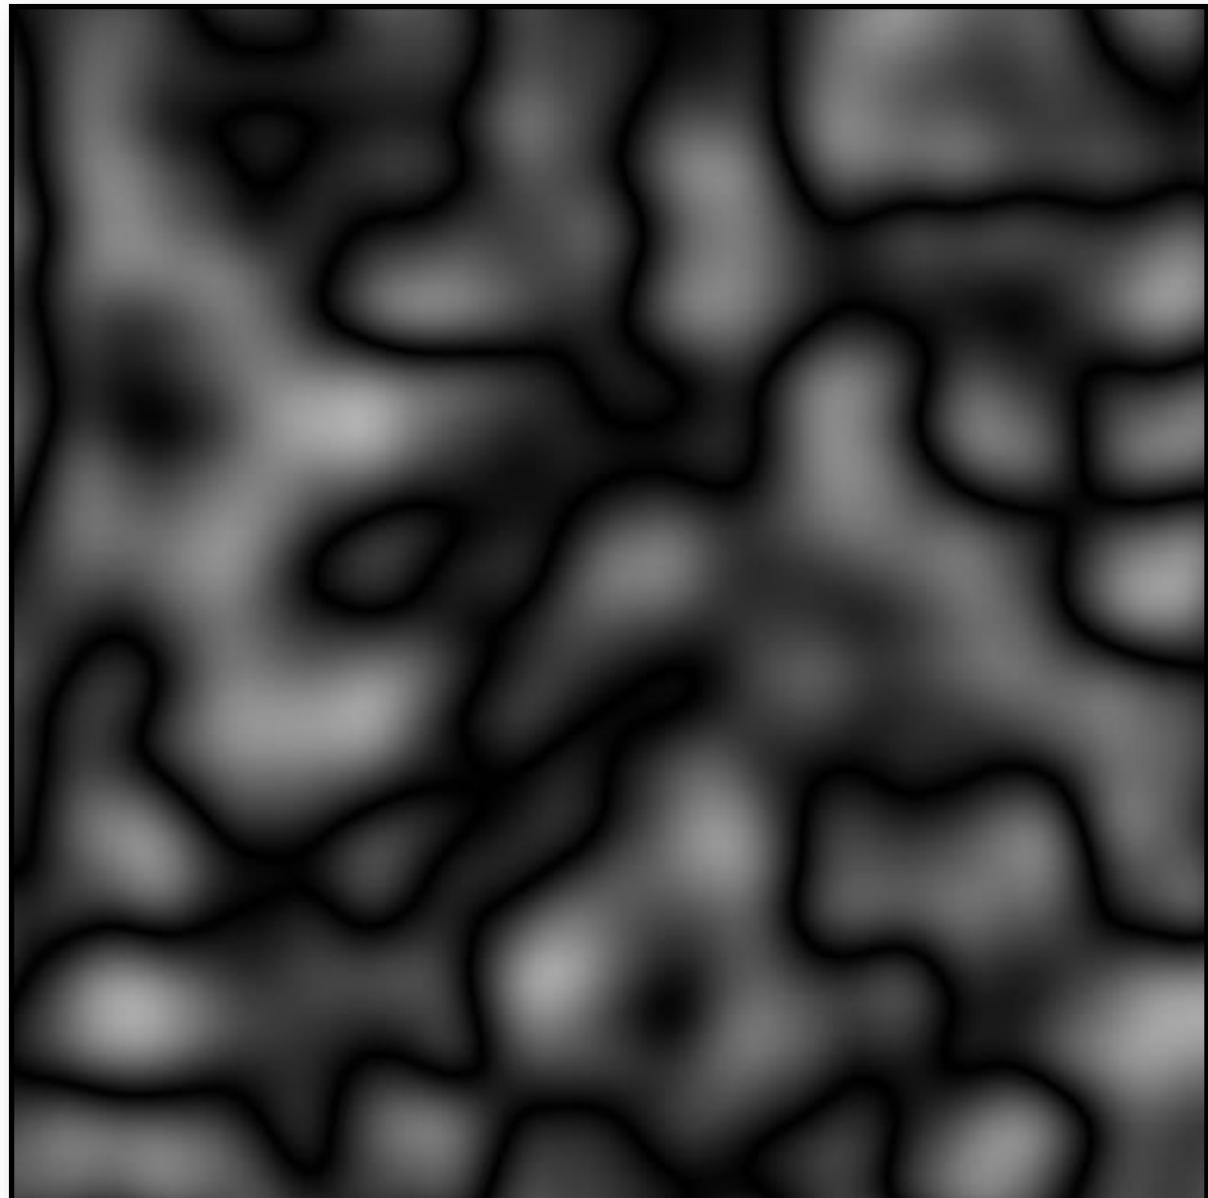
$(\text{noise}(\mathbf{p})+1)/2$



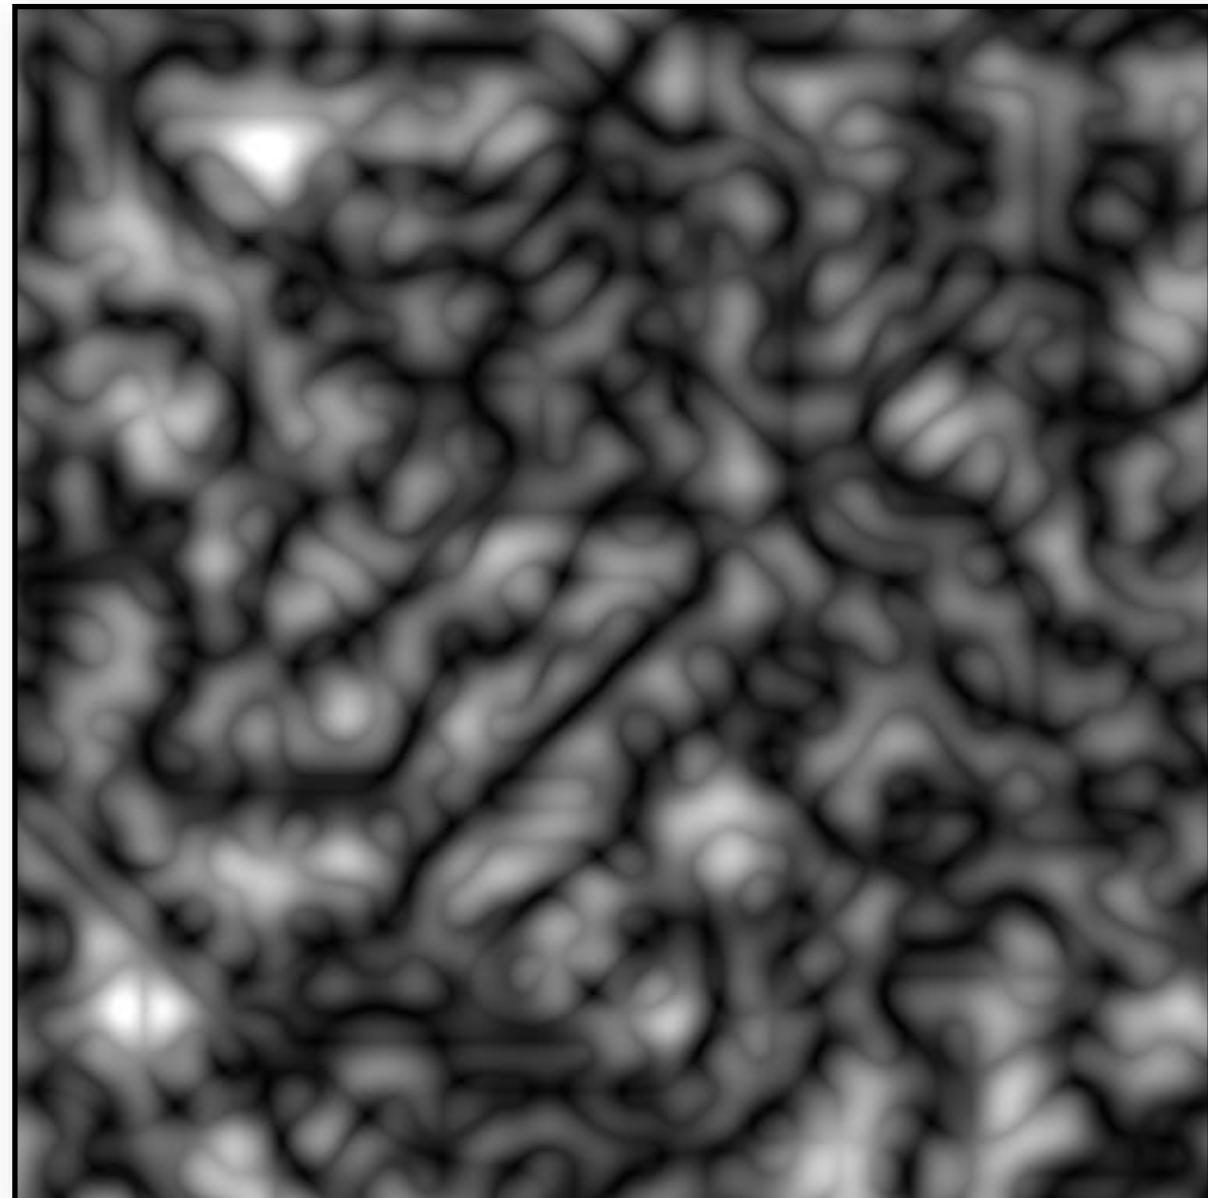
$|\text{noise}(\mathbf{p})|$

Turbulence

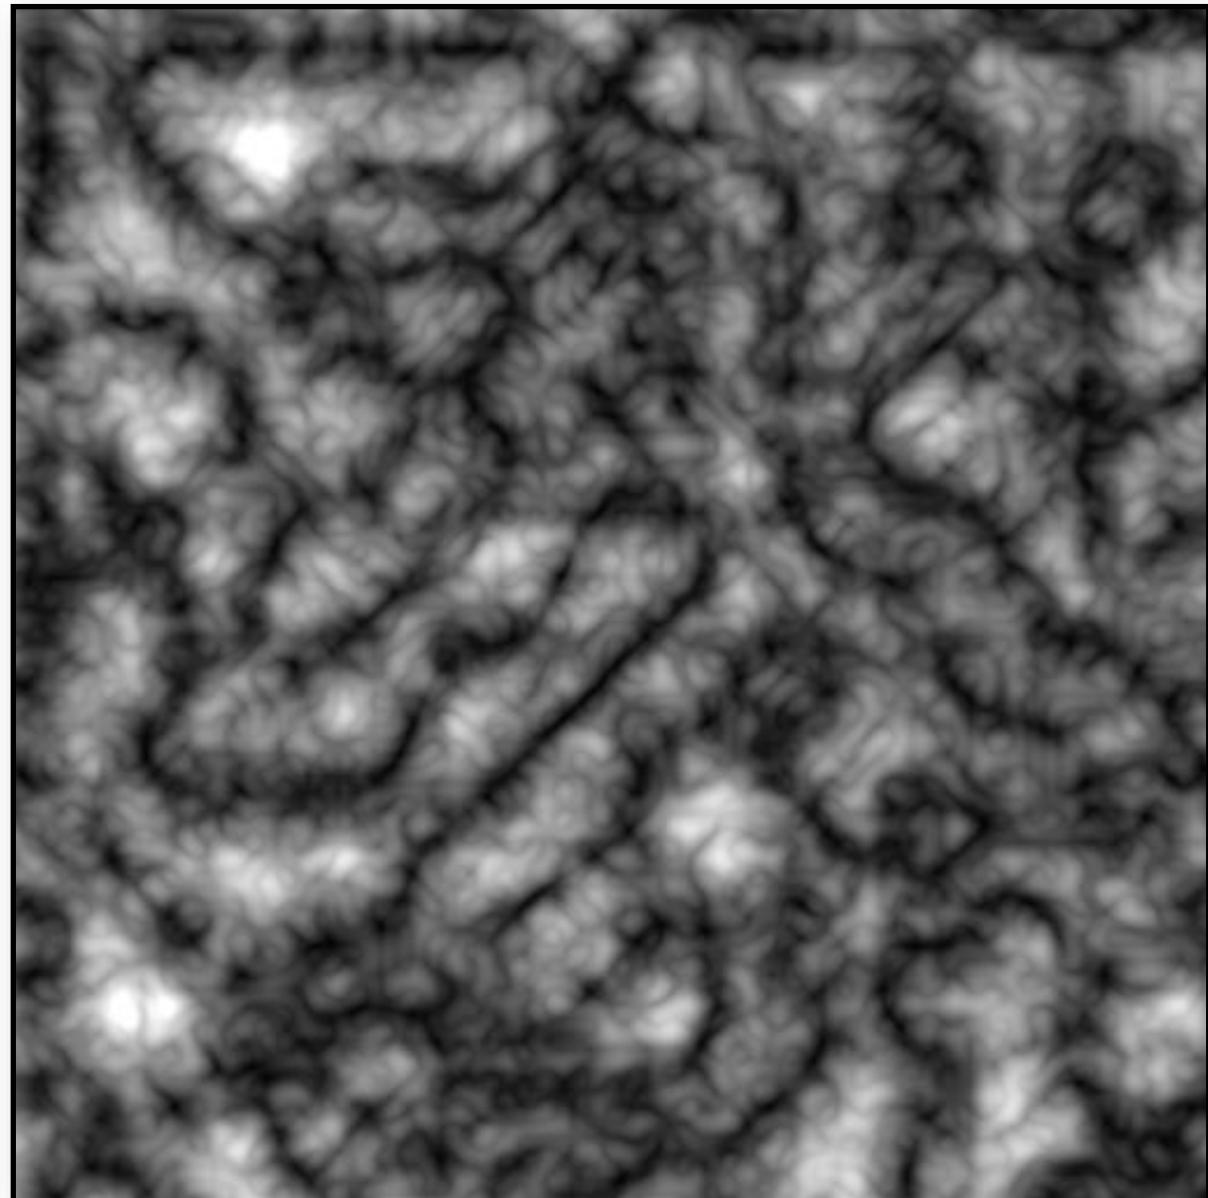
- Same as fBm, but sum **absolute value** of noise function



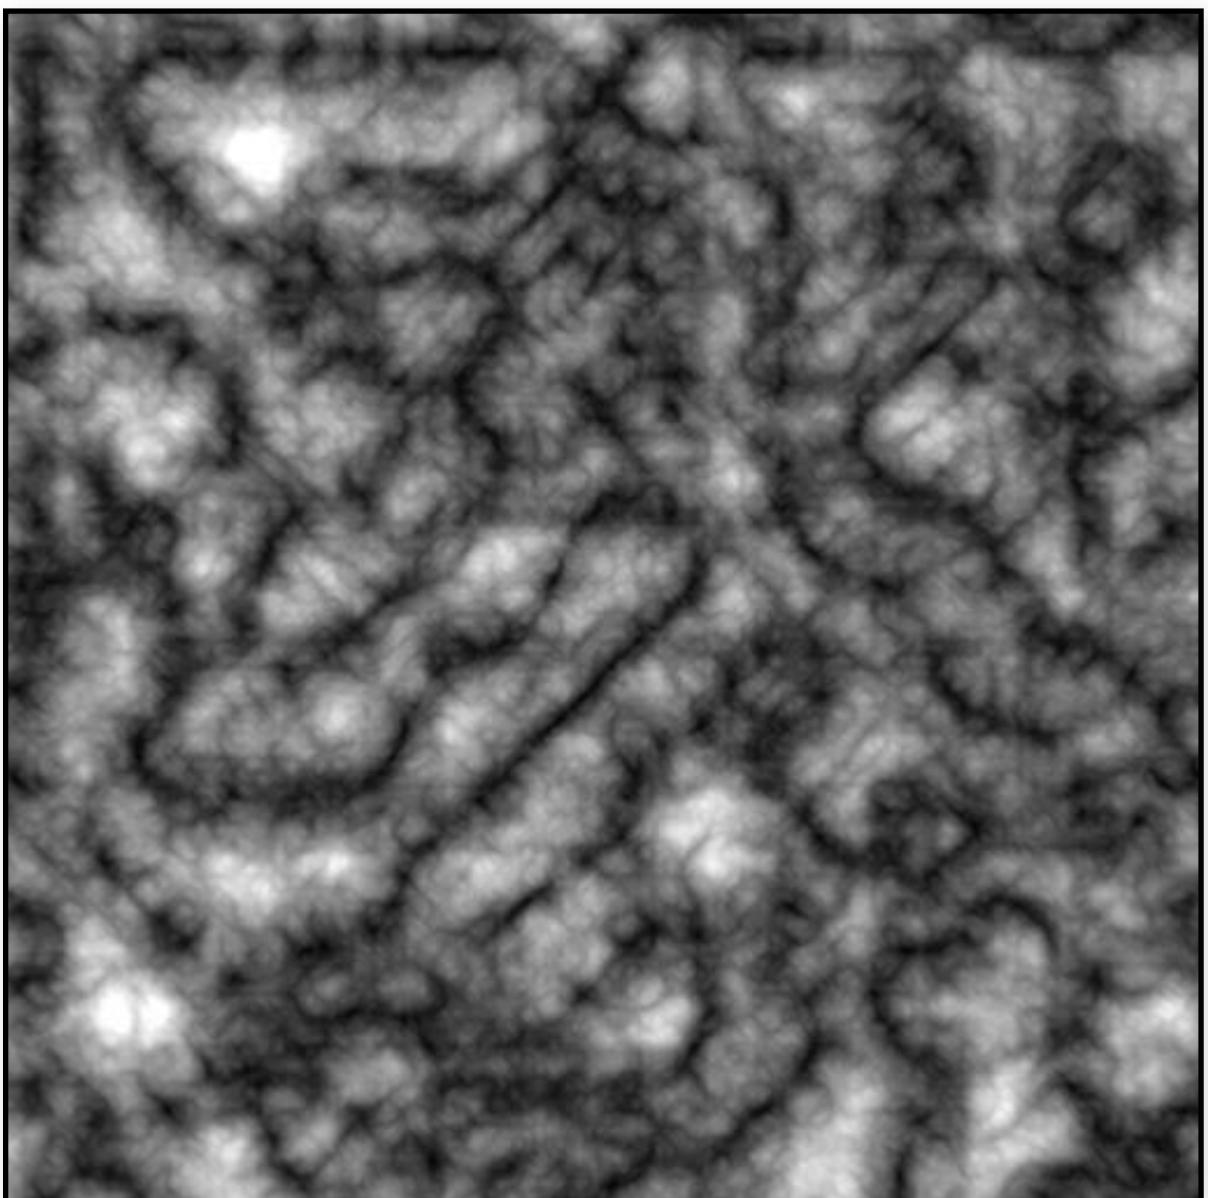
1 octaves



2 octaves



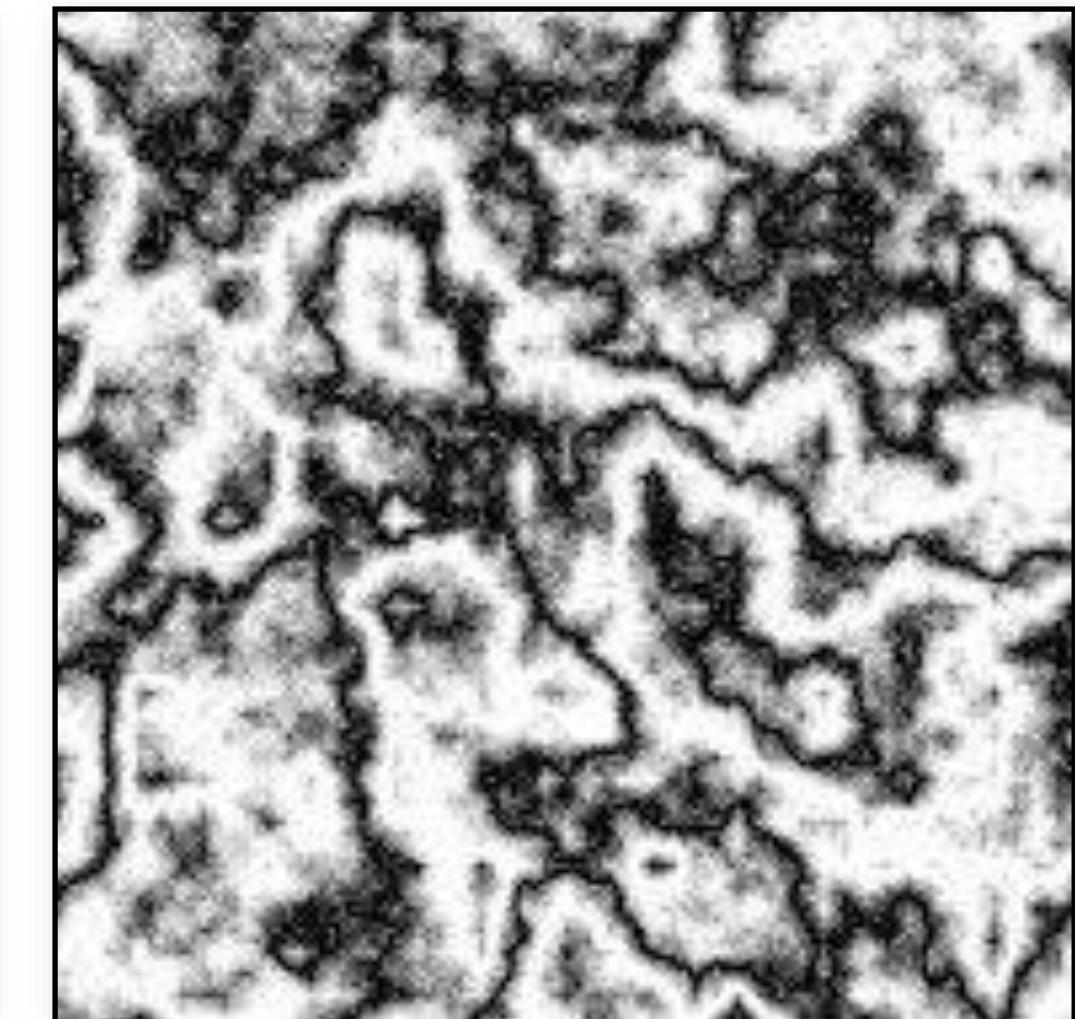
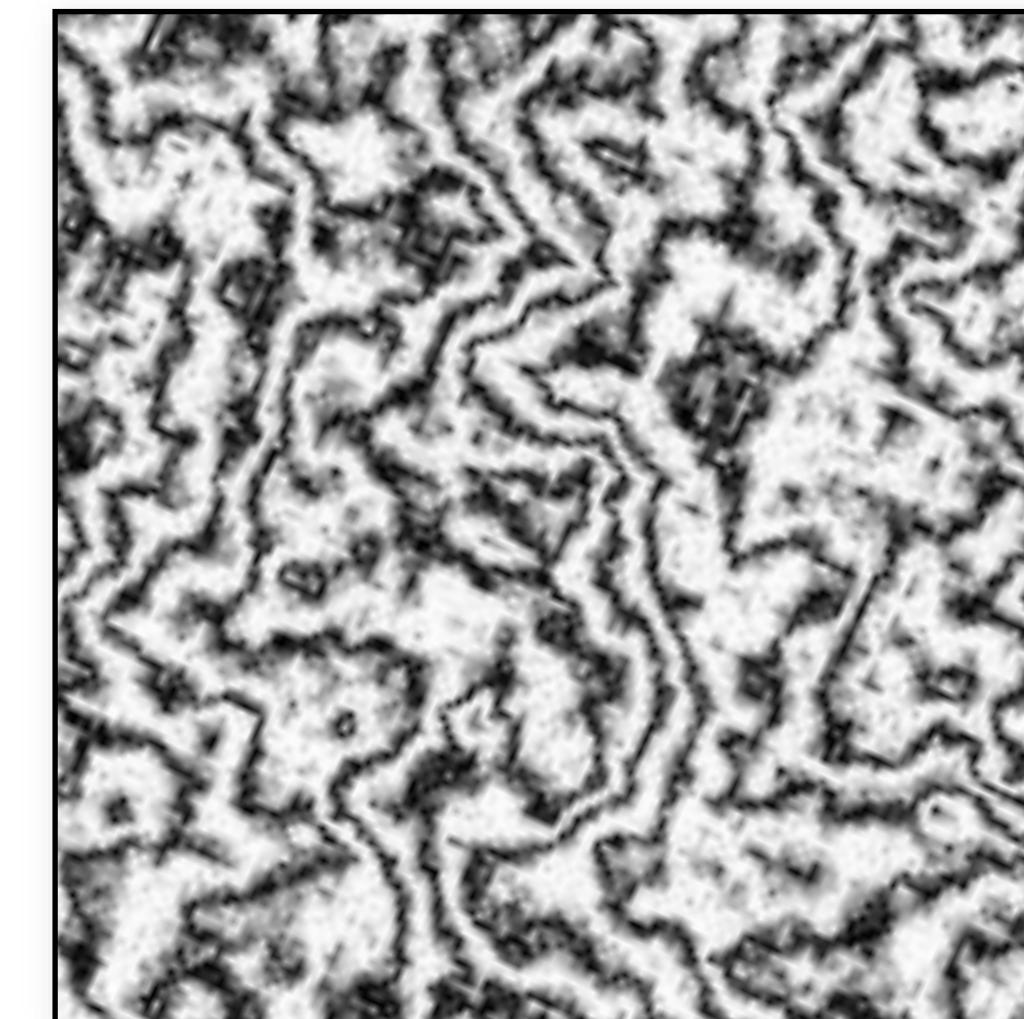
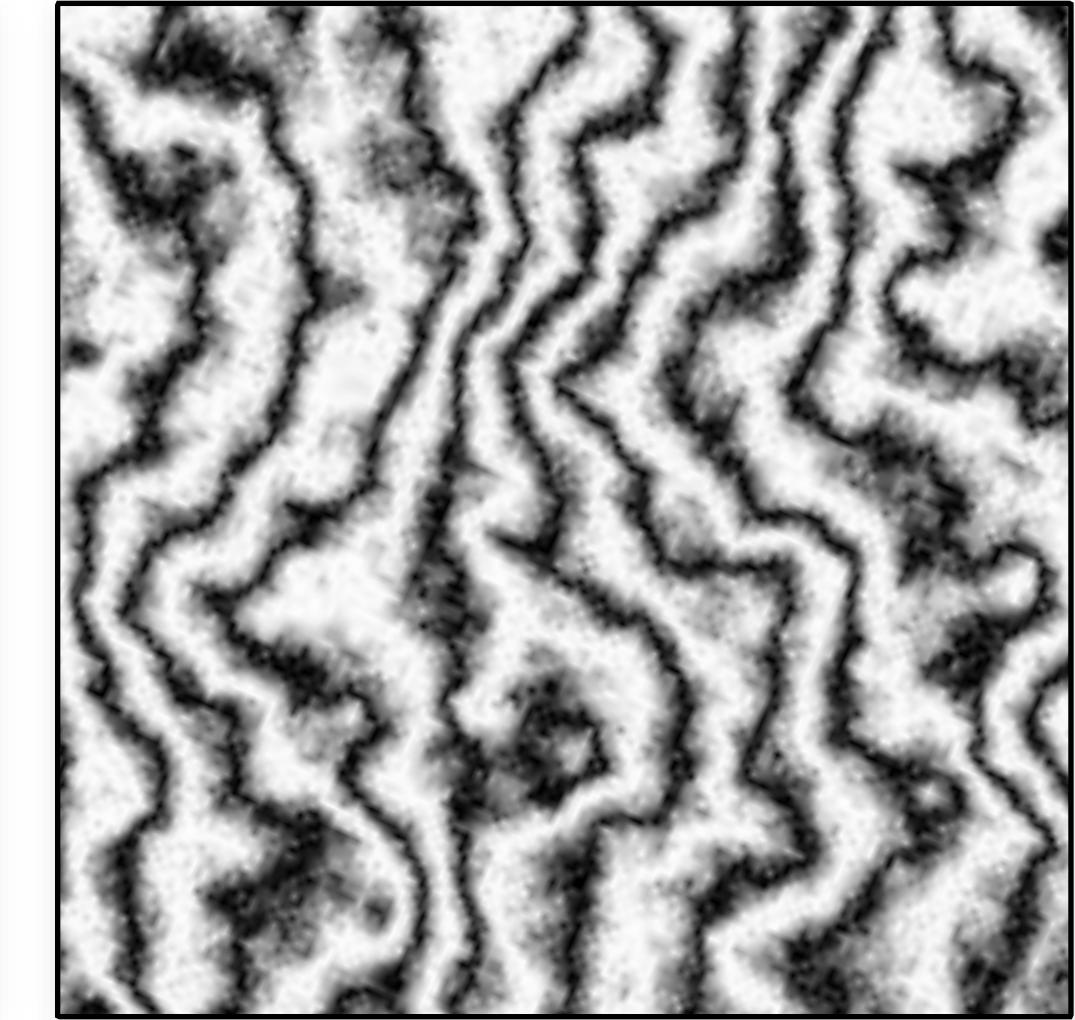
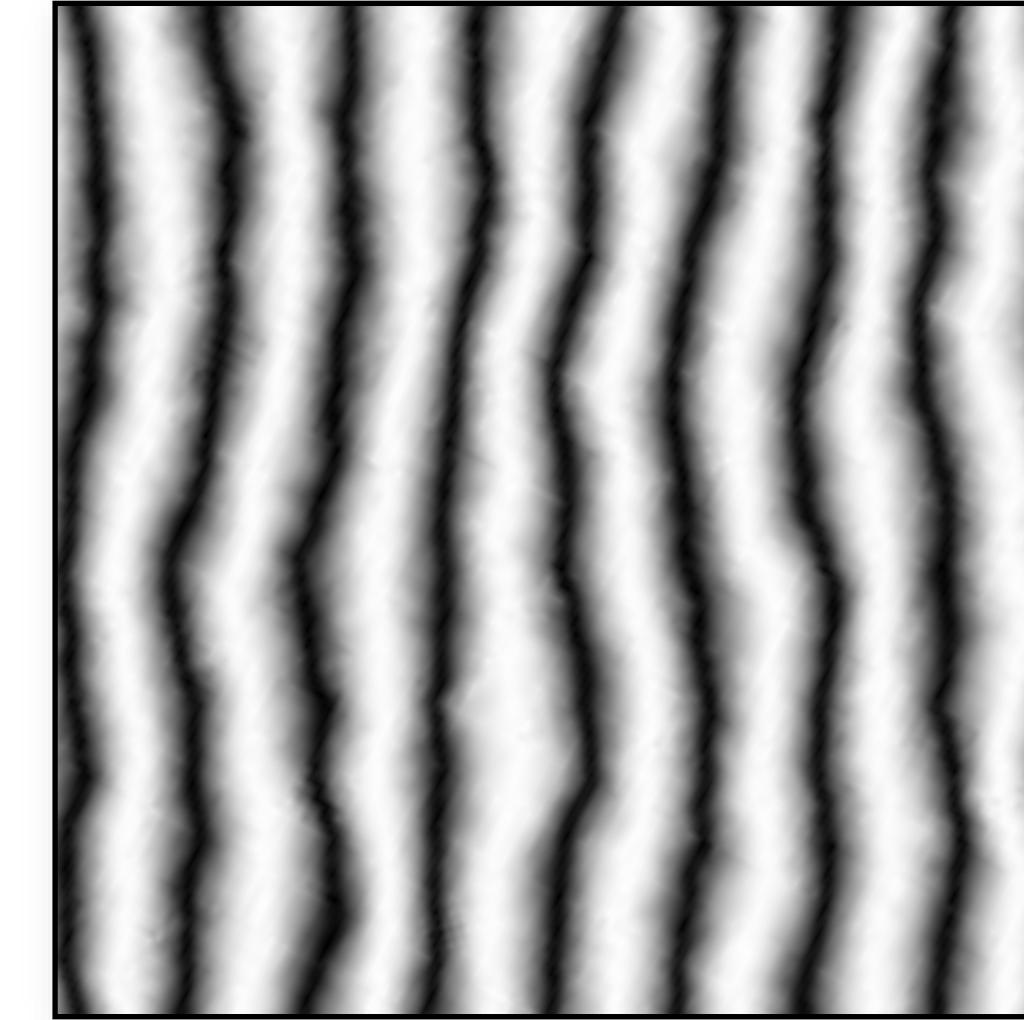
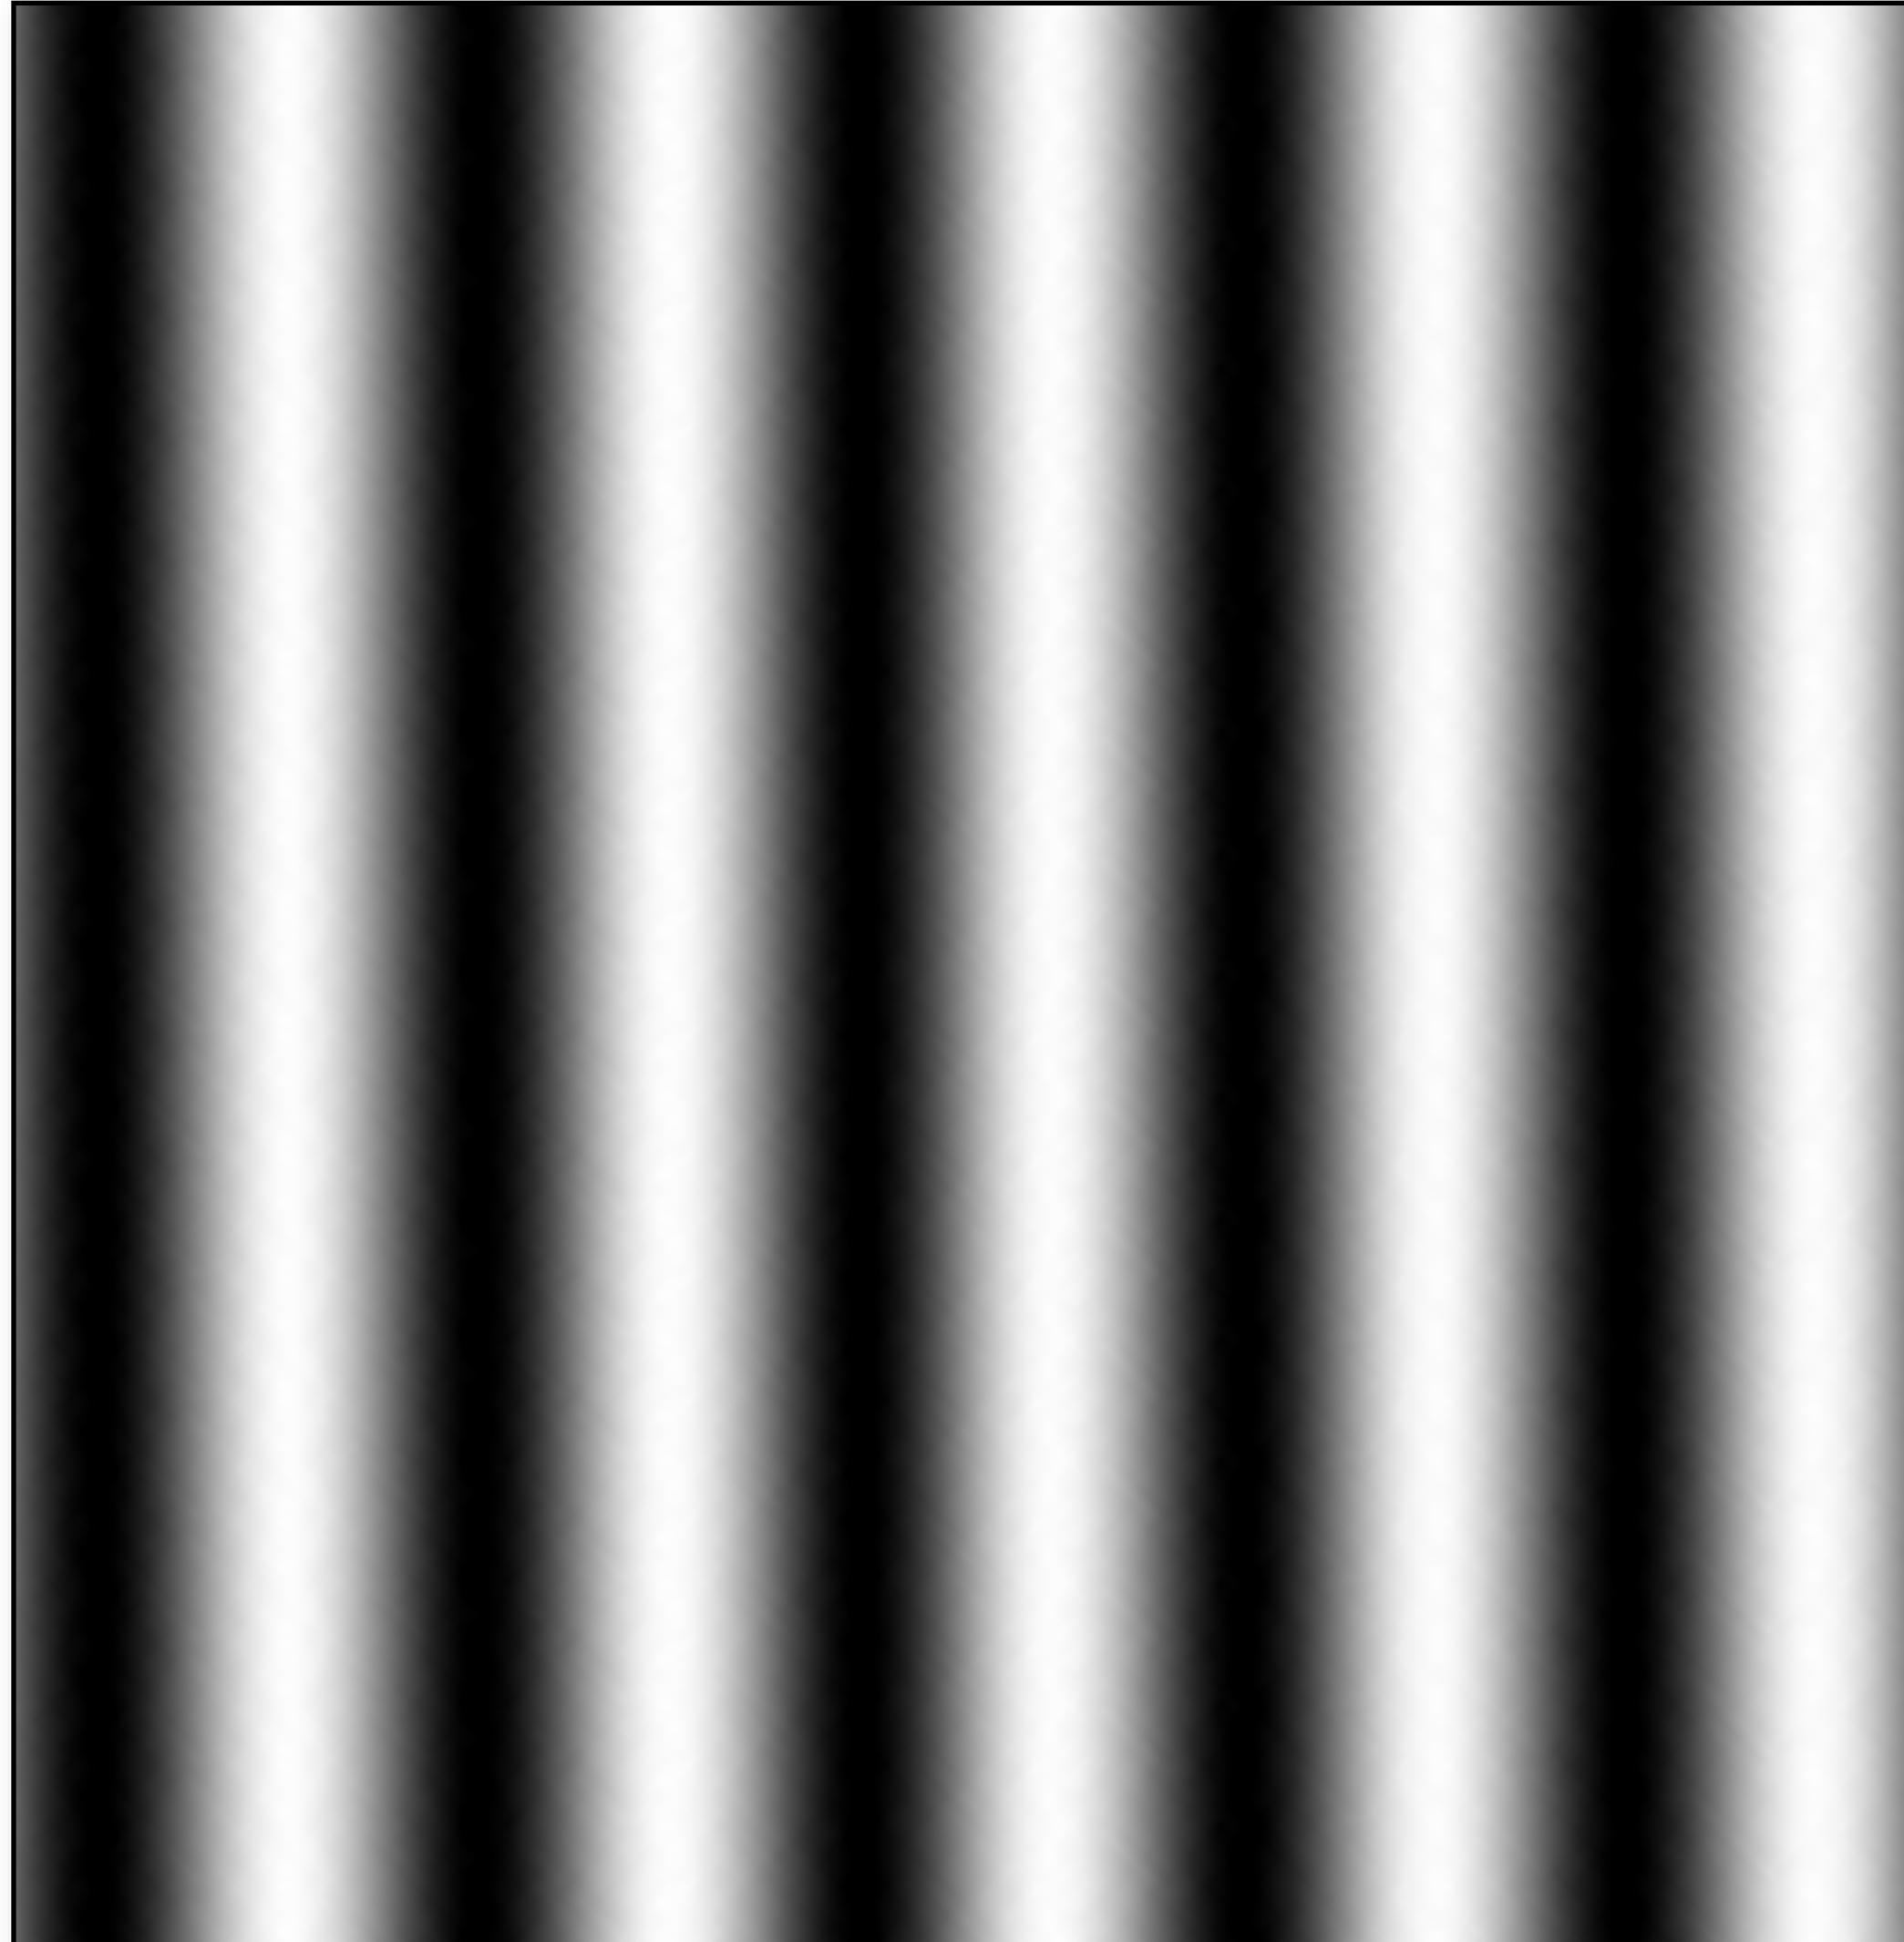
3 octaves



4 octaves

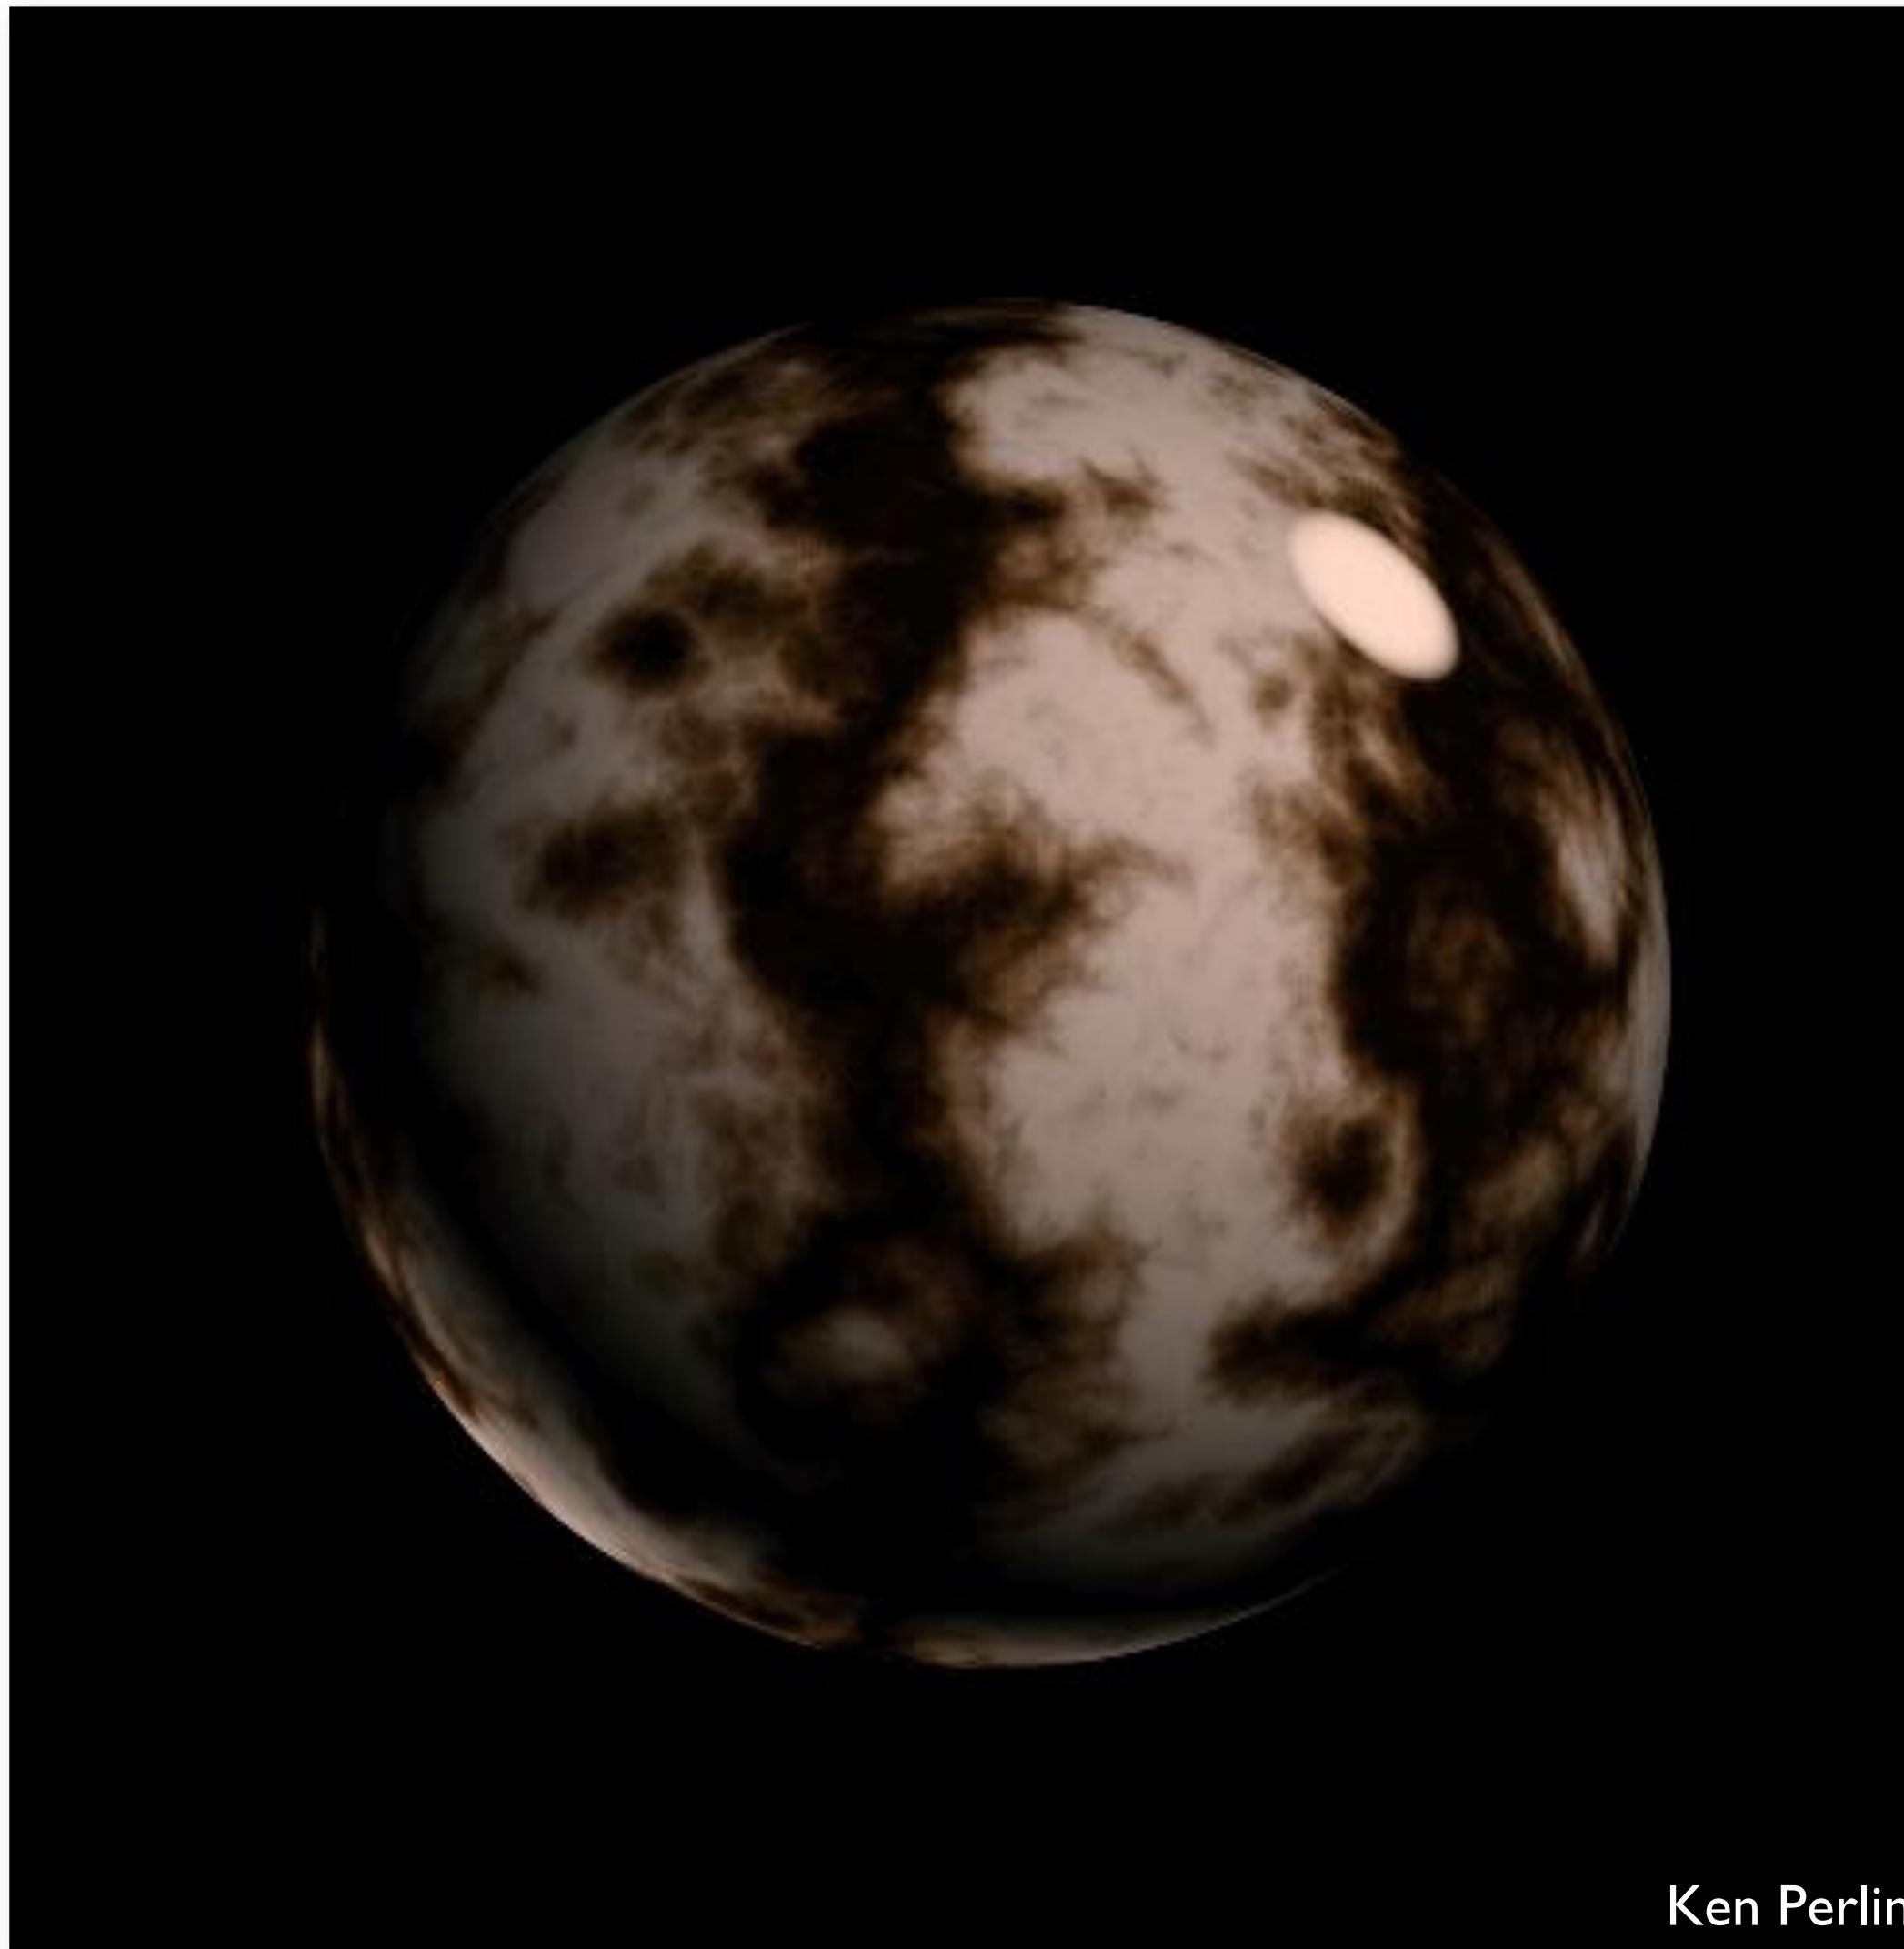
Marble

$(1 + \sin(k_1 p_x + \text{turbulence}(k_2 p))) / w) / 2$



Marble

$$(1 + \sin(k_1 p_x + \text{turbulence}(k_2 p))) / w) / 2$$



Ken Perlin

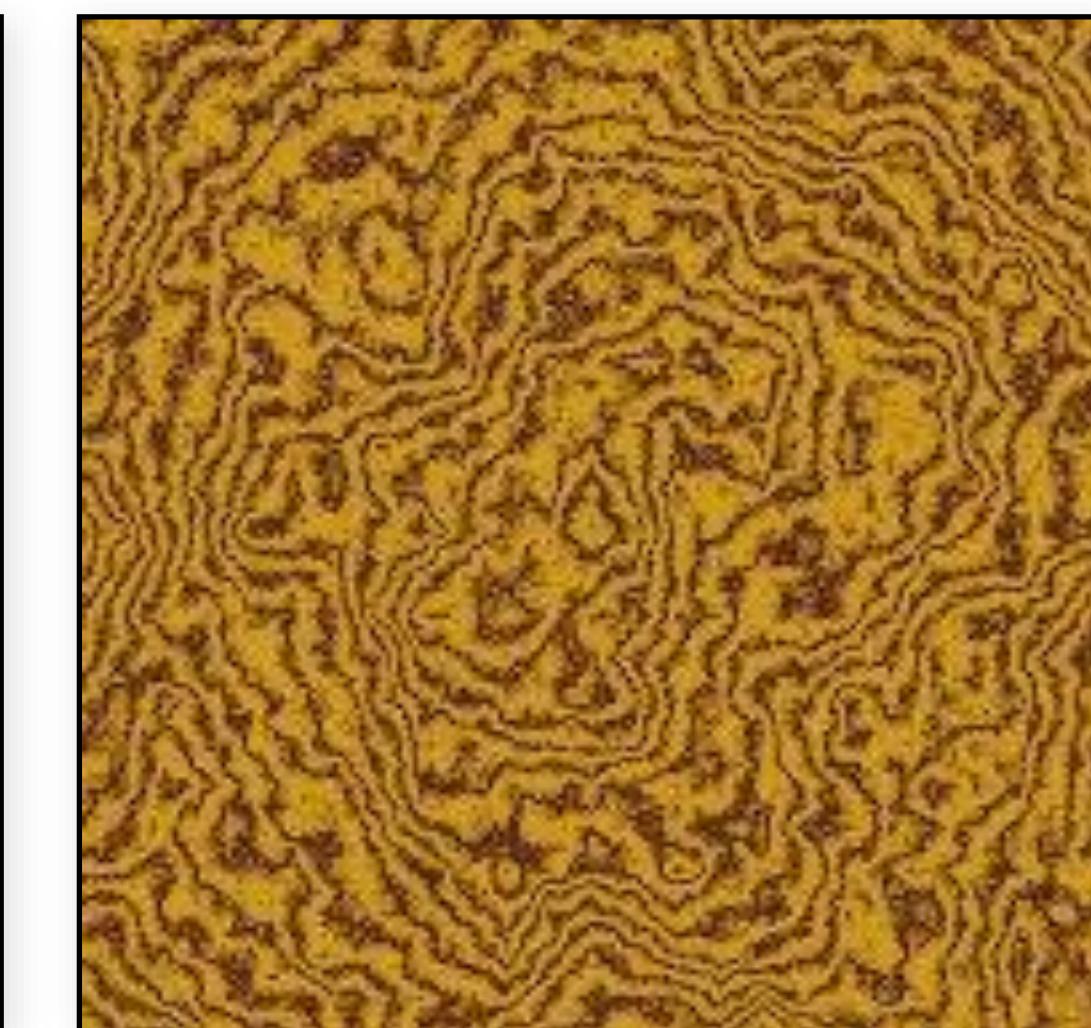
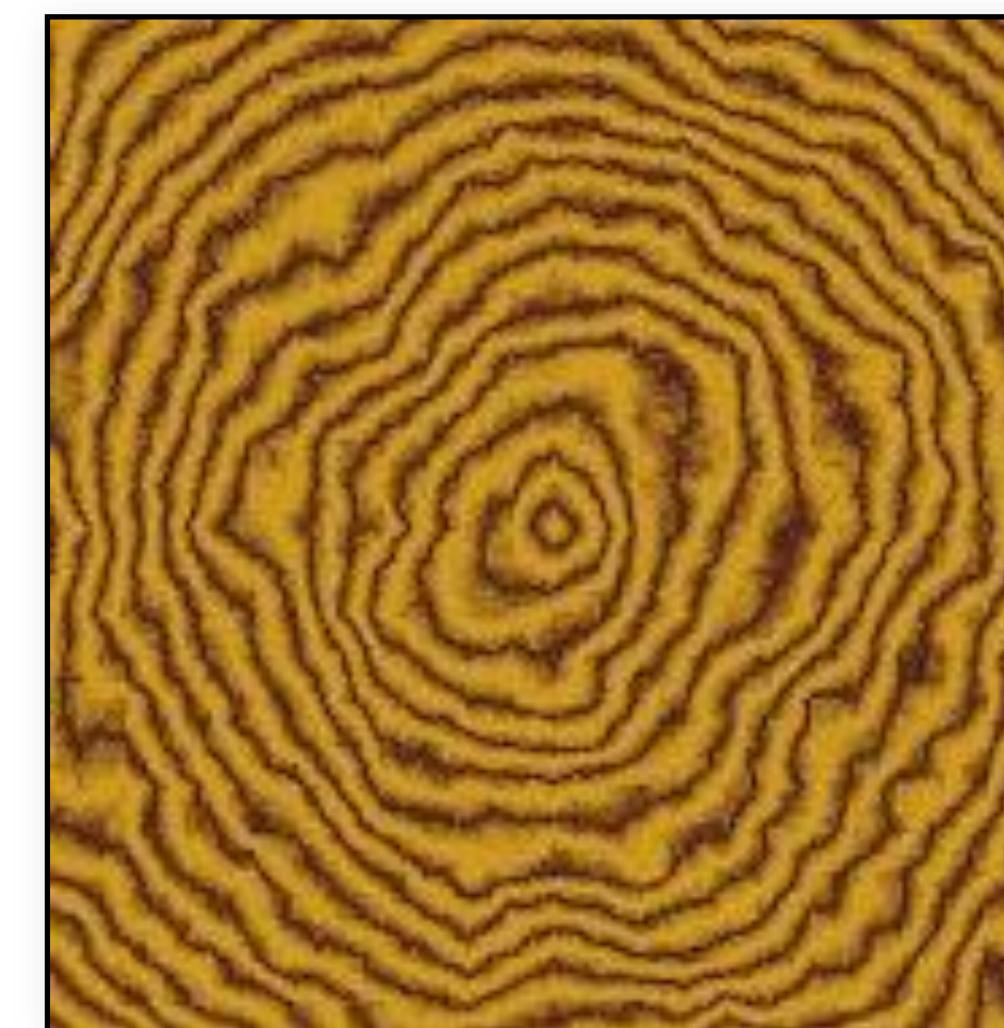
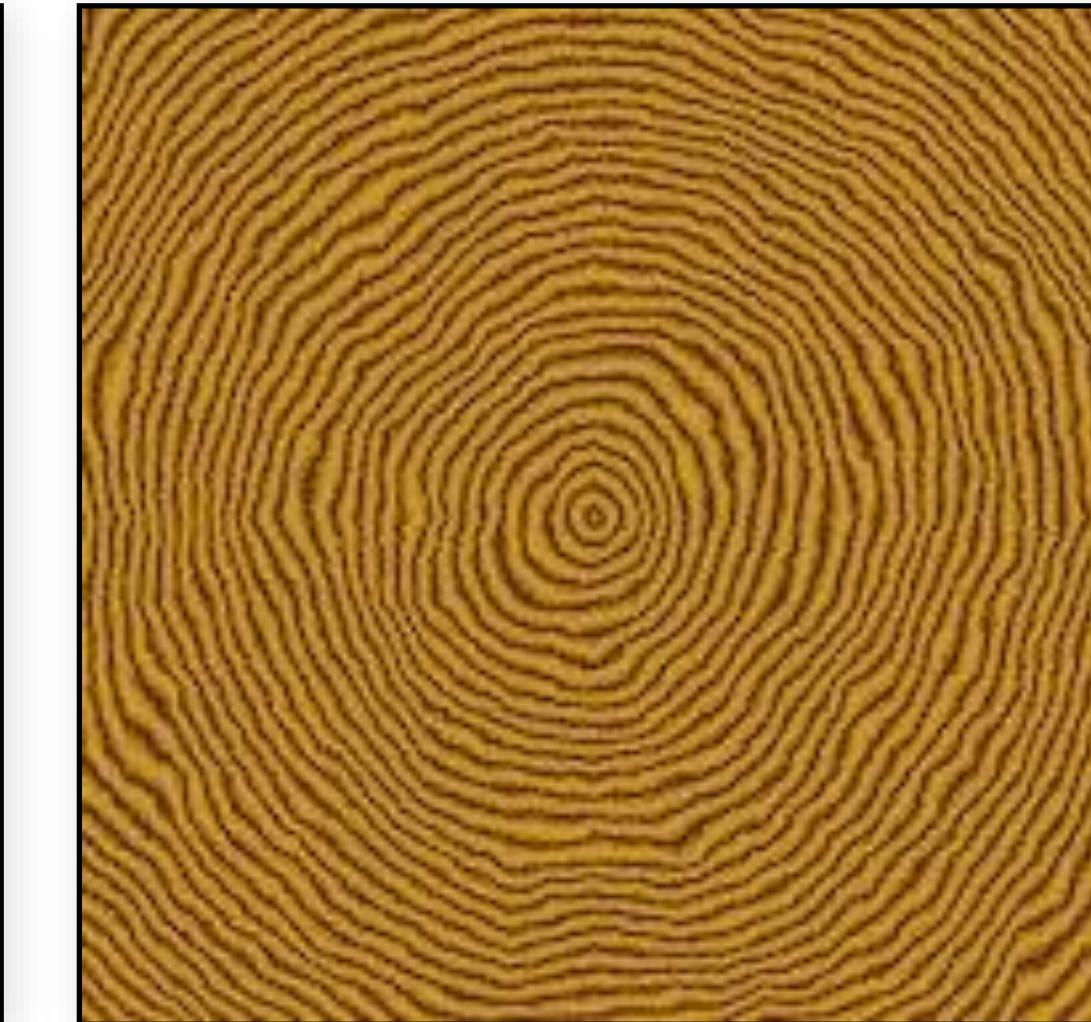
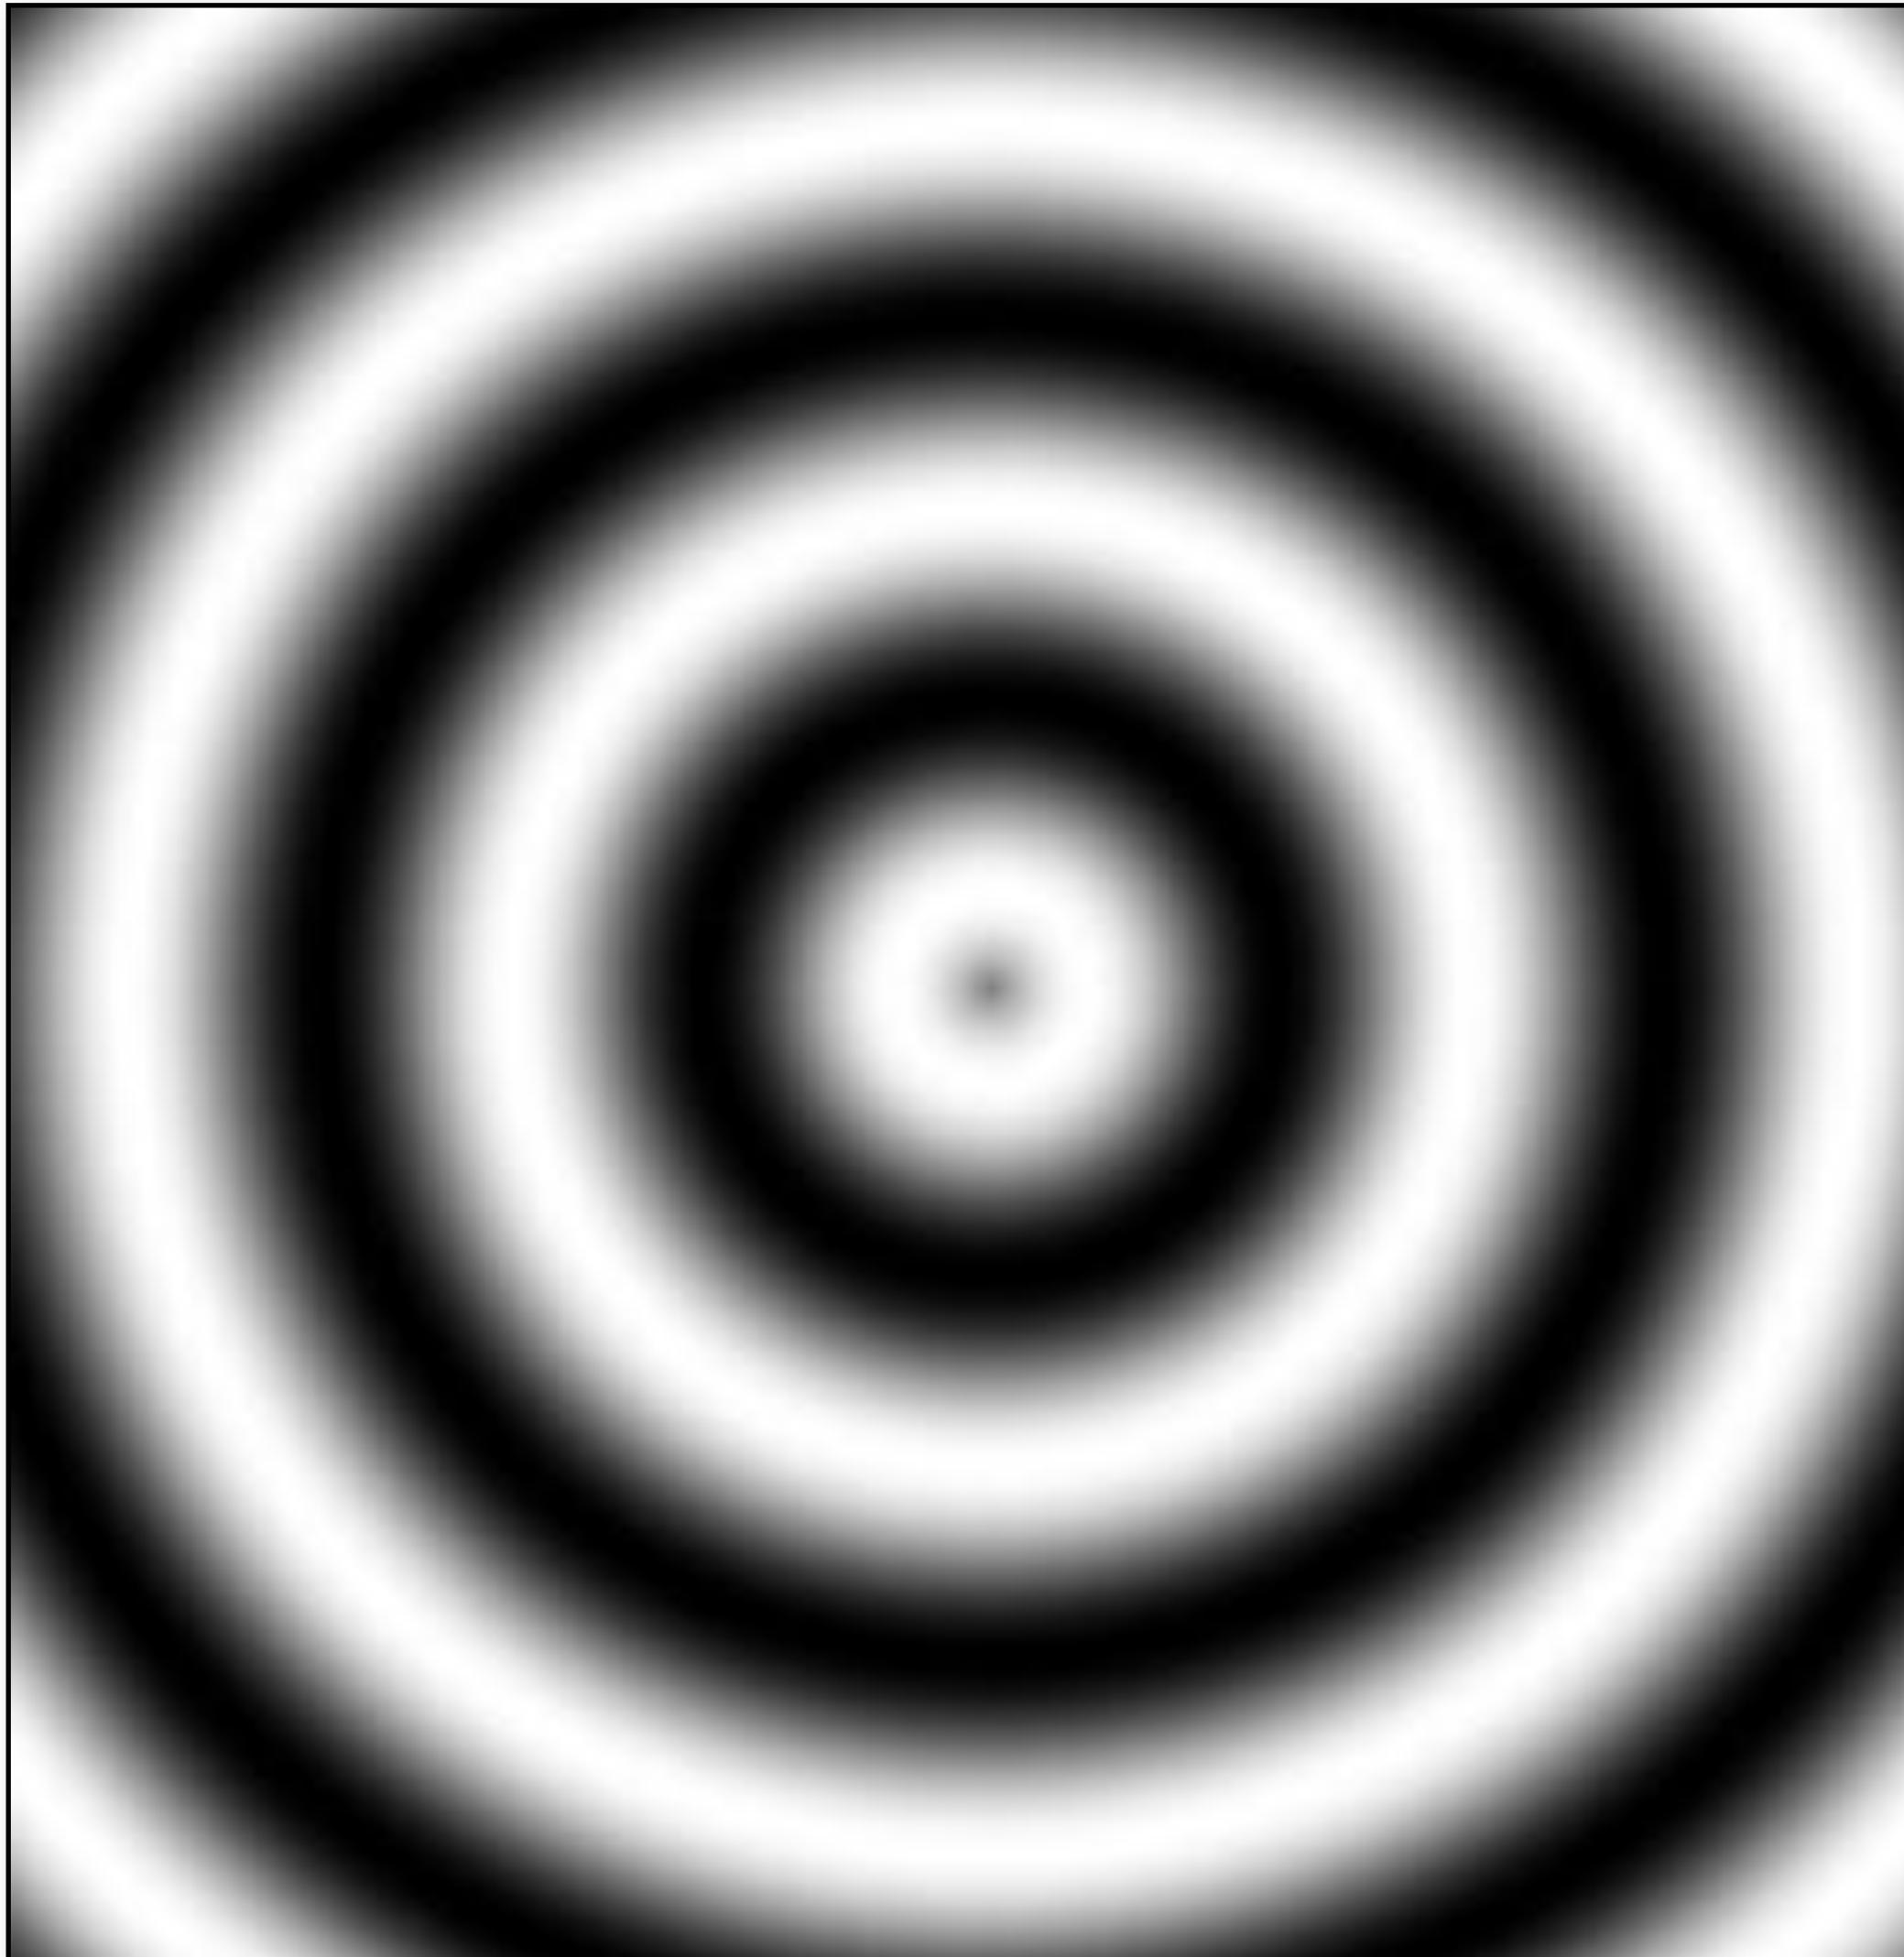


Ken Perlin



Wood

$(1 + \sin(\sqrt{p_x^2 + p_y^2}) + fBm(p))) / 2$

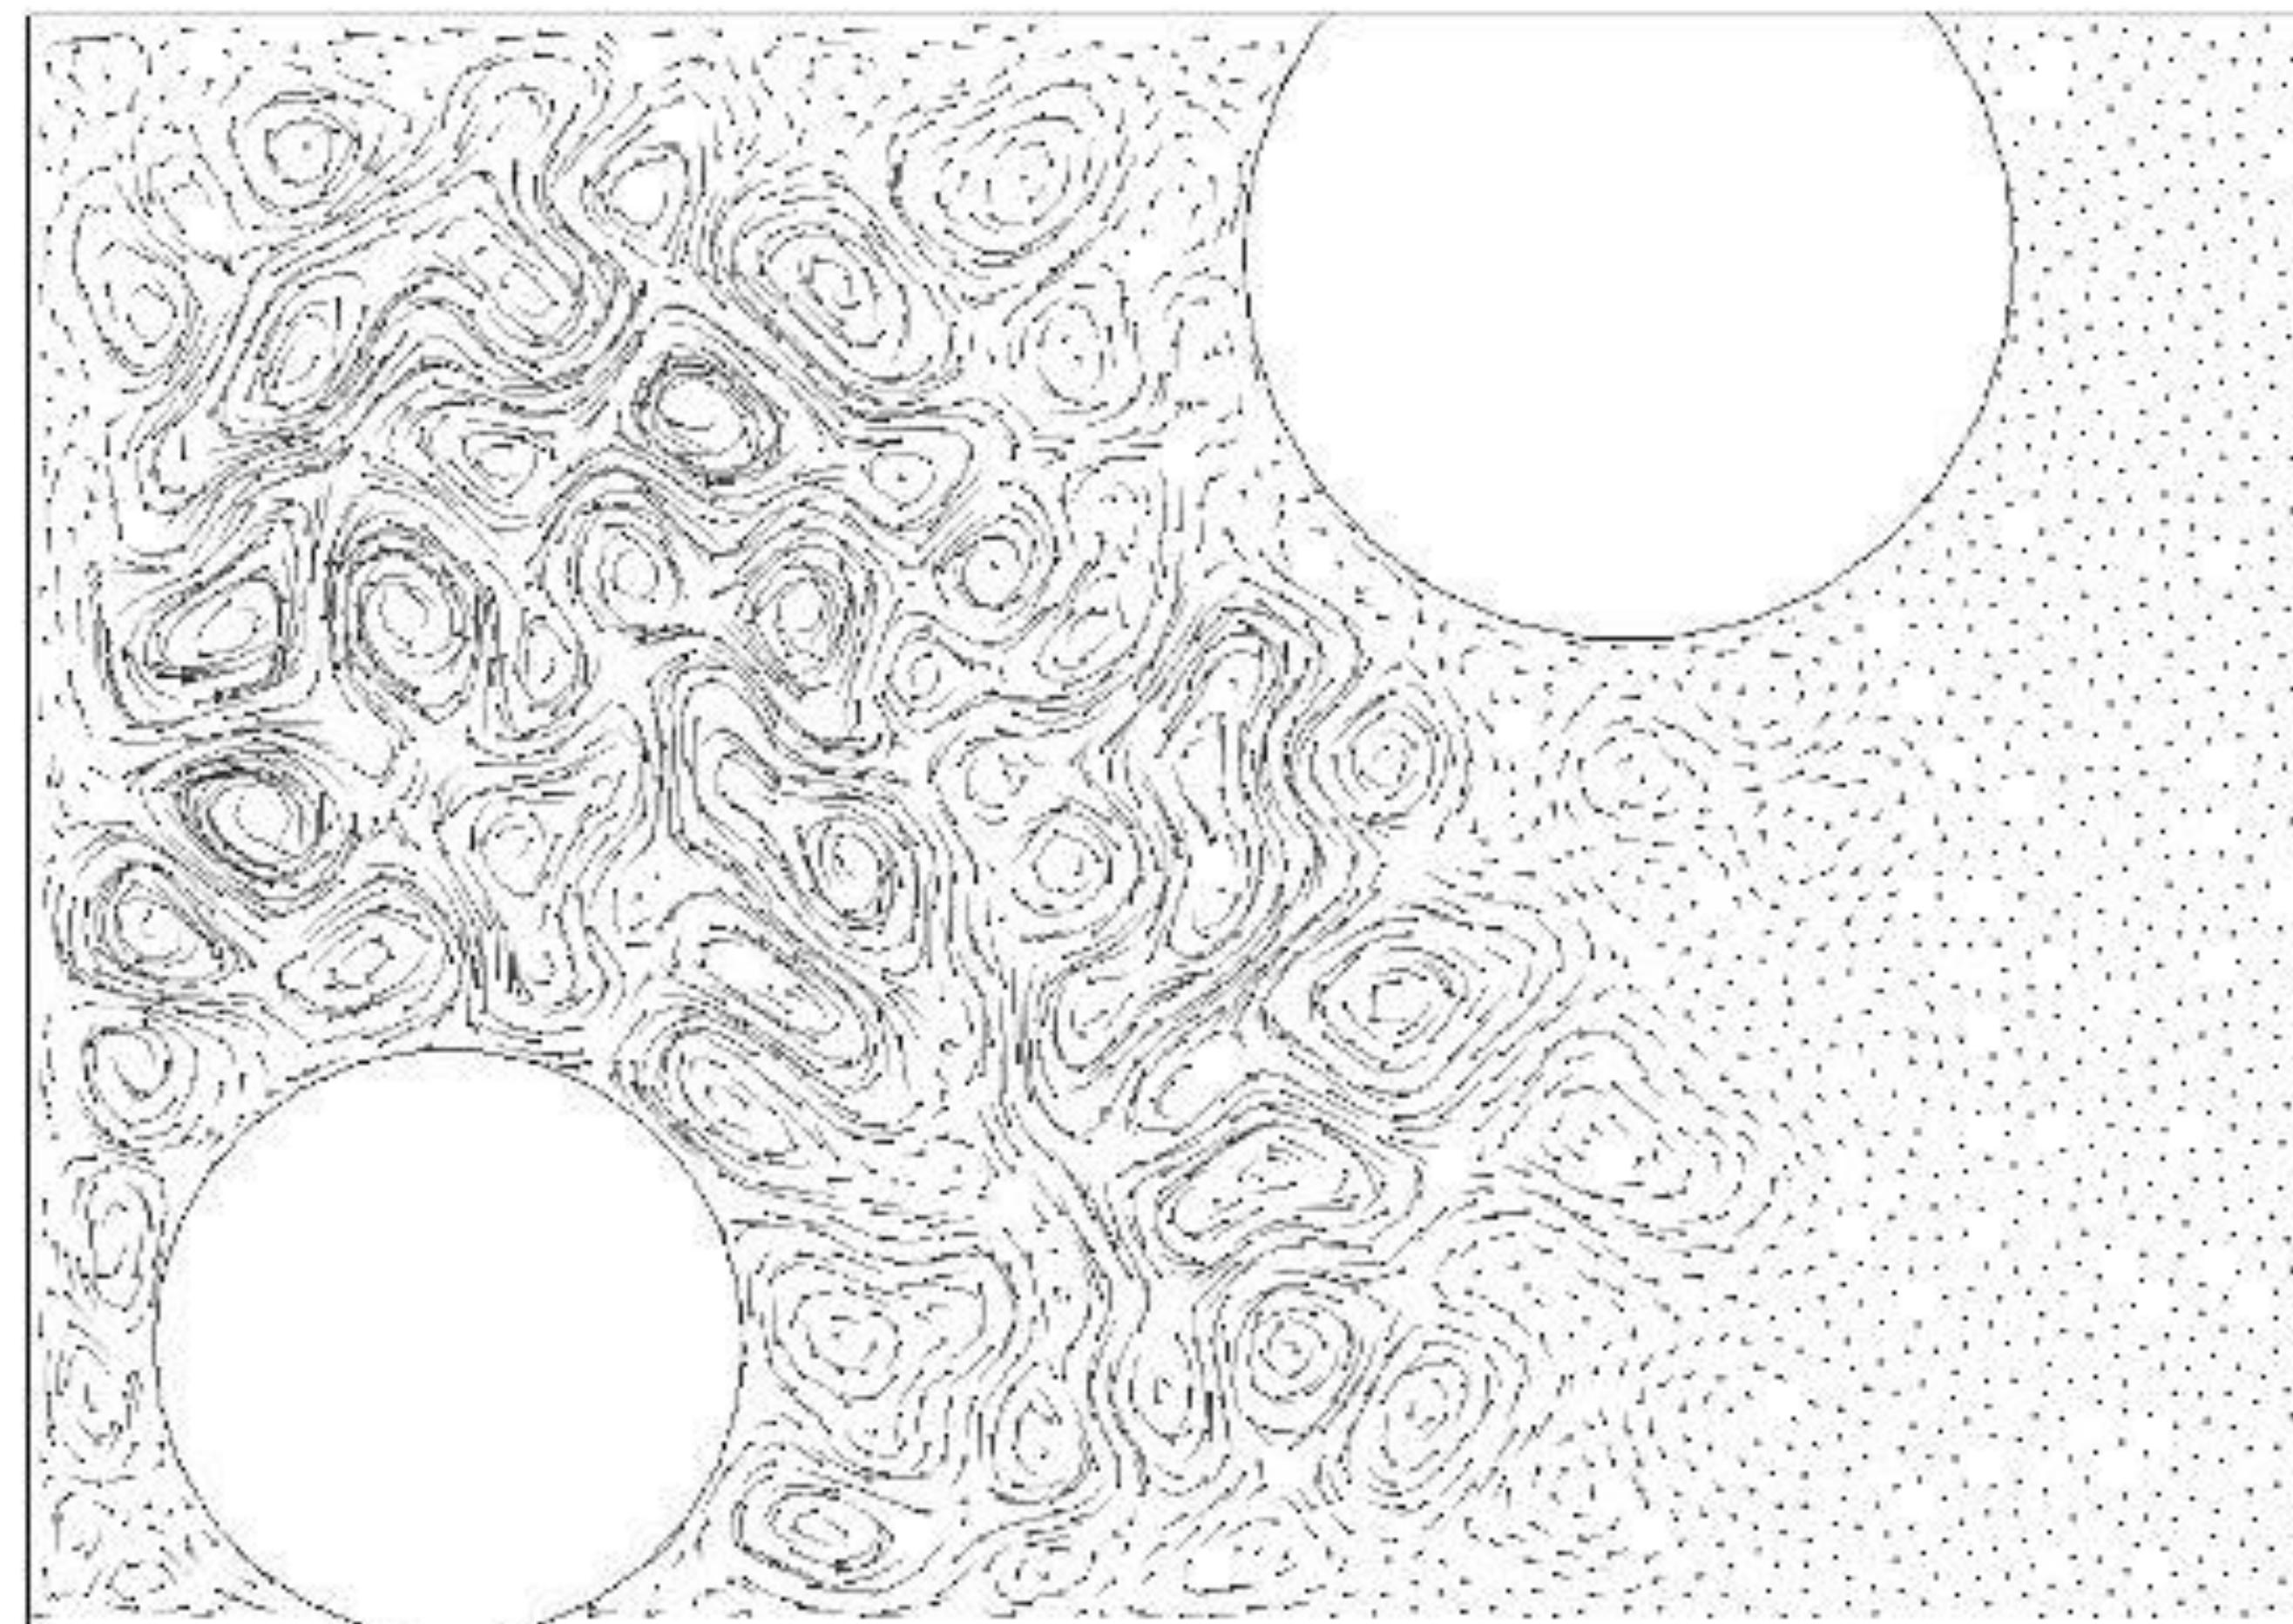


Wood

$$(1 + \sin(\sqrt{p_x^2 + p_y^2}) + fBm(p))) / 2$$



Curl Noise for Animation



- Curl noise for procedural fluid flow, R. Bridson, J. Hourihan, and M. Nordenstam, Proc. ACM SIGGRAPH 2007.

Worley Noise

- “Cellular texture” function
 - Introduced in 1996 by Steve Worley
- Randomly distribute “feature points” in space
$$f_n(x) = \text{distance to } n^{\text{th}} \text{ closest point to } x$$

